


Artificial Intelligence

Steve James

Search



Search

- Search = find solution to reach goal
- Automated theorem proving:
 - Start with axioms, find steps to reach theorem
- Checkers/Chess
 - Start at position, find moves to win

Definitions

- **Agent**: a **decision-making** entity
 - E.g. robot, software code
- **State**: a **representation** of agent's **world**/environment at a given time
- **State space**: the space of all possible states
- Start in an **initial state**
- **Successor function**: what **actions** agent can take at a given time, the **cost** of doing so, and **what happens next**
- **Goal test**: function that, given state, returns true if **a goal condition** has been met

Search problem


- Consists of:

- A state space



- Successor function

- Start state




- Goal test: $test(state) = \begin{cases} true & \text{if dots all eaten} \\ false & \text{otherwise} \end{cases}$

- Solution is sequence of actions (**plan**) that **transforms start state into a goal state**
- Optimal** solution: **minimises total cost** from start to goal

Defining the problem

- Challenge is taking problem to be solved and deciding what the **state space** should be
- A state need only keep details necessary for planning
 - i.e. can be **abstracted**
- Problem: navigation
 - States: **xy-location**; actions: NSEW; successor: update location; goal test: $(x, y) = \text{TARGET}$
- Problem: eat all dots:
 - States: **xy-location** + **dot booleans**; actions: NSEW; successor: update location and dot boolean; goal test: dots all false



State space = locations $\times 2^{\text{dots}}$

Formal Definition

- Set of states S
- Start state $s \in S$
- Set of actions A and rules $a(s) \rightarrow s'$
- Goal test $g(s) = \{0, 1\}$
- Cost function $C(s, a, s') \rightarrow \mathbb{R}^+$
- Search problem: $\langle S, s, A, g, C \rangle$

Problem statement

Find a sequence of actions a_1, \dots, a_n
and corresponding states s_1, \dots, s_n such that:

$$s_0 = s$$

$$s_i = a_i(s_{i-1}), i = 1, \dots, n$$

$$g(s_n) = 1$$

while minimising

$$\sum_{i=1}^n C(s_{i-1}, a_i, s_i)$$

Search tree

- A tree that shows **future outcomes** of actions!

- **Root** is **start** state



- **Children** are **successor** states


- Edges from root to node in tree

- Edges = **plan**!

- Sum of edge costs = **cost of plan**

- For most problems, we can't build the full tree!

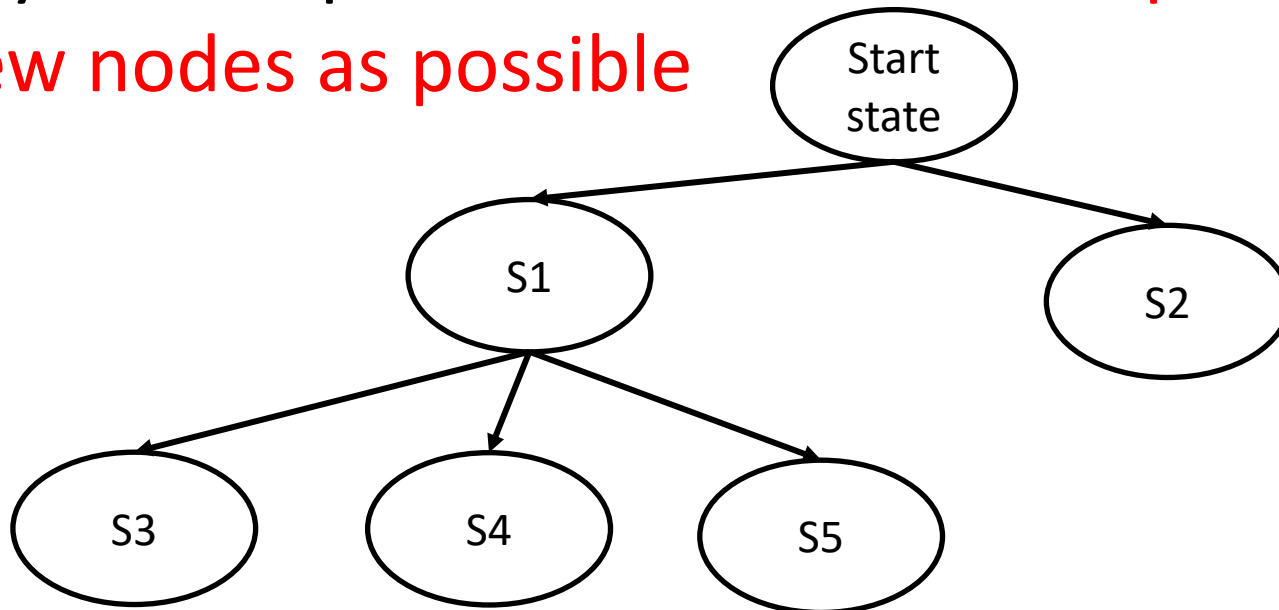
- $O(b^d)$ where d is the max plan length



This is the challenge!

Planning with search tree

- **Expand** tree nodes (potential plans)
- Maintain **frontier** of partial plans under consideration
- Try come up with solution **while expanding as few nodes as possible**



General tree search

```
function TREE-SEARCH( problem, strategy ) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy Here is where we can be clever!
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:

- Frontier
- Expansion
- Exploration strategy

*Implementation detail:
Remember visited states.
Don't add them more than
once to tree!*

- Which frontier nodes to explore?

Search strategies


- Tree search algorithms are about picking **order of node expansion**
- Things to consider:
 - **Completeness**: will it always find a **solution** if it exists?
 - **Optimality**: will it always find a **least-cost** solution?
 - **Time** complexity: **number** of **nodes** generated
 - **Space** complexity: max nodes in **memory**
- Complexity depends on max **branching factor**, **depth** of optimal solution, **max depth** of state space

Uninformed search

- **Uninformed** search: we use only the info in the **problem definition**
- If we had additional **domain knowledge**, we could use **informed/heuristic** search (later)
- General uninformed strategies:
 - **Breadth-first** search
 - **Uniform-cost** search
 - **Depth-first** search

Breadth-first search

- Expand **shallowest** unexpanded node
- Frontier as a **FIFO queue**
 - New successors go at **end**



Breadth-first search

```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)
```

BFS properties

- Let d be depth of **shallowest** solution
- Complete?
 - **Yes** (if b is finite). **Shallowest solution** returned
- Time?
 - $1 + b + b^2 + \dots + b^d = O(b^d)$
- Space?
 - Keeps all **frontier nodes in memory**: $O(b^d)$
- Optimal:
 - Only if costs are **constant**

BFS Issues

- Both time and space are $O(b^d)$
 - Assume:
 - we can generate 1m nodes/sec
 - 1 kbyte/node
 - $b = 10$
 - Then, for $d = 8$
 - Time is 2 mins, but memory is 102GB
 - For $d = 12$
 - Time is 13 days, but memory is 1000TB
- Memory is massive issue!*

Uniform-cost search

- BFS finds plan with **shortest length**
 - But what if cost of plan is **not optimal**? i.e. a **longer plan may have smaller cost** overall
- UCS almost same as BFS, but use **priority queue** instead of queue
 - Each node in queue ordered by **cost to node**
- BFS = UCS when costs are constant everywhere
- Very similar to **Dijkstra's** algorithm (but Dijkstra doesn't have a goal)

Uniform-cost search

function UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

node \leftarrow a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

frontier \leftarrow a priority queue ordered by PATH-COST, with *node* as the only element

explored \leftarrow an empty set

loop do

if EMPTY?(*frontier*) **then return** failure

node \leftarrow POP(*frontier*) /* chooses the lowest-cost node in *frontier* */

if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

add *node*.STATE to *explored*

for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

child \leftarrow CHILD-NODE(*problem*, *node*, *action*)

if *child*.STATE is not in *explored* or *frontier* **then**

frontier \leftarrow INSERT(*child*, *frontier*)


else if *child*.STATE is in *frontier* with higher PATH-COST **then**

replace that *frontier* node with *child*

*Priority queue instead
of queue*


*Extra check: if path
to child shorter than
previous one found,
update it!*

BFS vs UCS



BFS

Solution cost = C^ ,
 costs $\geq \epsilon$*




UCS

UCS properties


- Expands all nodes with **cost less than cheapest solution**
- If solution is costs C^* and each edge is ε , then **effective depth** is approx C^* / ε
- Complete?
 - Yes, if cost $\geq \varepsilon$ (**positive costs**) and best solution has finite cost
- Time?
 - $O(b^{\text{ceiling}(C^* / \varepsilon)})$
- Space?
 - Keeps all **frontier nodes in memory**: $O(b^{\text{ceiling}(C^* / \varepsilon)})$ *Again* 😞
- Optimal:
 - Yes! Nodes expanded in **increasing order of total cost**

Depth-first search

- Expand **deepest** unexpanded node
- Frontier as a **stack**
 - New successors go at **front**
 - Or **recursive** function



DFS



DFS properties

- We use special trick for **visited states**
 - Only remember states **along path from root to current node**
- Complete?
 - Yes, for finite spaces
- Time?
 - $O(b^m)$ - **terrible if $m \gg d$**
- Space?
 - Remember only **path from root to current node** (and **unexplored siblings** along path: $O(bm)$ *Finally!* 😊)
- Optimal:
 - No, finds “**leftmost**” solution!

Depth-limited DFS

- Solution might be at **finite depth**, but if $m = \infty$,
 - DFS will never find it!
- Solution: just **limit max depth** to l
 - Time is now $O(b^l)$ and space is $O(bl)$
- But **how to pick l** ? And what **if $l < d$** ?
 - Could use **domain knowledge** e.g. if you know solution is at most k , pick $l = k$
 - Or use **diameter**: max shortest distance between any 2 states

Iterative deepening

- Instead, we can just try **multiple depths!**
- i.e. RUN DFS with $l = 1$
 - If solution, great! If not, run it again with $l = 2$
 - **Keep going** until you find solution!
- Combines the benefits of **DFS + BFS**

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```

- Seems **wasteful!**
 - But is it?

IDS Analysis

- If we ran **depth-limited DFS** to depth d

$$N_{\text{DLS}} = b^0 + b^1 + b^2 + \dots + b^d$$

- If we run **IDS** $l = 0, 1, \dots, d$

$$N_{\text{IDS}} = (d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + 2b^{d-1} + b^d$$

- For $b = 10, d = 5$

- $N_{\text{DLS}} = 111\ 111$

- $N_{\text{IDS}} = 123\ 456$

- 11% difference!


*Most work happens
in lowest levels*

- Rule of thumb: if search space is **large** and **depth** of solution **unknown**, use **IDS**

IDS properties

- Complete?
 - Yes, it's like BFS
- Time?
 - $(d + 1)b^0 + db^1 + (d - 1)b^2 + b^d = O(b^d)$
- Space? *Like DFS!*
 - $O(bd)$
- Optimal:
 - Only if **step costs are same** everywhere (like BFS)

Bidirectional Search



Bidirectional Search

- Run searches **forward and backward** until intersection
 - $O(b^{d/2}) + O(b^{d/2}) \ll O(b^d)$
- Extra requirements:
 - **Invert** action rules (sometimes hard)
 - Not always **unique**
 - Need **single solution**
- When to stop:
 - Candidate solution when **frontier nodes intersect**
 - First solution may not be optimal – must check **possible shortcuts**

Summary of methods

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}


Figure 3.21 Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

Now

- What if we know something about the problem?
 - How can we incorporate knowledge?
 - **Task-specific** expansion strategy?

What does UCS suggest?

- Cost-so-far: how much it costs to get to a node
 - UCS says: cheapest node first!
- What remains?
 - Total-cost = Cost-so-far + Cost-to-go



From UCS to heuristics


- Recall: UCS is like BFS with **priority queue**
- Nodes ordered by priority from smallest to largest
 - Priority is **cost estimate** $f(n)$ *Evaluation function*
- In UCS, $f(n) = g(n)$, where g is cost of reaching n from root
- **Heuristic function** $h(n)$ = estimate of cost from n to goal
- g is **backward cost** (start to node), h is (**estimated**) **forward cost** (node to goal)

Greedy best-first search

- UCS is $f(n) = g(n)$
- GBFS is $f(n) = h(n)$
 - Order nodes by estimated cost to goal
- Where does this estimate come from?
 - Domain knowledge!
- e.g. In navigating task with location x

$$h(x) = |x - g|_2 \text{ Euclidean distance}$$


Example



*If allowing
repeated
nodes, will
be stuck!
(incomplete)*

*Euclidean
distance
heuristic*

Example




*Euclidean
distance
heuristic*

GBFS properties

- Incomplete for tree search (i.e. repeated nodes allowed)
 - Complete for graph search in finite space
- Space and time complexity are both $O(b^m)$
 - For max search depth
- But can be reduced substantially depending on heuristic and problem

A* search


- UCS: $f(n) = g(n)$
- GBFS: $f(n) = h(n)$



- Ideally: explore in order of $g(n) + CostToGo(n)$
- A*: explore in order $f(n) = g(n) + h(n)$
 - h is estimate of future cost
- Implement A* as UCS with different priority!

Optimality of A*

- Depends on heuristic



- Fails! Bad cost goal < estimated good cost

Heuristics in A*

- h admissible if $h(n) \leq \text{TrueFutureCost}(n)$
 - We **never overestimate** the cost!
 - $f(n) = g(n) + h(n)$ means we never overestimate cost from start to goal through n
 - **Optimistic**: estimates cost as less than it actually is
- h is consistent if $h(n) \leq c(n, a, n') + h(n')$
 - All **consistent heuristics are admissible**
- A* **tree search** is optimal if h is **admissible**
- A* **graph search** is optimal if h is **consistent**

Triangle inequality


Properties of A*

- A* is **complete, optimal**
- A* is **optimally efficient** w.r.t. heuristic
 - No other algorithm will expand fewer nodes
- Time complexity is $O(b^{\epsilon d})$
 - ϵ is **relative error** of heuristic, d is solution depth
- Alternate view, $O((b^*)^d)$ where b^* is **effective branching factor**
 - A* doesn't need to consider certain nodes (**prunes**) and so **branching factor is reduced!**
- Since exponential, **memory** is biggest **issue**
 - Can do **iterative deepening equivalent** (IDA*) *Like BFS*

Creating heuristics

- Main challenge for A*: how to make **good heuristics** that are **admissible/consistent**?
- For navigation, straight-line path is good
 - Can never be faster than that!
- Could **learn heuristics** from data (e.g. **machine learning**/precomputed database lookups)
- Use **relaxations** – solve an easier, **less constrained** version of a problem
 - E.g. Relaxed version of a maze → remove all the walls!

8-puzzle relaxation



- **Relaxation 1:** allow tiles to be placed anywhere directly
 - i.e. heuristic is number of misplaced tiles
 - Clearly underestimates true cost of moving tiles
- **Relaxation 2:** allow tiles to be slid around, ignoring other tiles
 - i.e. heuristic is sum of distances of each tile to its final location

Effect of heuristics

	Average nodes expanded when the optimal path has...		
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
Relax 1	13	39	227
Relax 2	12	25	73

- **Tradeoff** between quality of estimate and work per node
 - Closer heuristic is to true cost, fewer nodes are expanded...
 - But more work to compute heuristic per node!