

Artificial Intelligence


Steve James

Game Theory & Adversarial Search

What is game theory?

- Field involving games, answering such questions as:
 - How should you play games?
 - How do most people play games?
 - How can you create a game that has certain desirable properties?

What is a game?




What is a game?

- Game is defined by:
 - Initial state
 - $Player(s)$: decision-making entities
 - $Actions(s)$: available actions
 - $Result(s, a)$: successor function or transition model
 - $Terminal(s)$: is the game over/state terminal
 - $Utility(s, p)$: the value for the game ending in s for player p

Why study game theory in AI?

- Making good **decisions** \subseteq AI
- Making good decisions in games \subseteq Game Theory
- AI often created for situations that can be thought of as games



Types of games

- Sequential




- Simultaneous



Types of games

- Constant-sum



- Variable-sum



Classic game theory

- 2-player, one-turn, simultaneous-move games

	R	P	S
R	$\frac{1}{2}, \frac{1}{2}$	0, 1	1, 0
P	1, 0	$\frac{1}{2}, \frac{1}{2}$	0, 1
S	0, 1	1, 0	$\frac{1}{2}, \frac{1}{2}$

Strategy

- **Strategy:** A specification of what to do in every single non-terminal state of the game
 - Functions from states to (probability distributions over) legal actions
 - Pure vs. Mixed
- Examples:
 - Trading: I'll accept an offer of R2m or higher, but not lower
 - Chess: Full lookup table of moves and actions to make

Best response

- What's the **best strategy** in rock-paper-scissors?
 - Depends on what the **other player** is doing!
 - If we knew it, then we could choose the best strategy (**optimisation**)
 - But we don't know what they want!
 - How to reason when we don't know opponent's strategy

Dominated strategies

- Strategy s is dominated by s^* if s^* always gives higher payoff

	C	D
C	3, 3	0, 5
D	5, 0	1, 1

Iterated dominance

	L	C	R
U	6, 1	1, 0	6, 2
M	1, 4	0, 5	5, 5
D	3, 4	4, 3	2, 0

Iterated Dominance

- Iterated Elimination of Dominated Strategies (IEDS)
 - Won't always produce a unique solution
 - Common Knowledge of Rationality (CKR)
 - “Faithful Approach”

Conservative approach: Maximin

- Ensures **best worst-case** scenario

	L	C	R
U	6, 1	1 , 0	6, 2
M	1, 4	0 , 5	5, 5
D	3 , 4	4, 3	2 , 0

Nash equilibrium

- **Strategy profile**: specification of strategies for all players
- **Nash equilibrium**: strategy profile such that players are mutually best-responding
- In other words: From a NE, **no player** can do **better by switching** strategies alone

Stag hunt

- Strategy s is dominated by s^* if s^* always gives higher payoff

	B	S
B	2, 2	2, 0
S	0, 2	3, 3

Also: play B with prob $\frac{1}{3}$ is an NE!

Diplomacy/Society



E

R



E

R

Economy ↑ planet ↓	Economy ↓ planet ↓
Economy ↓ planet ↓	Economy → planet ↑

Properties

- There is always **at least one** NE
 - Might be **mixed**
- If **IEDS** produces **unique solution**, it is a **NE**

Now

- Let's consider finding pure strategies in
 - Sequential
 - Alternating
 - Constant-sum (zero-sum)
 - Many-turn
 - Perfect information

Games are big

- Tic-tac-toe $\sim 10^3$
- Connect Four $\sim 10^3$
- English draughts $\sim 10^{23}$
- Othello $\sim 10^{28}$
- Chess $\sim 10^{44}$
- Shogi $\sim 10^{71}$
- # atoms in observable universe $\sim 10^{82}$
- Twixt $\sim 10^{140}$
- Go (19x19 board) $\sim 10^{170}$

*Size of state space
(reachable states)*

Zero-sum games

- Assume two players, **competitive**
 - Player 1 wins, player 2 loses and vice versa
- Game is defined by:
 - Initial state
 - $Player(s)$: whose **turn** it is
 - $Actions(s)$: available **actions**
 - $Result(s, a)$: successor function or **transition** model
 - $Terminal(s)$: is the game over/state **terminal**
 - $Utility(s, p)$: the **value** for the game ending in s for player p
- Zero sum game: sum of utilities for all players is constant: e.g. Win = +1, Draw = 0, Loss = -1

Can extend to n

Two player, zero-sum


- Two players, MAX and MIN
 - We are **MAX**, try to maximise utility
 - Opponent is **MIN**, tries to minimise utility
 - Denote $V(s)$ as utility at a given state
 - Utility is known at terminal states
 - Start at root node, expand tree
 - Players **alternate turns**
 - At level 0, us to play. Level 1, them to play, etc
 - We want to compute **optimal play**, assuming our **opponent is also optimal**
- Each level is called a ply*

Example

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

Calculating minimax

- Want to calculate $V(s)$ for all s
- If s is terminal, use utility function directly
- else if player to play is MAX:
 - Value is best **maximising value** at state
- else player to play is MIN:
 - Value is best **minimising value** at state

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

def min-value(state):

initialize $v = +\infty$

for each successor of state:


$v = \min(v, \text{value}(\text{successor}))$

return v

MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Example



Game of chance

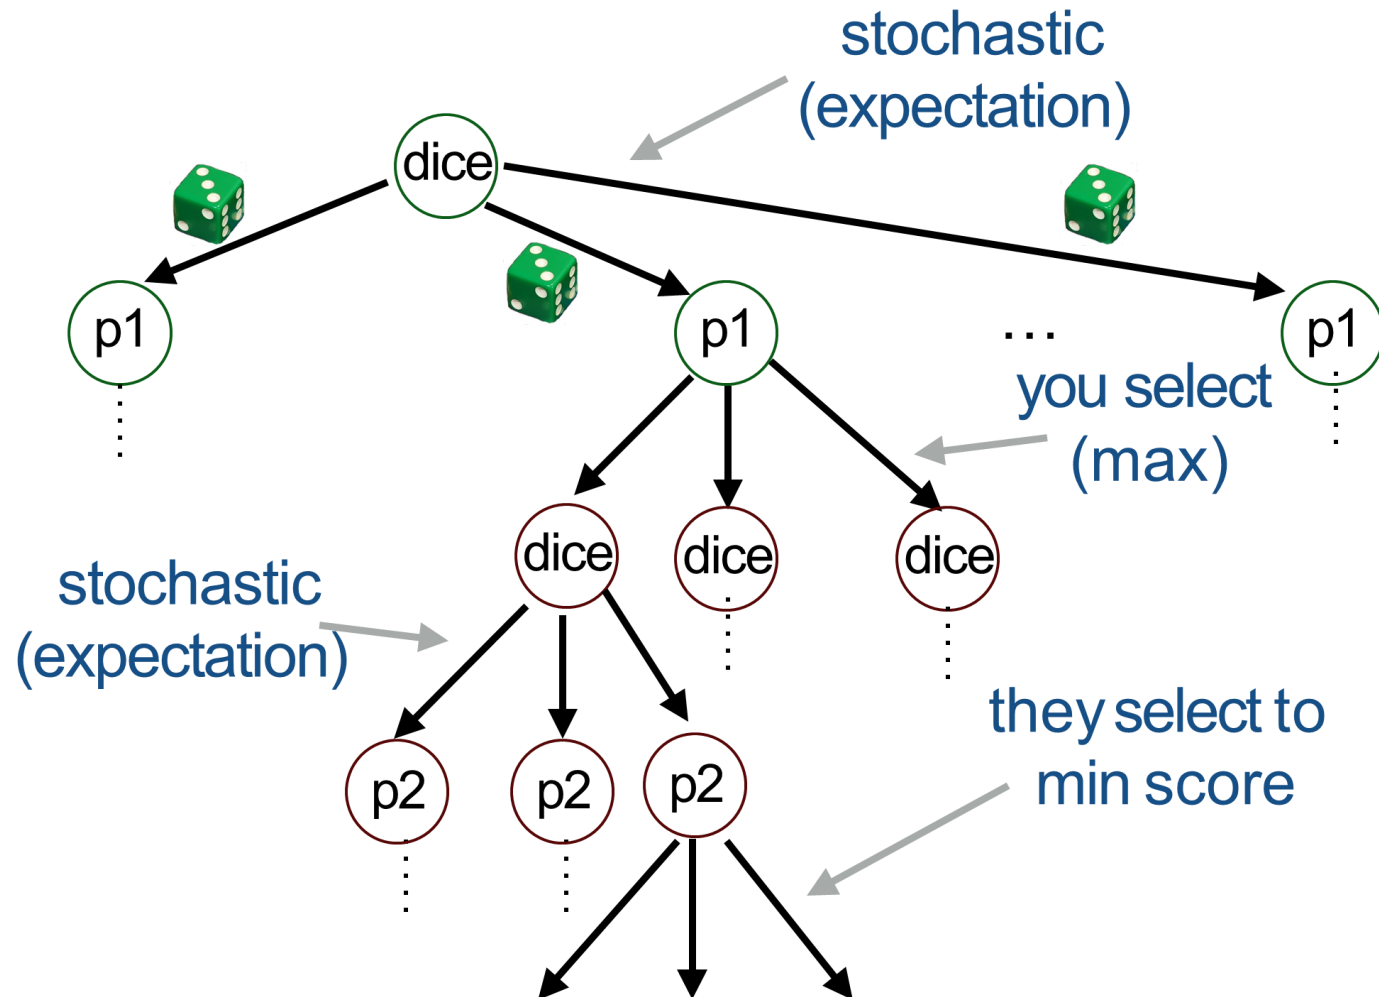
- What if there is **stochasticity** in the game?



Stochasticity

- Be aware of who is choosing at each level
- Sometimes it is:
 - You
 - Opponent
 - Random generator
- We already have min/max nodes
 - So now add chance node


Expectimax



Minimax properties


- Like **DFS**:
 - Time: $O(b^m)$
 - Space: $O(bm)$
- But instead of searching for single goal, we need to **exhaustively** try everything!
 - And we only get **value at leaf nodes**
- Chess, for e.g., $b \sim 20, m \sim 70$
 - Exact solution **infeasible**
 - So what do we do?

Pruning the tree



$\alpha\beta$ -pruning

- α – minimum score MAX is guaranteed of
 - β – maximum score MIN is guaranteed of
 - If at a given min node, $\beta < \alpha$
 - Then MIN can guarantee a score that makes MAX sad ☹️
 - So MAX will never go down this road
 - No need to expand rest of node's children!
 - Symmetric argument for other way around
 - If children are expanded in optimal order, complexity is halved: $O(b^{m/2})$
- Worst case guarantees*
- Could be pruning entire subtrees!*




α – minimum score MAX is guaranteed of
 β – maximum score MIN is guaranteed of

Depth-limited search

- Even with pruning, can't reach leaf nodes in real games
- So must **limit depth** of search
 - Must replace utility function with **estimate** (like heuristic in A*)
 - No longer optimal
- More plies = better performance
- Given time budget, use **IDS**!

Evaluation functions

- Estimate of utility of non-terminal state
- Ideally: want actual minimax value of state
 - But this is unknown
- In practice: use domain knowledge
 - E.g. $\text{eval}(s) = w_1(|\text{pawns}_w - \text{pawns}_b|) + w_2(|\text{bishops}_w - \text{bishops}_b|) + \dots$
- Tradeoff between complexity vs depth
 - More complex eval function may be more accurate, but longer to compute → less time to search deeper
 - Stockfish: fast eval function, huge depth
 - Komodo: slow, complex eval function, less depth



Other improvements

- Base algorithm of $\alpha\beta$ + IDS + eval function
- **Transposition tables**: stored previous states and their evals
- **Aspiration windows**: pretend that the $\alpha\beta$ window is smaller than it is
- **Evaluation functions optimised** from data (machine learning etc)
- **Move ordering**: try certain classes of moves first (e.g. captures, then regular moves)

Board games

- “ ... board games are more or less done and it's time to move on.”

