

# Artificial Intelligence

Steve James

Supervised and Unsupervised  
Learning

# Machine learning

- Subfield of AI concerned with learning from data
- Broadly, using:
  - Experience
  - To improve performance
  - On some task

*(Tom Mitchell, 1997)*

# Supervised learning

- Input:

- $X = \{x_1, \dots, x_n\}$

- $Y = \{y_1, \dots, y_n\}$

*Training data*


- Learning to predict **new labels**

- Given  $x$ ,  $y$ ?




# Classification vs regression

- If set of labels  $Y$  is **discrete**:
  - Classification
  - Minimise number of errors



- If  $Y$  is **real-valued**
  - Regression
  - Minimise sum squared error




# Supervised learning


- Formal definition:
  - Given training data:
    - $X = \{x_1, \dots, x_n\}$
    - $Y = \{y_1, \dots, y_n\}$
  - Produce **decision function**  $f: X \rightarrow Y$
  - That minimises error
    - $\sum_i \text{err}(f(x_i), y_i)$

# Test/train split

- Minimise error on what?
  - Don't get to see future data
- General principle: do not measure error on the data you train on



# Overfitting



# Test/train split

- Methodology:
  - Split data into **train** and **test** set
  - Fit  $f$  using training set
  - Measure error on test set
- Common alternative:  $k$ -fold cross validation
  - Repeat  $k$  times:
    - Partition data into train ( $n - n/k$ ) and test ( $n/k$ ) data sets
    - Train on training set, test on test set
  - Average results across  $k$  choices of test set.



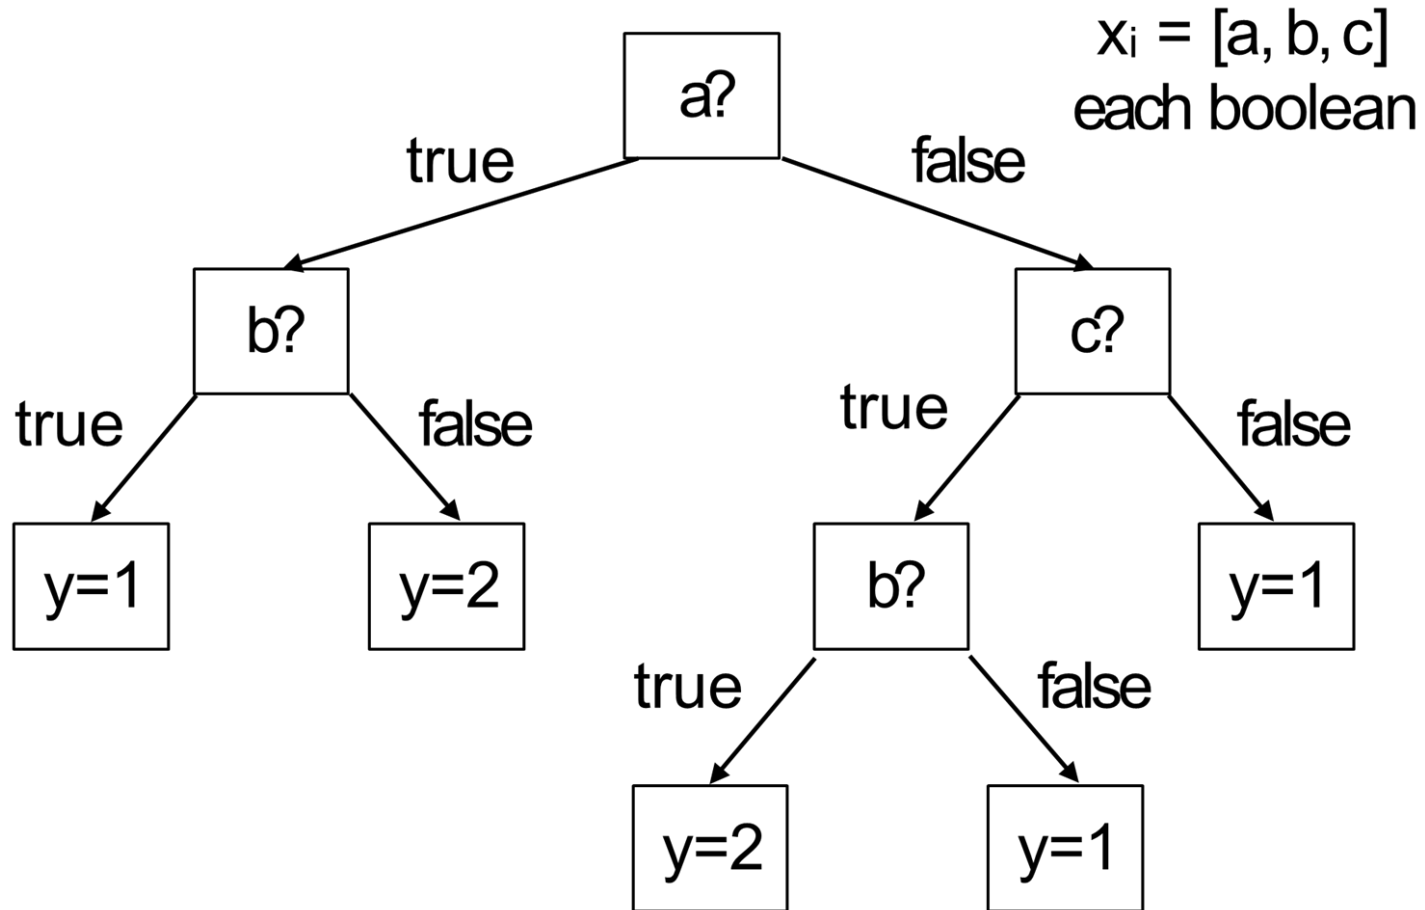
# Key idea: hypothesis space

- Typically **fixed representation** of classifier/regressor
  - Learning algorithm constructed to match
- Representation induces class of functions  $F$  from which to find  $f$ 
  - $F$  is known as the **hypothesis space**
  - Tradeoff: power vs expressibility vs data efficiency
  - Not every  $F$  can represent every function

# Decision trees

- Assume:
  - Two classes (true / false)
  - Input is vector of discrete values
- Simplest thing we could do?
  - How about if/else rules?
- Relatively simple classifier
  - Tree of **tests**
  - Evaluate test for each  $x_i$  and follow branch
  - Leaves are class labels

# Decision trees



# Decision trees

- How to make one?
- Given  $X = \{x_1, \dots, x_n\}$ ,  $Y = \{y_1, \dots, y_n\}$
- Repeat:
  - If all labels are the same, then we have a leaf node
  - Pick an attribute and split data based on its value
  - Recurse on each half
- If we run out of splits and data not perfectly in one class, then take a max

# Attribute picking

- But which attribute to split over?
- Information contained in a data set:


$$I(D) = -f_1 \lg(f_1) - f_2 \lg(f_2)$$

- How “bits” of information do we need to determine the label in a dataset?
- Pick attribute with **max information gain**:


$$Gain(E) = I(D) - \sum_i f_i I(E_i)$$

# Neural networks

- Single neuron:
  - Takes as input  $x$ , produces output of the form
$$y = a(w \cdot x + b)$$
- $w$  (weights) and  $b$  (bias) are **learnable parameters**
- $a$  is **activation function**. Many choices
  - $a(x) = \text{sgn } x$
  - $a(x) = \sigma(x)$
  - $a(x) = \text{ReLU}(x)$




# Neural networks



# Neural network learning

- Feed input to NN
  - Compute forward pass and get output
  - Compare output to label and compute error
  - Update weights/biases based on error.
- But how?
  - Compute derivative of weights with respect to error (chain rule)
  - Can be efficiently done using dynamic programming (**backpropagation** algorithm)






# Regression

- Broadly similar to classification
  - But predicting real valued numbers
  - Can modify decision trees, neural nets for regression (e.g. output layer is real valued, use MSE)
- Most ML methods are **parametric**:
  - Characterised by setting a few parameters
  - $y = f(x, w)$


# Nonparametric regression

- Alternative approach:
  - Let the data speak for itself
  - No finite-sized parameter vector
  - Usually more interesting decision boundaries
- Given  $X = \{x_1, \dots, x_n\}$ ,  $Y = \{y_1, \dots, y_n\}$ , distance metric  $D(x_i, x_j)$ 
  - For a new data point  $x_{new}$ :
    - Find  $k$  nearest points in  $x$  (measured by  $D$ )
    - Set  $y_{new}$  to the (weighted by  $D$ ) average  $y_i$  labels


# Nonparametric regression




# Nonparametric regression





# Nonparametric regression



# Nonparametric regression



# Gaussian processes



# Unsupervised learning

- Input:
  - $X = \{x_1, \dots, x_n\}$
- Try understand **structure of the data**



- E.g. How many types of cars? How can they vary?



# Clustering


- One particular type of unsupervised learning:
  - Split the data into **discrete clusters**
  - Assign new data points to each cluster
  - Clusters can be thought of as **types**
- Formally:
  - Given data points  $X = \{x_1, \dots, x_n\}$
  - Find number of clusters  $k$ 
    - Assignment function  $f(x) = \{1, \dots, k\}$

# K-means

- One approach
  - Pick  $k$
  - Place  $k$  points (“means”) in the data
  - Assign new point to  $i$ th cluster if nearest to  $i$ th mean
- Place  $k$  means at random
- Assign all points in the data to each “mean”
  - $f(x_j) = i$  such that  $d(x_j, \mu_i) \leq d(x_j, \mu_l) \forall l \neq i$
  - Move each “mean” to mean of assigned data and repeat


$$\mu_i = \sum_{v \in C_i} \frac{x_v}{|C_i|}$$

Start!



# Density estimation

- Clustering can answer *which cluster?*, but not *does this belong?*




# Density estimation

- Estimate the distribution the data is drawn from
- This allows us to evaluate the probability that a new point is drawn from the same distribution as the old data
- Formally:
  - Given data points  $X = \{x_1, \dots, x_n\}$
  - Find PDF  $P(X)$

# GMM

- Model data as **mixture of Gaussians**
- Each Gaussian has its own **mean and variance**
- Each has its own **weight** (sums to 1)
  - Weighted sum of Gaussians is still a PDF



# GMM


- Almost the same as  $k$ -means
- Place means at random, set variances high
- Assign all points to highest probability distribution

$$C_i = \{x_v | N(x_v | \mu_i, \sigma_i^2) > N(x_v | \mu_j, \sigma_j^2) \forall j\}$$

- Set mean, **variance and weights** to match assigned data and repeat

$$\mu_i = \sum_{v \in C_i} \frac{x_v}{|C_i|}; \sigma_i^2 = \text{variance}(C_i); w_i = \frac{|C_i|}{\sum_j |C_j|}$$

# GMM





# Nonparametric density estimation

- Parametric:
  - Define parameterised model (e.g. Gaussian)
  - Fit parameters
  - Done!
- Key assumptions
  - Data is distributed according to **parameterised form**
  - We know which parameterised form **in advance**
- What is shape of the distribution over human faces?



# Nonparametric density estimation

- Alternative:
  - Avoid fixed parameterised form
  - Compute density estimate directly from data

- Kernel density estimator


$$PDF(x) = \frac{1}{nb} \sum_{i=1}^n D\left(\frac{x_i - x}{b}\right)$$

- Where
  - $D$  is a special kind of distance metric called a **kernel**
    - Falls away from 0, integrates to 1
  - $b$  is **bandwidth**. Controls how fast kernel falls away

# Nonparametric density estimation

- $$PDF(x) = \frac{1}{nb} \sum_{i=1}^n D\left(\frac{x_i - x}{b}\right)$$
- Kernel: lots of choices, **Gaussian** often works in practice
- Bandwidth:
  - High: distant points have higher “contribution” to sum
  - Low: distant points have lower

# Nonparametric density estimation



# Dimensionality reduction

- $X = \{x^1, \dots, x^n\}$ , each  $x^i$  has  $m$  dimensions:  $x^i = [x_1, \dots, x_m]$
- If  $m$  is high, data can be hard to deal with
  - High-dimensional decision boundary
  - Need more data
  - But data is often not really high-dimensional
- **Dimensionality reduction:**
  - Reduce or **compress** the data
  - Try not to lose too much!
  - Find intrinsic dimensionality

# Dimensionality reduction

- Often can be phrased as a projection:

$$f : X \rightarrow X'$$

- Where
  - $|X'| \ll |X|$
  - Our goal: retain as much **sample variance** as possible
- Variance captures what varies within the data

# Principal component analysis

- Gather data  $x^1, \dots, x^n$ 
  - Adjust data to be zero-mean
  - Compute covariance matrix  $C$
  - Compute unit **eigenvectors**  $V_i$  and **eigenvalues**  $v_i$  of  $C$
- Each  $V_i$  is a direction, and each  $v_i$  is its importance - the amount of the data's variance it accounts for
- New data points:

*Compressed  
data point*

$$\hat{x}^i = [V_1, \dots, V_p] x^i$$

*Compression  
matrix  $V$*

*Original  
data point*

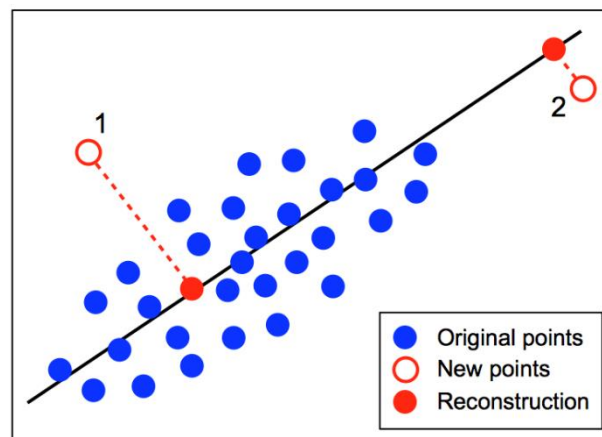
# PCA

- To recover original data point:

$$\bar{x}^i = V^{-1} \hat{x}^i$$
$$\bar{x}^i = V^T \hat{x}^i$$

*V is orthonormal  
so  $V^{-1} = V^T$*

- Every data point is expressed as a **linear combination** of basis (eigenvector) functions

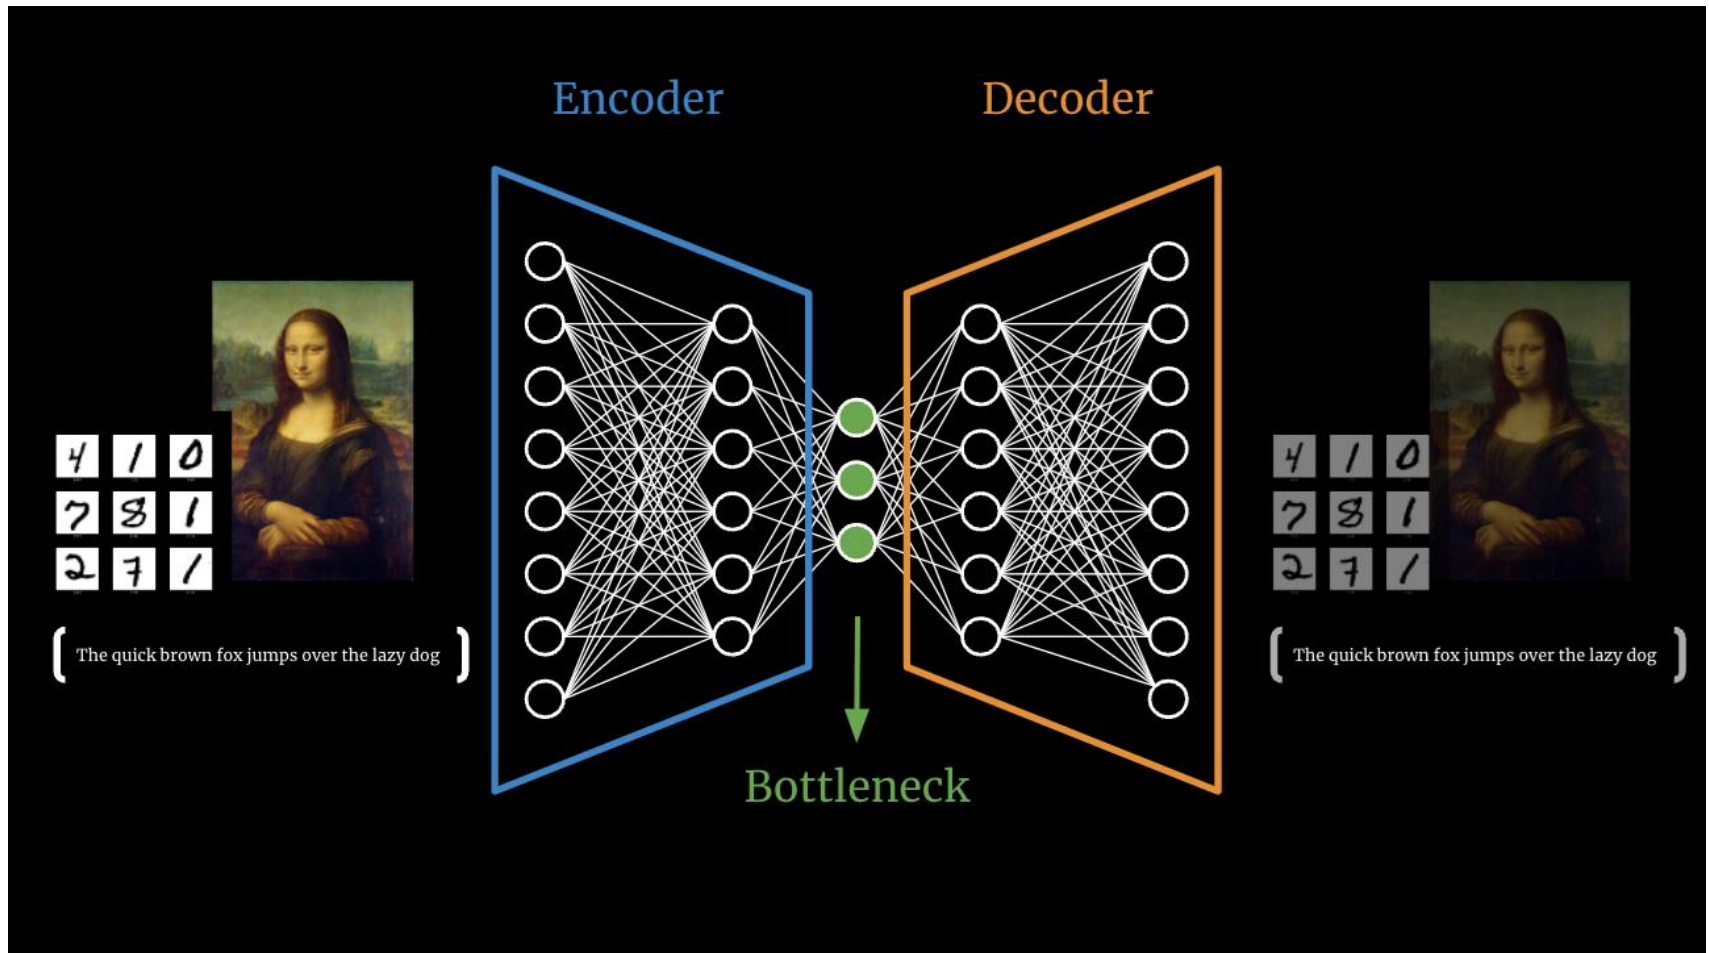





# Autoencoders

- Fundamental issue with PCA
  - Linear reconstruction
- Can we use a nonlinear method for construction?
  - Extract more complex relationships within the data.
  - Remove “linear reconstruction” property.
- One idea: train neural network to reproduce output

# Autoencoders



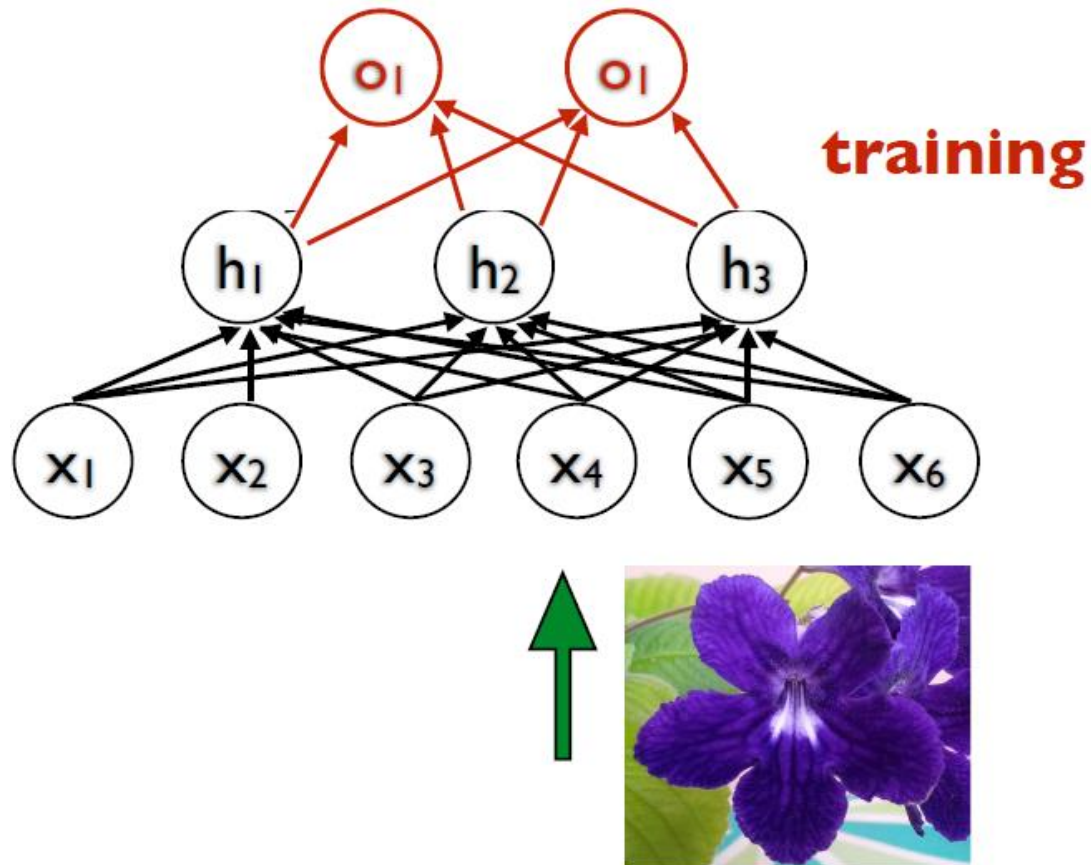
# Autoencoders for classification



**pretraining**



# Autoencoders for classification



# Data mining

- Most common application of unsupervised learning
  - Given large corpus of data, what can be learned?



# Data mining

“As Pole’s computers crawled through the data, he was able to identify about 25 products that, when analyzed together, allowed him to assign each shopper a “pregnancy prediction” score. More important, he could also estimate her due date to within a small window, so Target could send coupons timed to very specific stages of her pregnancy.

One Target employee I spoke to provided a hypothetical example. Take a fictional Target shopper named Jenny Ward, who is 23, lives in Atlanta and in March bought cocoa-butter lotion, a purse large enough to double as a diaper bag, zinc and magnesium supplements and a bright blue rug. There’s, say, an 87 percent chance that she’s pregnant and that her delivery date is sometime in late August.”


# Spurious correlations

## Per capita cheese consumption

correlates with

## Number of people who died by becoming tangled in their bedsheets

Correlation: 94.71% ( $r=0.947091$ )




Data sources: U.S. Department of Agriculture and Center

## People who drowned after falling out of a fishing boat

correlates with

## Marriage rate in Kentucky

Correlation: 95.24% ( $r=0.952407$ )



Data sources: Centers for Disease Control & Prevention and National Vital Statistics Reports

tylervigen.com

<https://www.tylervigen.com/spurious-correlations>