


Control Theory: Introduction and PID

Robotics – COMS4045A / COMS7049A

Benjamin Rosman

Over the next few weeks...

- You will learn how to:
 - determine where the robot is (forward kinematics)
 - determine where you'd like it to be (inverse kinematics)
 - determine how it may move (dynamics)
- But now:
 - How to get it to move
 - From where it is
 - To where I want it
 - (Control theory)




States

- State = robot's representation of the configuration of the world (and itself)
 - Ideally unique
 - Typically don't model things that don't change (e.g. obstacles)
- Why states?
 - We want to know (or learn) what actions (a or u) to take at each state (x or s or q)
 - In continuous systems, state is a function of time:
 - $x = x(t)$
 - Then we get velocity as a derivative:
 - $x'(t) = \dot{x}(t) = \frac{dx}{dt}$

Examples of state

- $x = \{\phi_1, \phi_2, \phi_3\}$
- $x = \{\phi_1, \phi_2, \phi_3, \dot{\phi}_1, \dot{\phi}_2, \dot{\phi}_3\}$
- $x = \{\phi_1, \phi_2, \phi_3, isHandOpen\}$




- $x = \{xPos, yPos\}$
- $x = \{xPos, yPos, direction\}$
- $x = \{xPos, yPos, colour\}$
- $x = \{xPos, yPos, hasKey\}$
- $x = \{xAll, yAll\}$



System dynamics

- System dynamics describe how the state changes:
 - Modelled as differential equations
 - Passively: $\dot{\mathbf{x}} = f(\mathbf{x})$
 - With time: $\dot{\mathbf{x}} = f(\mathbf{x}, t)$
 - Controlled: $\dot{\mathbf{x}} = f(\mathbf{x}, u)$
- Examples:
 - Newton's 2nd law: $\ddot{x}(t) = F/m$
 - Falling under gravity: $\dot{x}(t) = gt$
 - Population growth: $\dot{x}(t) = Ax(t)(1 - \frac{x(t)}{B})$



Aside: systems of DEs

- General trick:
 - Any higher-order DE can be written as a first-order SYSTEM
- E.g.
 - $\ddot{x} = 3x$
 - Now let $x_1 = x$, and $x_2 = \dot{x}$
 - So: $\dot{x}_1 = x_2$
 - Then: $\dot{x}_2 = 3x_1$
 - As a system: $\begin{pmatrix} \dot{x}_1 \\ \dot{x}_2 \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 3 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$
- In general: $\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x})$
 - Question: what about $x''' = 2x'' - x + 1$?

And now...

- What we really want is to **tell the robot how to function in a specified manner**
 - When this would not happen naturally
 - System subject to perturbations – maintain stability
- This is control theory
 - In general, find control input u to drive the system to desired states x
 - $\dot{x} = f(x) + g(x, u)$
- Much of the core of robotics

Issues to think about: the system

- System may be **linear**, which makes the maths easy:
 - $\dot{\mathbf{x}}(t) = A(t)\mathbf{x}(t) + B(t)\mathbf{u}(t)$
- System may be **time invariant**:
 - $A(t) = A; \quad B(t) = B$
- System is not always **observable**. Often we only have a partial view of the state \mathbf{x} through some observables \mathbf{y} :
 - i.e. it can't see everything (very common)
 - $\mathbf{y}(t) = C\mathbf{x}(t) + D\mathbf{u}(t)$
- Does the system have **delays** in updates, sensing, or control?
 - $\dot{\mathbf{x}}(t) = A\mathbf{x}(t) + B\mathbf{u}(t - 10)$

Issues to think about: stability

- It is worth considering the **steady state behaviour** of the system
 - i.e. what is the asymptotic behaviour: $x(t \rightarrow \infty)$
- This relates to the notion of **stability**:
 - Informally, a system is stable if the behaviour converges (or in a weaker case, stays similar)
- In linear systems, consider the **eigenvalues** of A :
 - $Re(\lambda_A) < 0$?

Controllability

- **Controllability:** is it *possible* to fully control the system (affect each of the n state variables of the system)?
 - E.g. $\mathbf{x} = [x, \dot{x}, \ddot{x}]^T$ (position, velocity, acceleration)
 - $\mathbf{u} = [u, 0, 0]^T$ (we can directly affect the position)
 - $\mathbf{u} = [0, 0, u]^T$ (we can directly affect the acceleration)
- Consider continuous linear time-invariant system:
 - $\dot{\mathbf{x}}(t) = \mathbf{A}\mathbf{x}(t) + \mathbf{B}\mathbf{u}(t)$
 - $\mathbf{y}(t) = \mathbf{C}\mathbf{x}(t) + \mathbf{D}\mathbf{u}(t)$
 - Compute controllability matrix:
 - $\mathbf{R} = [\mathbf{B} \ \mathbf{A}\mathbf{B} \ \mathbf{A}^2\mathbf{B} \ \dots \ \mathbf{A}^{n-1}\mathbf{B}]$
 - System is controllable if full rank: $\text{rank}(\mathbf{R}) = n$

Intuition: Does repeatedly applying the control policy allow you to change every state independently?

Example

- $\begin{pmatrix} \dot{x}_1 \\ \dot{x}_2 \end{pmatrix} = \overset{\text{A}}{\begin{pmatrix} 1 & 0 \\ 0 & 3 \end{pmatrix}} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \overset{\text{B}}{\begin{pmatrix} 0 \\ 1 \end{pmatrix}} u ; y = \overset{\text{C}}{(1 \quad 1)} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$
- Controllable?
 $R = [B \ AB \ A^2B \ \dots A^{n-1}B]$
- $R = \begin{pmatrix} 0 & 0 \\ 1 & 3 \end{pmatrix}$: $rank(R) = 1 < n$ ($n = 2$): No!
- $\begin{pmatrix} \dot{x}_1 \\ \dot{x}_2 \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 3 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} u ; y = (1 \quad 1) \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$
- Controllable?
- $R = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$: $rank(R) = 2 = n$: Yes!

Observability

- **Observability:** is it *possible* to fully observe the system (observe each state variable)?
 - E.g. $\mathbf{x} = [x, \dot{x}, \ddot{x}]^T$
 - $\mathbf{x} = [x, \dot{x}, \ddot{x}]^T, \mathbf{y} = x$ (we only observe position)
 - $\mathbf{x} = [x, \dot{x}, \ddot{x}]^T, \mathbf{y} = \ddot{x}$ (we only observe acceleration)
- Consider continuous linear time-invariant system:
- Compute observability matrix:

$$\bullet O = \begin{pmatrix} C \\ CA \\ \vdots \\ CA^{n-1} \end{pmatrix}$$

Intuition: Does repeatedly observing the system allow you to see every state variable?

- System is observable if full rank: $\text{rank}(O) = n$

Example

- $\begin{pmatrix} \dot{x}_1 \\ \dot{x}_2 \end{pmatrix} = \overset{\text{A}}{\begin{pmatrix} 3 & 0 \\ 0 & 1 \end{pmatrix}} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \overset{\text{B}}{\begin{pmatrix} 0 \\ 1 \end{pmatrix}} u ; y = \overset{\text{C}}{(0 \quad 1)} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$

- Observable?

- $O = \begin{pmatrix} 0 & 1 \\ 0 & 1 \end{pmatrix} : \text{rank}(O) = 1 < n \ (n = 2) : \text{No!}$

$$O = \begin{pmatrix} C \\ CA \\ \vdots \\ CA^{n-1} \end{pmatrix}$$

- $\begin{pmatrix} \dot{x}_1 \\ \dot{x}_2 \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 3 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \end{pmatrix} u ; y = (0 \quad 1) \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$

- Observable?


- $O = \begin{pmatrix} 0 & 1 \\ 3 & 0 \end{pmatrix} : \text{rank}(O) = 2 = n : \text{Yes!}$

Issues to think about: the controller

- What kind of control strategy to use?
- Is the strategy **efficient**?
 - Optimal control theory (later)
- Can the controller **sense** some aspect of the system?
Can it **respond** to changes?
 - Open loop vs closed loop control
- Can the controller make **predictions** about the environment/system?
 - Model free vs model based control

Open loop control


- Open loop control:
 - Policy set and pre-determined (function of time, NOT state)



- Pros:
 - Cheap, simple
 - Fast (feedback and sensing may be too slow)
 - Can be calibrated
- Cons:
 - Policy doesn't take disturbances into account

Closed loop (feedback) control

- Feedback control:
 - Measure disturbance (error) and use this to adjust control




Closed loop (feedback) control

- Pros:
 - Simple
 - Doesn't require process model
 - Robust output for unpredictable disturbances
- Cons:
 - Requires sensors to measure output
 - Requires tuning: low gain slow, high gain unstable
 - Delays in feedback produce oscillations

Model based control

- Model predictive control:
 - Maintain a process model, make predictions and act accordingly



Model based control

- Pros:
 - Anticipate disturbances and effects of actions, plan to account for these
 - Plan long action sequences
 - Prevent potential problems
 - Robust responses to disturbances
- Cons:
 - Slow
 - Difficult to acquire and maintain a model
 - Traditionally provided by expert
 - Learning system (more on this later)



e.g. Room heating

- Heater to increase room temperature
- Open loop:
 - Switch heater on, after some pre-defined time, switch off
- Feedback:
 - Use thermometer in room to switch off at desired temperature
- Model based:
 - Make predictions based on the number of people in the room, outside temperature, doors and windows



Example – line following

- Bang-bang control (on-off control)
 - Switch between extremes
- Controller:
 - Turn left if line is to the left
 - Turn right if line is to the right




Proportional error control

- Idea: move proportionally left or right
 - Proportional control
- Let the desired state (setpoint) be x_{goal}
- Then, error $e(t) = x_{goal} - x(t)$
- Response output: $K_P e(t)$
- K_P is the proportional gain
 - Too low: slow response
 - Too high: overshoot (unstable)
 - Problem: undershoots (steady state error) when there is a steady disturbance!

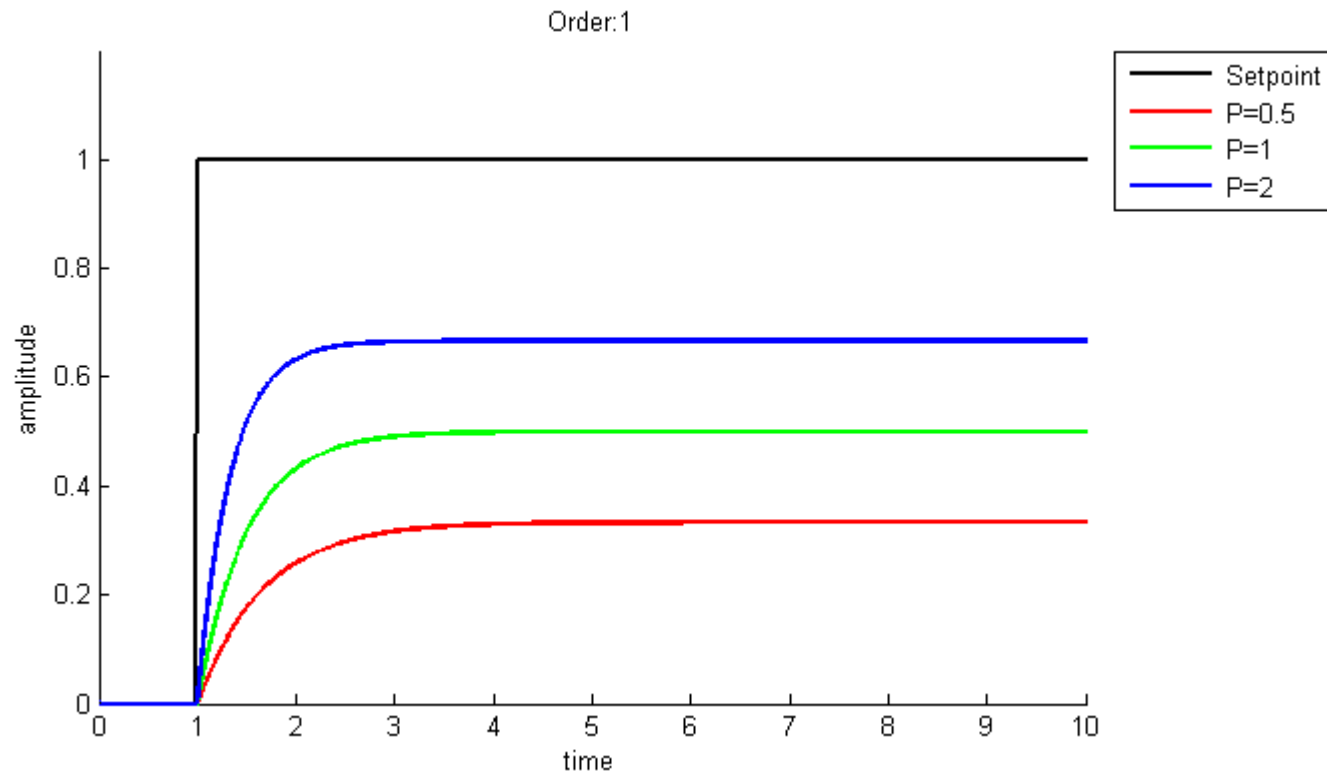
Proportional error control

- Simple first order system




Proportional error control

- Steady constant disturbance (droop problem)
 - (think of “wind”)



Proportional error control

- Second order system (with damping)
- Droop and oscillations




Proportional integral control

- Account for undershoot of P controller?
- Accumulate history of error
 - Increase response as history increases
- Then, error history $\int_0^t e(\tau) d\tau$
- Response output: $K_I \int_0^t e(\tau) d\tau$
- K_I is the integral gain
 - Tends to overshoot and cause oscillatory behaviour!

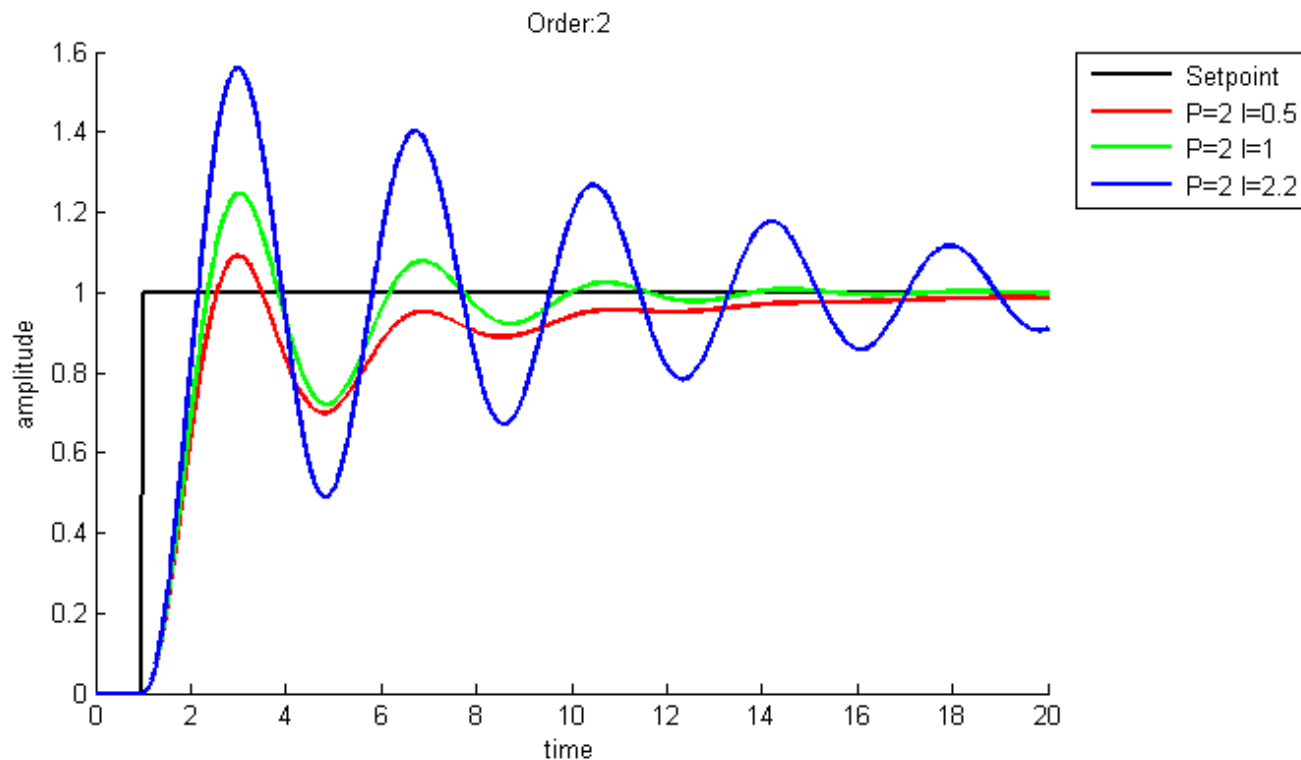
Proportional integral control

- Accounts for steady state error



Proportional integral control

- Introduce or amplify oscillations




Proportional derivative control

- Correct for (damp) oscillations?
 - Often caused by inertia, delays, and rapid overcorrecting
- Want to resist change: the faster the system changes, the harder it resists (artificial friction)
 - “Predictive” corrections
- Rate of change of error $\frac{d}{dt} e(t)$
- Response output: $K_D \frac{d}{dt} e(t)$
- K_D is the derivative gain
 - Improves settling time and stability
 - Can amplify noise, so not always used
 - Can slow control

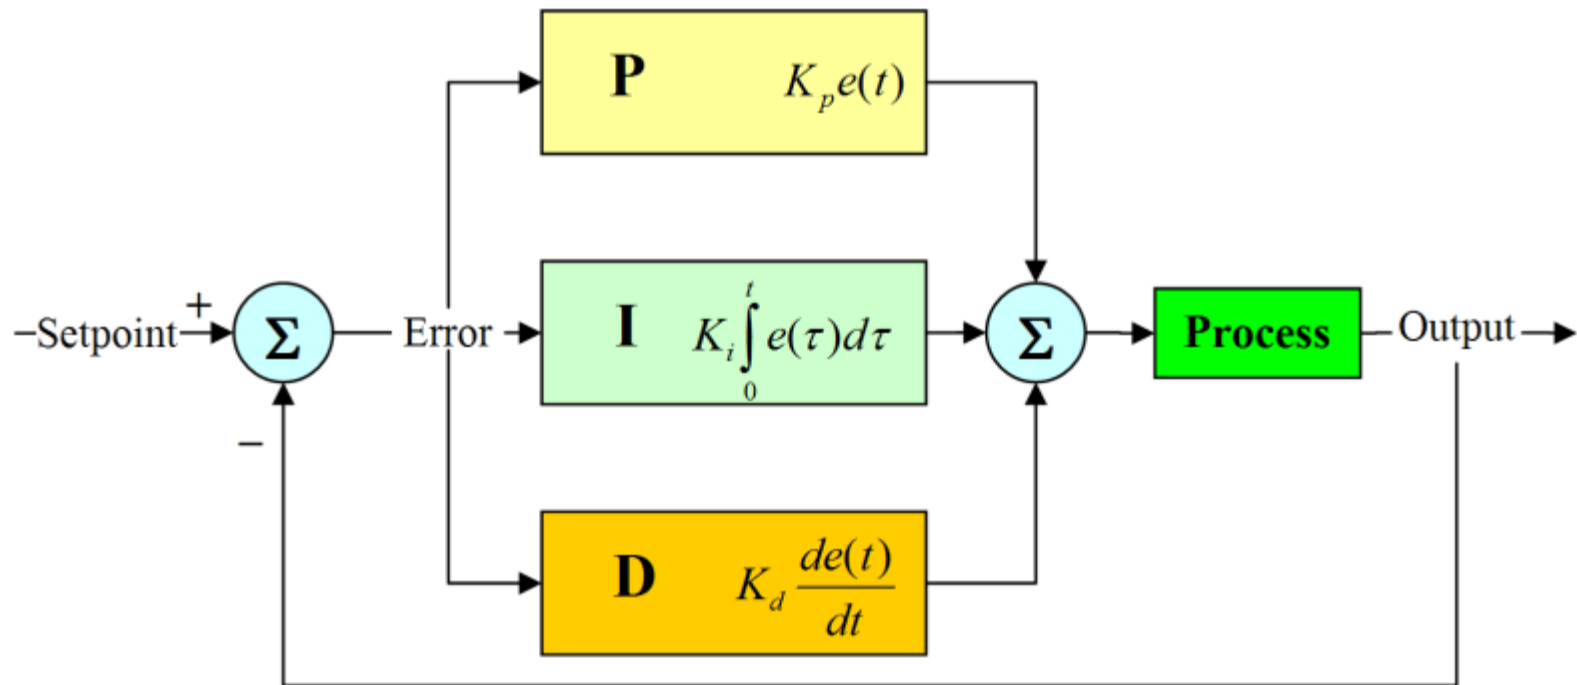
PID control

- Oscillation damping




PID control

- Combine: PID control
- Control law: $U = K_P e(t) + K_I \int_0^t e(\tau) d\tau + K_D \frac{d}{dt} e(t)$



PID control



PID control

- No knowledge of the plant dynamics
 - Only: current and desired behaviours
- Factors to tune for:
 - Responsiveness
 - Setpoint offset
 - Oscillation
- Tuning is often a fine art
- Different parameters may be needed in different regions of state space (tracking)
- D-term not often used in practice (noise amplifying)
- Not optimal

Tuning the parameters

- Manually:
 1. Set K_I and K_D to 0
 2. Increase K_P until oscillations
 3. Halve K_P
 4. Increase K_I until any offset corrected sufficiently fast (too high will be unstable)
 5. Increase K_D until quick recovery from perturbation (too high will overshoot)

Tuning the parameters

- Ziegler-Nichols method:

1. Set K_I and K_D to 0
2. Increase K_P until oscillations (at ultimate gain K_U) with oscillation period P_U
3. Then set:

Ziegler–Nichols method

Control Type	K_p	K_i	K_d
P	$0.50K_u$	-	-
PI	$0.45K_u$	$1.2K_p/P_u$	-
PID	$0.60K_u$	$2K_p/P_u$	$K_pP_u/8$

Illustrated parameter tuning

