

Unit 5: Advanced Data Modeling

Learning Objectives

- In this chapter, you will learn:
 - About the extended entity relationship (EER) model
 - How entity clusters are used to represent multiple entities and relationships
 - The characteristics of good primary keys and how to select them
 - How to use flexible solutions for special data-modeling cases

Extended Entity Relationship Model (EERM)

- Result of adding more semantic constructs to the original entity relationship (ER) model
- Question:
 - What is the meaning of semantics?
- **EER diagram (EERD):** Uses the EER model

Use Case: An Aviation Company

- Employee (common characteristics – firstname, hiredate etc.)
 - Pilots
 - License, ratings, flight_hours (unique to only pilots)
 - Accountants
 - Clerks
 - Managers
- Saving the details of all employees in a table will lead to lots of nulls or the need for dummy entries

Entity Supertypes and Subtypes

- **Entity supertype:** Generic entity type related to one or more entity subtypes
 - Contains common characteristics
- **Entity subtype:** Contains unique characteristics of each entity subtype
- Criteria to determine the usage
 - There must be different, identifiable kinds of the entity in the user's environment
 - The different kinds of instances should each have one or more attributes that are unique to that kind of instance

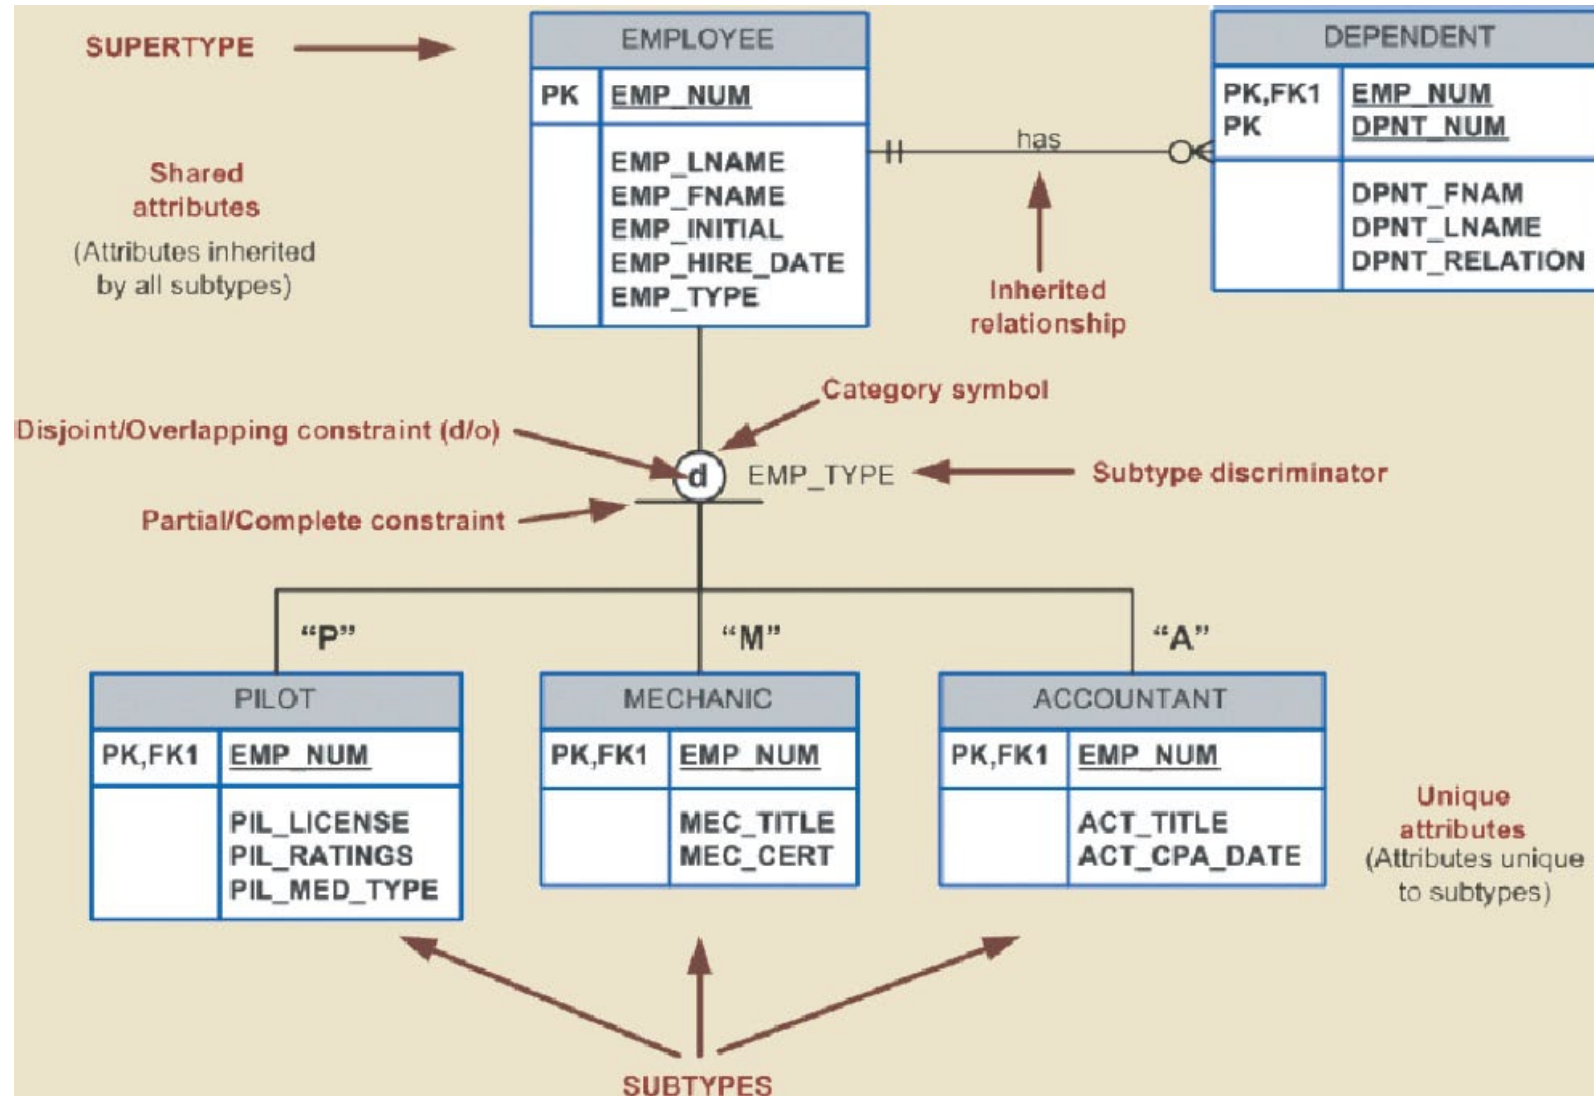
Specialization Hierarchy (1 of 2)

- Depicts arrangement of higher-level entity supertypes and lower-level entity subtypes
- Relationships are described in terms of “is-a” relationships
- Subtype exists within the context of a supertype
- Every subtype has only one supertype to which it is directly related
- Supertype can have many subtypes

Specialization Hierarchy (2 of 2)

- Provides the means to:
 - Support attribute inheritance
 - Define a special supertype attribute known as the subtype discriminator
 - Define **disjoint/overlapping** constraints and **complete/partial** constraints

Figure 5.2 - Specialization Hierarchy



Inheritance

- Enables an entity subtype to inherit attributes and relationships of the supertype
- All entity subtypes inherit their primary key attribute from their supertype
- At the implementation level, supertype and its subtype(s) maintain a 1:1 relationship
- Entity subtypes inherit all relationships in which supertype entity participates
- Lower-level subtypes inherit all attributes and relationships from its upper-level supertypes

Subtype Discriminator

- Attribute in the supertype entity that determines to which entity subtype the supertype occurrence is related

Disjoint and Overlapping Constraints

- **Disjoint subtypes:** Contain a unique subset of the supertype entity set
 - Known as **nonoverlapping subtypes**
 - Implementation is based on the value of the subtype discriminator attribute in the supertype
 - Employee can either be a pilot or mechanic; not both
- **Overlapping subtypes:** Contain nonunique subsets of the supertype entity set
 - Implementation requires the use of one discriminator attribute for each subtype
 - Staff can be both a professor and an administrator

Figure 5.4 - Specialization Hierarchy with Overlapping Subtypes

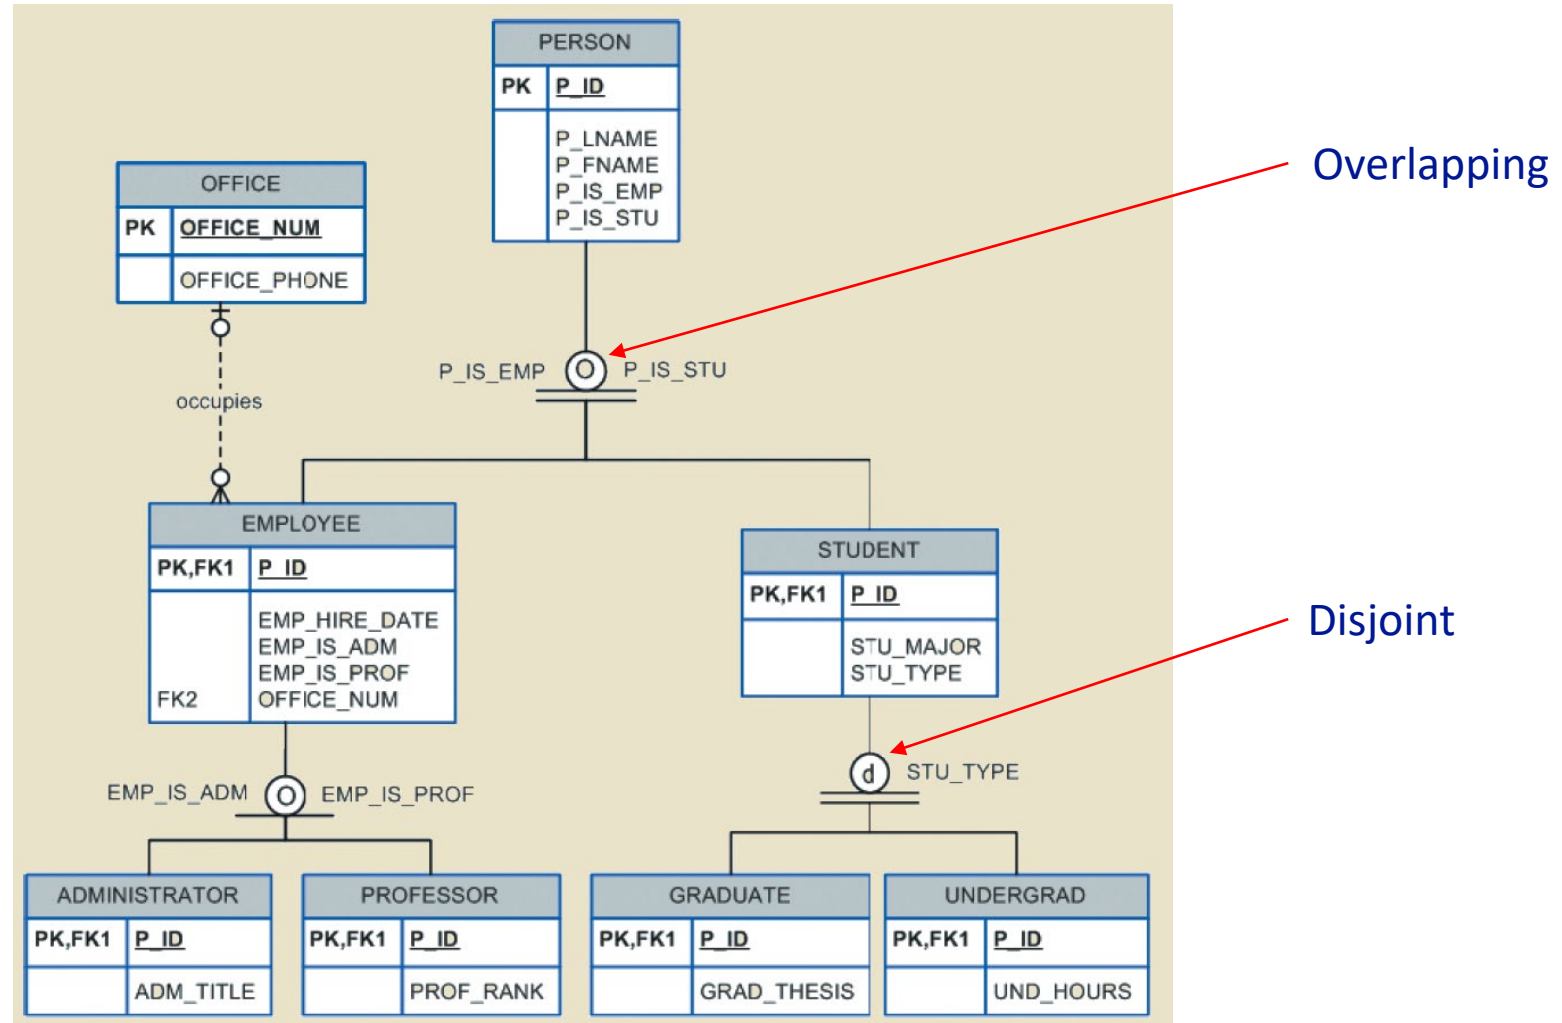




Table 5.1 - Discriminator Attributes with Overlapping Subtypes

| DISCRIMINATOR ATTRIBUTES OF PROFESSOR | DISCRIMINATOR ATTRIBUTES OF ADMINISTRATOR | COMMENT |
|---------------------------------------|---|--|
| Y | N | The Employee is a member of the Professor subtype. |
| N | Y | The Employee is a member of the Administrator subtype. |
| Y | Y | The Employee is both a Professor and an Administrator. |

Completeness Constraint

- Specifies whether each supertype occurrence must also be a member of at least one subtype
- Types
 - **Partial completeness:** Not every supertype occurrence is a member of a subtype
 - **Total completeness:** Every supertype occurrence must be a member of a subtype

Table 5.2 - Specialization Hierarchy Constraint Scenarios

| TYPE | DISJOINT CONSTRAINT | OVERLAPPING CONSTRAINT |
|--|--|---|
| Partial  | Supertype has optional subtypes. Subtype discriminator can be null. Subtype sets are unique. | Supertype has optional subtypes. Subtype discriminators can be null. Subtype sets are not unique. |
| Total  | Every supertype occurrence is a member of only one subtype. Subtype discriminator cannot be null. Subtype sets are unique. | Every supertype occurrence is a member of at least one subtype. Subtype discriminators cannot be null. Subtype sets are not unique. |

Specialization and Generalization

Specialization

- Top-down process
- Identifies lower-level, more specific entity subtypes from a higher-level entity supertype
- Based on grouping **unique** characteristics and relationships of the subtypes

Generalization

- Bottom-up process
- Identifies a higher-level, more generic entity supertype from lower-level entity subtypes
- Based on grouping **common** characteristics and relationships of the subtypes

Entity Cluster

- Virtual entity type used to represent multiple entities and relationships in ERD
- Combines multiple interrelated entities into a single, abstract entity object.
- It is not actually an entity in the final ERD.
- Helps to simplify the ERD; enhances readability.
- Avoid the display of attributes to eliminate complications that result when the inheritance rules change

FIGURE 4.35 THE COMPLETED TINY COLLEGE ERD

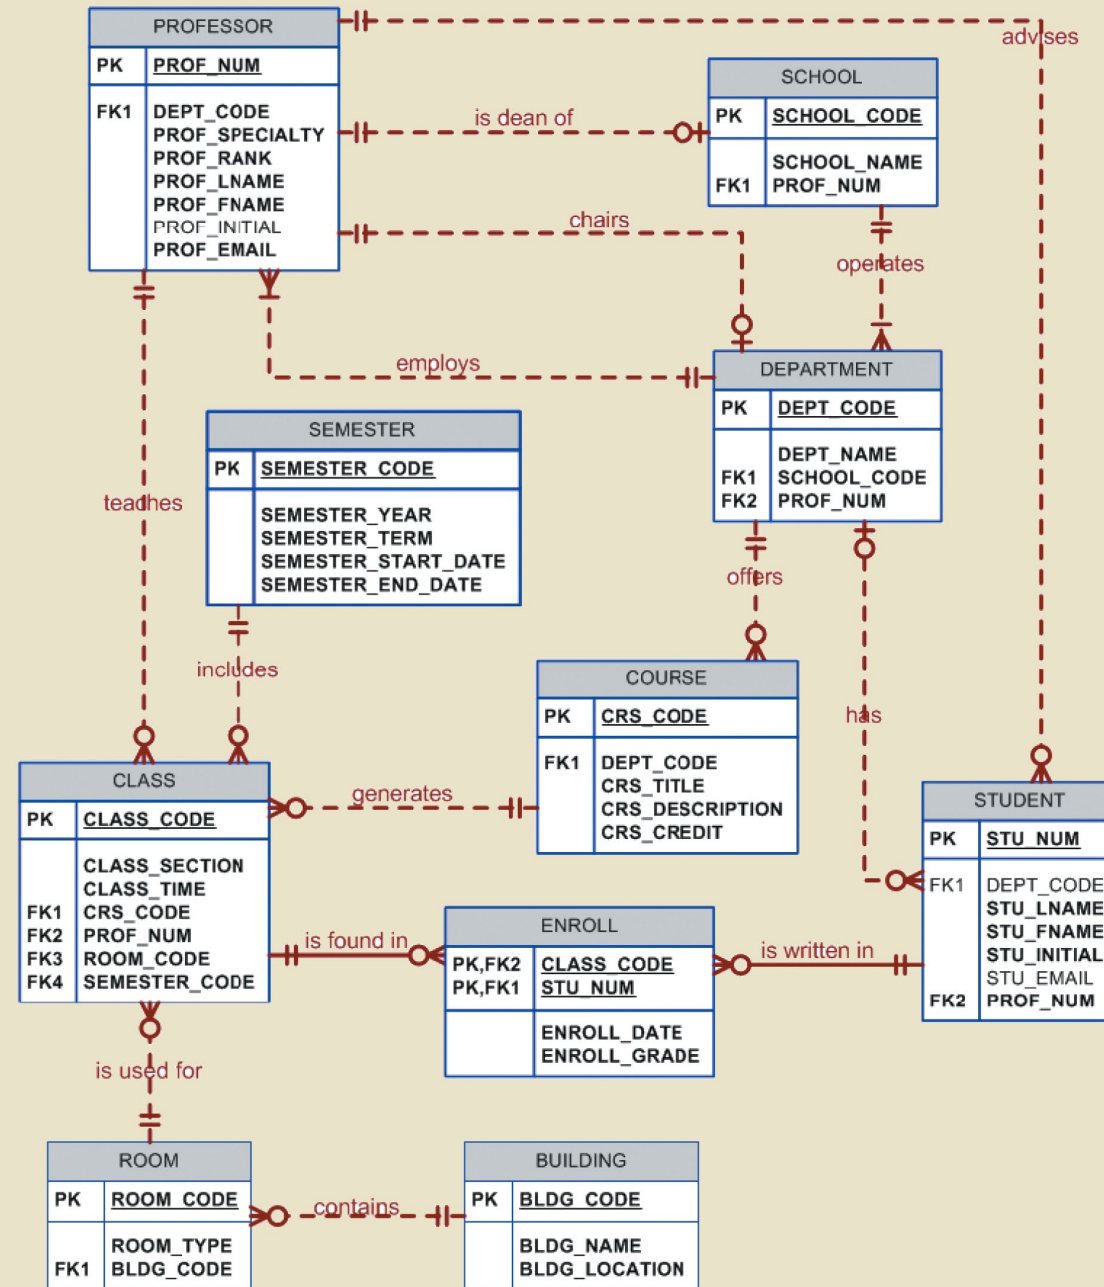
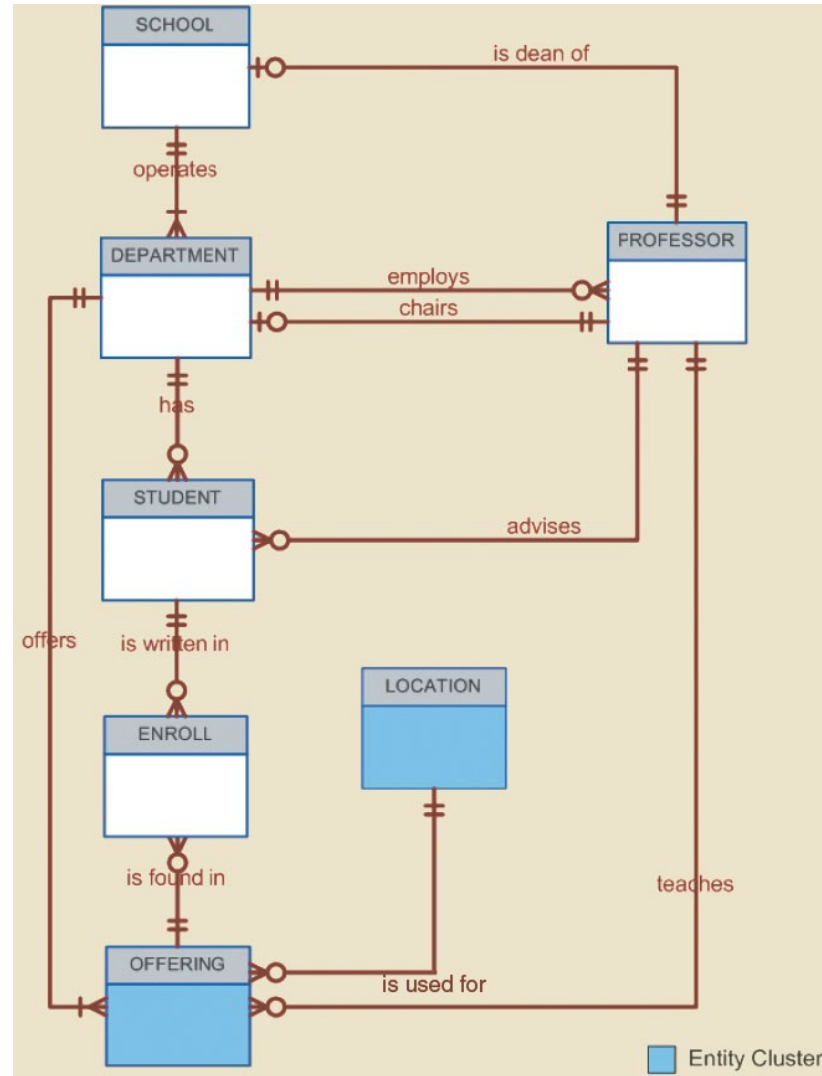


Figure 5.5 - Tiny College ERD Using Entity Clusters



NEXT: Selecting a Primary Key

What is a Primary Key?

- Question:
- Which of the following is TRUE:
 - A Primary key MUST be made up of a combination of attributes
 - A Primary key MUST be a single attribute
 - A Primary key can be a single attribute or be made up of a combination of attributes

Primary Keys

- Single attribute or a combination of attributes, which uniquely identifies each entity instance
 - Guarantees entity integrity
 - Works with foreign keys to implement relationships

Natural Keys or Natural Identifier

- Real-world identifier used to uniquely identify real-world objects
 - Familiar to end users and forms part of their day-to-day business vocabulary
 - Also known as natural identifier
 - Used as the primary key of the entity being modeled
 - E.g Student_number, South African ID number

Desirable Primary Key Characteristics - 1

- Unique values
- Non intelligent
 - should not have embedded semantic meaning
 - eg. STD_NO preferred to STD_NAME, use STD_NAME as a descriptive characteristic of an entity
- No change over time
 - descriptive attributes tend to change, choose attributes that are permanent and unchangeable
- Preferably single-attribute
 - simplifies implementation of foreign keys

Desirable Primary Key Characteristics - 2

- Preferably numeric
 - better managed, easily implemented using auto-increment functionality
- Security-compliant
 - void of any attribute(s) that might be considered a security risk. eg: South African ID as PK

Use of Composite Primary Keys (1 of 2)

- Although, they are not preferred, they are sometimes useful/needed:
 - Identifiers of composite (bridge) entities
 - Each primary key combination is allowed once in M:N relationship
 - Identifiers of weak entities
 - Weak entity has a strong identifying relationship with the parent entity

Use of Composite Primary Keys (2 of 2)

- When used as identifiers of weak entities, represent a real-world object that is:
 - Existence-dependent on another real-world object
 - Represented in the data model as two separate entities in a strong identifying relationship

Figure 5.6 - The M:N Relationship between STUDENT and CLASS

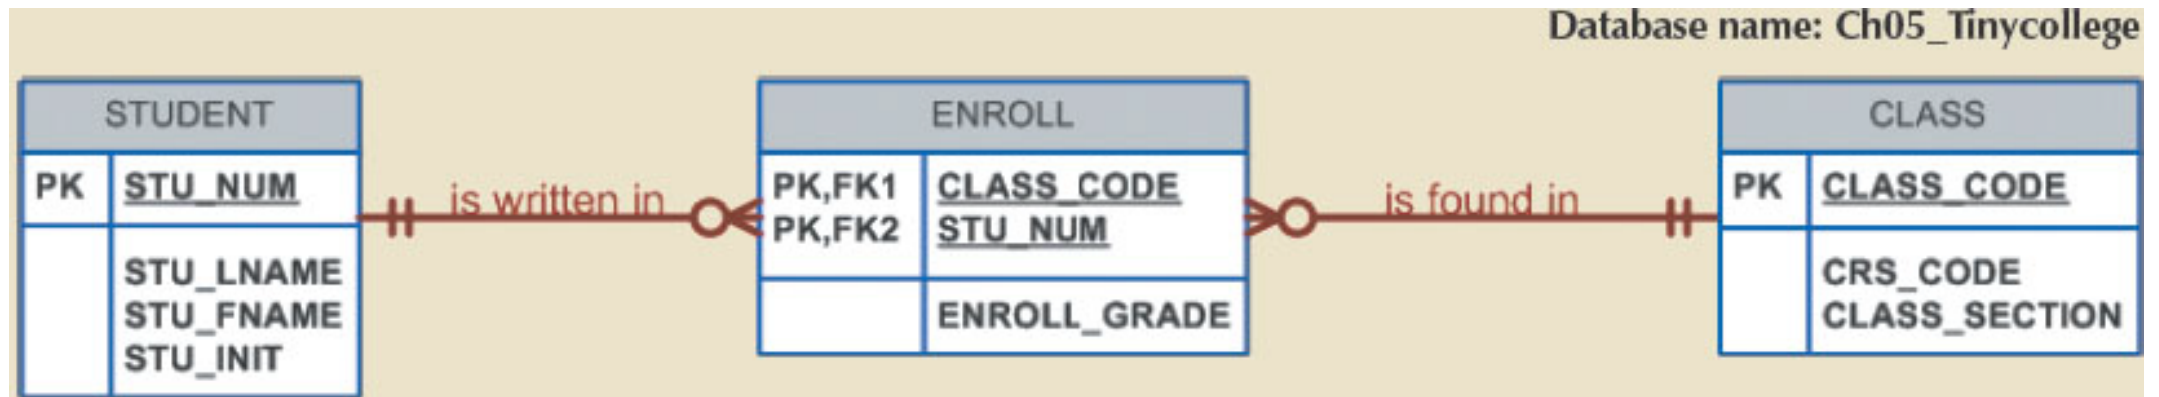


Table name: STUDENT
(first four fields)

| STU_NUM | STU_LNAME | STU_FNAME | STU_INIT |
|---------|-----------|-----------|----------|
| 321452 | Bowser | William | C |
| 324257 | Smihson | Anne | K |
| 324258 | Brewer | Juliette | |
| 324269 | Oblonski | Walter | H |
| 324273 | Smith | John | D |
| 324274 | Katinga | Raphael | P |
| 324291 | Robertson | Gerald | T |
| 324299 | Smith | John | B |

Table name: ENROLL

| CLASS_CODE | STU_NUM | ENROLL_GRADE |
|------------|---------|--------------|
| 10014 | 321452 | C |
| 10014 | 324257 | B |
| 10018 | 321452 | A |
| 10018 | 324257 | B |
| 10021 | 321452 | C |
| 10021 | 324257 | C |

Table name: CLASS
(first three fields)

| CLASS_CODE | CRS_CODE | CLASS_SECTION |
|------------|----------|---------------|
| 10012 | ACCT-211 | 1 |
| 10013 | ACCT-211 | 2 |
| 10014 | ACCT-211 | 3 |
| 10015 | ACCT-212 | 1 |
| 10016 | ACCT-212 | 2 |
| 10017 | CIS-220 | 1 |
| 10018 | CIS-220 | 2 |
| 10019 | CIS-220 | 3 |
| 10020 | CIS-420 | 1 |
| 10021 | QM-261 | 1 |
| 10022 | QM-261 | 2 |
| 10023 | QM-362 | 1 |
| 10024 | QM-362 | 2 |
| 10025 | MATH-243 | 1 |

15 Minutes break

Surrogate Primary Keys

- A surrogate key is a primary key created by the database designer to simplify the identification of entity instances.
- It has no meaning in the user's environment—it exists only to distinguish one entity instance from another (just
- like any other primary key).
- Can be generated by the DBMS to ensure that unique values are always provided.

Surrogate Primary Keys

- Primary key used to simplify the identification of entity instances are useful when:
 - There is no natural key
 - Selected candidate key has embedded semantic contents or is too long
- Require ensuring that the candidate key of entity in question performs properly
 - Use unique index and not null constraints

Case Study: Data Used to Keep Track of Events

What would the primary key be?

| DATE | TIME_START | TIME_END | ROOM | EVENT_NAME | PARTY_OF |
|-----------|------------|-----------|---------|----------------|----------|
| 6/17/2016 | 11.00a.m. | 2.00p.m. | Allure | Burton Wedding | 60 |
| 6/17/2016 | 11.00a.m. | 2.00p.m. | Bonanza | Adams Office | 12 |
| 6/17/2016 | 3.00p.m. | 5.30p.m. | Allure | Smith Family | 15 |
| 6/17/2016 | 3.30p.m. | 5.30p.m. | Bonanza | Adams Office | 12 |
| 6/17/2016 | 1.00p.m. | 3.00p.m. | Bonanza | Boy Scouts | 33 |
| 6/17/2016 | 11.00a.m. | 2.00p.m. | Allure | March of Dimes | 25 |
| 6/17/2016 | 11.00a.m. | 12.30p.m. | Bonanza | Smith Family | 12 |

Options for the primary key

- EVENT (DATE, TIME_START, TIME_END, ROOM, EVENT_NAME, PARTY_OF)
 - (DATE, TIME_START, ROOM) or
 - (DATE, TIME_END, ROOM)
- Let's select: (DATE, TIME_START, ROOM)

The CASE STUDY continues...

- There's also another entity – RESOURCE; to capture resources such as projectors, speakers that an event might use
- The RESOURCE entity would be represented by the following attributes:
 - RESOURCE (RSC_ID, RSC_DESCRIPTION, RSC_TYPE, RSC_QTY, RSC_PRICE)

M:N Relationship?

- Now, there's a M:N relationship between RESOURCE and EVENT
- Let's resolve this via the EVNTRSC composite entity with a composite primary key
 - EVNTRSC (DATE, TIME_START, ROOM, RSC_ID, QTY_USED)
- You now have a lengthy, four-attribute composite primary key.
- What would happen if the EVNTRSC entity's primary key were inherited by another existence-dependent entity?

Problems with the Selected Primary Key

- The EVENT entity's selected primary key does not fare well
 - EVENT entity's selected primary key contains embedded semantic information and is
 - formed by a combination of date, time, and text data columns.
 - the selected primary key would cause lengthy primary keys for existence-dependent entities.
- The preferred alternative is to use a numeric, single-attribute surrogate primary key.

NEXT: Design Cases

- This section presents four special design cases that highlight the importance of flexible designs, proper identification of primary keys, and placement of foreign keys.

RECAP: Example of a 1:1 Relationship

- A 1:1 relationship between EMPLOYEE and DEPARTMENT based on the business rule
 - “one EMPLOYEE is the manager of one DEPARTMENT, and one
 - DEPARTMENT is managed by one EMPLOYEE.”

RECAP: Example of a 1:1 Relationship

- The 1:1 relationship is a preferred relationship in the relational data model?
- TRUE or FALSE

Figure 5.7 - The 1:1 Relationship between Department and Employee

A One-to-One (1:1) Relationship:

An EMPLOYEE manages zero or one DEPARTMENT;
each DEPARTMENT is managed by one EMPLOYEE.



Design Case 1: Implementing 1:1 Relationships

- Foreign keys work with primary keys to properly implement relationships in relational model
- Rule (1:M)
 - Put the primary key of the “one” side (the parent entity) on the “many” side (the dependent entity) as a foreign key.
- Options for selecting and placing the foreign key (1:1):
 - Place a foreign key in both entities
 - Place a foreign key in one of the entities

Options for selecting and placing the foreign key (1:1)

- Place a foreign key in both entities
 - Place EMP_NUM as a foreign key in DEPARTMENT, and
 - Place DEPT_ID as a foreign key in EMPLOYEE.
 - Not recommended
 - Duplicates,
 - Can cause conflict with other existing relationships.
 - E.g. DEPARTMENT and EMPLOYEE also participate in a 1:M relationship—one department employs many employees.

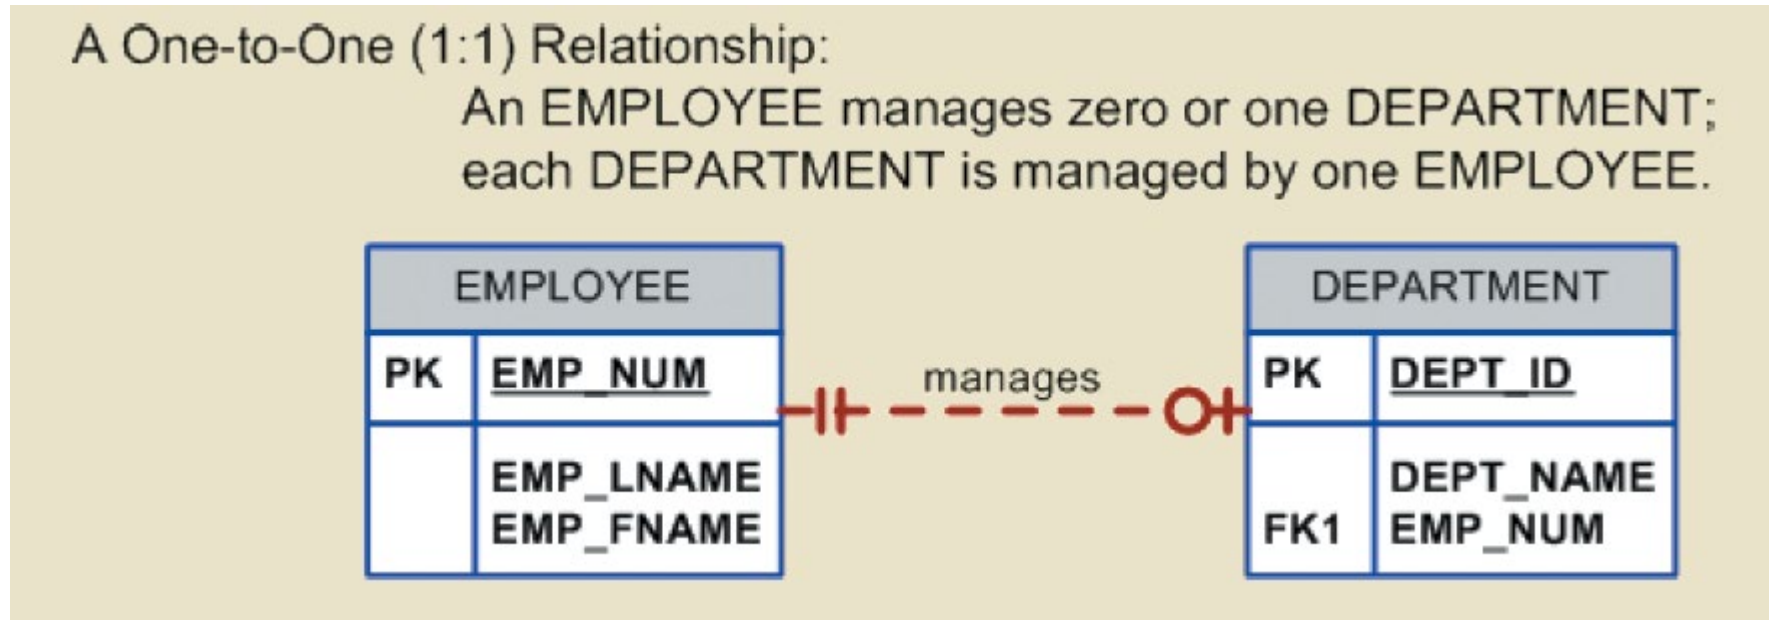
Options for selecting and placing the foreign key (1:1)

- Place a foreign key in one of the entities
 - the primary key of one of the two entities appears as a foreign key in the other entity.
 - The preferred solution
- Which primary key should be used as a foreign key?

Table 5.5 - Selection of Foreign Key in a 1:1 Relationship

| CASE | ER RELATIONSHIP CONSTRAINTS | ACTION |
|------|---|---|
| I | One side is mandatory and the other side is optional. | Place the PK of the entity on the mandatory side in the entity on the optional side as a FK, and make the FK mandatory. |
| II | Both sides are optional. | Select the FK that causes the fewest nulls, or place the FK in the entity in which the (relationship) role is played. |
| III | Both sides are mandatory. | See Case II, or consider revising your model to ensure that the two entities do not belong together in a single entity. |

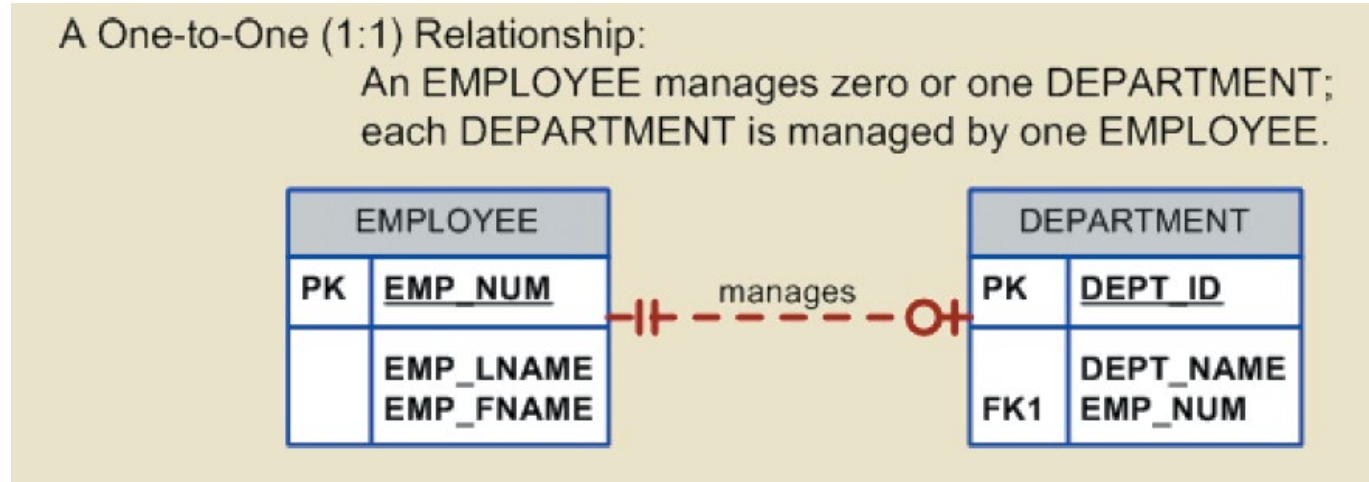
Figure 5.7 - The 1:1 Relationship between Department and Employee



EMP_NUM is placed as the foreign key in DEPARTMENT.

Why?

Figure 5.7 - The 1:1 Relationship between Department and Employee



Two Reasons:

- EMPLOYEE is mandatory to DEPARTMENT.
- “manager” role is played by the EMPLOYEE in the DEPARTMENT.

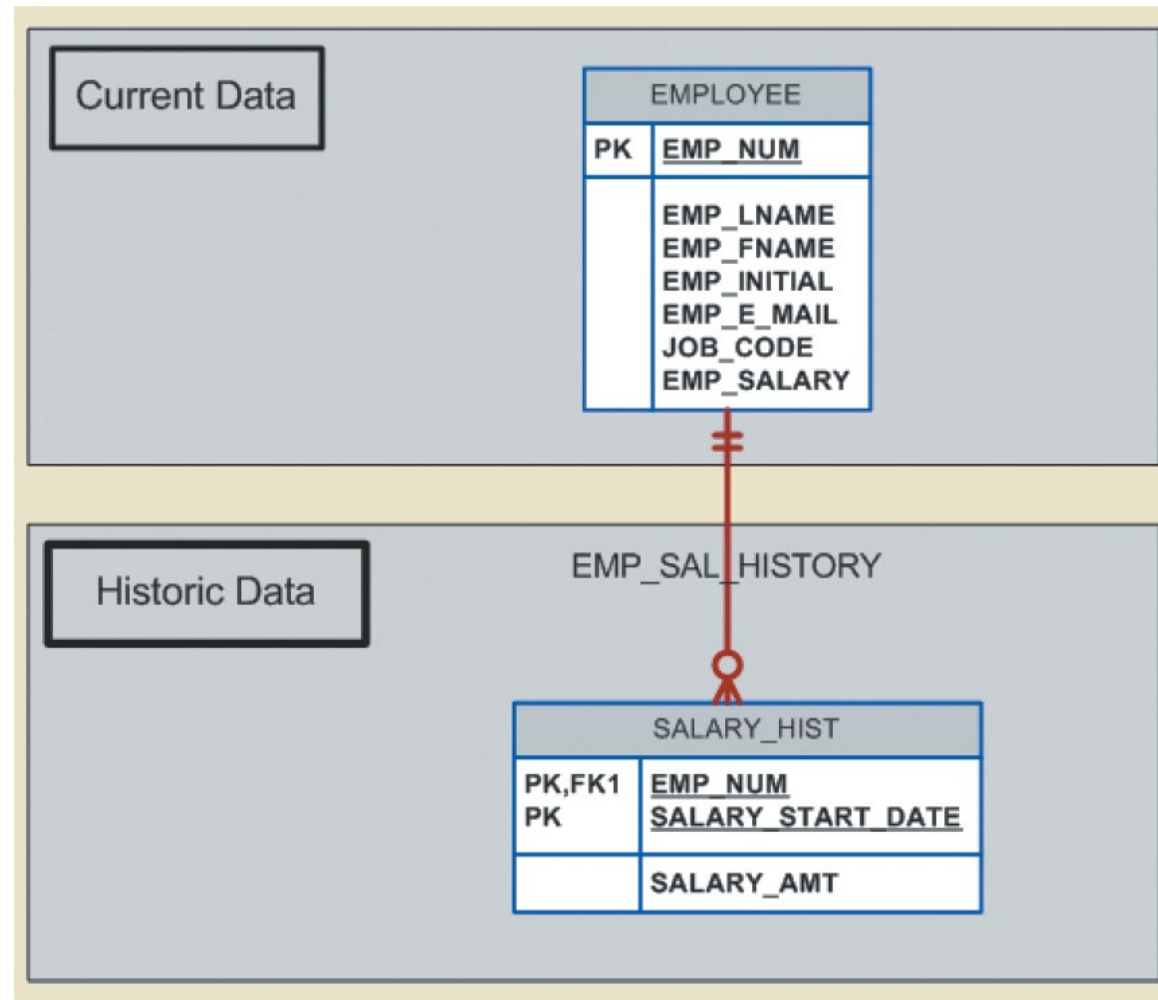
Design Case 2: Maintaining History of Time-Variant Data

- **Time-variant data:** Data whose values change over time and for which a history of the data changes must be retained
- The storage of time-variant data requires changes in the data model;
 - the type of change depends on the nature of the data
- Ex:
 - an employee's salary over the years
 - a department's manager over the years

Design Case 2: Maintaining History of Time-Variant Data

- Some time-variant data is equivalent to having a multivalued attribute in your entity.
 - An employee's salary over the years (multivalued-attribute);
 - Requires creating a new entity in a 1:M relationship with the original entity
 - New entity contains the new value, date of the change, and other pertinent attribute

Figure 5.8 - Maintaining Salary History



Design Case 2: Maintaining History of Time-Variant Data

- Other time-variant data can turn a 1:M relationship into an M:N relationship.
 - Each department is managed by only one employee
 - Each employee can manage one department at most
 - 1:1 relationship exist between EMPLOYEE and DEPARTMENT.
 - If you want to keep track of the history of all department managers as well as the current manager?

Figure 5.9 - Maintaining Manager History

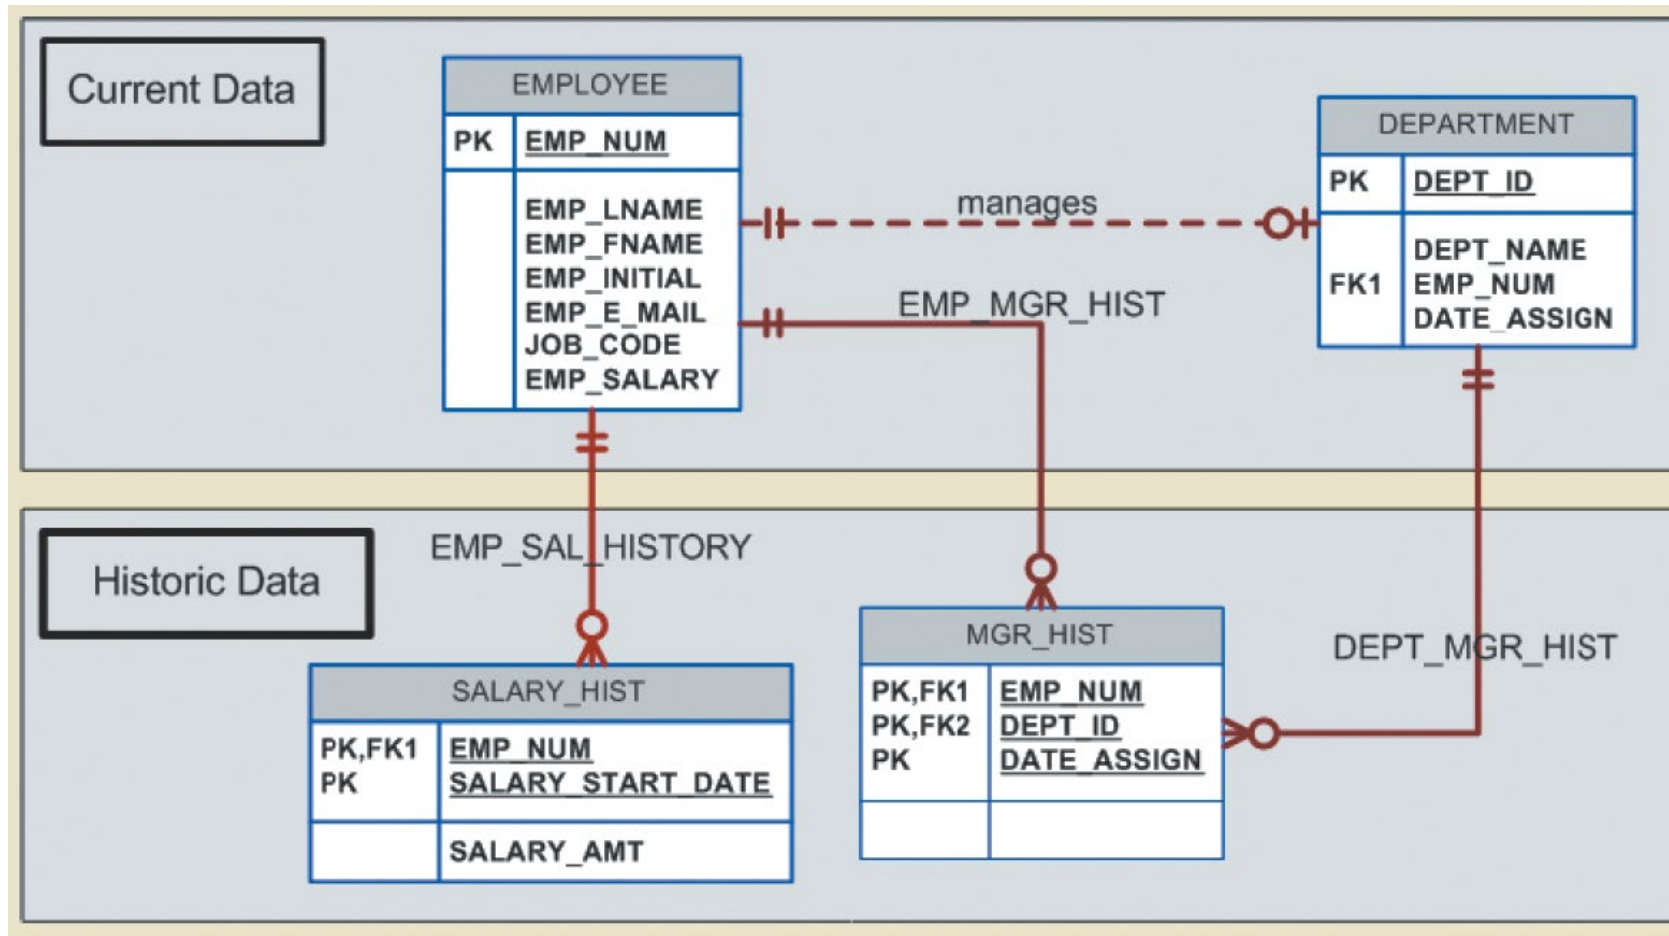
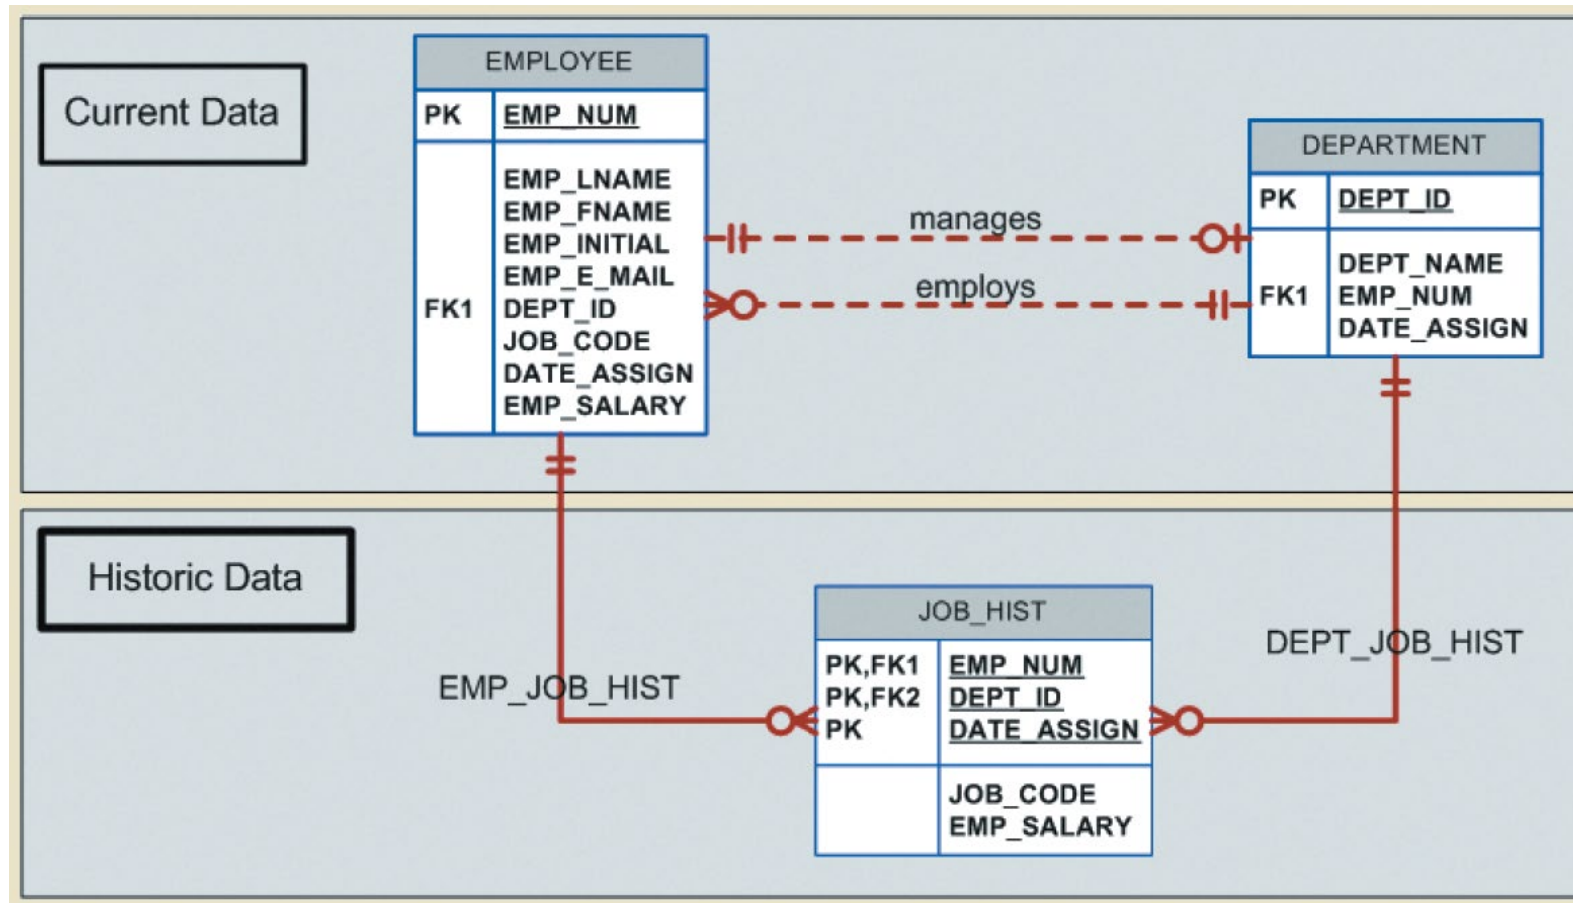


Figure 5.10 - Maintaining Job History



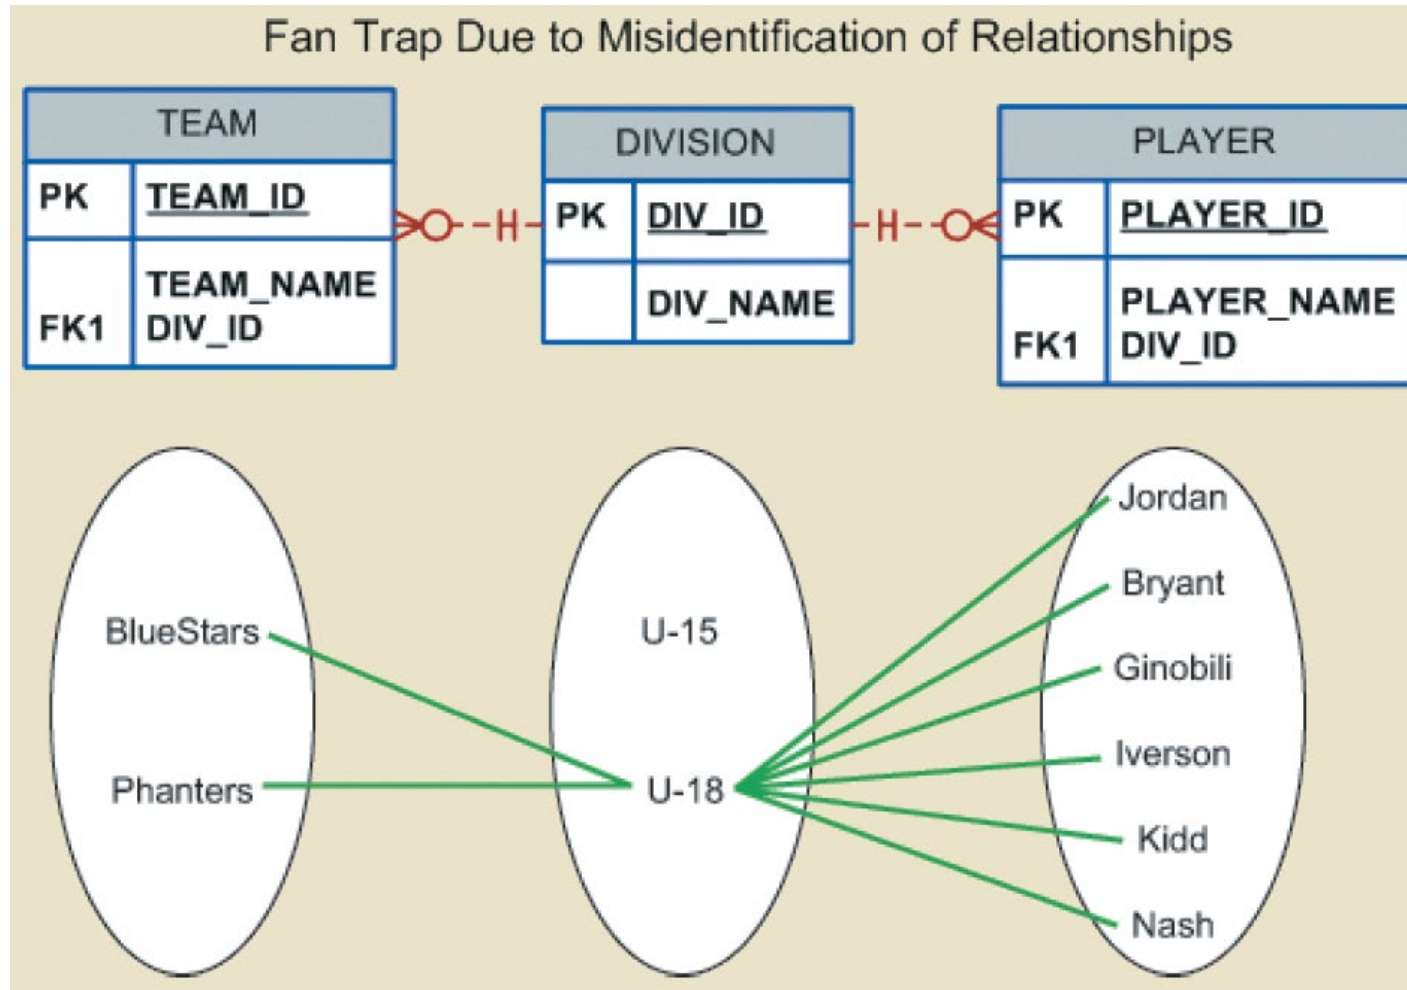
Design Case 3: Fan Traps

- **Design trap:** Occurs when a relationship is improperly or incompletely identified
 - Represented in a way not consistent with the real world
- **Fan trap:**
 - The most common type of design trap
 - Occurs when one entity is in two 1:M relationships to other entities
 - Produces an association among other entities not expressed in the model

Example

- A basketball league has many divisions.
- Each division has many players,
- Each division has many teams

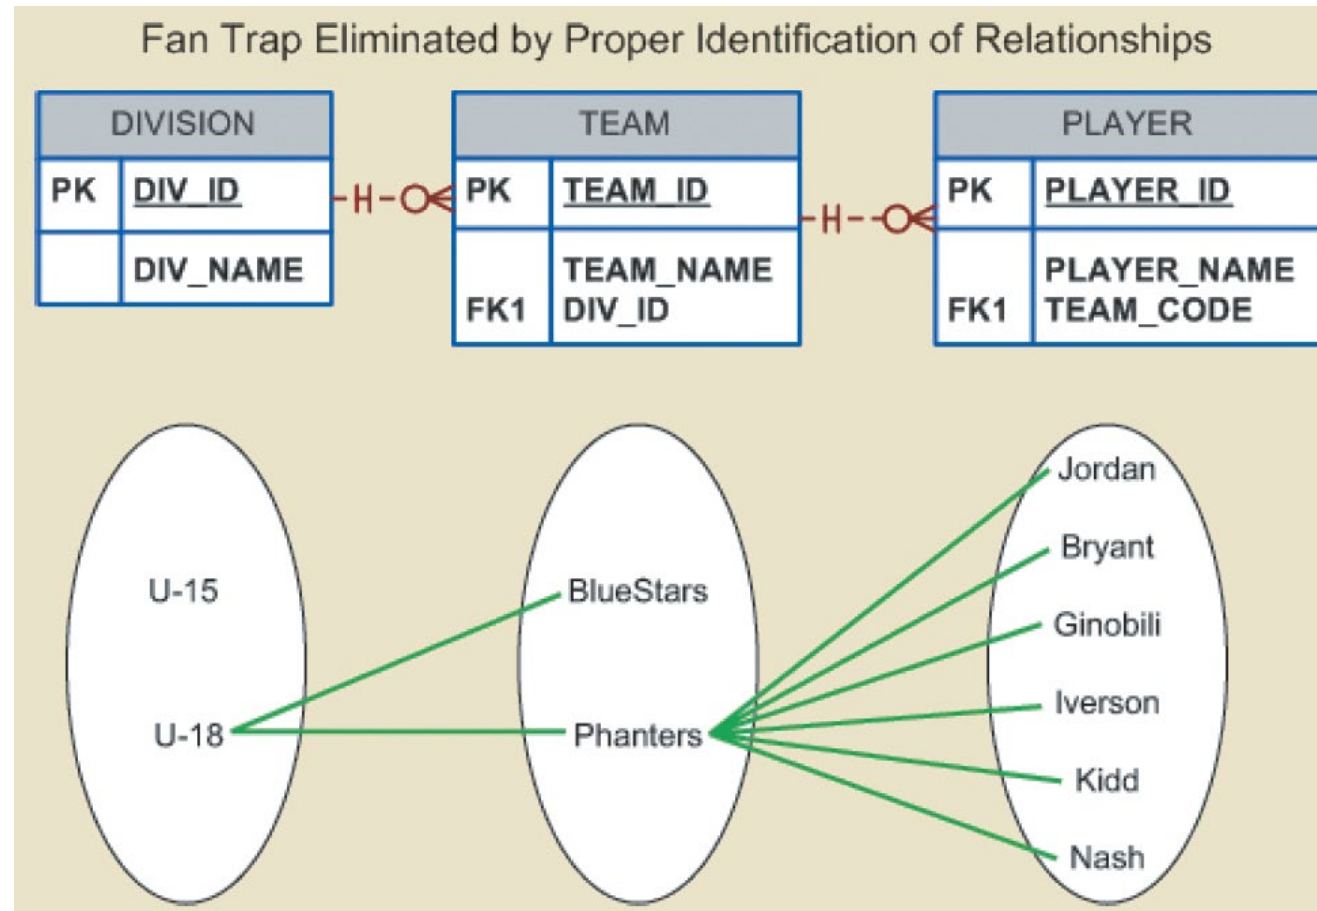
Figure 5.11 - Incorrect ERD with Fan Trap Problem



Although that representation is semantically correct, the relationships are not properly identified.

How do you identify which players belong to which team?

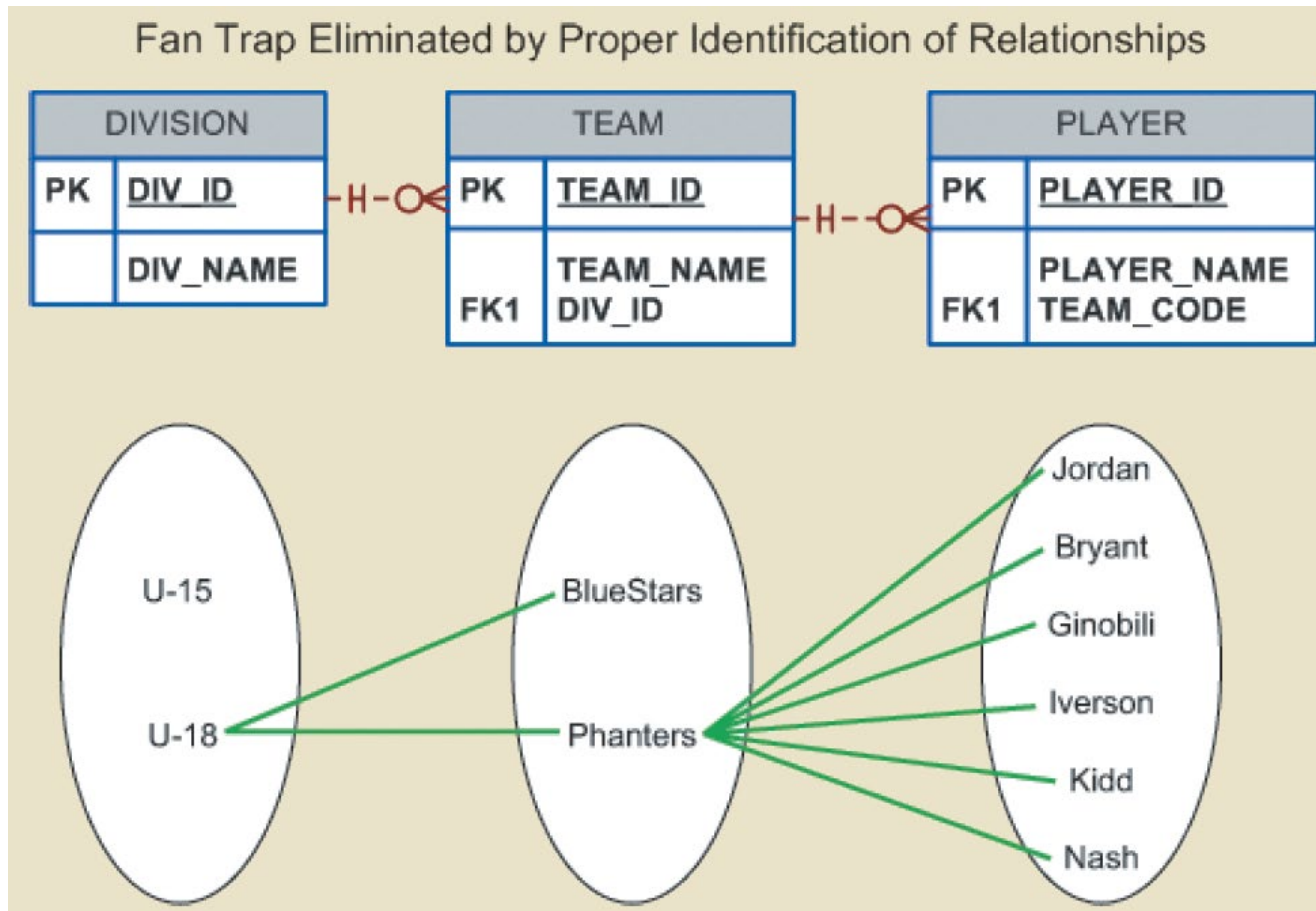
Figure 5.12 - Corrected ERD After Removal of the Fan Trap



Design Case 4: Redundant Relationships

- Occur when there are multiple relationship paths between related entities
- Need to remain consistent across the model
- Help simplify the design

Figure 5.12 - Corrected ERD After Removal of the Fan Trap



Which division
does a player
belong to ?

Figure 5.13 - A Redundant Relationship

