

Informed Search

Ian Sanders

Second Semester, 2024



Where to now?

- ▶ Previously *uninformed strategy* for node expansion
 - ▶ Shallowest
 - ▶ Deepest



Where to now?

- ▶ Previously *uninformed strategy* for node expansion
 - ▶ Shallowest
 - ▶ Deepest – depth first search
 - ▶ Smallest total cost from root
- ▶ Now
 - ▶ What if we know something about the problem?
 - ▶ How can we incorporate knowledge?
 - ▶ Task-specific expansion strategy?
 - ▶ What if we have an adversary?
 - ▶ Can't just find a goal
 - ▶ Opponent can prevent us from doing so!




From UCS to heuristics

- ▶ Recall: UCS is like BFS with priority queue
- ▶ Nodes ordered by priority from smallest to largest
- ▶ Priority is cost estimate $f(n)$
- ▶ In UCS, $f(n) = g(n)$, where g is cost of reaching n from root
- ▶ Now we include a *Heuristic function*
- ▶ $h(n)$ = estimate of cost from n to goal
- ▶ $g(n)$ is backward cost (start to node), h is (estimated) forward cost (node to goal)
- ▶ So $f(n) = g(n) + h(n)$
- ▶ This leads to the idea of *best first search* – a generalisation of “normal” tree and graph searches.



Greedy best first search

- ▶ UCS is $f(n) = g(n)$ – no heuristic
- ▶ GBFS is $f(n) = h(n)$ – $h(n)$ is the heuristic, the estimate of the future cost to the goal
Uses the heuristic in the hope that it is a good guide to the solution.



GBFS

- ▶ Order nodes in the priority queue by estimated cost to goal
- ▶ Where does this estimate come from?
- ▶ Domain knowledge!
- ▶ For example, in choosing the shortest path from one town to another in a graph of towns, we could choose the *direct distance* (“as the crow flies”) between two towns as the heuristic to use.




GBFS properties

- ▶ Incomplete for tree search (i.e. repeated nodes allowed)
- ▶ Complete for graph search in finite space
- ▶ Space and time complexity are both $O(b^m)$ for max search depth
- ▶ But can be reduced substantially depending on the heuristic and problem



A* search

- ▶ A* search is the most well known of all *best first searches*
- ▶ Recall
UCS: $f(n) = g(n)$
GBFS: $f(n) = h(n)$
- ▶ For A* Search we use $g(n)$ the known path cost and $h(n)$ which is an estimate of *future cost*.




- ▶ Implement A* as UCS with different priority!



A* vs UCS vs GBFS

- ▶ Uniform-cost orders by path cost, or backward cost $g(n)$
- ▶ Greedy orders by goal proximity, or forward cost $h(n)$
- ▶ A* Search orders by the sum: $f(n) = g(n) + h(n)$




A* vs UCS vs GBFS

Search	expand	frontier
UCS	S	$a, g = 1$
GBFS	S	$a, h = 6$
A*	S	$a, f = 1 + 6 = 7$
UCS	a	$d, g = 4$ $b, g = 5$ $e, g = 9$
GBFS	a	$e, h = 1$ $d, h = 3$ $b, h = 5$
A*	a	$d, f = 4 + 2 = 6$ $b, f = 2 + 6 = 8$ $e, f = 9 + 1 = 10$



Optimality of A*

- Depends on heuristic



- Expand S
To $G - f$ is 5
To $A - f$ is 7
- First choice is $S \rightarrow G$
Not the the best of options
A better heuristic should have made A a better option.



Heuristics in A*

- ▶ h admissible if $h(n) \leq \text{TrueFutureCost}(n)$
We never overestimate the cost!
- ▶ $f(n) = g(n) + h(n)$ means we never overestimate cost from start to goal through n
- ▶ Optimistic: estimates cost as less than it actually is
- ▶ h is consistent if $h(n) \leq c(n, a, n') + h(n')$
 n' is a successor of n , a is the action, c step cost to reach n' from n
The triangle inequality
- ▶ All consistent heuristics are admissible
- ▶ A* tree search is optimal if h is admissible
- ▶ A* graph search is optimal if h is consistent



Properties of A*

- ▶ A* is complete, optimal
- ▶ A* is optimally efficient w.r.t. heuristic
No other algorithm will expand fewer nodes
- ▶ Time complexity is $O(b^{\epsilon d})$
 ϵ is relative error of heuristic, d is solution depth
- ▶ Alternate view, $O(b^{*d})$ where b^* is effective branching factor
- ▶ A* doesn't need to consider certain nodes (prunes) and so branching factor is reduced!
- ▶ Since exponential, memory is biggest issue
- ▶ Can do iterative deepening equivalent (IDA*)



Creating heuristics

- ▶ Main challenge for A*: how to make good heuristics that are admissible/consistent?
- ▶ For navigation, straight-line path is good
- ▶ Can never be faster than that!
- ▶ Could learn heuristics from data (e.g. machine learning/precomputed database lookups)
- ▶ Use relaxations – solve an easier, less constrained version of a problem
- ▶ E.g. Relaxed version of a maze – remove all the walls
- ▶ Tradeoff between quality of estimate and work per node
- ▶ Closer heuristic is to true cost, the fewer nodes are expanded...



Adversarial search

- ▶ So far, aim was to reach goal with min cost
- ▶ But what if another agent is trying to stop us?
- ▶ Must take into account their actions



Games are big

The size of the state space in some games....

- ▶ Tic-tac-toe $\approx 10^3$
- ▶ Connect Four $\approx 10^3$
- ▶ English draughts $\approx 10^{23}$
- ▶ Othello $\approx 10^{28}$
- ▶ Chess $\approx 10^{44}$
- ▶ Shogi $\approx 10^{71}$
- ▶ # atoms in observable universe $\approx 10^{82}$
- ▶ Twixt $\approx 10^{140}$
- ▶ Go (19x19 board) $\approx 10^{170}$



Games are hard

- ▶ NP-Hard or harder

If a problem is NP-hard, then it is at least as difficult to solve as the problems in NP

A Problem X is NP-Hard if there is an NP-Complete problem Y, such that Y is reducible to X in polynomial time.

NP-Hard problems are as hard as NP-Complete problems.

NP-Hard Problem need not be in NP class.

NP-Complete is a subset of NP-Hard



Zero-sum games

- ▶ Assume two players, competitive
- ▶ Player 1 wins, player 2 loses and vice versa
- ▶ Game is defined by:
 - ▶ Initial state
 - ▶ $Player(s)$: whose turn it is
 - ▶ $Actions(s)$: available actions
 - ▶ $Result(s, a)$: successor function or transition model
 - ▶ $Terminal(s)$: is the game over/state terminal
 - ▶ $Utility(s, p)$: the value for the game ending in s for player p
- ▶ Zero sum game: sum of utilities for all players is constant:
e.g. Win = +1, Draw = 0, Loss = -1



Two player, zero-sum

- ▶ Two players, MAX and MIN
- ▶ We are MAX, try to maximise utility
- ▶ Opponent is MIN, tries to minimise utility
- ▶ Denote $V(s)$ as utility at a given state
- ▶ Utility is known at terminal states
- ▶ Start at root node, expand tree
- ▶ Players alternate turns
- ▶ At level 0, us to play. Level 1, them to play, etc.
Each level is called a *ply*.
- ▶ We want to compute optimal play, assuming our opponent is also optimal

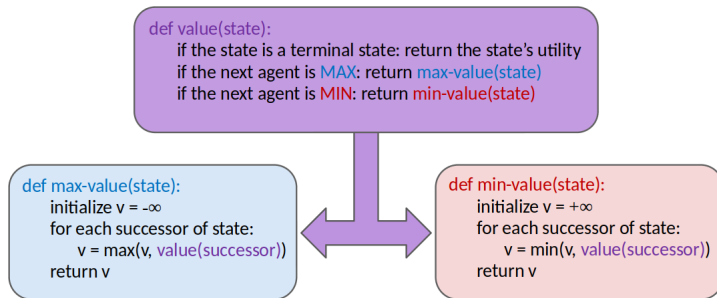


Calculating minimax

- ▶ Want to calculate $V(s)$ for all s
- ▶ If s is terminal, use utility function directly
- ▶ else
 - ▶ if player to play is MAX:
 - ▶ Value is best maximising value at state
 - ▶ else player to play is MIN:
 - ▶ Value is best minimising value at state



Calculating minimax



$MiniMax(s)$


$= Utility(s), \text{ If Terminal} - Test(s)$

$= \max_{a \in Actions(s)} MiniMax(Result(s, a)), \text{ If Player}(s) = Max$


$= \min_{a \in Actions(s)} MiniMax(Result(s, a)), \text{ If Player}(s) = Min$




Example



Example



Example




Minimax properties

- ▶ Like DFS:
- ▶ Time: $O(b^m)$
- ▶ Space: $O(b^m)$
- ▶ But instead of searching for single goal, we need to *exhaustively* try everything!
And we only get value at leaf nodes
- ▶ Chess, for e.g., $b \approx 20$, $m \approx 70$
Exact solution *infeasible*
So what do we do?




Pruning



Max to play.



Pruning



Max gets at least 3 – from the minimum of the left subtree

Min would consider *H* from this can get 2 for node *C* but this would not be chosen anyway as Max already has 3

If *I* and/or *J* where greater than 2 then Min would not choose them.

If *I* and/or *J* where less than 2 then Min could choose them but it would make no difference up the tree

I and/or *J* can be pruned!




Where to now?

- ▶ $\alpha\beta$ -pruning
- ▶ α – minimum score MAX is guaranteed of
- ▶ β – maximum score MIN is guaranteed of
- ▶ If at a given node, $\beta < \alpha$
- ▶ Then MIN can guarantee a score that makes MAX sad
- ▶ So MAX will never go down this road
- ▶ No need to expand rest of node's children!
- ▶ Symmetric argument for other way around
- ▶ If children are expanded in optimal order, complexity is halved: $O(b^{m/2})$



Pruning



$\alpha\beta$ Search Algorithm

```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```

Figure 5.7 The alpha-beta search algorithm. Notice that these routines are the same as the MINIMAX functions in Figure 5.3, except for the two lines in each of MIN-VALUE and MAX-VALUE that maintain α and β (and the bookkeeping to pass these parameters along).



Tracing $\alpha\beta$ pruning

1. A: max – $v = -\infty, \alpha = -\infty, \beta = \infty$ expand B
2. B: min – $v = \infty, \alpha = -\infty, \beta = \infty$, expand E
3. E: max – terminal state true, returns utility value, ie returns 3
4. B: min – was $v = \infty, \alpha = -\infty, \beta = \infty$
 v min of v and returned value
 $v = MIN(\infty, 3) = 3$
 v *not* $\leq \alpha$ so do not return now
 $v = \cancel{\infty} 3, \alpha = -\infty, \beta = \cancel{\infty} 3$, expand F
5. F: max – terminal state true, returns utility value, ie returns 12
6. B: min – $v = 3, \alpha = -\infty, \beta = 3$, expand G
7. F: max – terminal state true, returns utility value, ie returns 8
8. B: min – $v = 3, \alpha = -\infty, \beta = 3$, returns $v = 3$ to A



9. A: $\max - v = -\infty, \alpha = -\infty, \beta = \infty$

$$v = \text{MAX}(v, 3) = 3$$

v not $\geq \beta$ so do not return

$$\alpha = \text{MAX}(\alpha, v) = 3$$

$$v = \cancel{-\infty} 3, \alpha = \cancel{-\infty} 3, \beta = \infty$$

Expand C

10. C: $\min - v = \infty, \alpha = 3, \beta = \infty$, expand H

11. H: \max – terminal state true, returns utility value, ie returns 2

12. C: \min

before expanding H – $v = \infty, \alpha = 3, \beta = \infty$

v min of v and returned value

$$v = \text{MIN}(\infty, 2) = 2$$

now $v \leq \alpha$ so return 2 to A

This effectively prunes nodes I and J from the search tree.



13. A: $\max - v = 3, \alpha = 3, \beta = \infty$ when calling C
 $v = \text{MAX}(v, 2) = 2$
 v not $\geq \beta$
 $\alpha = \text{MAX}(\alpha, v) = \text{MAX}(3, 2) = 3$
So $v = 3, \alpha = 3, \beta = \infty$
Expand D
14. D: $\min - v = \infty, \alpha = 3, \beta = \infty$, Expand K
15. Continue as an exercise.



Depth-limited search

- ▶ Even with pruning, can't reach leaf nodes in real games
- ▶ So must limit depth of search
- ▶ Must replace utility function with estimate (like heuristic in A*)
- ▶ No longer optimal
- ▶ More plies = better performance
- ▶ Given time budget, use IDS!



Evaluation functions

- ▶ Estimate of utility of non-terminal state
- ▶ Ideally: want actual minimax value of state
- ▶ But this is unknown
- ▶ In practice: use domain knowledge
- ▶ Tradeoff between complexity vs depth
- ▶ More complex evaluation function may be more accurate, but take longer to compute therefore less time to search deeper
- ▶ Stockfish: fast eval function, huge depth
- ▶ Komodo: slow, complex eval function, less depth



Other improvements

- ▶ Base algorithm of $\alpha\beta$ + IDS + eval function
- ▶ Transposition tables: stored previous states and their evals
- ▶ Aspiration windows: pretend that the $\alpha\beta$ window is smaller than it is
- ▶ Evaluation functions optimised from data (machine learning, etc.)
- ▶ Move ordering: try certain classes of moves first (e.g. captures, then regular moves)



Summary

- ▶ For single-agent search:
- ▶ Uninformed node expansion (DFS/BFS/UCS)
- ▶ Informed with heuristics (GBFS/A*)
- ▶ A* optimal, but need admissible/consistent heuristics
- ▶ Heuristics from problem relaxation
- ▶ Adversarial search:
- ▶ Minimax for optimal play
- ▶ Can use pruning to reduce nodes
- ▶ In practice, must cut-off depth
- ▶ Use evaluation function

Primary focus of AI for 60 years!

