

Operating Systems COMS(3010A)

Kernels and Processes 2




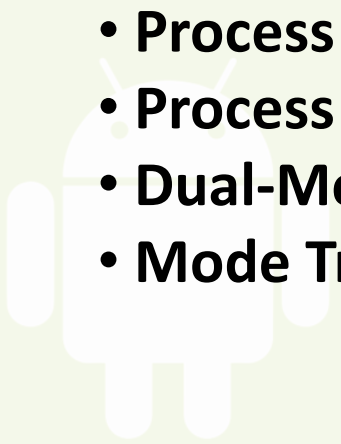
Branden Ingram

branden.ingram@wits.ac.za

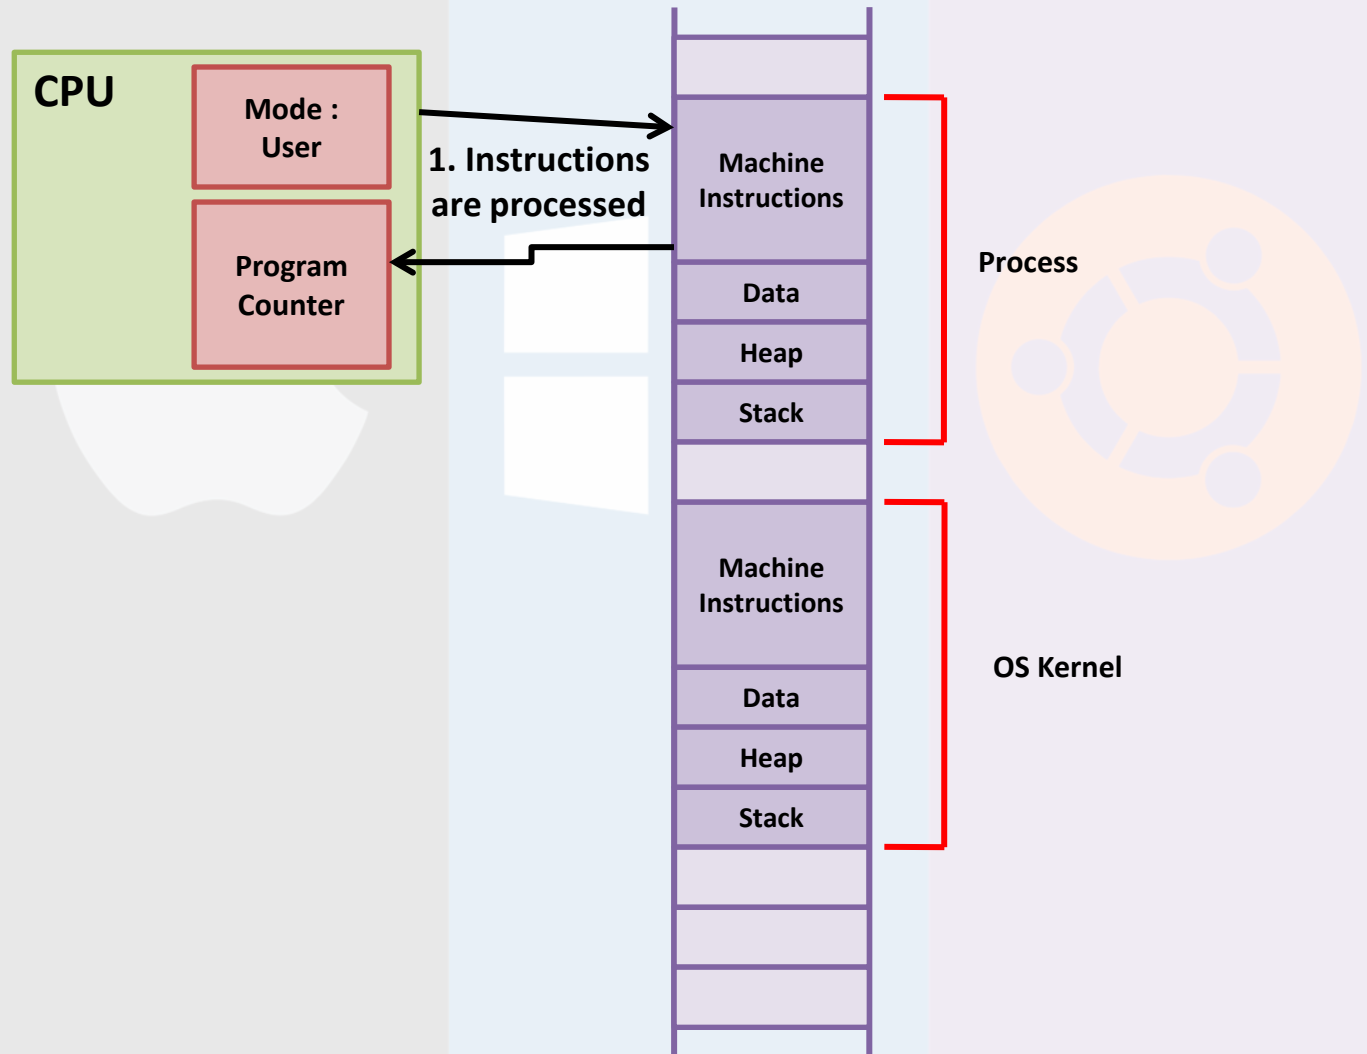
Office Number : ???

Recap

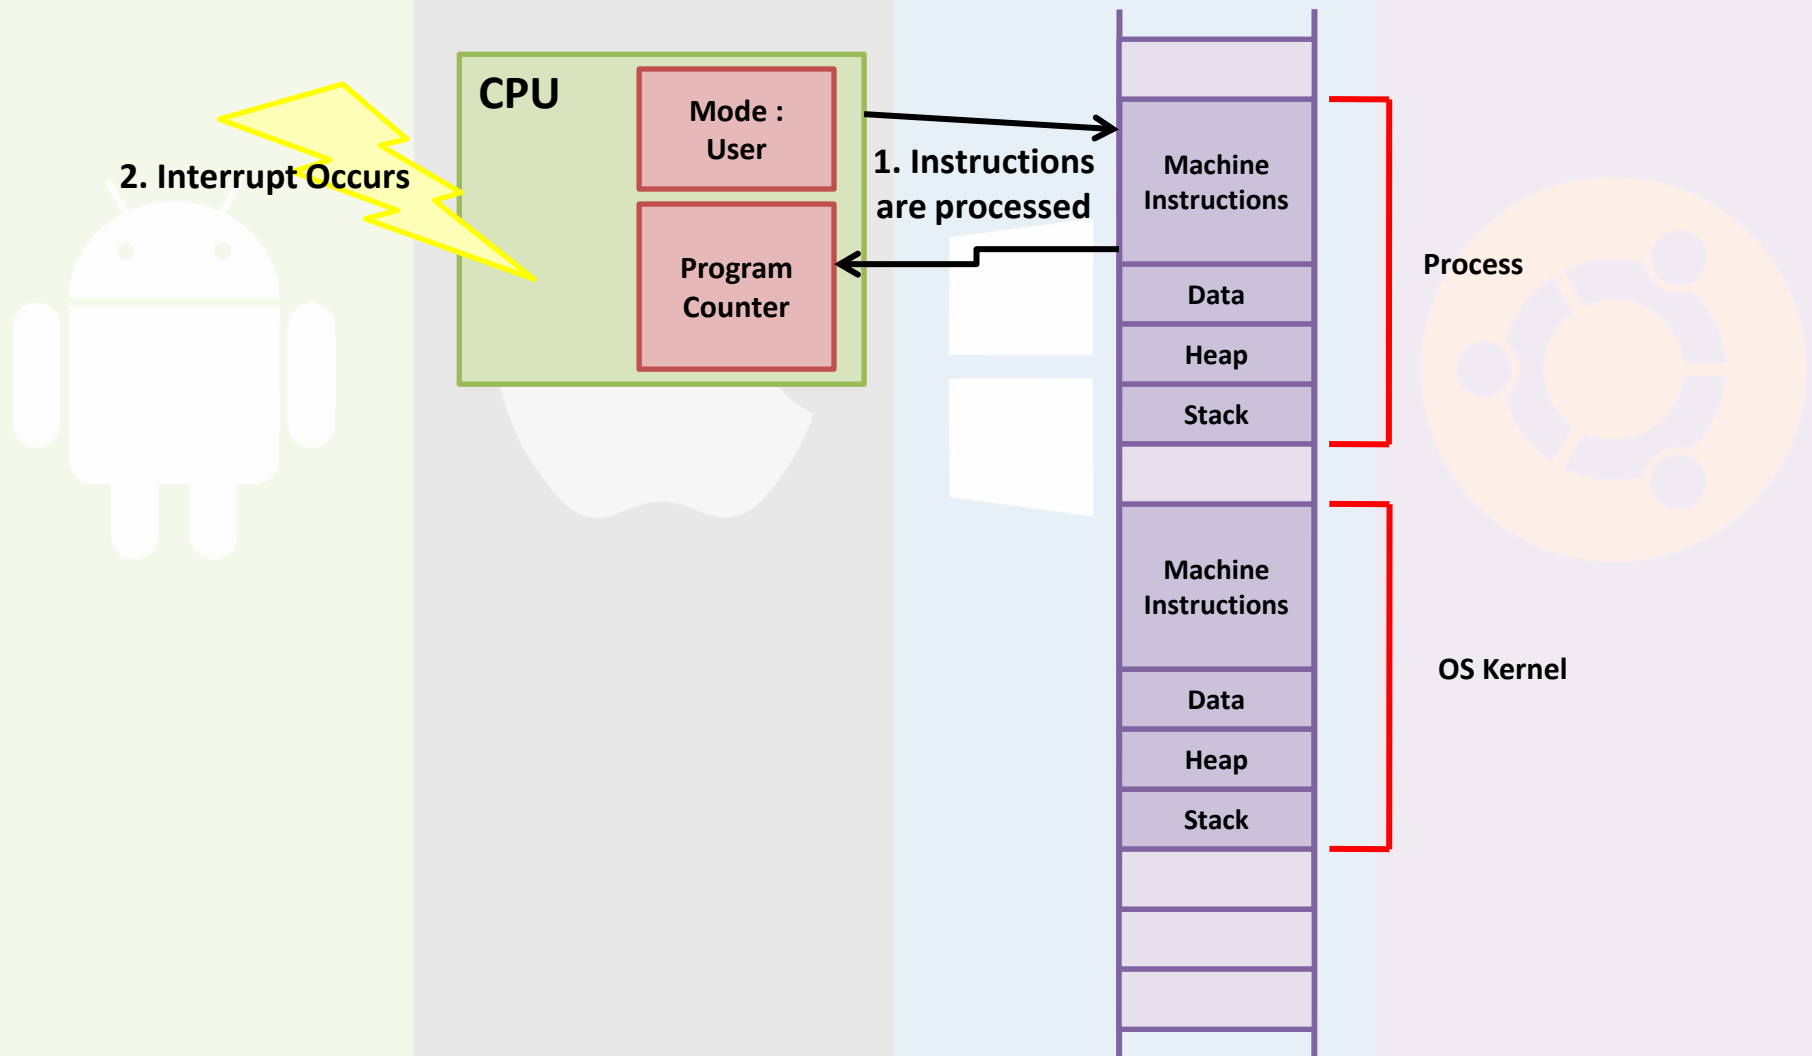
- Kernel
- Process Abstraction
- Process Control Block
- Dual-Mode Operation
- Mode Transfer



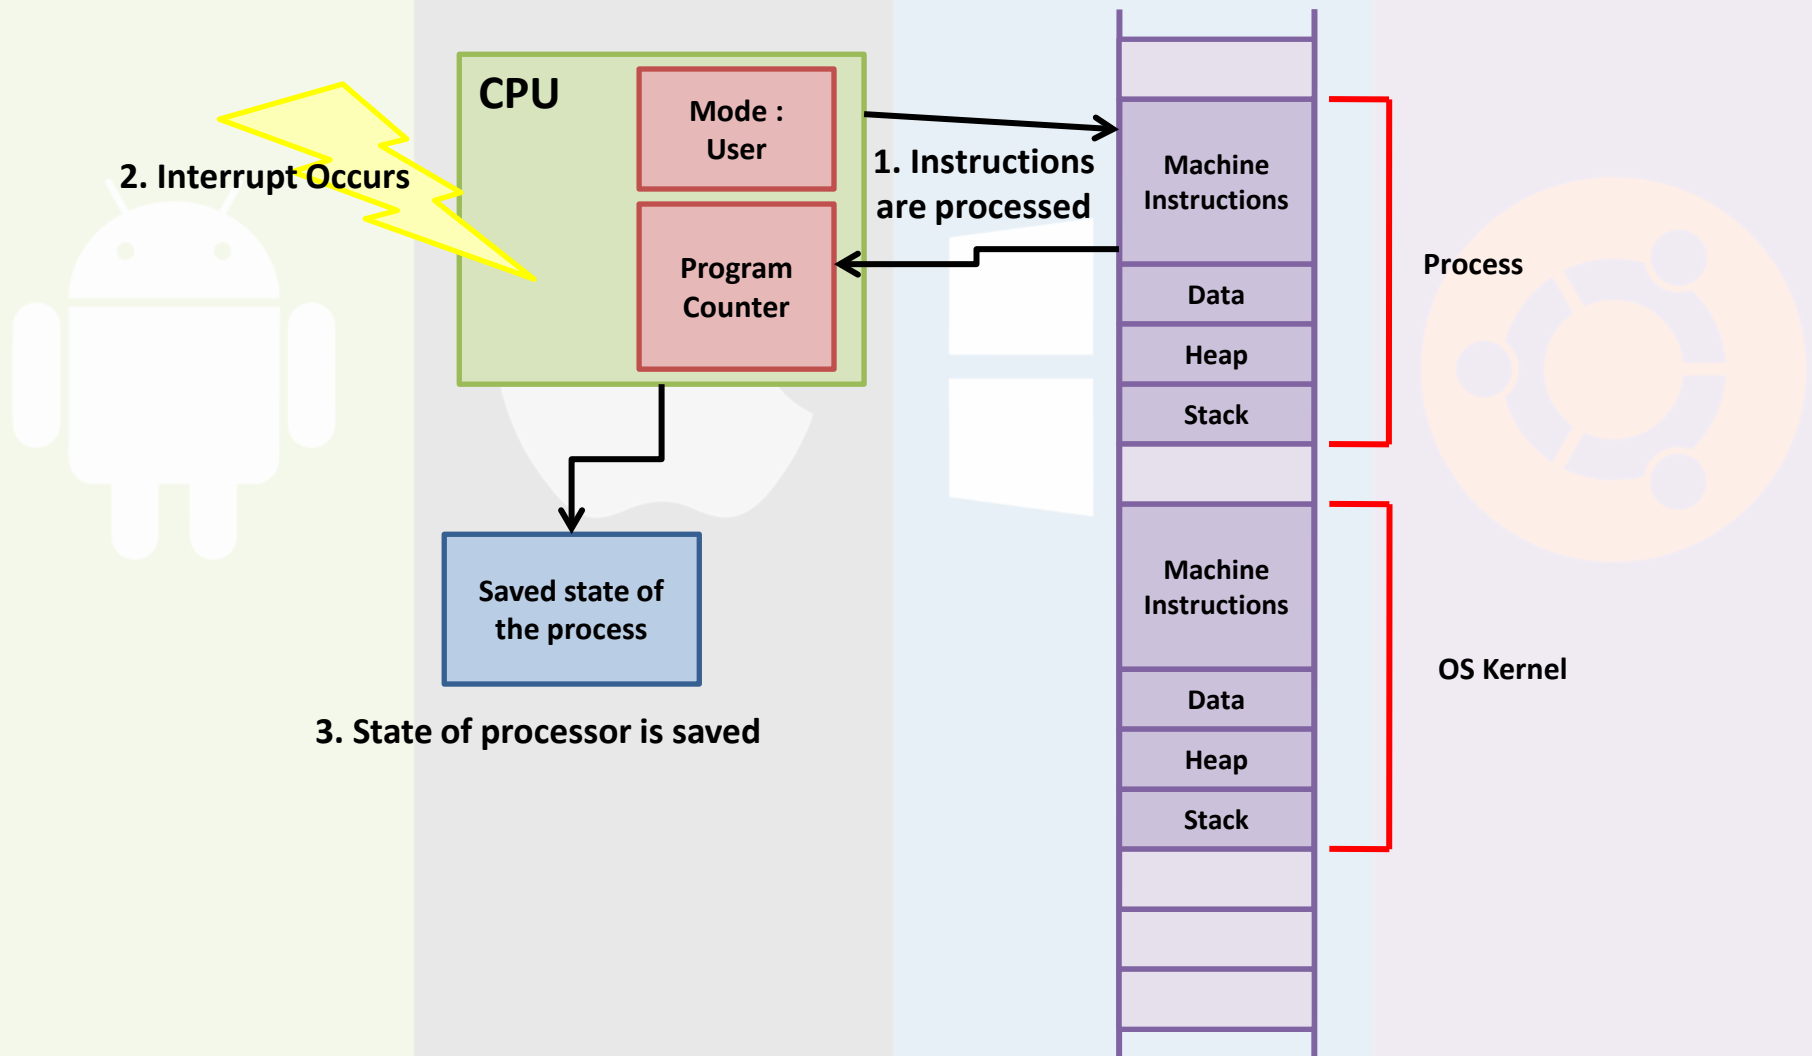
Summary



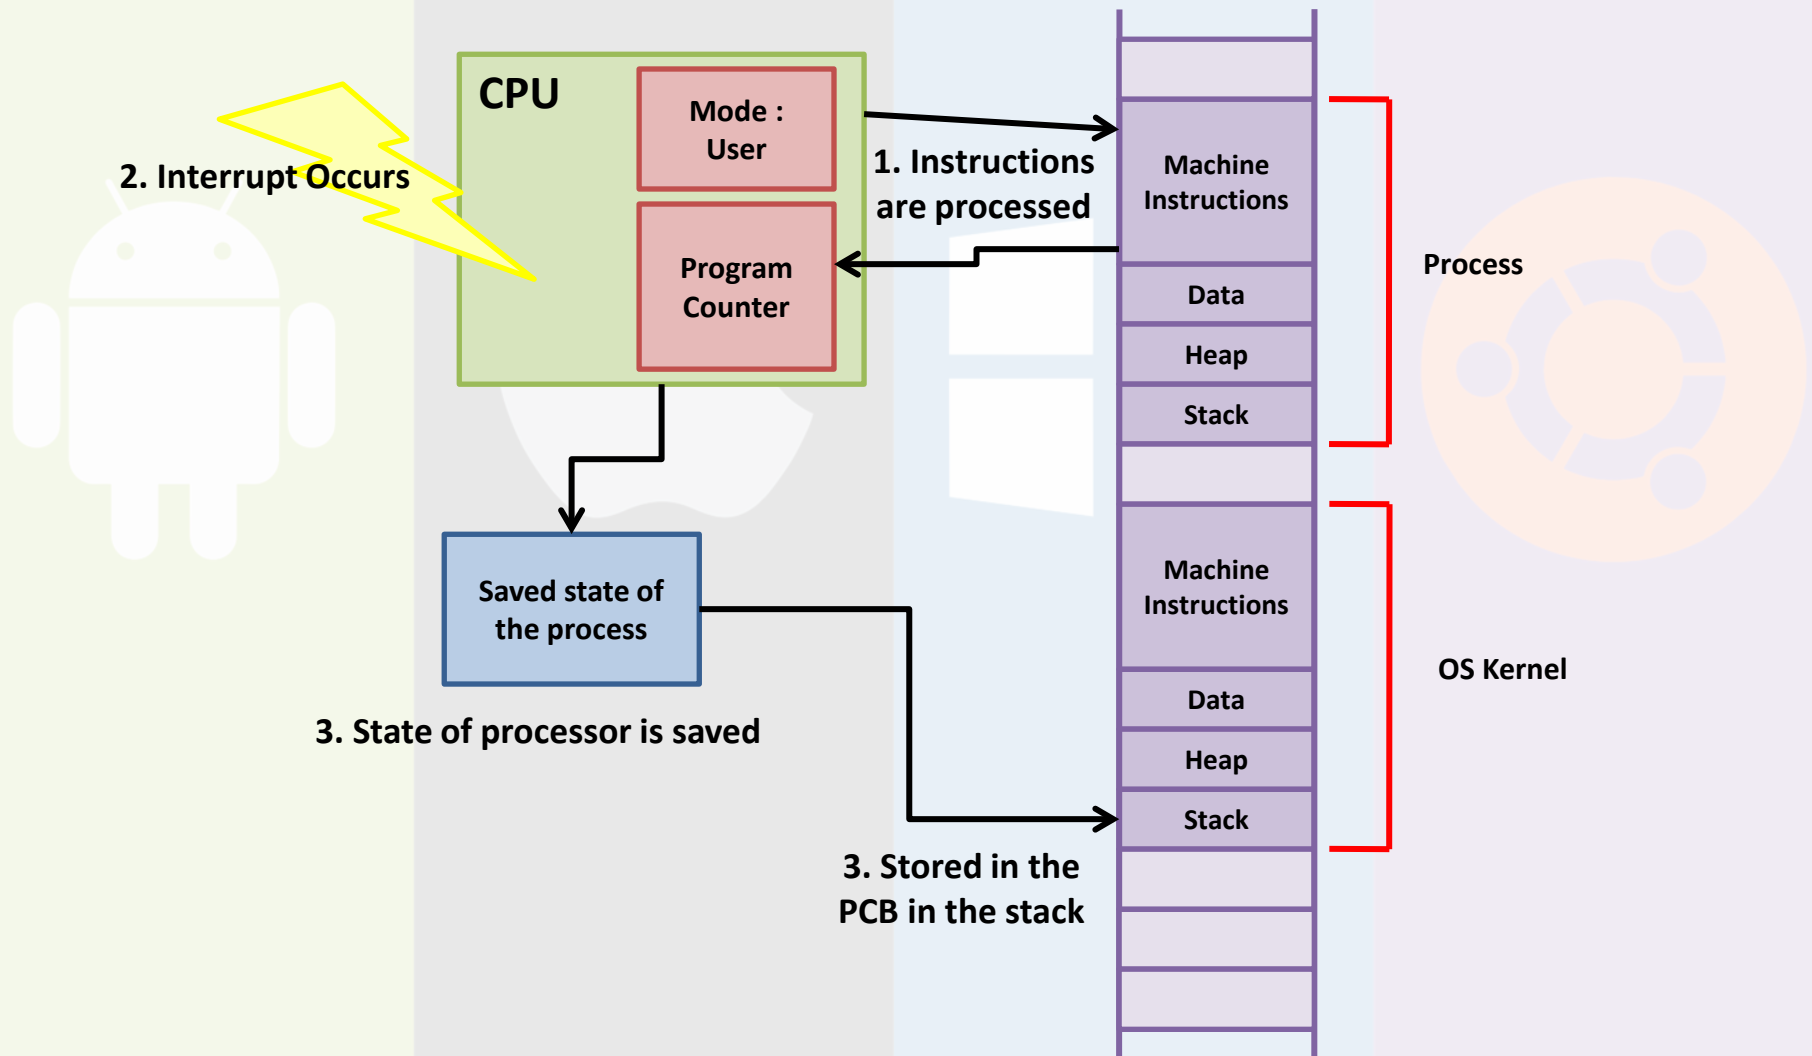
Summary



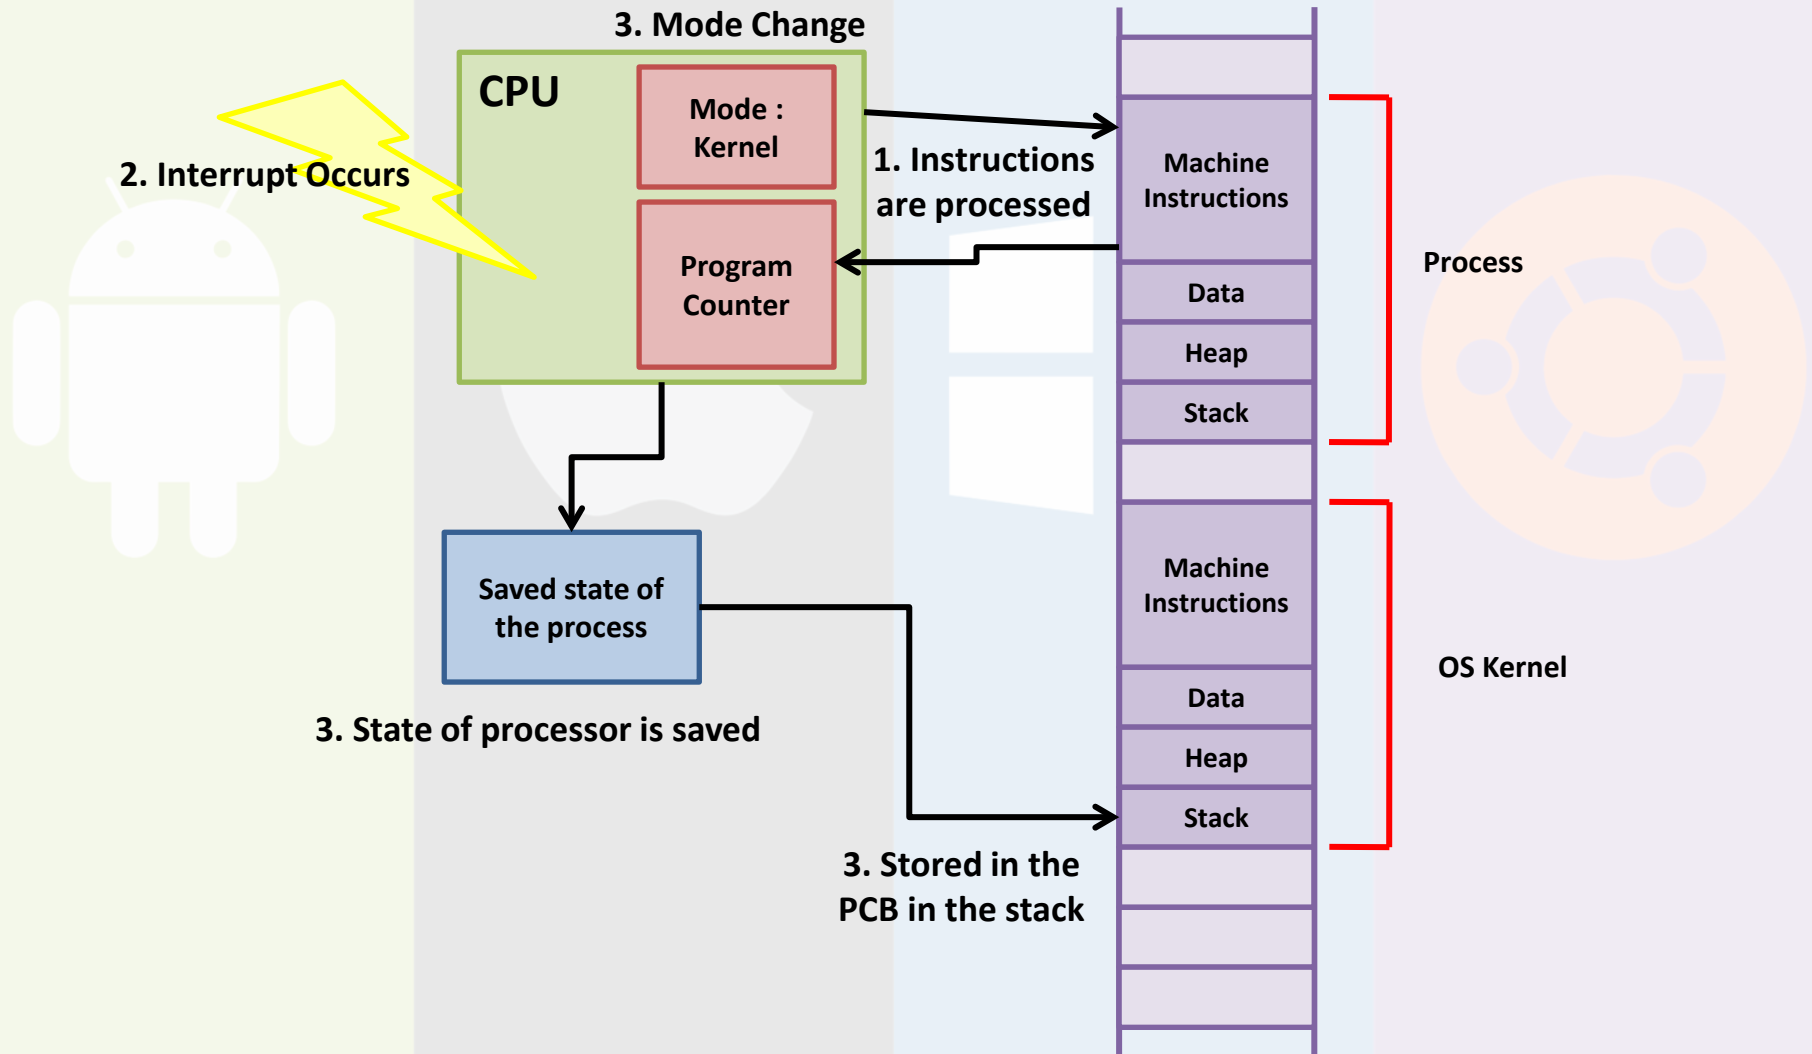
Summary



Summary




Summary



Summary

1. Interrupt Processed



CPU

Mode :
Kernel

Program
Counter

Saved state of
the process

2. State of processor is loaded

2. Loaded from
the PCB in the
stack

Machine
Instructions

Data

Heap

Stack

Machine
Instructions


Data

Heap

Stack


Process

OS Kernel




Summary

1. Interrupt Processed




2. Mode Change



Saved state of the process

2. State of processor is loaded

2. Loaded from the PCB in the stack




Process


OS Kernel

Summary

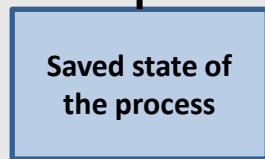
1. Interrupt Processed



2. Mode Change




3. Instructions are processed



2. State of processor is loaded

2. Loaded from the PCB in the stack




Process

OS Kernel

Implementing Safe Mode Transfer

- Care needs to be taken when implementing mode transfer to ensure malicious programs cannot gain access to the kernel

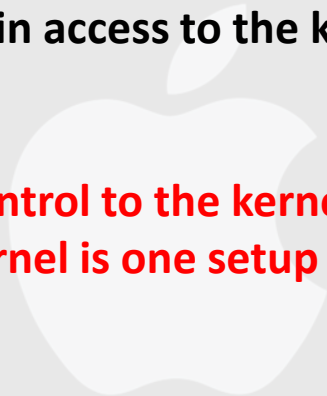



Implementing Safe Mode Transfer

- Care needs to be taken when implementing mode transfer to ensure malicious programs cannot gain access to the kernel

- **Limited entry**

- To transfer control to the kernel, the hardware must ensure that the entry point to the kernel is one setup by the kernel





Implementing Safe Mode Transfer

- Care needs to be taken when implementing mode transfer to ensure malicious programs cannot gain access to the kernel

- **Atomic changes to processor state**

- Changing of program counter, mode, stack and memory protection are all changed at the same time



Implementing Safe Mode Transfer

- Care needs to be taken when implementing mode transfer to ensure malicious programs cannot gain access to the kernel

- **Transparent, restartable execution**

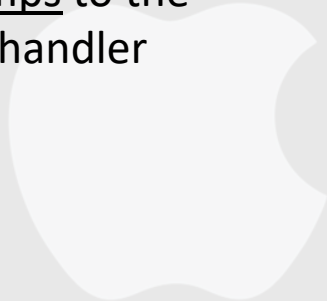
- **Interrupts should be invisible to the user processes**

```
1  #include <stdio.h>
2
3  int main()
4  {
5
6      int i = 0;
7      i = i/0;
8      printf("Hello World");
9
10     return 0;
11 }
```

Stop and Continue at the same point

Implementing Safe Mode Transfer

- On an **interrupt** the **processor** saves its **current state** to memory, changes to kernel mode and jumps to the exception/interrupt handler



System Calls

OS @ boot
(kernel mode)

initialize trap table

Hardware

remember address of ...
syscall handler

OS @ run
(kernel mode)

Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from -trap

Hardware

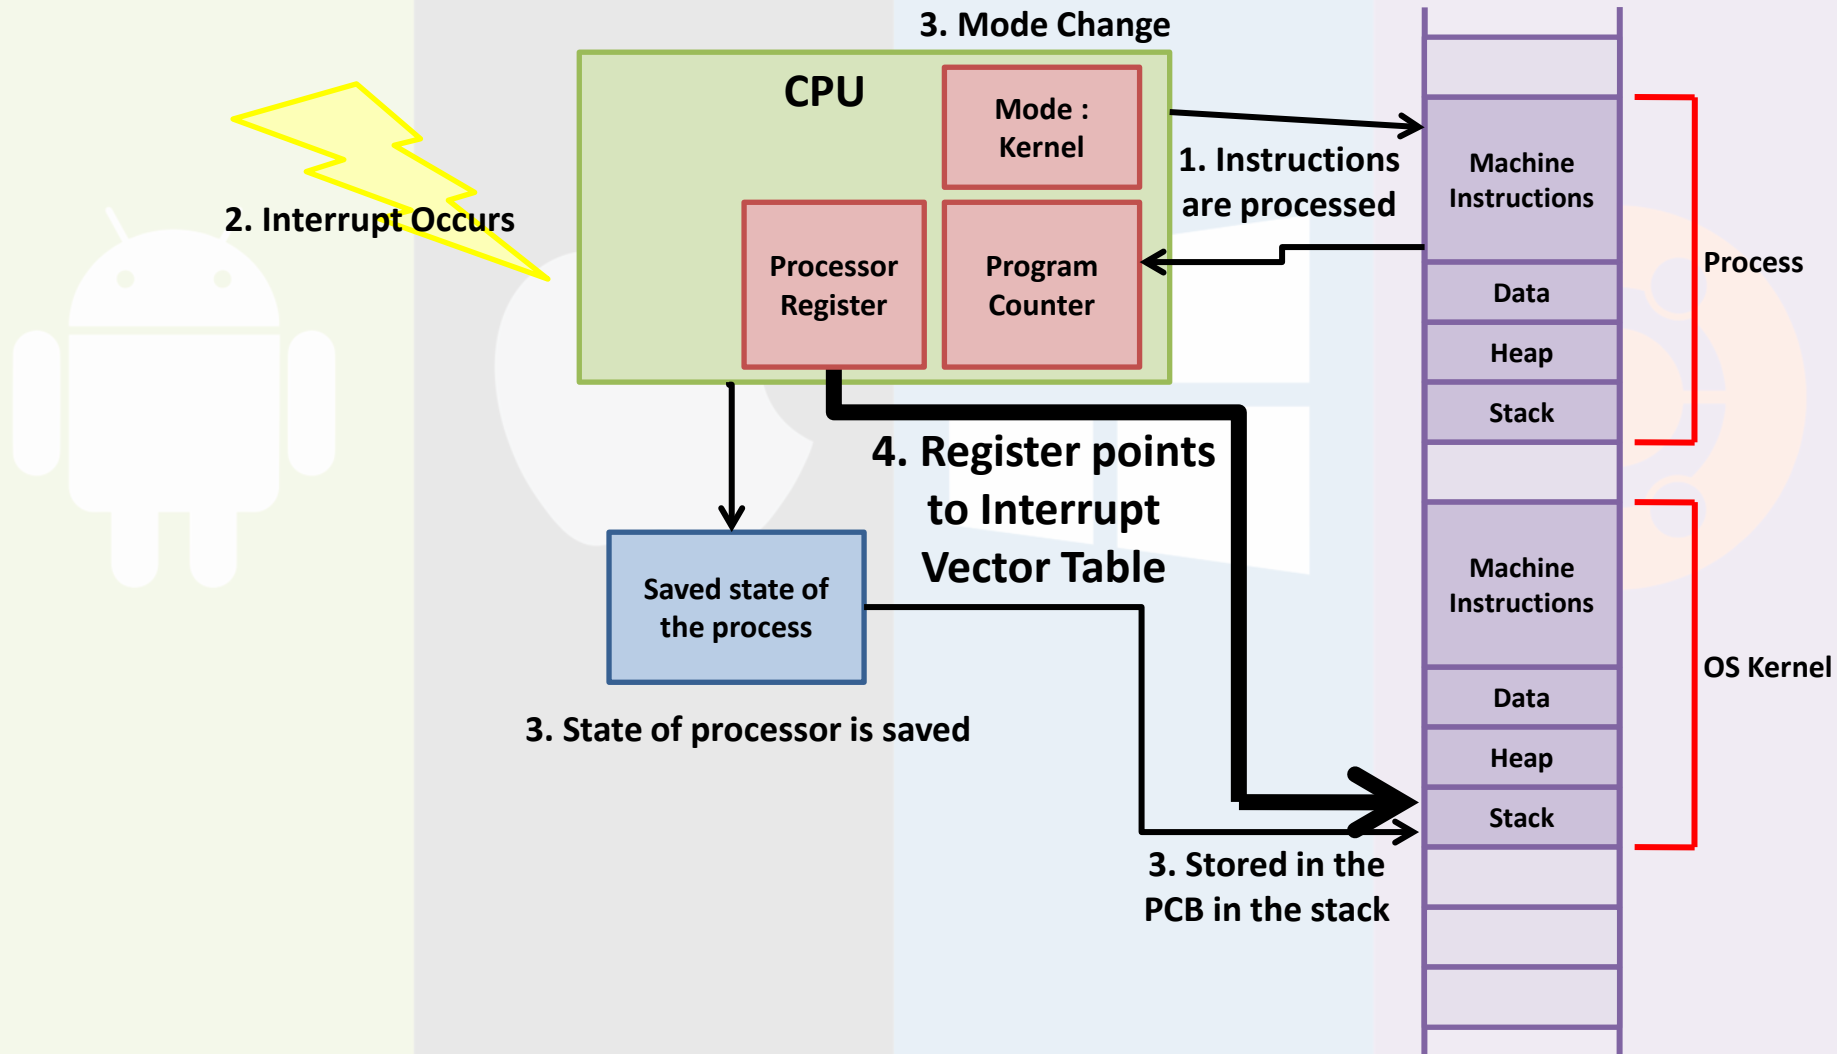
restore regs from kernel stack
move to user mode
jump to main

Program
(user mode)

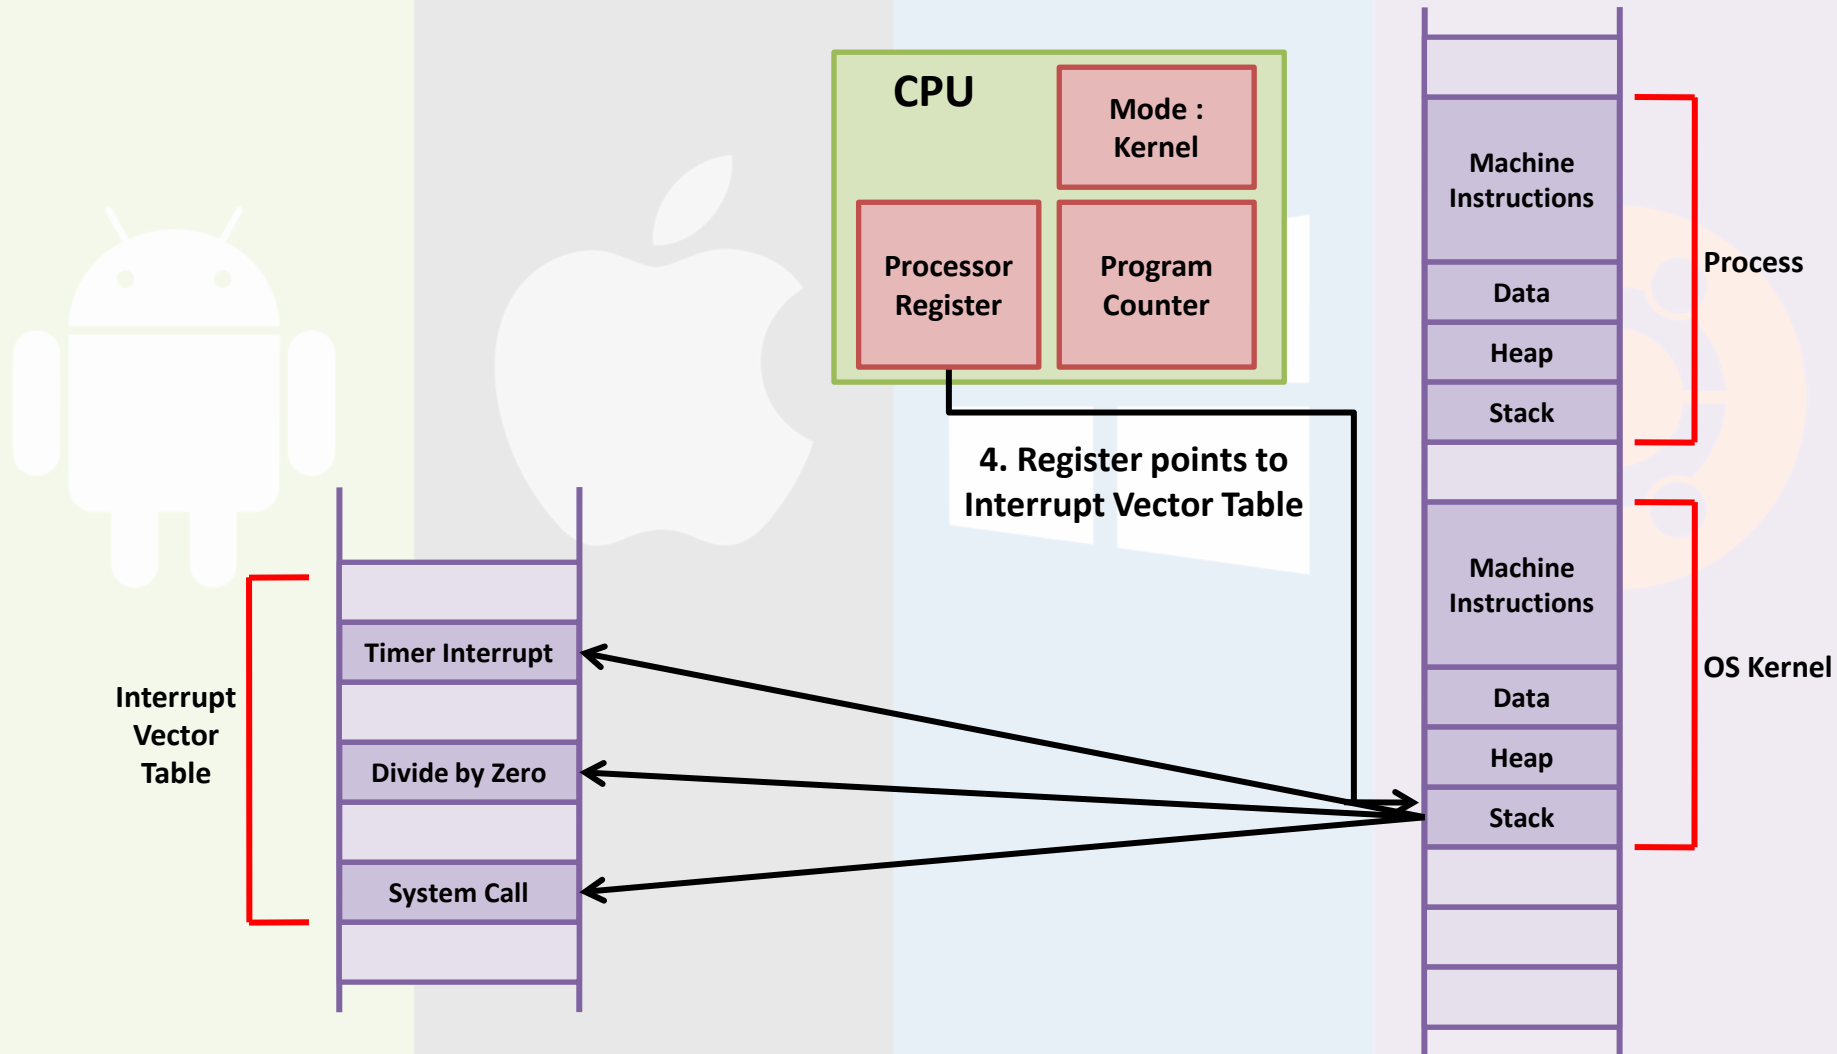
Run main()

...
Call system
trap into OS

Implementing Safe Mode Transfer

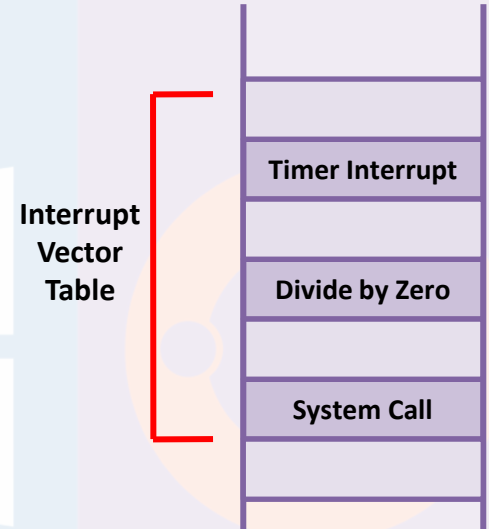


Implementing Safe Mode Transfer

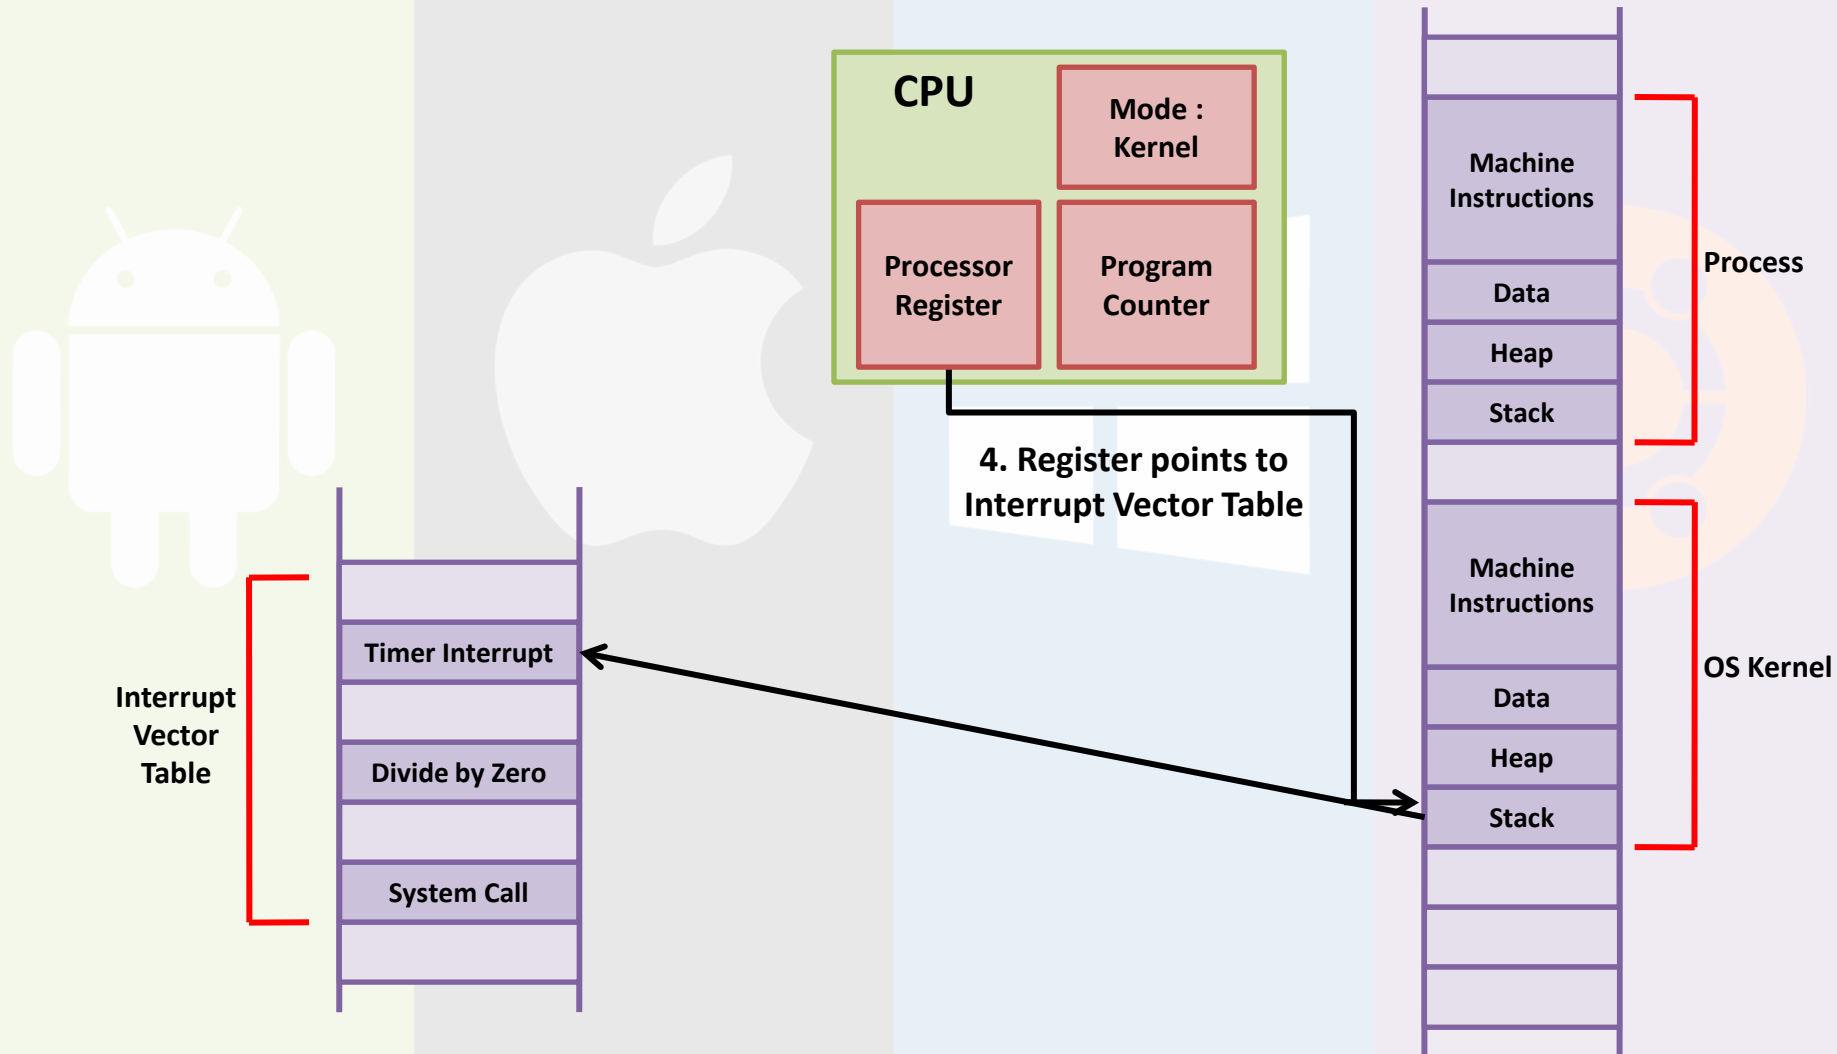


Interrupt Vector Table

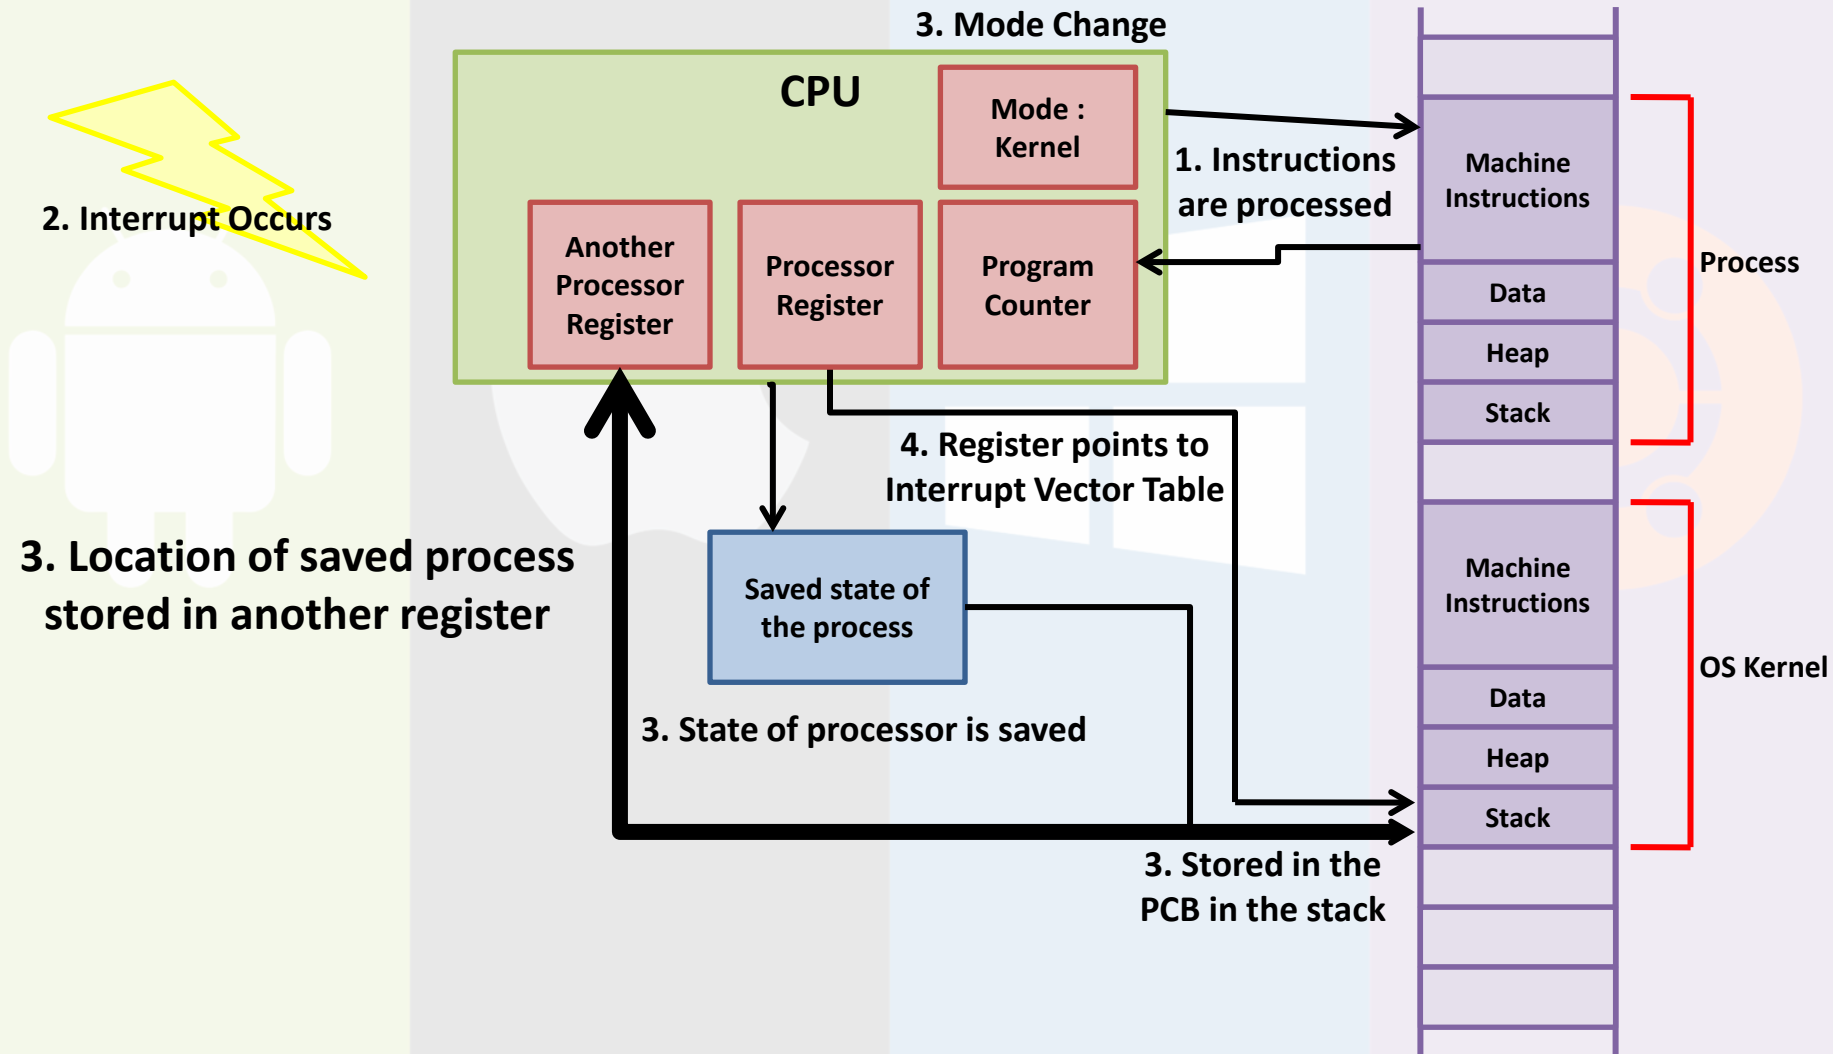
- The processor has a **special register** that in the event of an interrupt **points** to the corresponding **interrupt handler**.
- The handlers are **found** in a special part of a **kernel's stack** memory called the **Interrupt vector table**
- The format of the vector is processor specific
- x86 Architecture
 - 0 - 31 exception handlers
 - 32 – 255 interrupt handlers
- Since the hardware determines which device caused the exception the hardware will be able to use the correct handler



Implementing Safe Mode Transfer



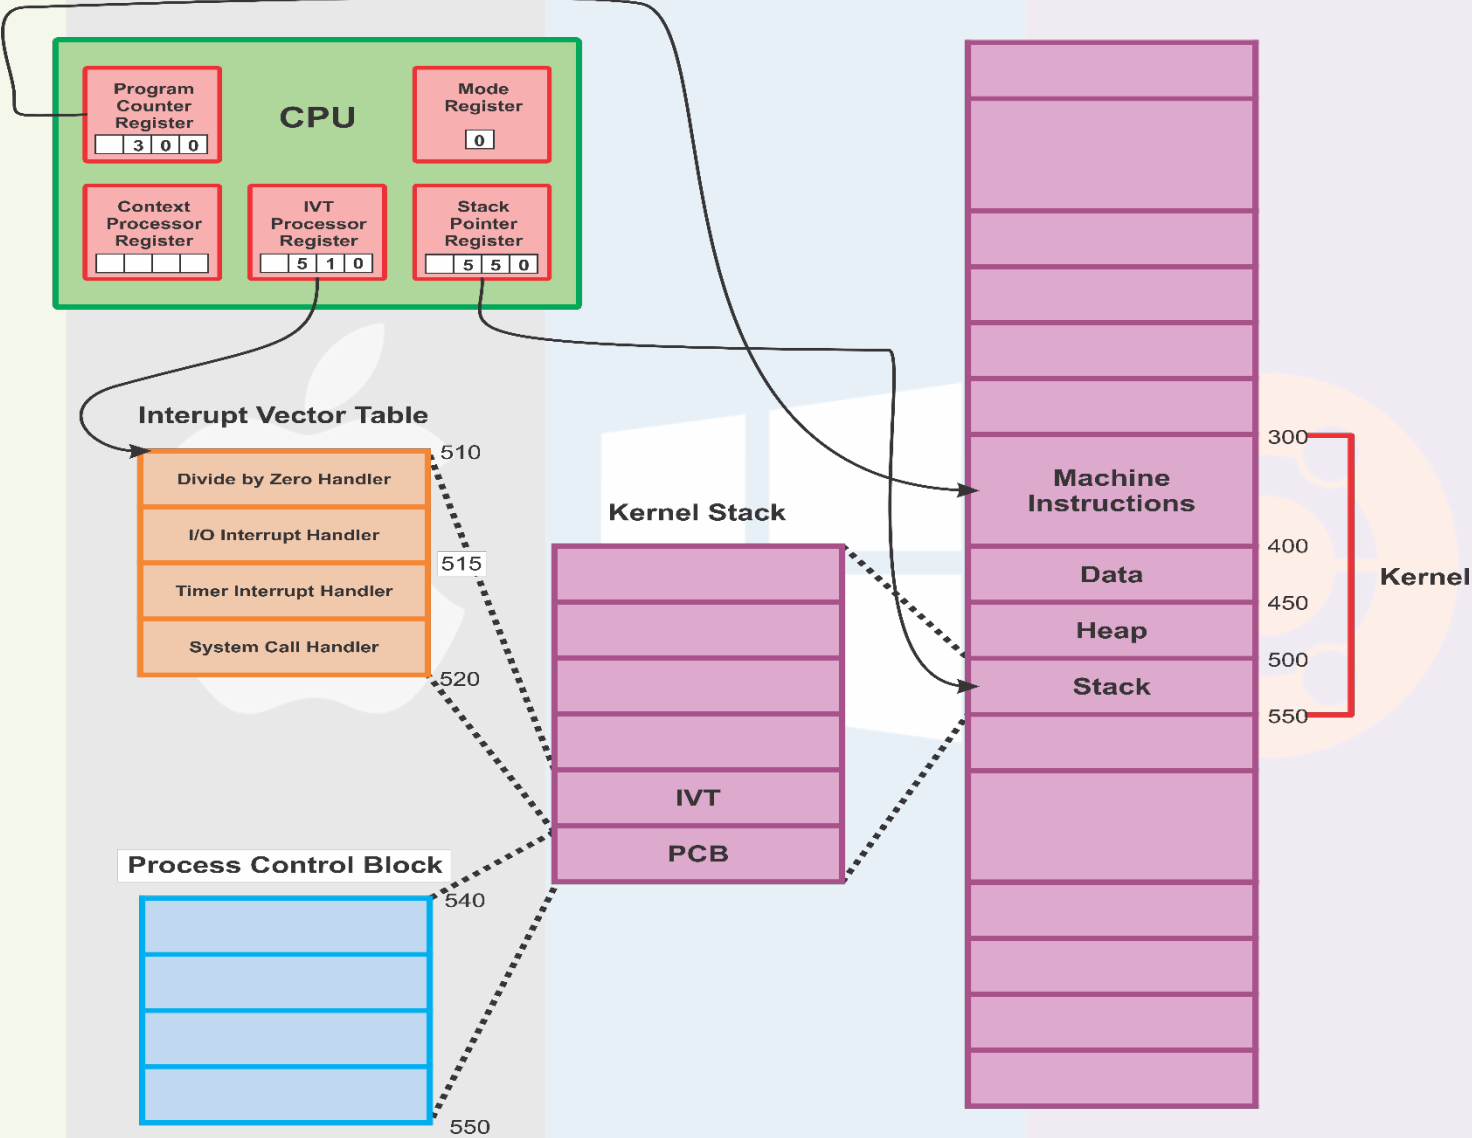
Interrupt Stack



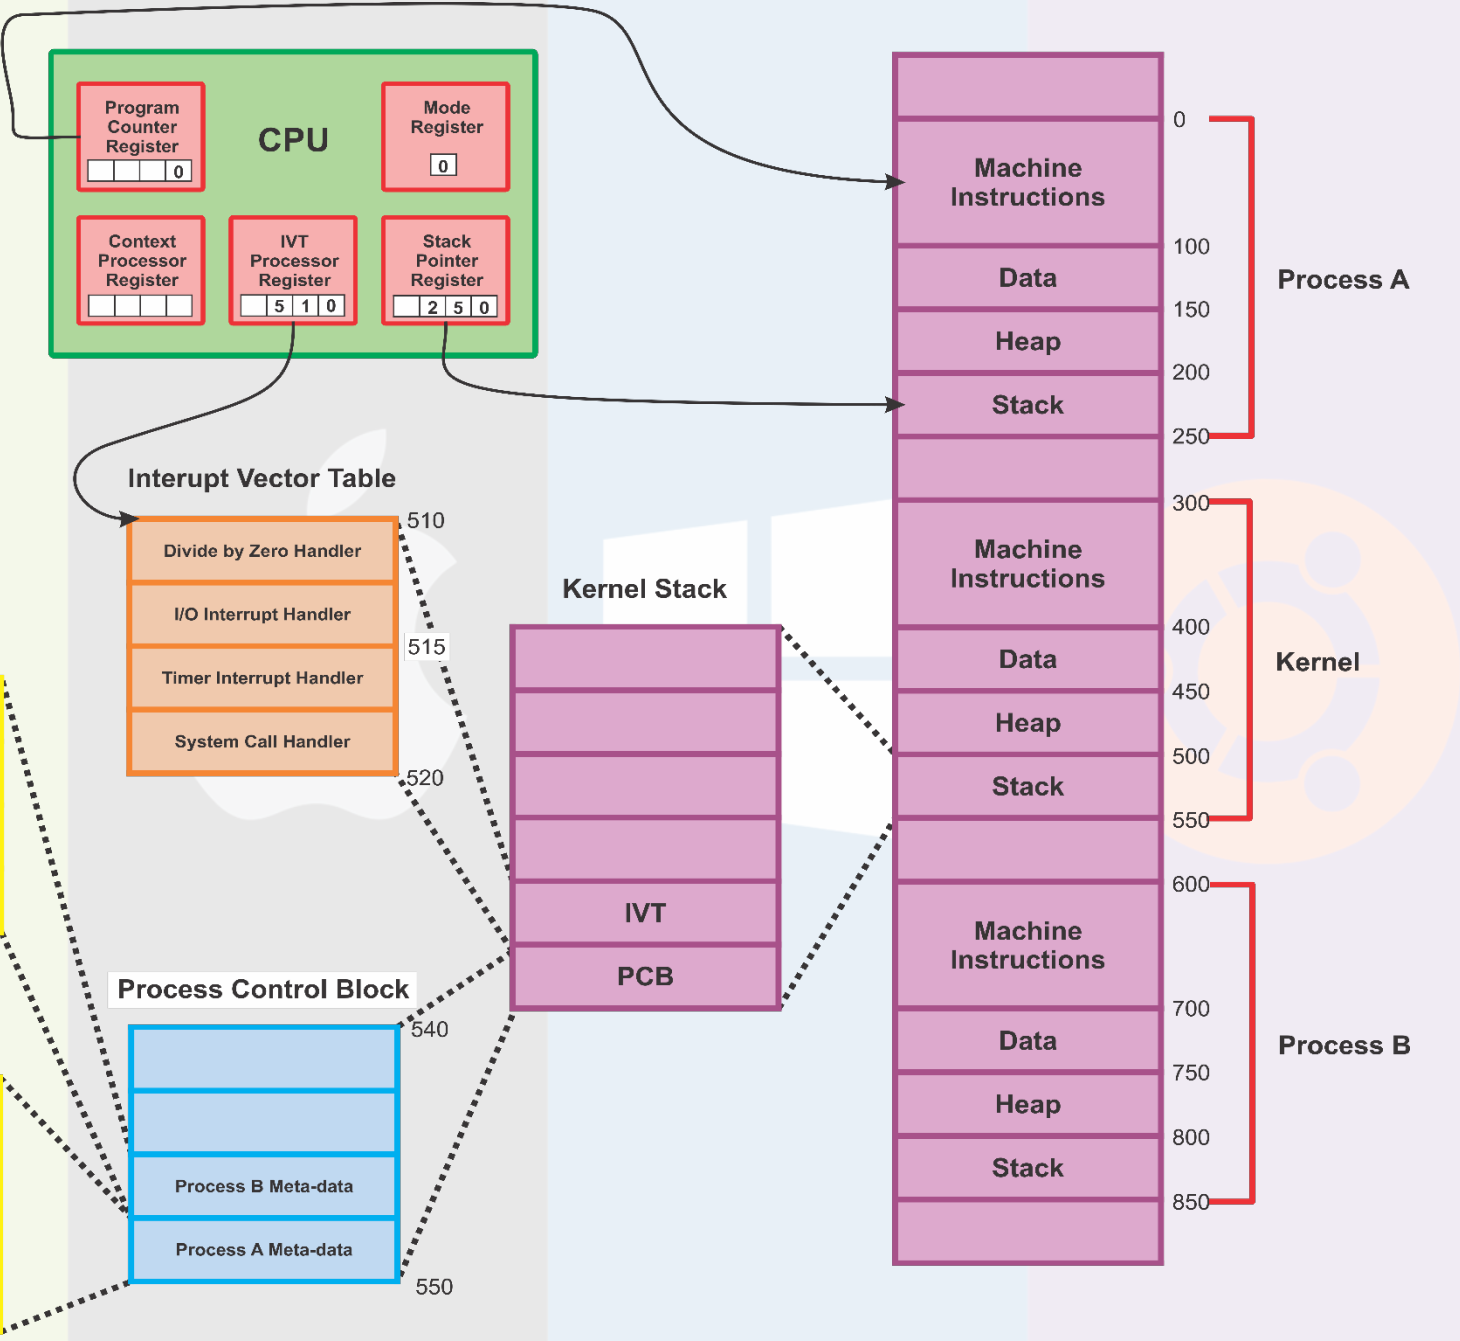
Interrupt Stack

- Interrupt/Exception/System Call occurs
- Process halted
- Stack pointer set to base of kernel stack
- Hardware saves(pushes) process's registers to interrupt stack in kernel
- Control switched to kernel
- Kernel does its job and handle event
- When returning from Interrupt/Exception/System Call
- The saved registers are loaded(popped) **using the address stored in the CPU**
- Process continues from where it left off

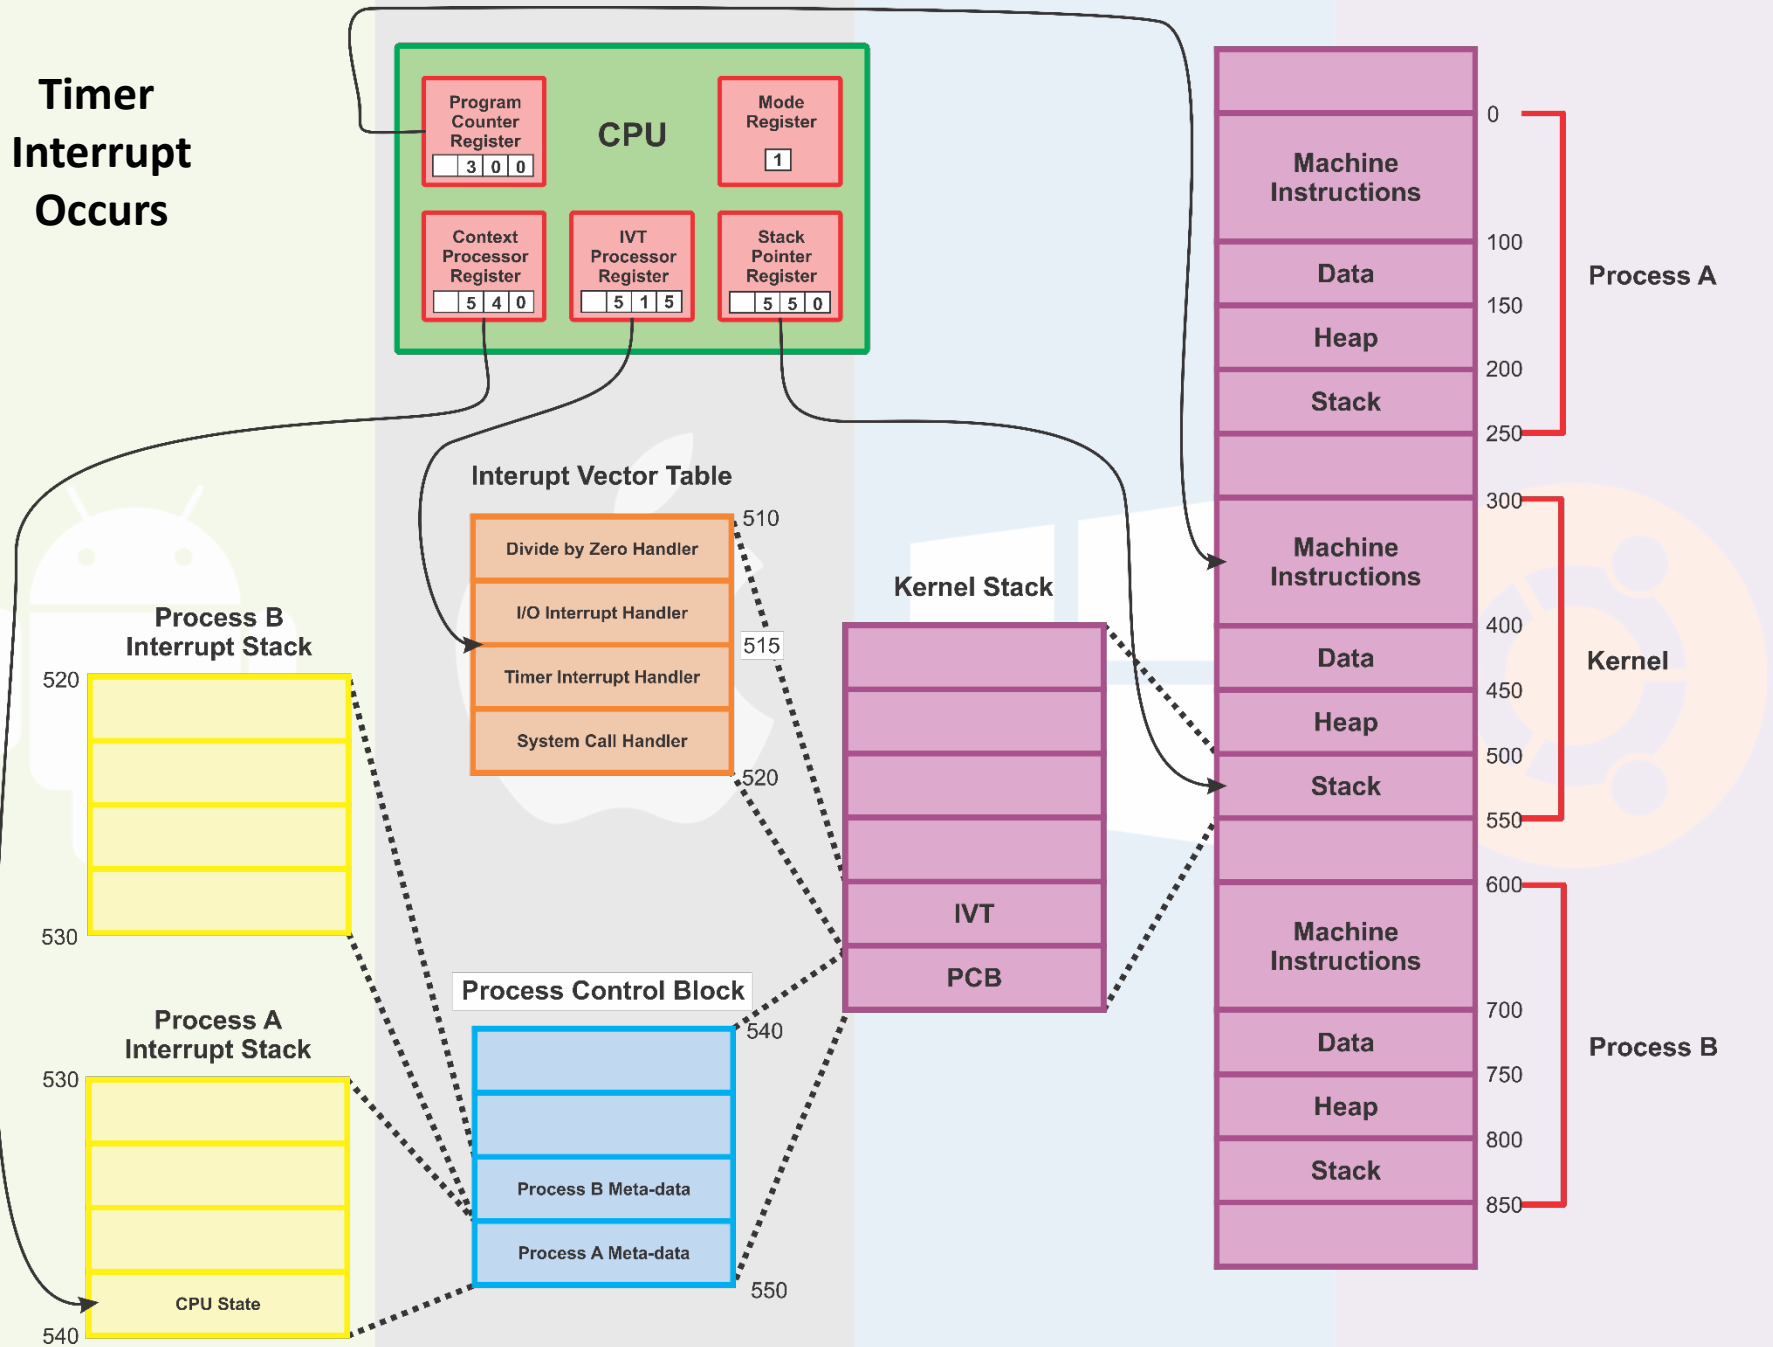
On Boot



Process A
Is Running



Timer Interrupt Occurs



Why not just store it on the process stack?



- Two reasons



Why not just store it on the process stack?

- Two reasons

- Reliability – The process's user-level stack might be invalid due to some bug


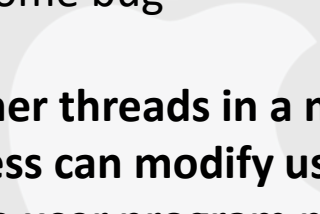
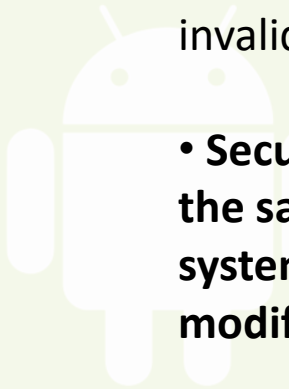


Why not just store it on the user-level stack?

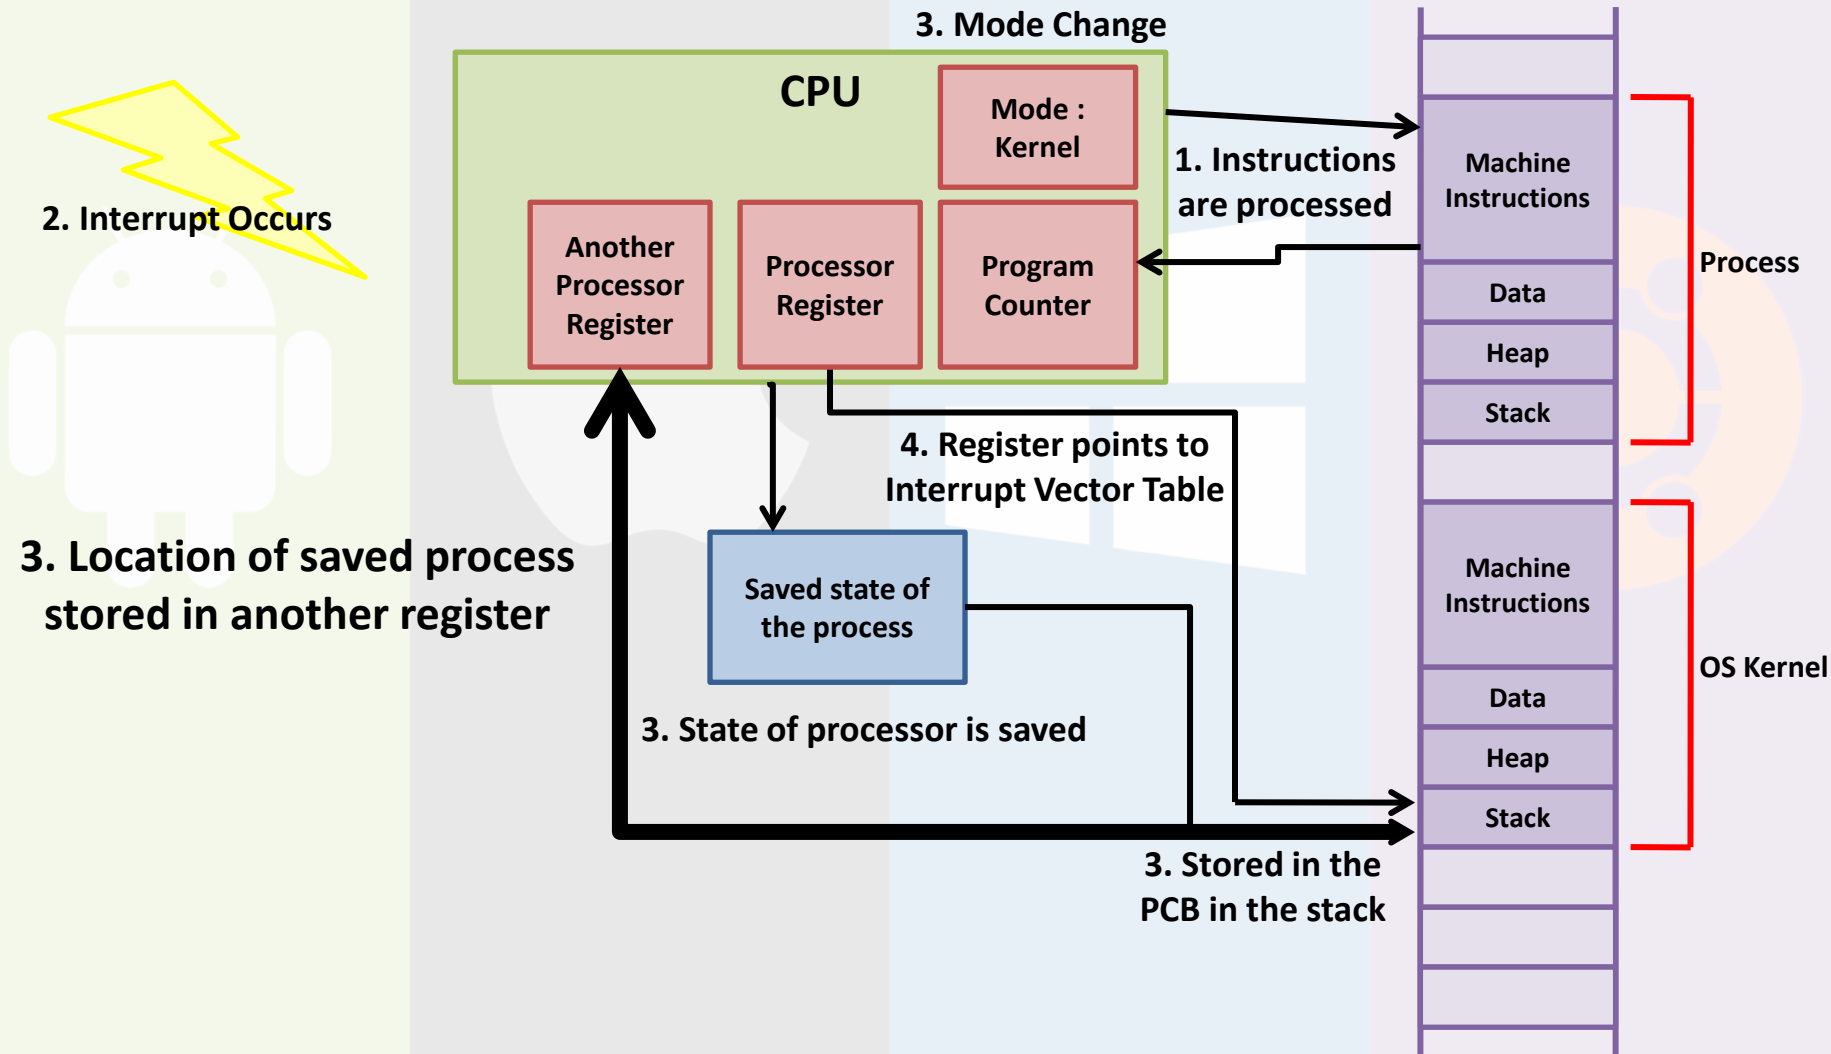
- **Two reasons**

- **Reliability** – The process's user-level stack might be invalid due to some bug

- **Security** – Other threads in a multiprocessor running the same process can modify user memory during a system call. The user program might then be able to modify the kernel's return address



Interrupt Stack



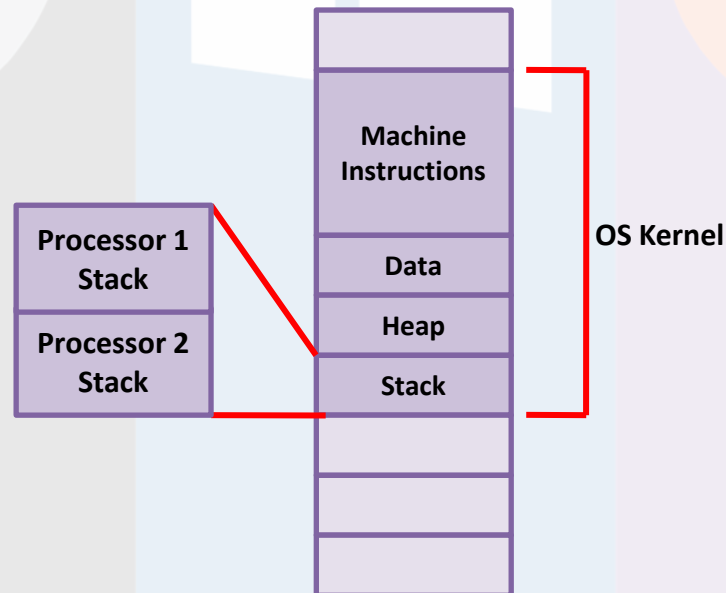
What happens when we have multiple processors?

- How can the kernel handle simultaneous system calls and exceptions across multiple processors?



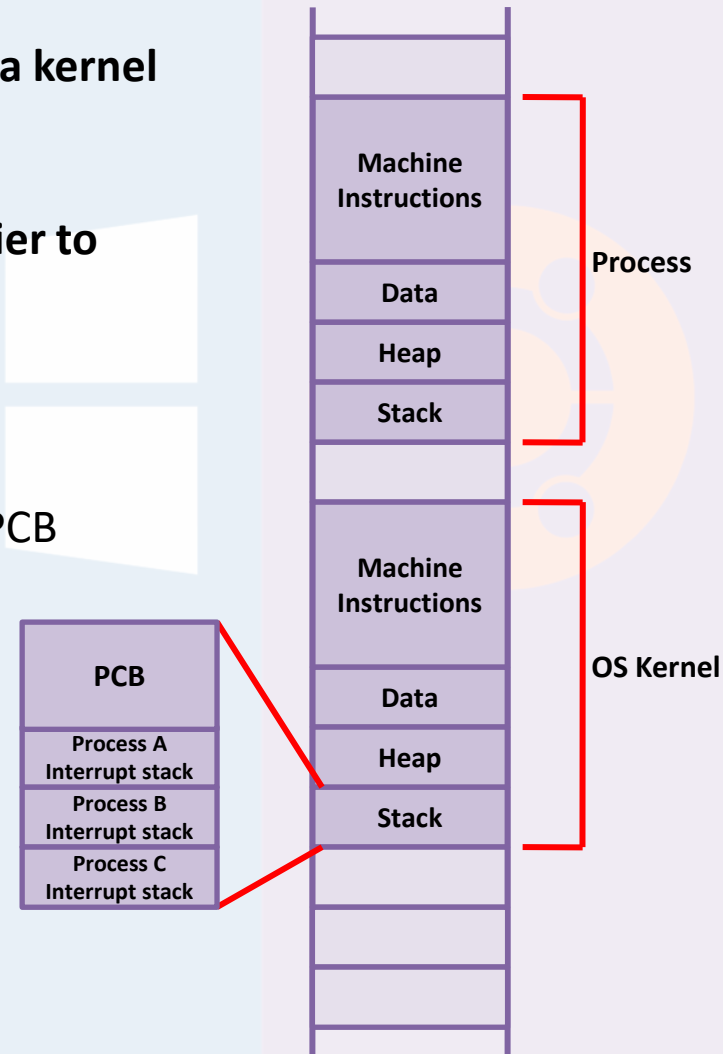
What happens when we have multiple processors?

- How can the kernel handle simultaneous system calls and exceptions across multiple processors?
- **For each processor the kernel allocates a separate region of memory as that processors interrupt stack.**



Two Stacks per Process

- Most OS kernels go one step further and allocate a kernel interrupt stack per user process
- Allocating a kernel stack per process makes it easier to switch between processes
- Now to switch from process A to B
 - Stop processing A
 - Store a pointer to process A's kernel stack in PCB
 - Save state of CPU in kernel stack A
 - Clear registers
 - Read process B's kernel stack pointer found in PCB
 - Load state of kernel stack B into CPU
 - Process B's instructions



Two Stacks per Process



- If the process is running on the processor in user mode
 - Kernel stack empty
 - Ready for an interrupt

```
#include <stdio.h>
void method1(){
    method2();
}
void method2(){
    Printf("Nothing");
}
int main() {
    method1();
    return 0;
}
```

User Stack

Kernel Stack

Running



Two Stacks per Process



- If the process is ready to run on the processor but is awaiting its turn
 - Kernel stack contains registers and state to be restored

```
#include <stdio.h>
void method1(){
    method2();
}
void method2(){
    Printf("Nothing");
}
int main() {
    method1();
    return 0;
}
```

User Stack

Kernel Stack

Ready to Run

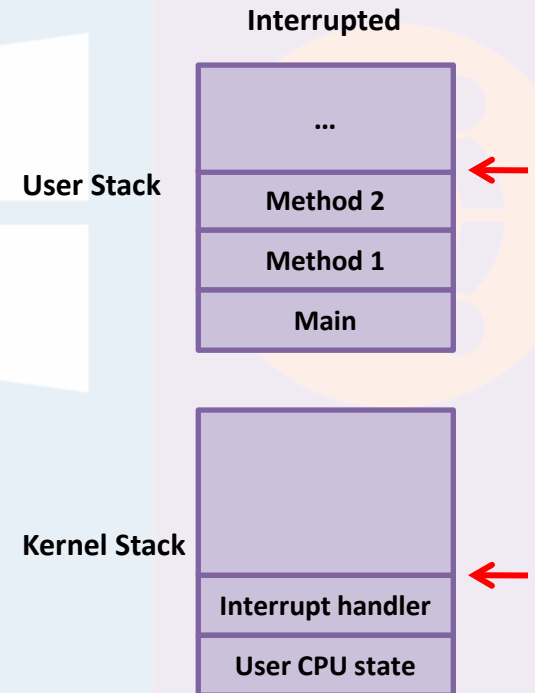


Two Stacks per Process

• If the process is running on the processor in kernel mode due to an Interrupt/System call/Exception

- Kernel stack is in use
- Saved registers and state of suspended CPU
- Current state of kernel handler

```
#include <stdio.h>
void method1(){
    method2();
}
void method2(){
    Printf("Nothing");
}
int main() {
    method1();
    return 0;
}
```




Two Stacks per Process


- If the process is waiting for I/O event
 - Saved registers and state of suspended CPU
 - Current state of kernel handler

```
#include <stdio.h>
void method1(){
    method2();
}
void method2(){
    char c;
    scanf("%c",&c);
}
int main() {
    method1();
    return 0;
}
```

User Stack

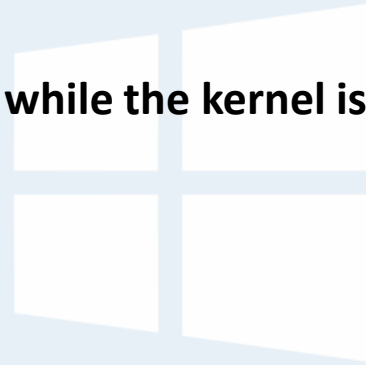
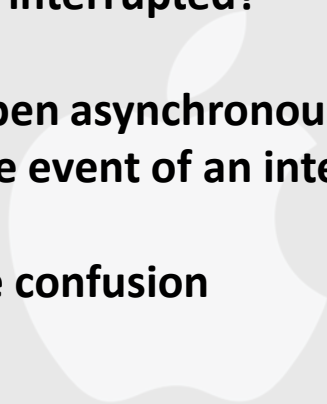
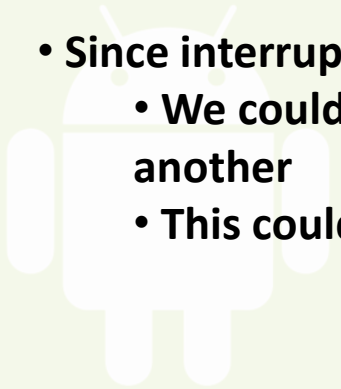


Kernel Stack



Interrupt Masking

- Can we interrupt the interrupted?
- Since interrupts happen asynchronously
 - We could get the event of an interrupt while the kernel is processing another
 - This could cause confusion


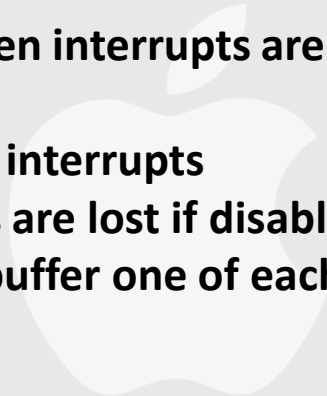



Interrupt Masking

- Can we interrupt the interrupted?
- Since interrupts happen asynchronously
 - We could get the event of an interrupt while the kernel is processing another
 - This could cause confusion
- **To simplify the kernel design provides a privileged instruction**
 - **This instruction temporarily defers deliveries of interrupts**
 - **on x86 infrastructure = “Disable Interrupts”**
- **Deferred not ignored**
 - **Once the corresponding “Enable Interrupts” instruction is executed, any pending interrupts are delivered**


What happens when multiple interrupts occur when disabled

- Stored in a buffer
- Delivered in turn when interrupts are re-enabled
- Limited buffering for interrupts
 - Some interrupts are lost if disabled too long
 - Generally only buffer one of each type



How do we prevent the same instruction from causing an exception ?

- If the handler returns back the instruction that caused the exception, the exception would immediately reoccur



```
1  #include <stdio.h>
2
3  int main()
4  {
5
6      int i = 0;
7      i = i/0;
8      printf("Hello World");
9
10     return 0;
11 }
```

Division by 0

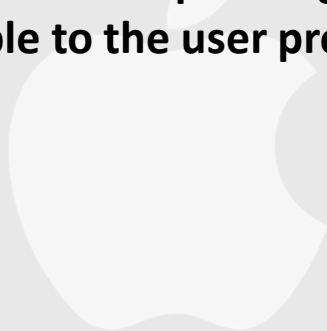
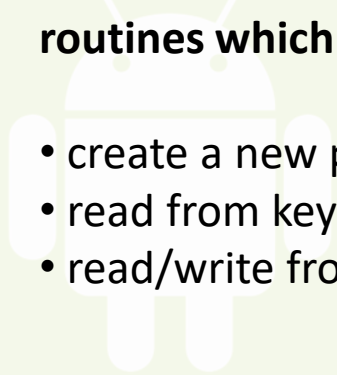
How do we prevent the same instruction from causing an exception ?

- If the handler returns back the instruction that caused the exception, the exception would immediately reoccur
- **To prevent this infinite loop**
 - **The exception handler modifies the program counter stored at the base of the stack to point to the instruction immediately after the one which caused the exception**

Implementing Secure System Calls

- **Voluntary mode switches from user to kernel**
- **Provide the illusion that the operating system kernel is simply a set of library routines which available to the user program**

- create a new process
- read from keyboard
- read/write from disk



Implementing Secure System Calls

- **To implement these system calls requires defining a calling convention**
 - how to name system calls
 - pass arguments
 - and receive return values across the user-kernel boundary
- **Once the arguments are in the correct format**
 - The user-level program can issue a system call by executing the trap instruction to transfer to the kernel mode

Implementing Secure System Calls

- **To implement these system calls requires defining a calling convention**
 - how to name system calls
 - pass arguments
 - and receive return values across the user-kernel boundary
- **Once the arguments are in the correct format**
 - The user-level program can issue a system call by executing the trap instruction to transfer to the kernel mode
- **Once inside the kernel**
 - A handler handles each system call
 - These handlers are implemented in a way that protects the kernel from any errors or attacks

Implementing Secure System Calls

- **To implement these system calls requires defining a calling convention**
 - how to name system calls
 - pass arguments
 - and receive return values across the user-kernel boundary
- **Once the arguments are in the correct format**
 - The user-level program can issue a system call by executing the trap instruction to transfer to the kernel mode
- **Once inside the kernel**
 - A handler handles each system call
 - These handlers are implemented in a way that protects the kernel from any errors or attacks
- **We bridge the divide between user calling a system call and the kernel implementing the system call with a pair of stubs**
- **A pair of stubs is a pair of procedures which mediate between two environments**

Implementing Secure System Calls

- **Step 1** : The user process makes a normal procedure call to a stub linked

User-Level Process

```
main(){  
    file_open(arg1,arg2)  
}
```



User Stub

```
File_open(arg1,arg2){  
    push #SYSCALL_OPEN  
    trap  
    return  
}
```

Kernel

```
File_open(arg1,arg2){  
    //do operation  
}
```

Kernel Stub

```
File_open(arg1,arg2){  
    // copy arguments from user memory  
    // check validity of arguments  
    file_open(arg1,arg2)  
    //copy return value into memory  
    return;  
}
```

Implementing Secure System Calls

- **Step 2:** The stub executes the trap instruction. This transfers control to the kernel trap handler. The trap handler copies and checks its arguments.

User-Level Process

```
main(){  
    file_open(arg1,arg2  
}
```

1



User Stub

```
File_open(arg1,arg2){  
    push #SYSCALL_OPEN  
    trap  
    return  
}
```

2 Hardware Trap



Kernel

```
File_open(arg1,arg2){  
    //do operation  
}
```

Kernel Stub

```
File_open(arg1,arg2){  
    // copy arguments from user memory  
    // check validity of arguments  
    file_open(arg1,arg2)  
    //copy return value into memory  
    return;  
}
```

Implementing Secure System Calls

- **Step 3: Kernel stub calls the kernel implementation of the system call, to do the operation**

User-Level Process

```
main(){  
    file_open(arg1,arg2  
}
```

1
↓

User Stub

```
File_open(arg1,arg2){  
    push #SYSCALL_OPEN  
    trap  
    return  
}
```

2 Hardware Trap →

Kernel

```
File_open(arg1,arg2){  
    //do operation  
}
```

3
↑

Kernel Stub

```
File_open(arg1,arg2){  
    // copy arguments from user memory  
    // check validity of arguments  
    file_open(arg1,arg2)  
    //copy return value into memory  
    return;  
}
```


Implementing Secure System Calls

- **Step 4:** The code returns to the trap handler which copies the return value into the user memory

User-Level Process

```
main(){  
    file_open(arg1,arg2  
}
```

1
↓

User Stub

```
File_open(arg1,arg2){  
    push #SYSCALL_OPEN  
    trap  
    return  
}
```

2 Hardware Trap
→

Kernel

```
File_open(arg1,arg2){  
    //do operation  
}
```

3
↑

4
↓

Kernel Stub

```
File_open(arg1,arg2){  
    // copy arguments from user memory  
    // check validity of arguments  
    file_open(arg1,arg2)  
    //copy return value into memory  
    return;  
}
```

Implementing Secure System Calls

- **Step 5:** The handler returns to the user level at the next instruction in the stub

User-Level Process

```
main(){  
    file_open(arg1,arg2  
}
```

1
↓

User Stub

```
File_open(arg1,arg2){  
    push #SYSCALL_OPEN  
    trap  
    return  
}
```

2 Hardware Trap
→

←

5 Trap Return

Kernel

```
File_open(arg1,arg2){  
    //do operation  
}
```

3
↑

4
↓

Kernel Stub


```
File_open(arg1,arg2){  
    // copy arguments from user memory  
    // check validity of arguments  
    file_open(arg1,arg2)  
    //copy return value into memory  
    return;  
}
```

Implementing Secure System Calls

- Step 6: The user stub returns to the user-level caller

User-Level Process

```
main(){  
  file_open(arg1,arg2  
}
```



User Stub

```
File_open(arg1,arg2){  
  push #SYSCALL_OPEN  
  trap  
  return  
}
```

2 Hardware Trap




5 Trap Return



Kernel

```
File_open(arg1,arg2){  
  //do operation  
}
```



Kernel Stub

```
File_open(arg1,arg2){  
  // copy arguments from user memory  
  // check validity of arguments  
  file_open(arg1,arg2)  
  //copy return value into memory  
  return;  
}
```

Kernel Stub

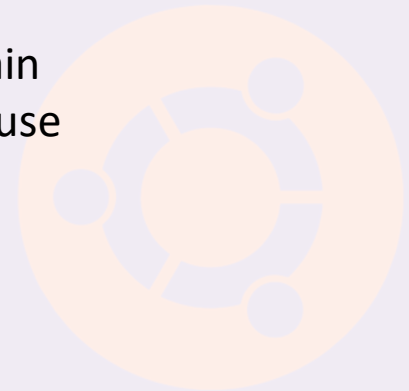
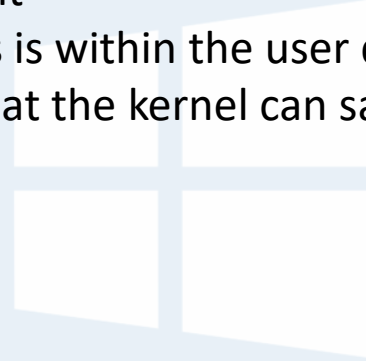
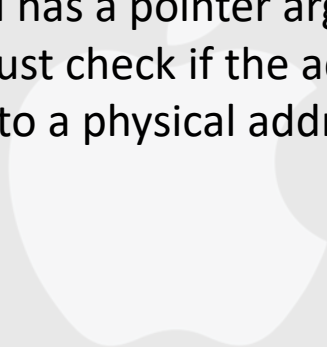
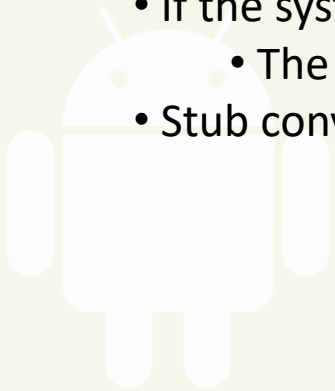
- The Kernel Stub has **four** tasks to perform before the kernel can do the operation



Kernel Stub

- **Locate System Call arguments**

- The arguments are stored in user memory unlike a regular kernel procedure
- If the system call has a pointer argument
 - The stub must check if the address is within the user domain
- Stub converts it to a physical address that the kernel can safely use



Kernel Stub

- **Copy Before Check**

- Kernel copies system call parameters into kernel memory before performing the necessary checks
- This is to ensure that the application can't modify the parameters after the stub checks but before the parameter is actually used
- **Time of use vs time of check attack (TOCTOU)**
- This happens when multiple processes share memory
 - One process traps into the kernel
 - The other modifies the parameters

Kernel Stub


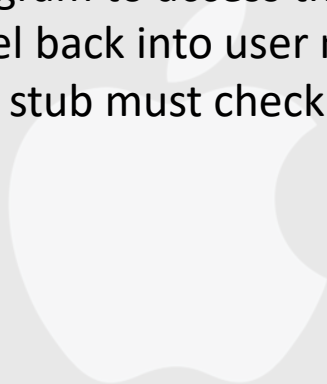

- **Validate Parameters**

- The kernel must also protect itself against malicious or accidental errors in format or content of its arguments
- A filename is typically a zero-terminated string, however, the kernel can't rely on user code to always work correctly
- The filename might point to regions outside the applications region
- Half the file might be stored within and the other half might exceed beyond
- File may not even exist
- If an error occurs it is returned to the user
- If not the kernel performs the operation

Kernel Stub


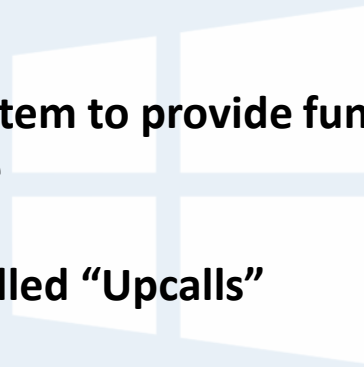


- **Copy Back Any Results**

- For the user program to access the results of the system call, the stub must copy the result of the kernel back into user memory
- Again the kernel stub must check the address



Implementing Upcalls


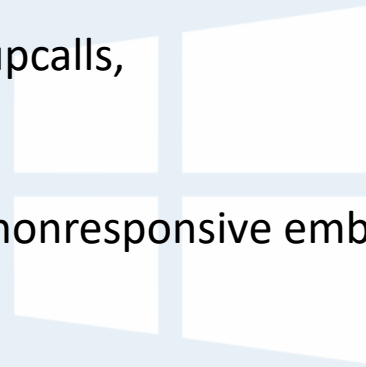
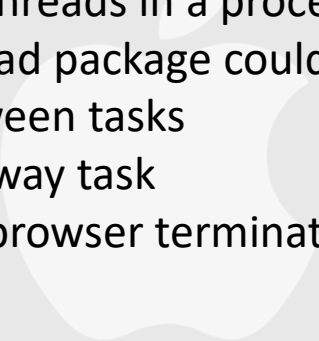
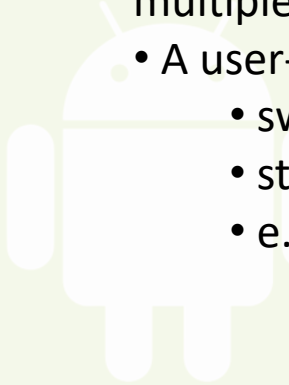
- For many of the same reasons that kernels need interrupt based event delivery, applications can benefit from being told when events that need their immediate attention occur
- We have virtualized a component of the system to provide functionality to user programs, we apply the same technique here
- Virtualized interrupts and exceptions are called “Upcalls”



Uses of Upcalls

- **Preemptive user-level threads**



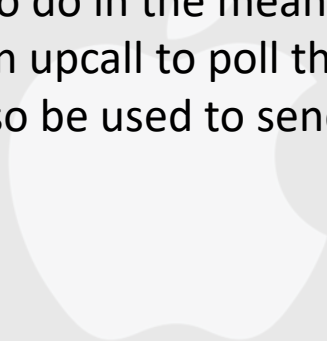
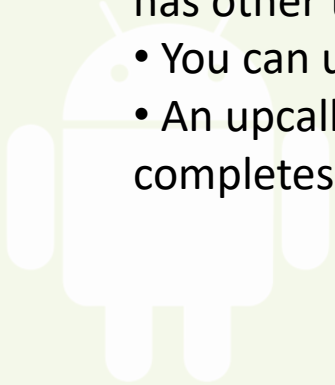
- Just as OS runs multiple processes on a single processor so too can an application run multiple tasks or threads in a process
- A user-level thread package could use upcalls,
 - switch between tasks
 - stop a runaway task
 - e.g. A web browser terminating a nonresponsive embedded script



Uses of Upcalls

- **Asynchronous I/O Notification**

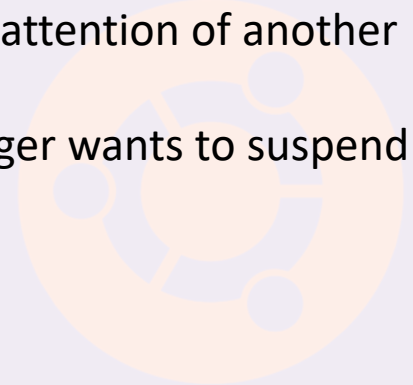
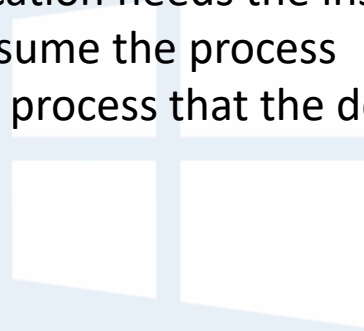
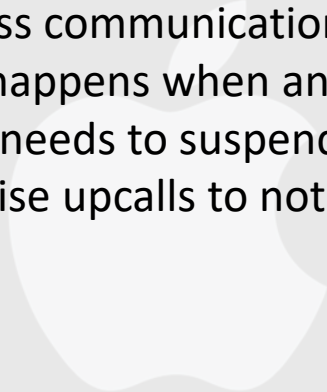
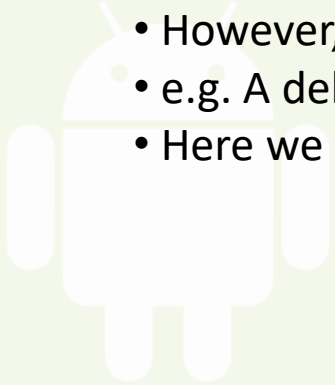
- Most system calls wait until the operation is completed, what happens if the process has other things to do in the meantime
- You can utilise an upcall to poll the kernel for I/O completion
- An upcall can also be used to send a notification to the application when the I/O completes



Uses of Upcalls

- **Interprocess communication**

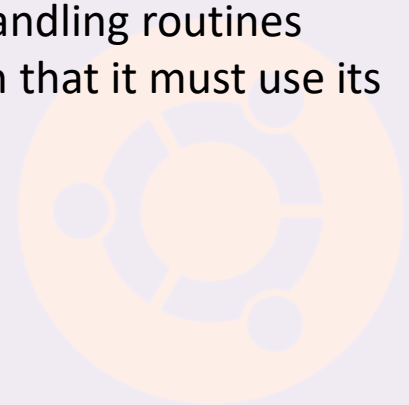
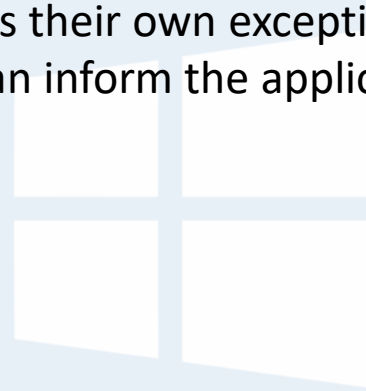
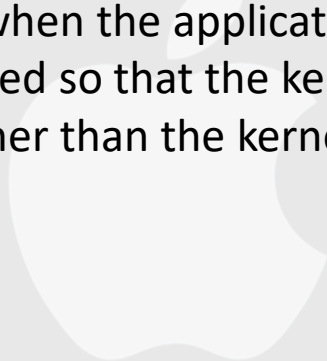
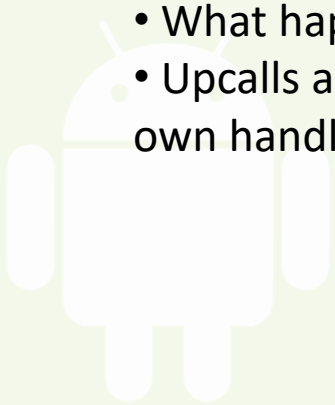
- Most interprocess communication can be handled with system calls
- However, what happens when an application needs the instant attention of another
- e.g. A debugger needs to suspend or resume the process
- Here we can utilise upcalls to notify the process that the debugger wants to suspend



Uses of Upcalls

- **User-Level exception handling**


- We have described processor exceptions such as divide by zero
- What happens when the application has their own exception handling routines
- Upcalls are utilised so that the kernel can inform the application that it must use its own handlers rather than the kernel



Uses of Upcalls

- **User-Level Resource Allocation**

- OS allocates resources, deciding which process/user gets how much
- Many applications are also resource adaptive, able to optimize their behaviour based on the resources available
- Java garbage collector, the more resources available the fewer amount of times the garbage collector is run



Upcalls

- **The virtualized interrupts share similarities to the hardware interrupts**
 - **Types of signals** : in place of hardware exceptions the kernel defines a limited number of signal types
 - **Handlers** : Each process defines its own handlers for each signal in a similar way as the interrupt vector table works
 - **Signal Stack** : Special stack for event handling is similar to that found in the kernel
 - **Signal Masking** : signals are deferred in the same way as in hardware interrupts by disabling interrupts
 - **Processor State** : The kernel copies onto the signal stack the saved state of the program counter, stack pointer and all other registers


Upcalls

- The state of a user program and signal stack **before a UNIX signal(Upcall)**
- Signals behave analogously to processor exceptions but at user level

User-Level Process

```
Foo(){  
  while(...){  
    x = x + 1  
    y = y - 2  
  }  
}
```

User Stack



Registers


Program Counter

Stack Pointer

User-Level event handling

```
Signal_Handler(){  
  ....  
}
```

Signal Stack




Upcalls

- The state of a user program and signal stack **during a UNIX signal(Upcall)**
- The signal stack stores the state of the registers at point when process interrupted
- With room for the signal handler to operate on the signal stack


User-Level Process

```
Foo(){  
  while(...){  
    x = x + 1  
    y = y - 2  
  }  
}
```

User Stack



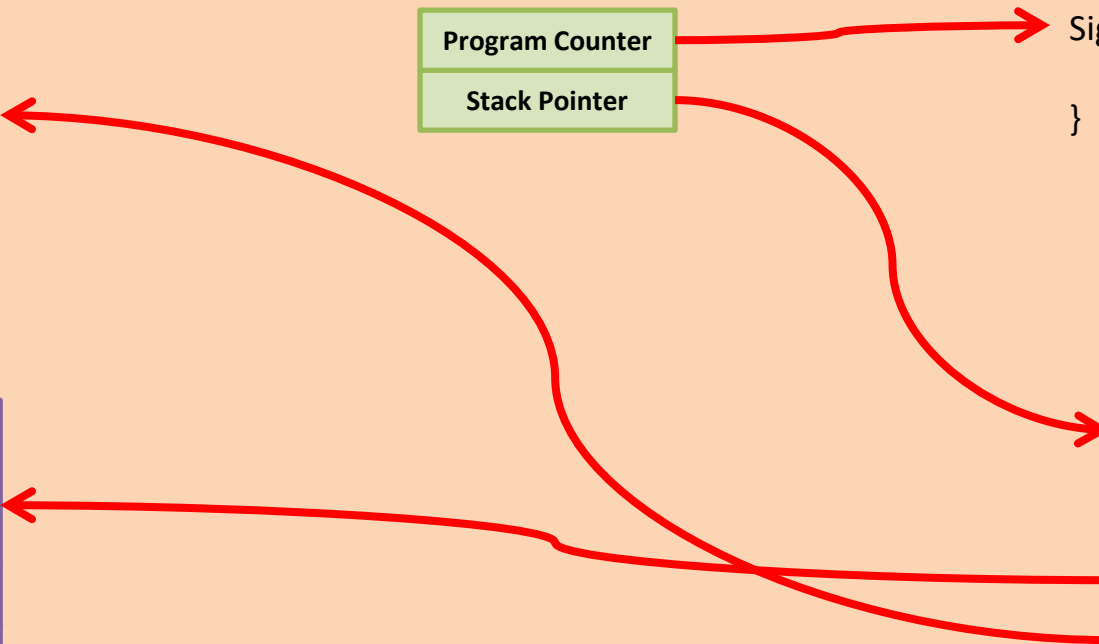

Registers



User-Level event handling

```
Signal_Handler(){  
  ....  
}
```

Signal Stack




Upcalls

- The state of a user program and signal stack **as Signal Processing finishes**


User-Level Process

```
Foo(){  
  while(...){  
    x = x + 1  
    y = y - 2  
  }  
}
```

User Stack



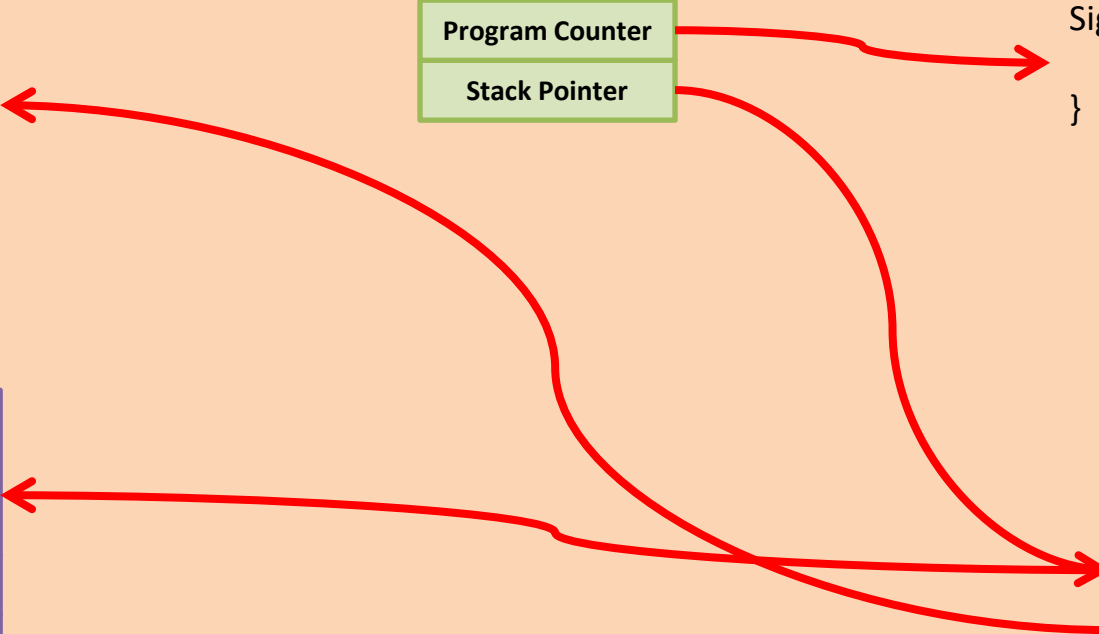
Registers



User-Level event handling

```
Signal_Handler(){  
  ....  
}
```

Signal Stack




Upcalls

- The state of a user program and signal stack **after** a UNIX signal(Upcall)

User-Level Process

```
Foo(){  
  while(...){  
    x = x + 1  
    y = y - 2  
  }  
}
```

User Stack




Registers

Program Counter
Stack Pointer

User-Level event handling




```
Signal_Handler(){  
  ....  
}
```

Signal Stack



UNIX

- Unix is a family of multitasking, multiuser computer operating systems that derive from the original AT&T Unix
- Formed the foundation of Linux/ MacOS
- We will be looking at a subset of the programming interface of UNIX
 - Process Management
 - I/O Control

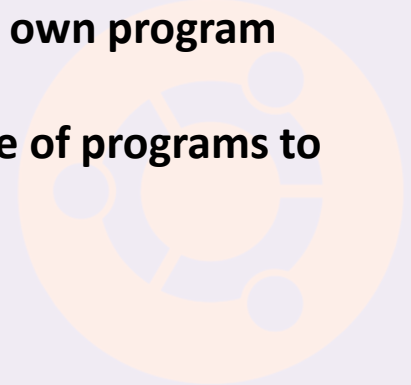
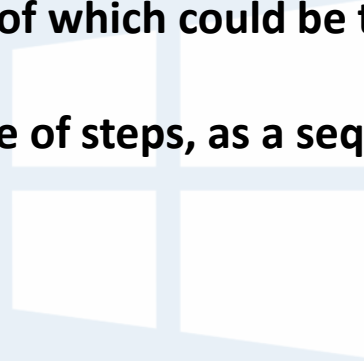
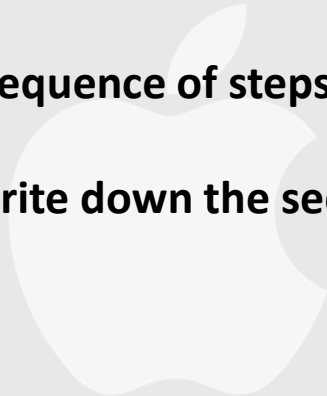
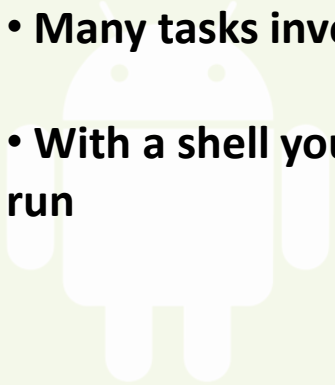


Process Management

- Traditionally in the early batch processing systems, the kernel handled process management by necessity
- A different approach that was developed was one that allowed users programs to create and manage their own processes
 - Web browsers managing embedded scripts
 - Window Manages managing various windows
- An early motivation for user-level process management was to allow developers to write there own shell command line interpreters

Shell

- A shell is a job control system
 - Windows and Unix both have a shell
- Many tasks involve a sequence of steps each of which could be there own program
- With a shell you can write down the sequence of steps, as a sequence of programs to run

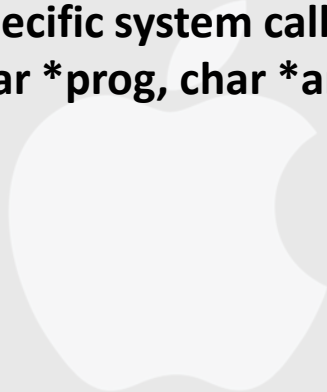
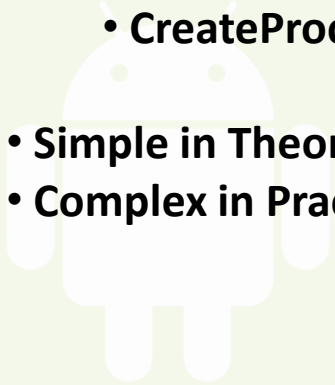


Shell

- A shell is a job control system
 - Windows and Unix both have a shell
- Many tasks involve a sequence of steps each of which could be there own program
- With a shell you can write down the sequence of steps, as a sequence of programs to run
- For example Makefiles are utilised to compile multiple C programs
- Makefiles are an example of a shell
- The C compiler itself is a shell program
 - The compiler first invokes a process to expand header include files
 - Separate process parses the output
 - Another process to convert to Assembly code
 - Lastly a process to convert Assembly into executable machine instructions




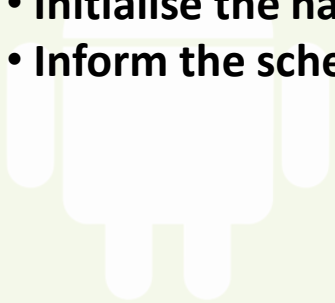
How does Windows handle process management?

- Windows simply has system calls to handle process operations
- For example there is specific system call to handle the creation of processes
 - `CreateProcess(char *prog, char *args)`
- Simple in Theory
- Complex in Practice



What steps does CreateProcess() take

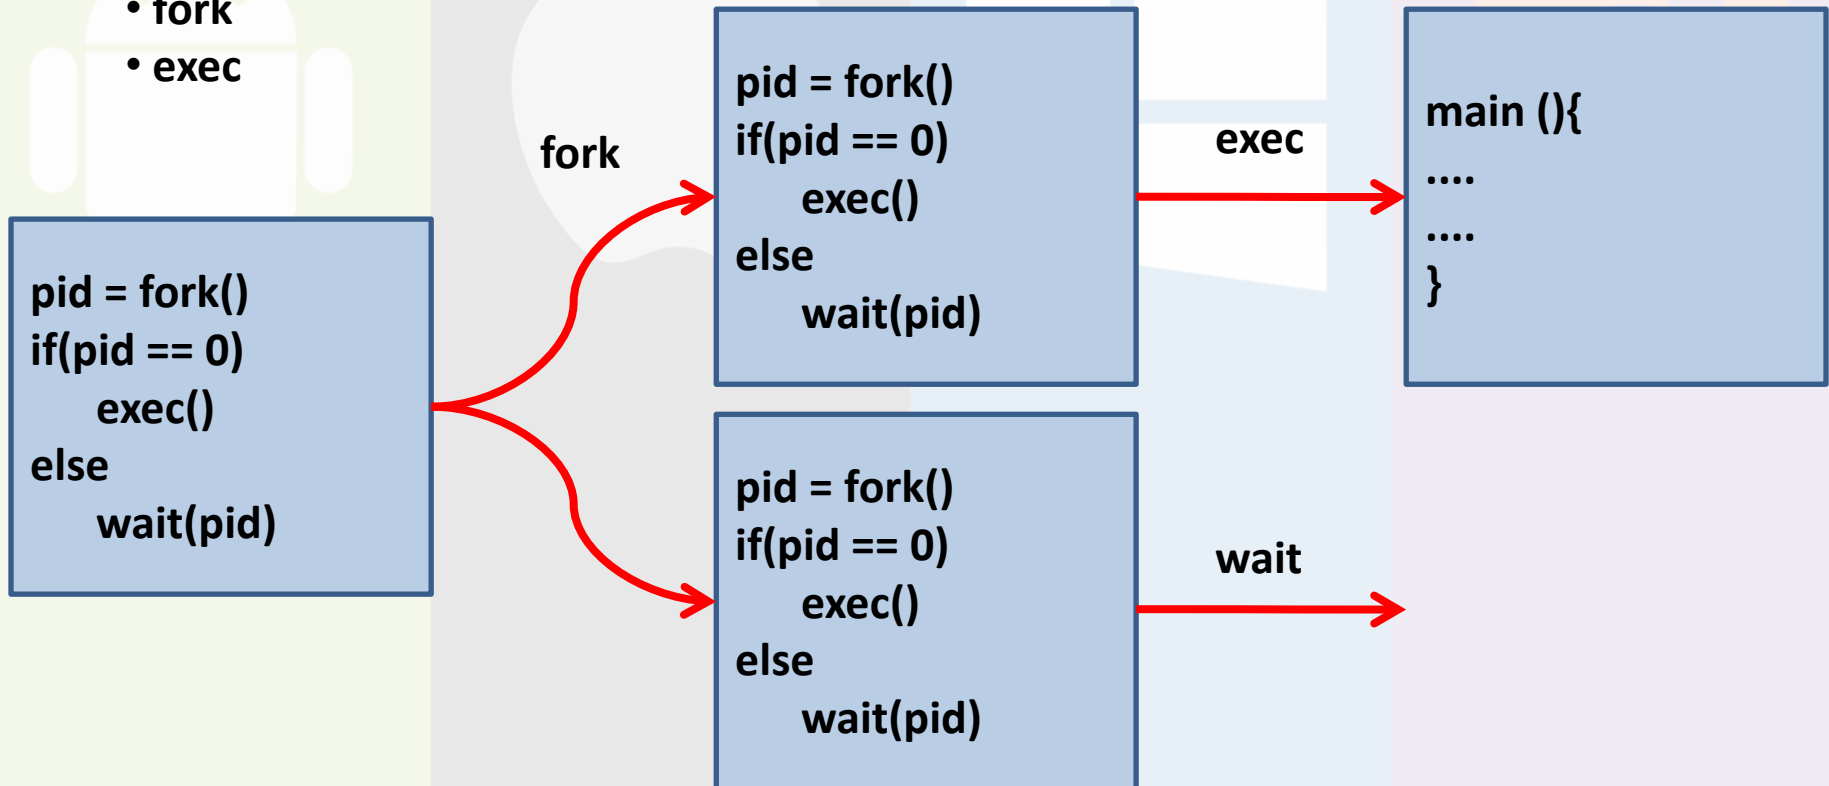
- Create and initialise PCB in the kernel
- Create and initialise a new address space
- Load the program into the address space
- Copy the arguments into the memory in the address space
- Initialise the hardware context to start execution at the start of the program
- Inform the scheduler that a new process is ready



How does UNIX handle process management?

- Complex in Theory
- Simple in Practice
- UNIX splits CreateProcess into two steps called

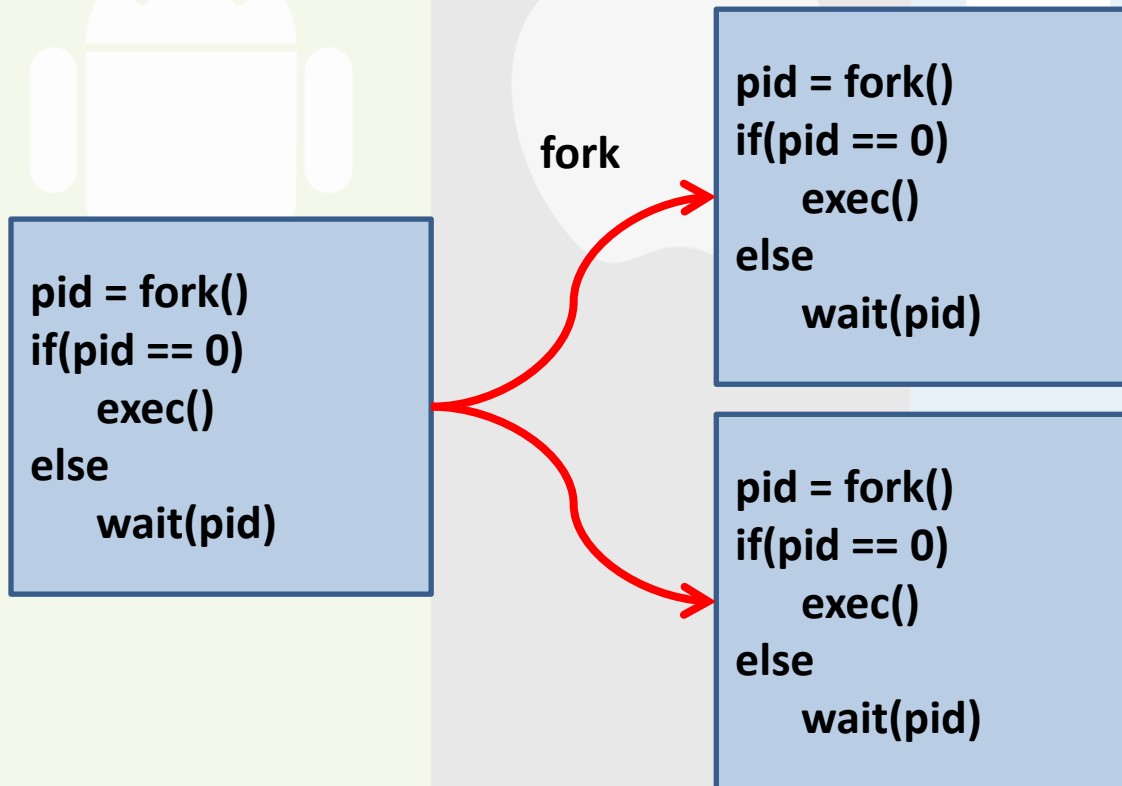
- fork
- exec



How does UNIX handle process management?


- **UNIX fork**

- Creates a complete copy of the parent process
- The only difference is the id which is used to identify it
- Child process sets up the same privileges, priorities and I/O the parent would



What steps does UNIX fork take

- Create and initialise the PCB in the kernel
- Create a new address space
- Initialise the address space with a copy of the entire contents of the address space of the parent
- Inherit the execution context of the parent
 - if any files have been opened
- Inform the scheduler that the new process is ready to run



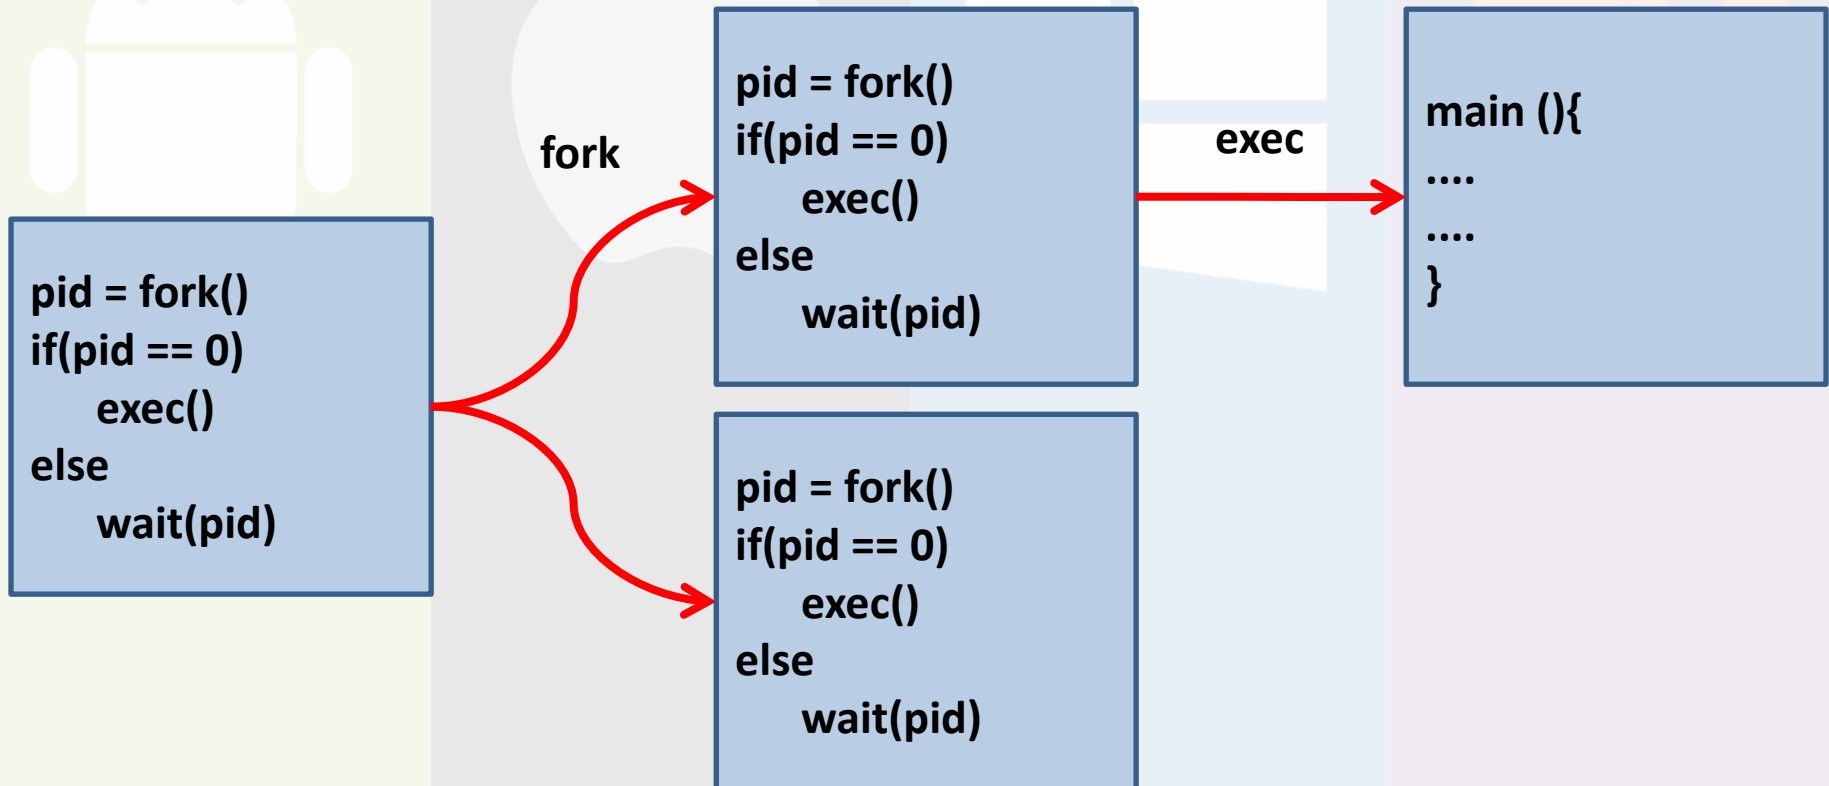
What steps does UNIX fork take

- Create and initialise the PCB in the kernel
- Create a new address space
- Initialise the address space with a copy of the entire contents of the address space of the parent
- Inherit the execution context of the parent
 - if any files have been opened
- Inform the scheduler that the new process is ready to run
- A strange aspect of Unix fork is that the system call returns twice once for the parent and once for the child
- Parent receives the process ID of the child
- Child receives 0 indicating success

How does UNIX handle process management?


- UNIX exec

- Once the context is set the child process calls UNIX exec
- Brings the new executable image into memory and runs it
- exec takes in as parameters the name of the program and the arguments for it



What steps does UNIX exec take

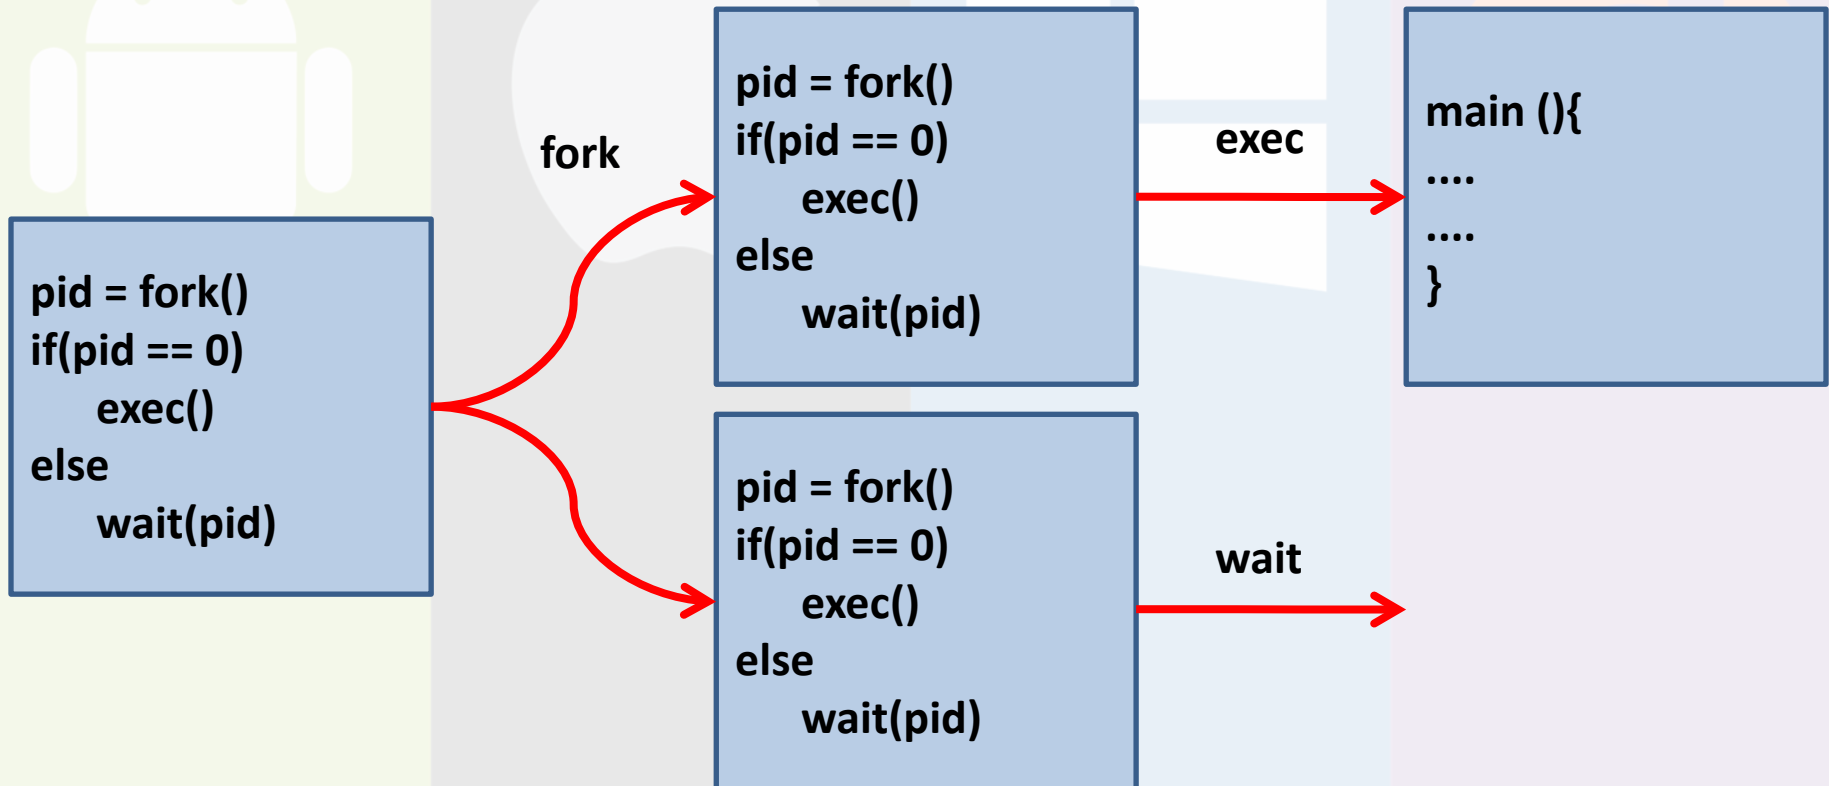
- Loads the program specified into the current address space
- Copy the arguments into memory in the address space
- Initialise the hardware context to start the execution at the start



How does UNIX handle process management?

- UNIX wait

- Often parent processes need to pause for child processes for this we use wait
- Pauses the parent process until the child process is finished
- UNIX wait takes in the process id for which the parent must wait for



How does UNIX handle process management?

- UNIX signal

- UNIX provides a facility for one process to send another an instant notification (Upcall)
- This notification is sent using signal

