


Operating Systems

COMS(3010A)

Memory



Branden Ingram

branden.ingram@wits.ac.za

Office Number : ???

Recap

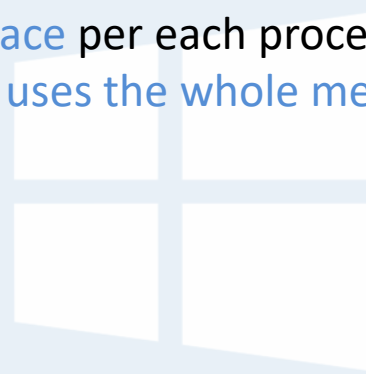
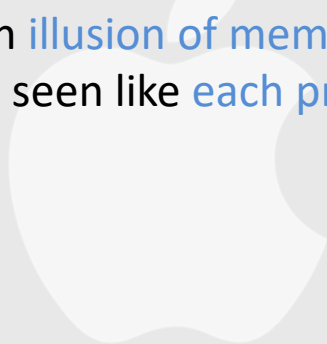

- Interrupt Stack
- System Calls
- UpCalls
- UNIX



Memory Virtualization




- **What is memory virtualization?**

- OS virtualizes its physical memory.
- OS provides an **illusion of memory space** per each process
- It seems to be seen like **each process uses the whole memory**



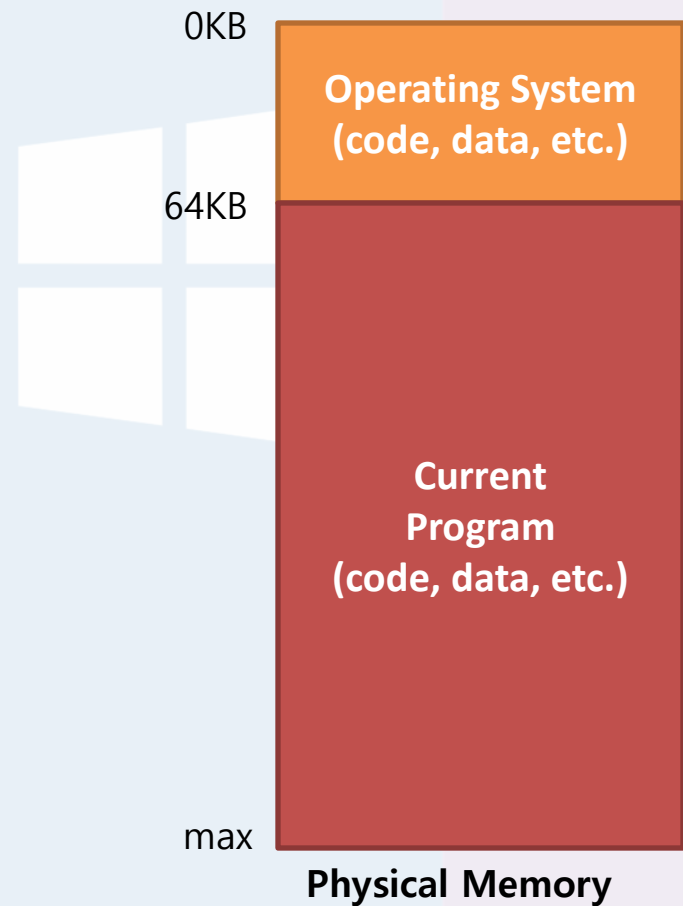

What are the benefits of Memory Virtualization

- Ease of use in programming
- Memory efficiency in terms of times and space
- The guarantee of isolation for processes as well as OS
 - Protection from errant accesses of other processes



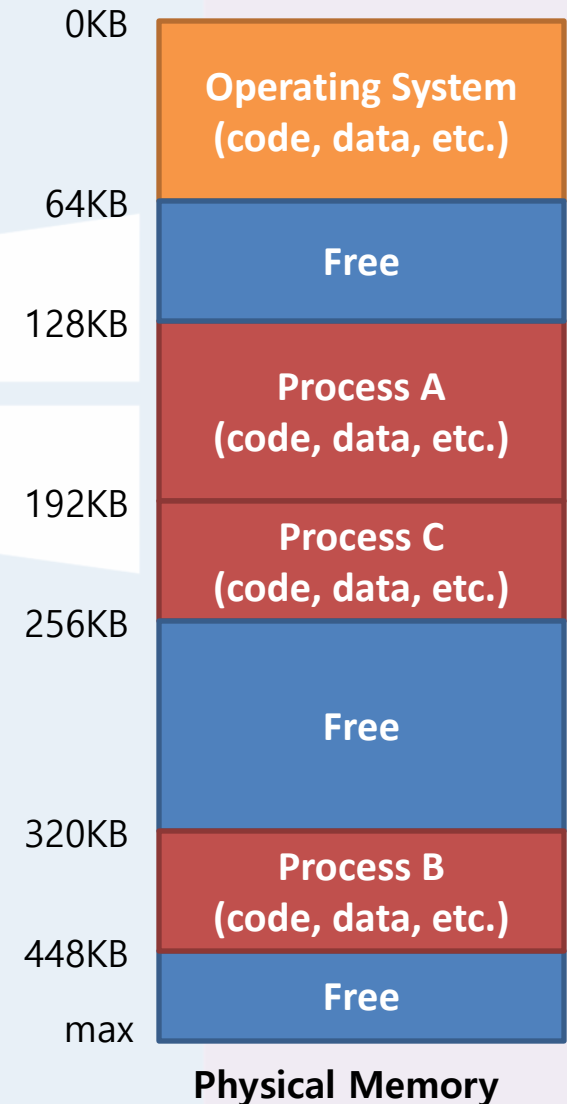
Memory Management in Early OS's

- Load only one process in memory
- Poor utilization and efficiency




Multiprogramming and Time Sharing

- **Load multiple processes in memory**
 - Execute one for a short while
 - Switch processes between them in memory
 - Increase utilization and efficiency
- **Cause an important protection issue**
 - Errant memory accesses from other processes



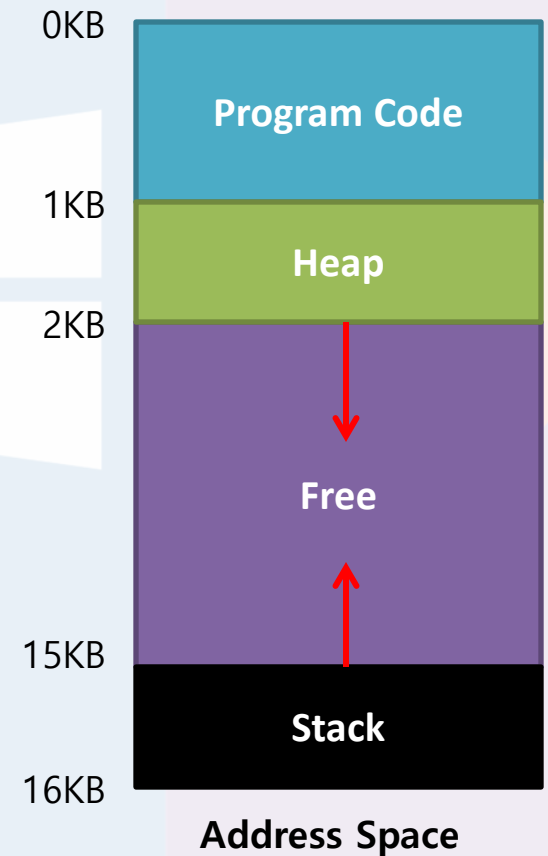
Address Space

- **OS creates an abstraction of physical memory**
 - This is called the Address Space
 - The address space contains all the info about a running process
 - That is it consists of program code, heap, stack and etc

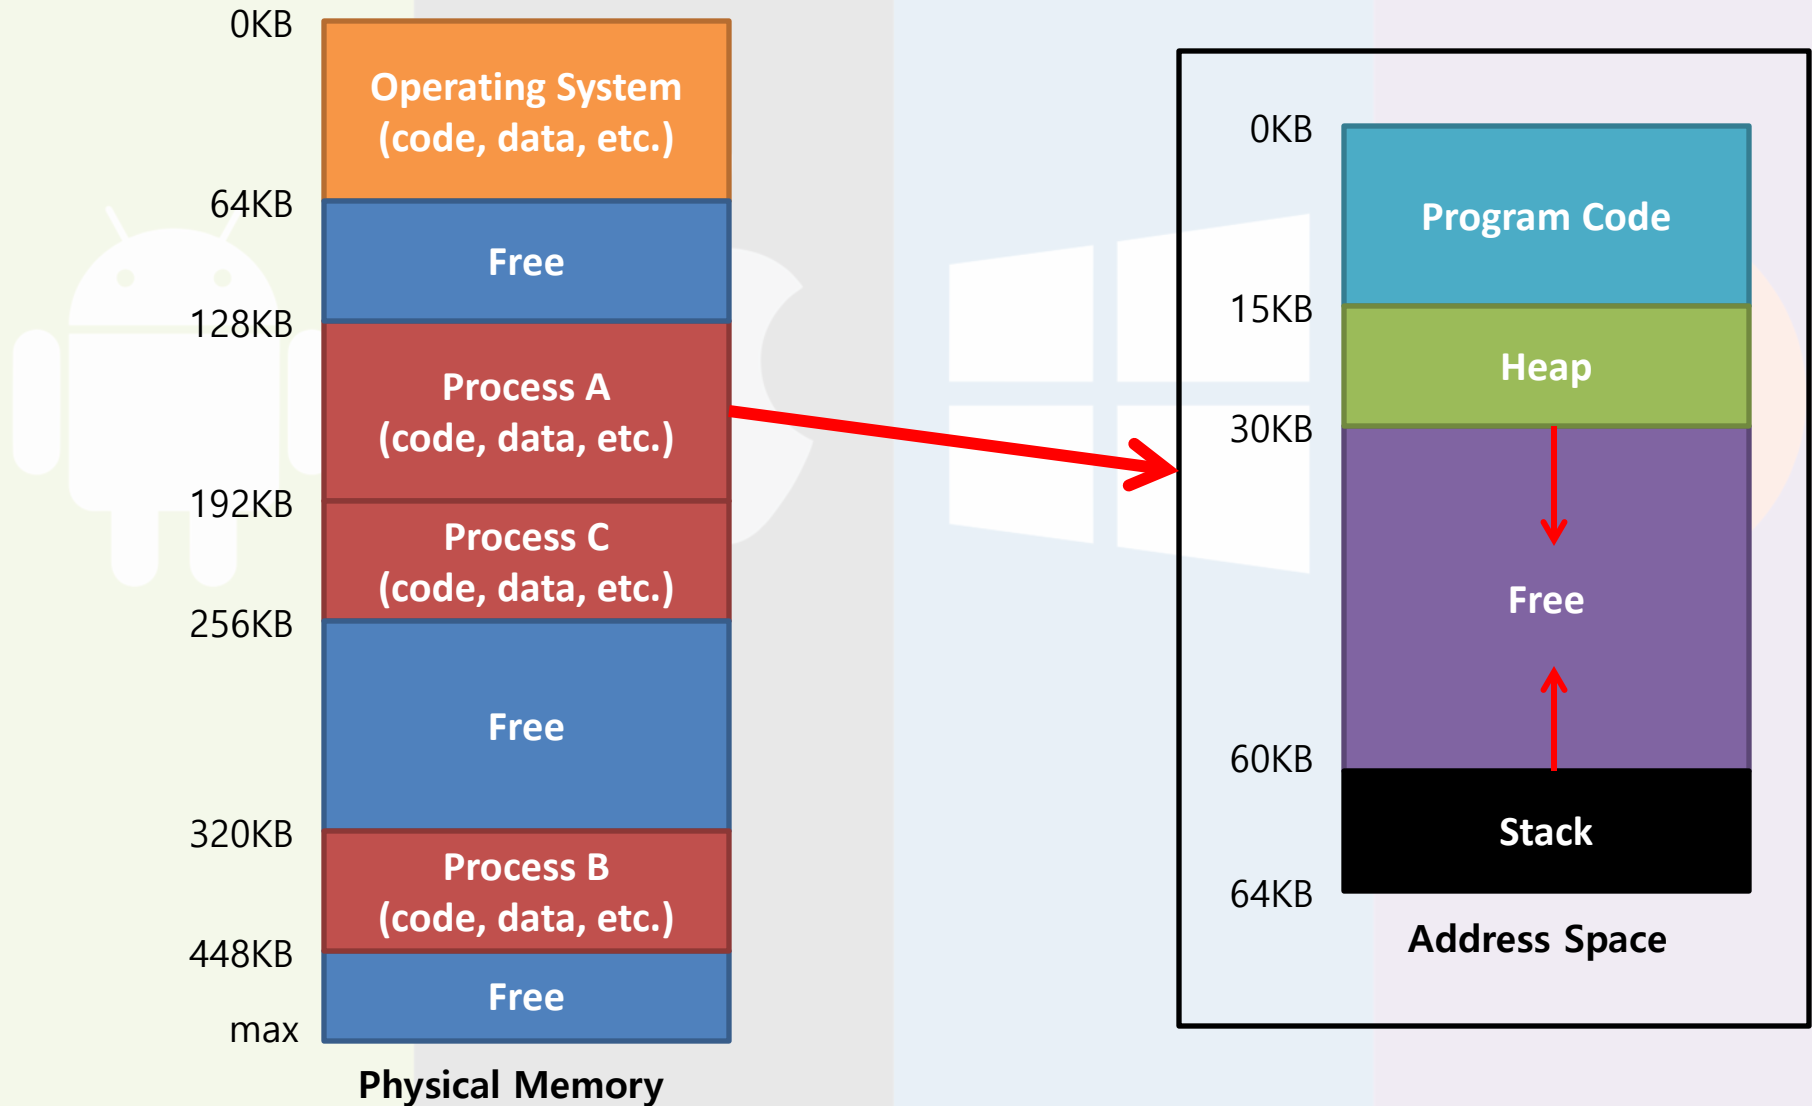


Address Space

- **Code**
 - Where instructions live
- **Heap**
 - Dynamically allocate memory
 - “malloc” in C language
 - “new” in object-oriented language
- **Stack**
 - Store return addresses or values
 - Contain local variables arguments to routines



Address Space



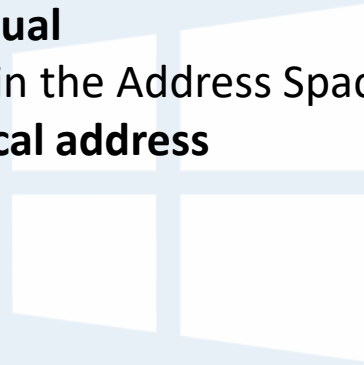
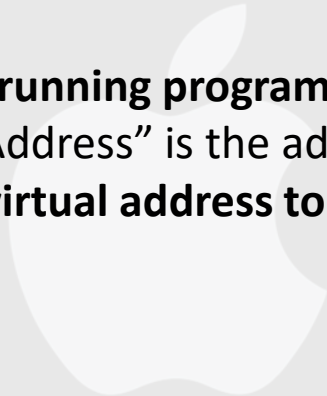
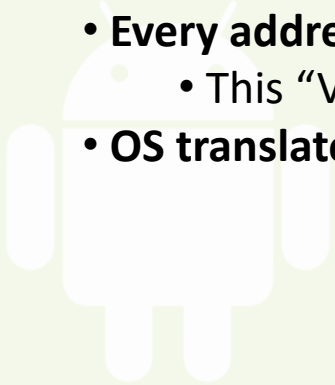
Well how do we map between Physical Memory and the Address Space?

- Well the OS ofcourse



Well how do we map between Physical Memory and the Address Space?

- Well the OS of course
- Every address in a running program is virtual
 - This “Virtual Address” is the address in the Address Space
- OS translates the virtual address to physical address



Virtual Addresses

- The virtualized address in address space

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]){

    printf("location of code   : %p\n", (void *) main);
    printf("location of heap   : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack  : %p\n", (void *) &x);

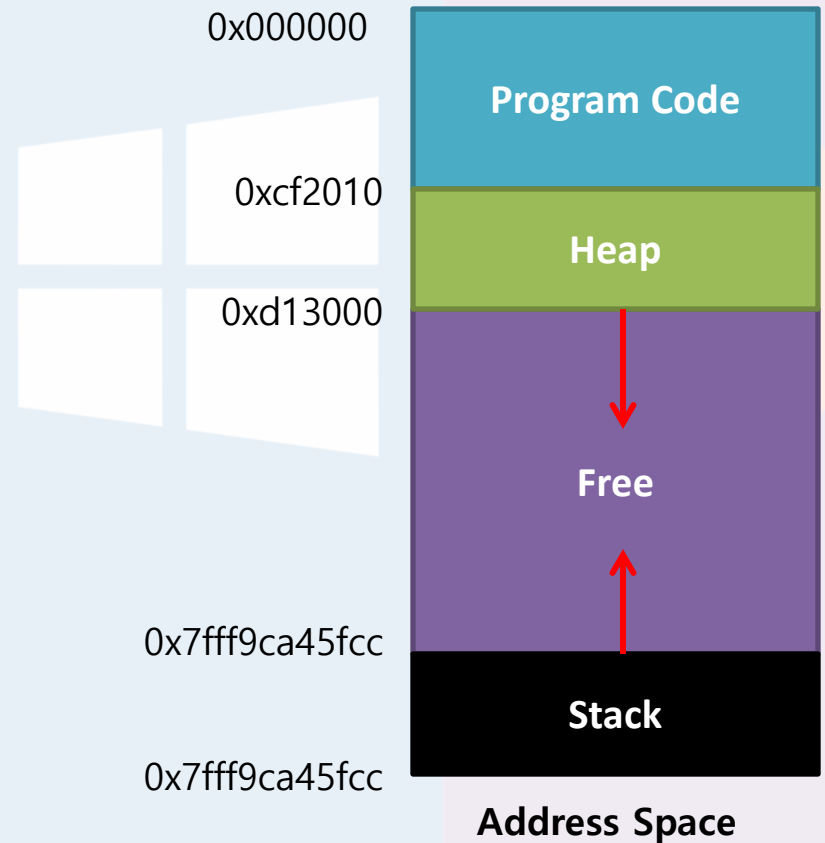
    return x;
}
```

A simple program that prints out addresses

Virtual Addresses

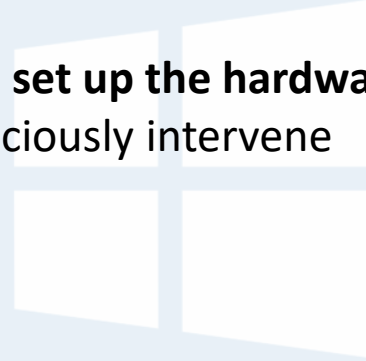

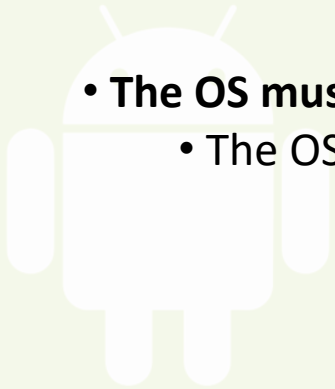
- The output in 64-bit Linux machine

```
location of code   : 0x000000  
location of heap   : 0xcf2010  
location of stack  : 0x7fff9ca45fcc
```



Address Translation

- **Hardware transforms a virtual address to a physical address**
 - The desired information is actually stored in a physical address
- **The OS must get involved at key points to set up the hardware**
 - The OS must manage memory to judiciously intervene



Address Translation - Example

- C - Language code

```
void func()  
    int x;  
    ...  
    x = x + 3; // this is the line of code we are interested in
```

- Load a value from memory
- Increment it by three
- Store the value back into memory

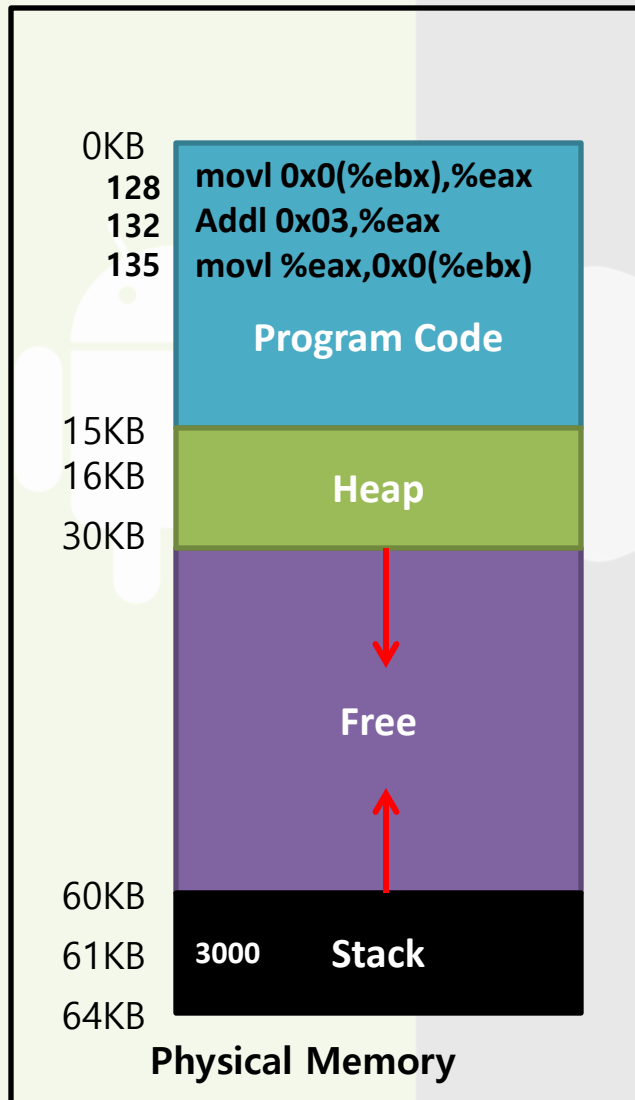
Address Translation - Example

- Assembly

128	:	movl	0x0(%ebx), %eax	; load 0+ebx into eax
132	:	addl	\$0x03, %eax	; add 3 to eax register
135	:	movl	%eax, 0x0(%ebx)	; store eax back to mem

- Presume that the address of 'x' has been place in ebx register
- Load the value at that address into eax register
- Add 3 to eax register
- Store the value in eax back into memory

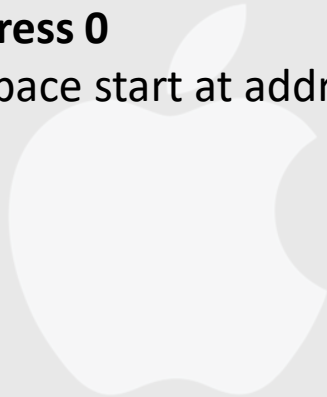

Address Translation - Example




- Fetch instruction at address 128
- Execute this instruction (load from address 61KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 61 KB)

What happens if the OS needs to relocate a process?

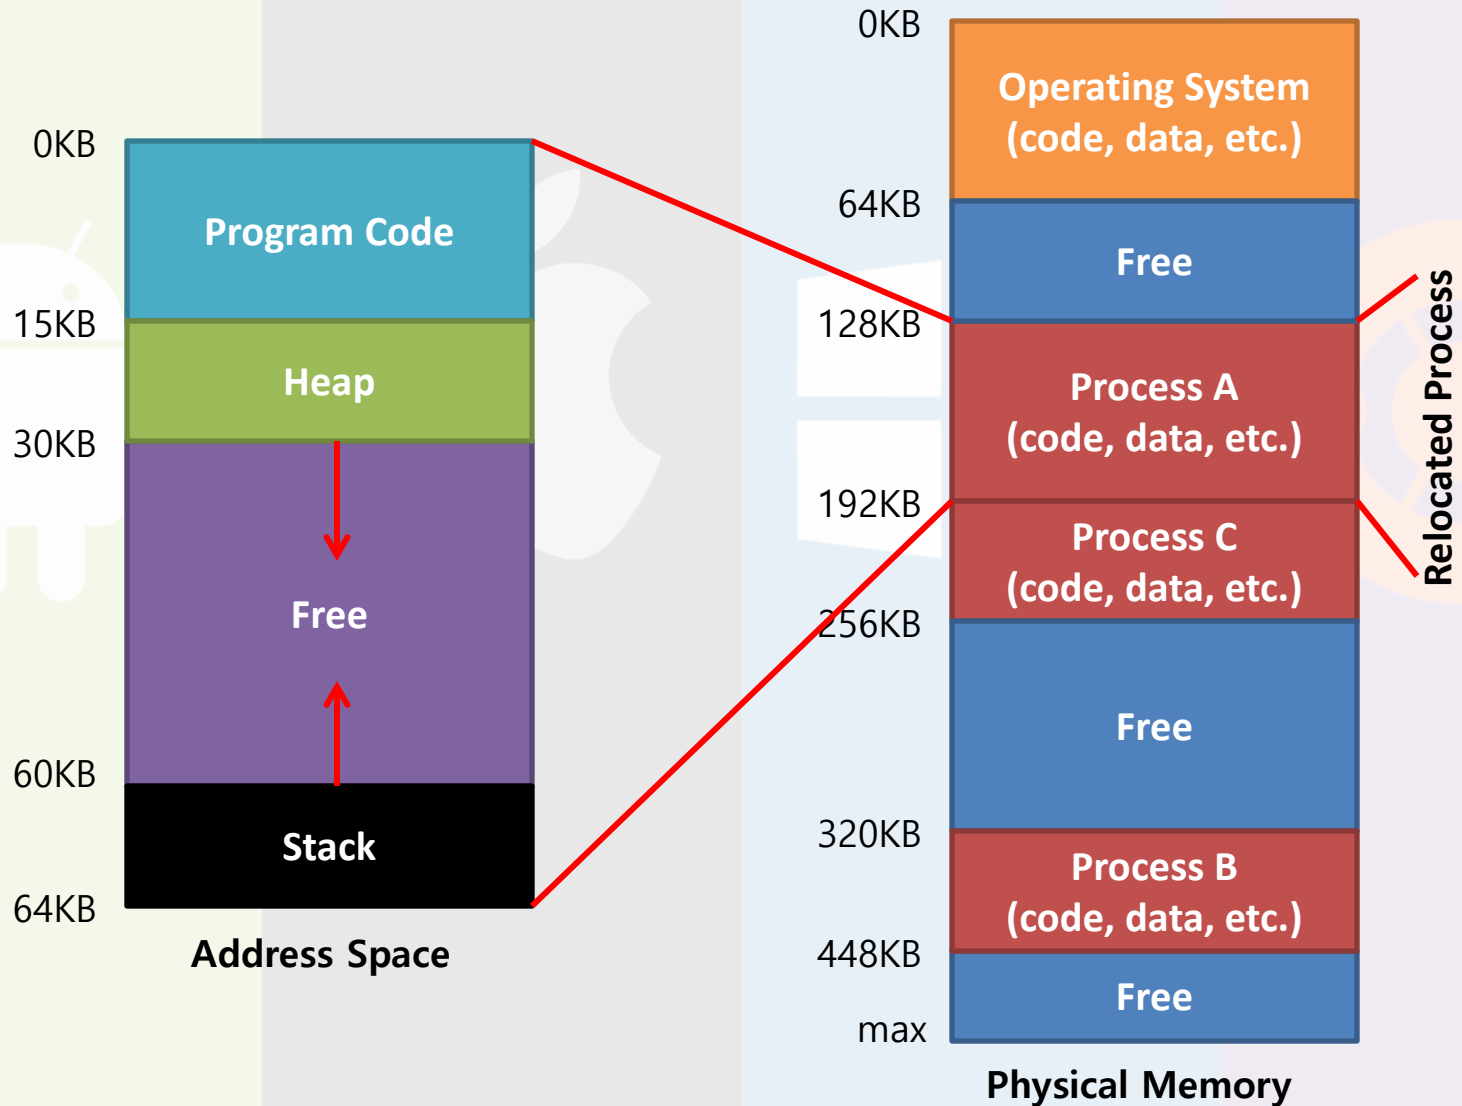
- The OS wants to place the process somewhere else in physical memory, not at address 0
 - The address space start at address 0



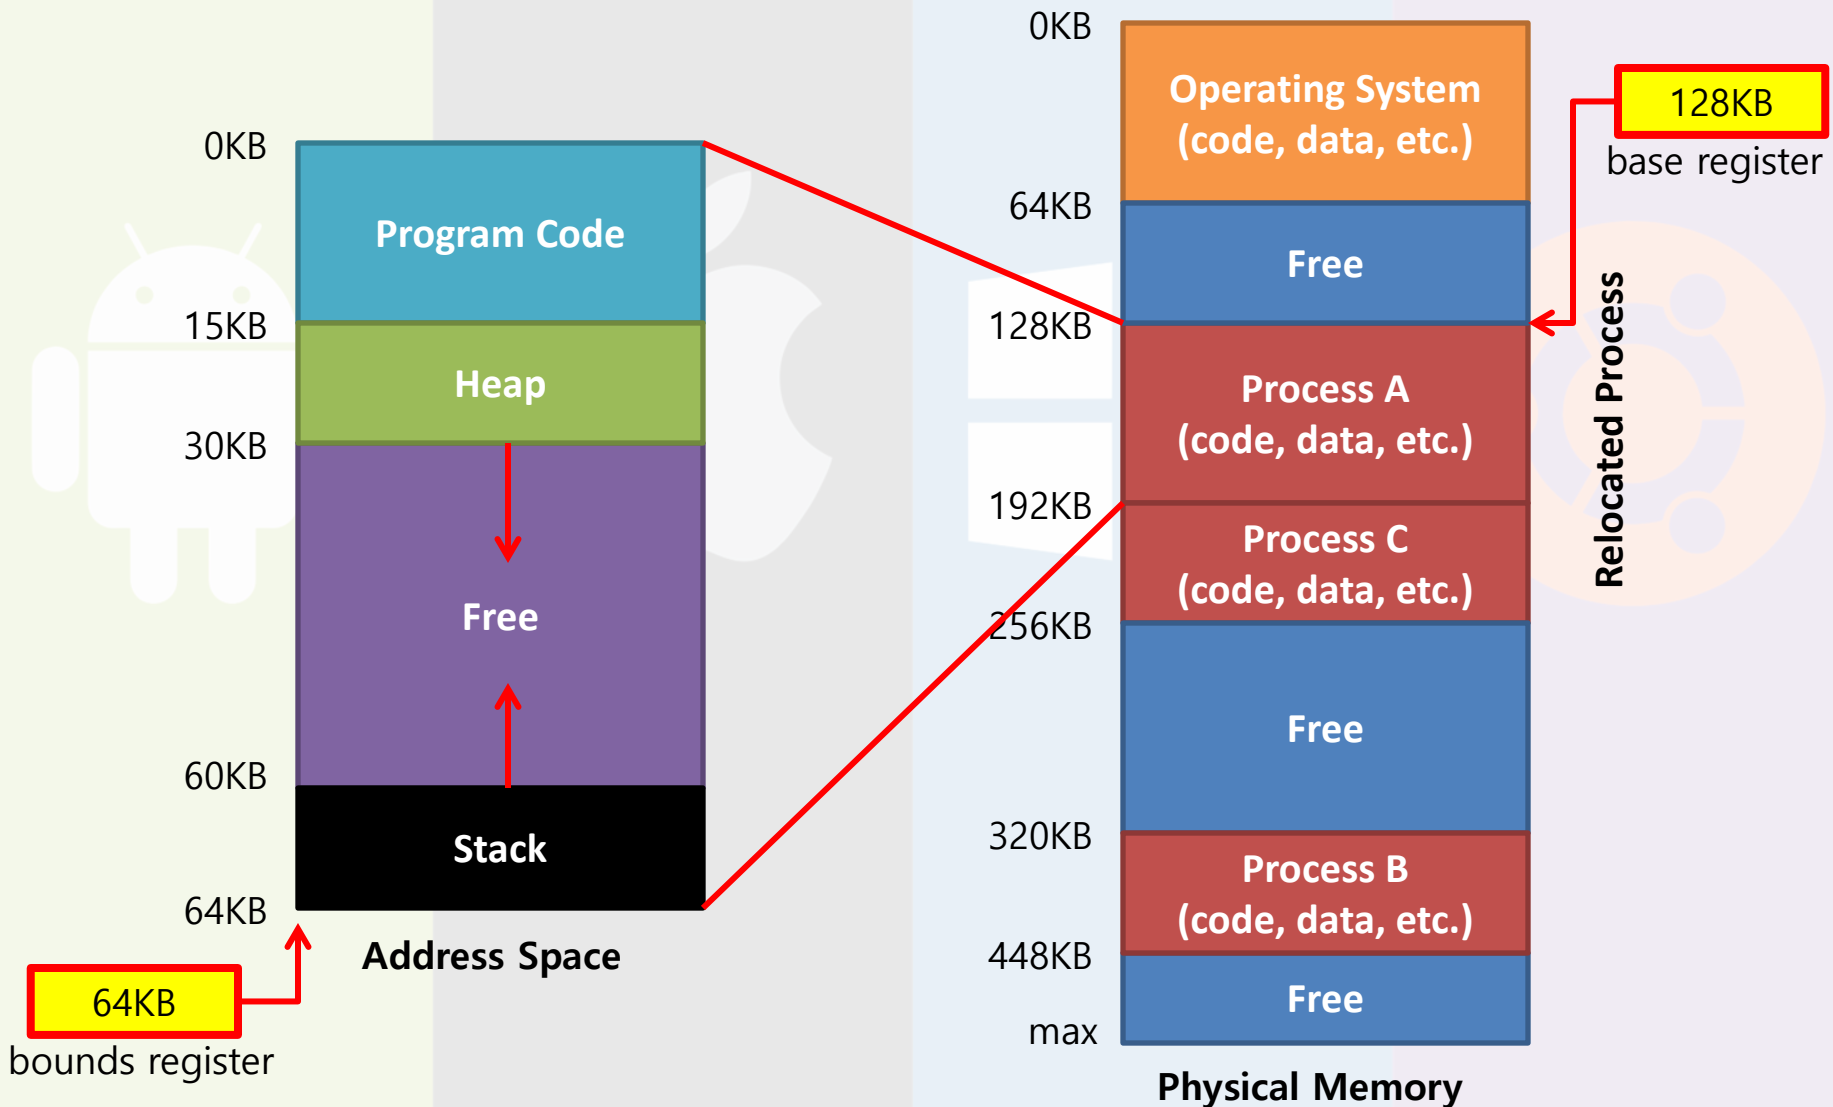
Memory Relocation



Memory Relocation



Base and Bounds Registers



Hardware base Relocation

- When a program starts running, the OS decides where in physical memory a process should be loaded
- Set the **base** register value

Physical address = virtual address + base

- Every virtual address must **not be greater than bound** and **negative**

$0 \leq \text{virtual address} < \text{bounds}$

Address Translation

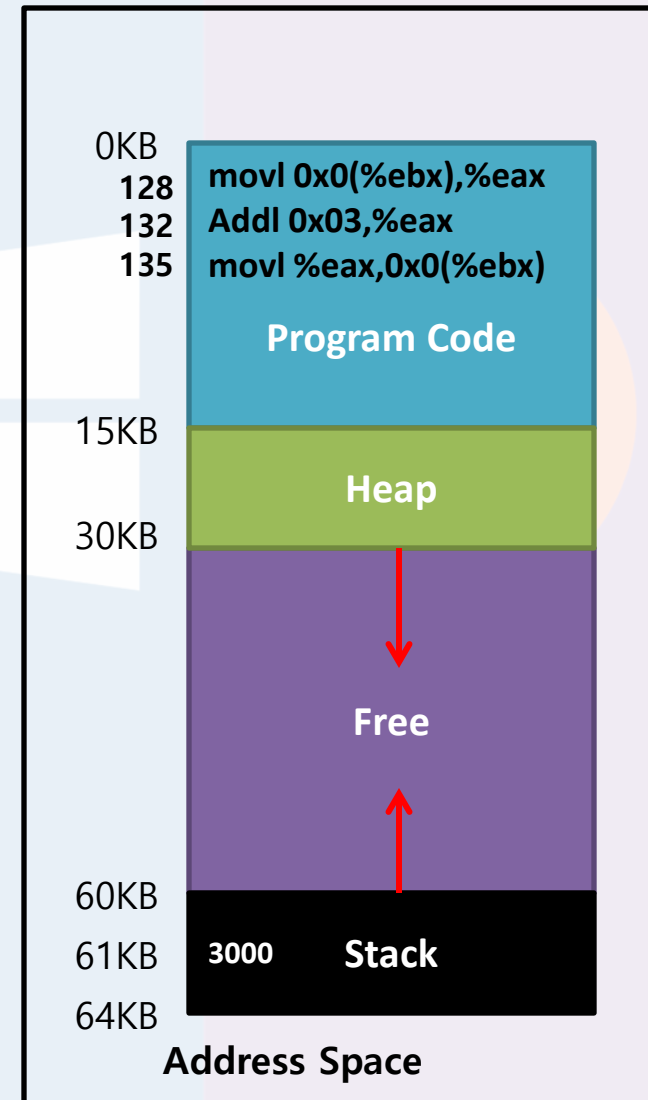
128 : `movl 0x0(%ebx), %eax`

- Fetch instruction at address 128

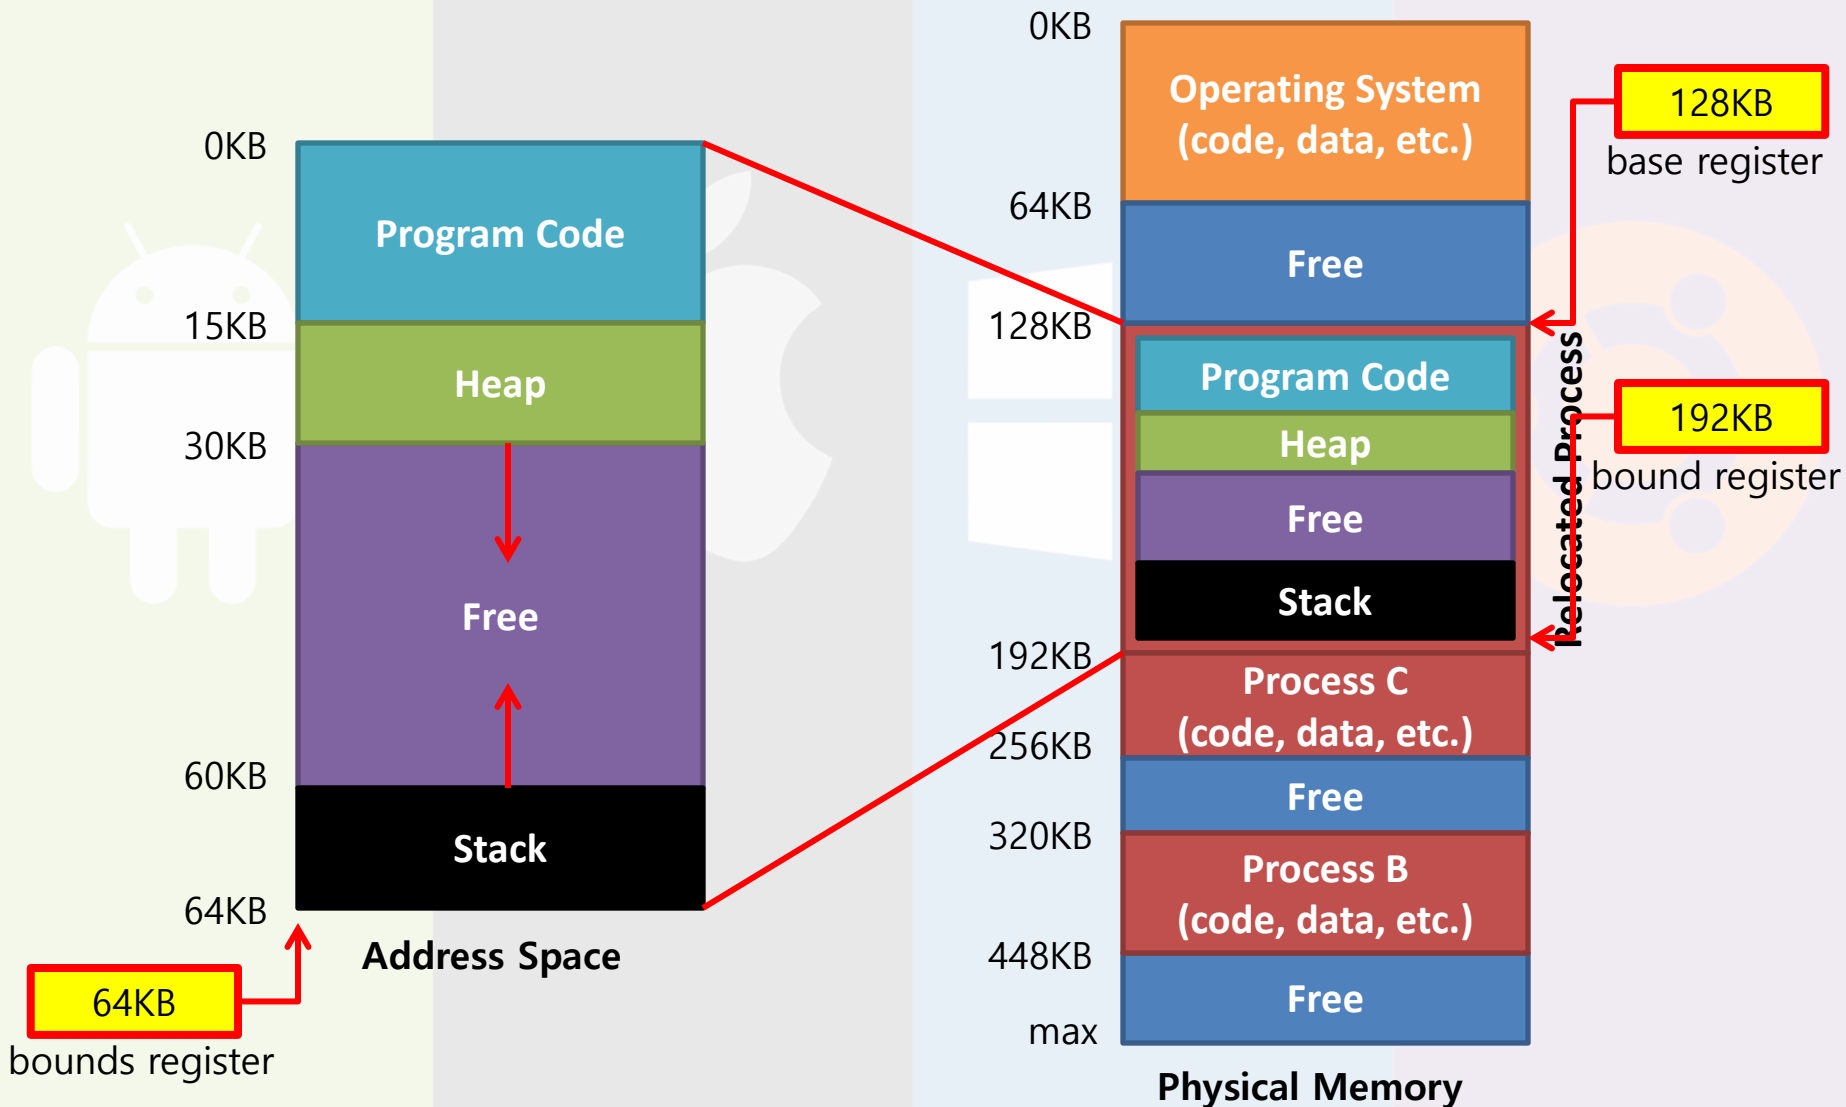
$$131200 = 128 + 128\text{KB}(\text{base})$$

- Execute this instruction
 - Load from address 61KB

$$189\text{KB} = 61\text{KB} + 128\text{KB}(\text{base})$$


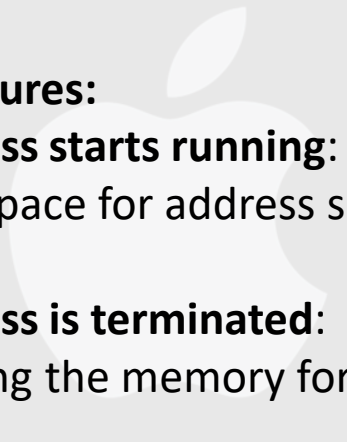
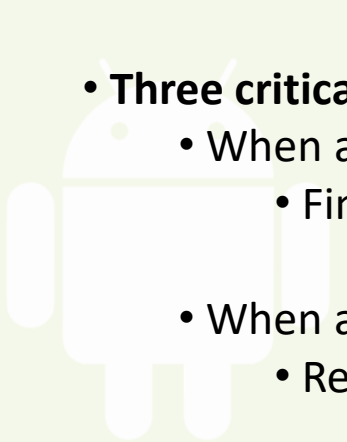


Two ways of Bounds Register



What are the Issues an OS must handle for Memory Virtualizing?

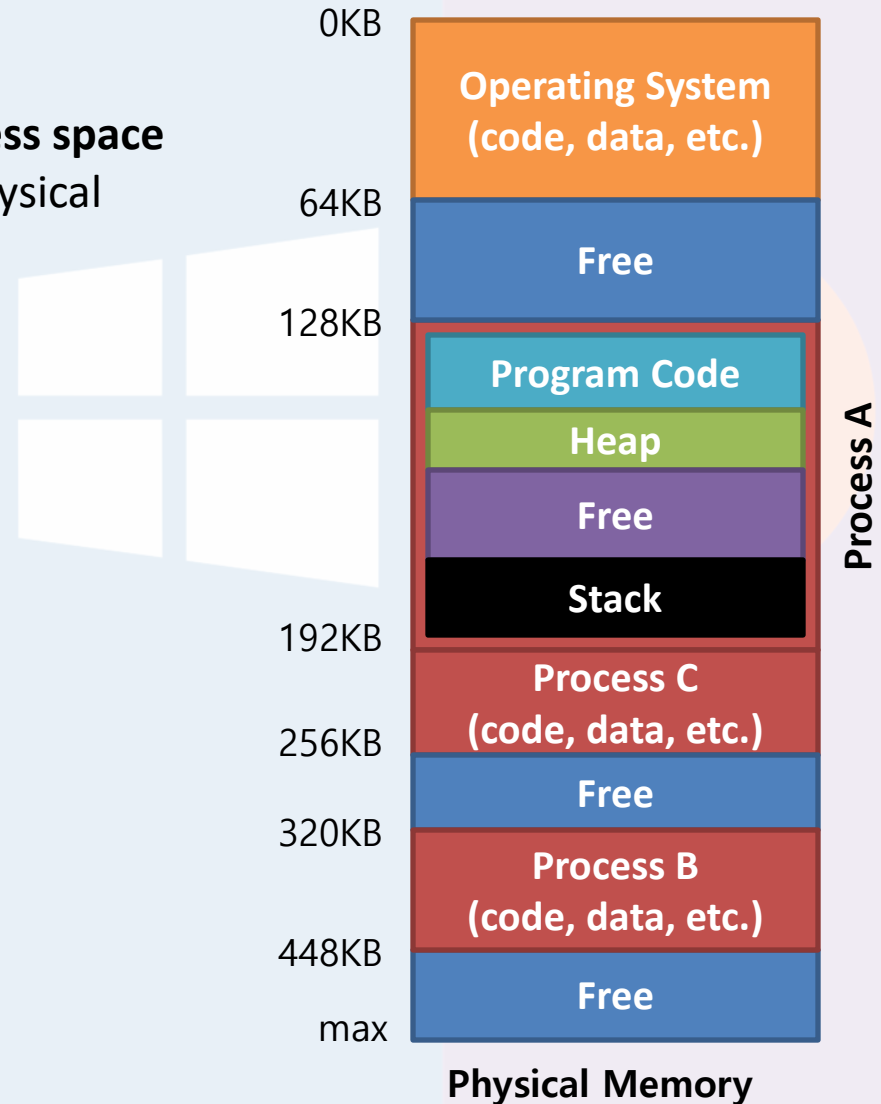

- The OS must take action to implement base-and-bounds approach
- Three critical junctures:
 - When a **process starts running**:
 - Finding space for address space in physical memory
 - When a **process is terminated**:
 - Reclaiming the memory for use
 - When **context switch occurs**:
 - Saving and storing the base-and-bounds pair



OS Issues: When a Process Starts Running

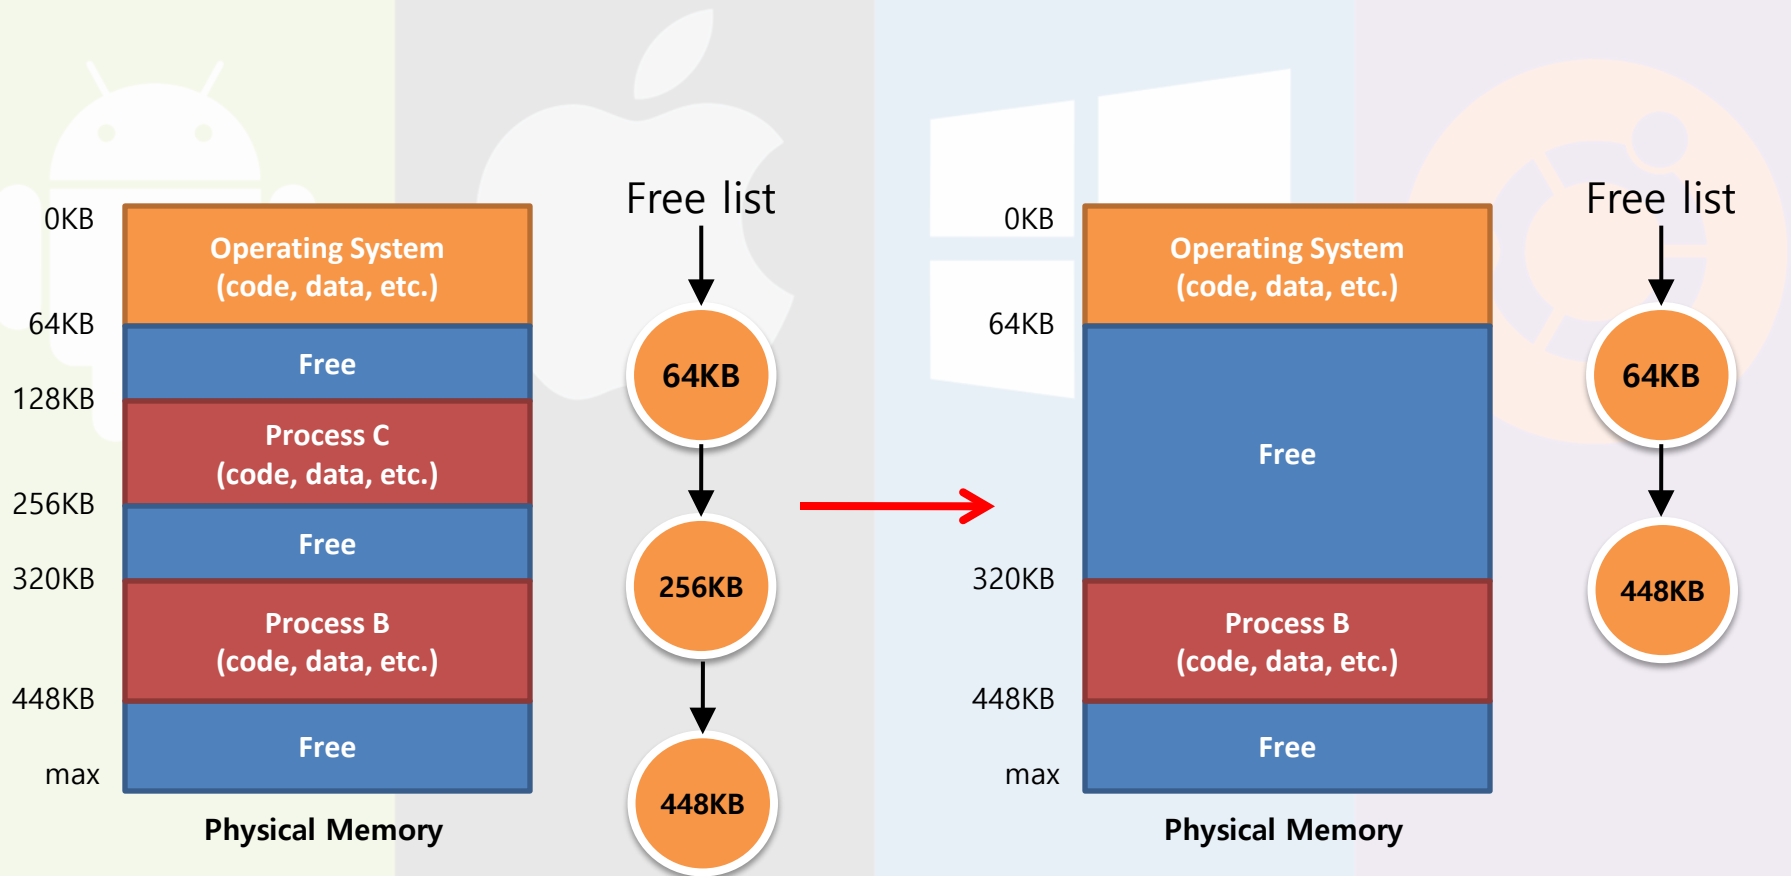
- The OS must find a room for a new address space
 - **free list** : A list of the range of the physical memory which are not in use

The OS lookup the free list



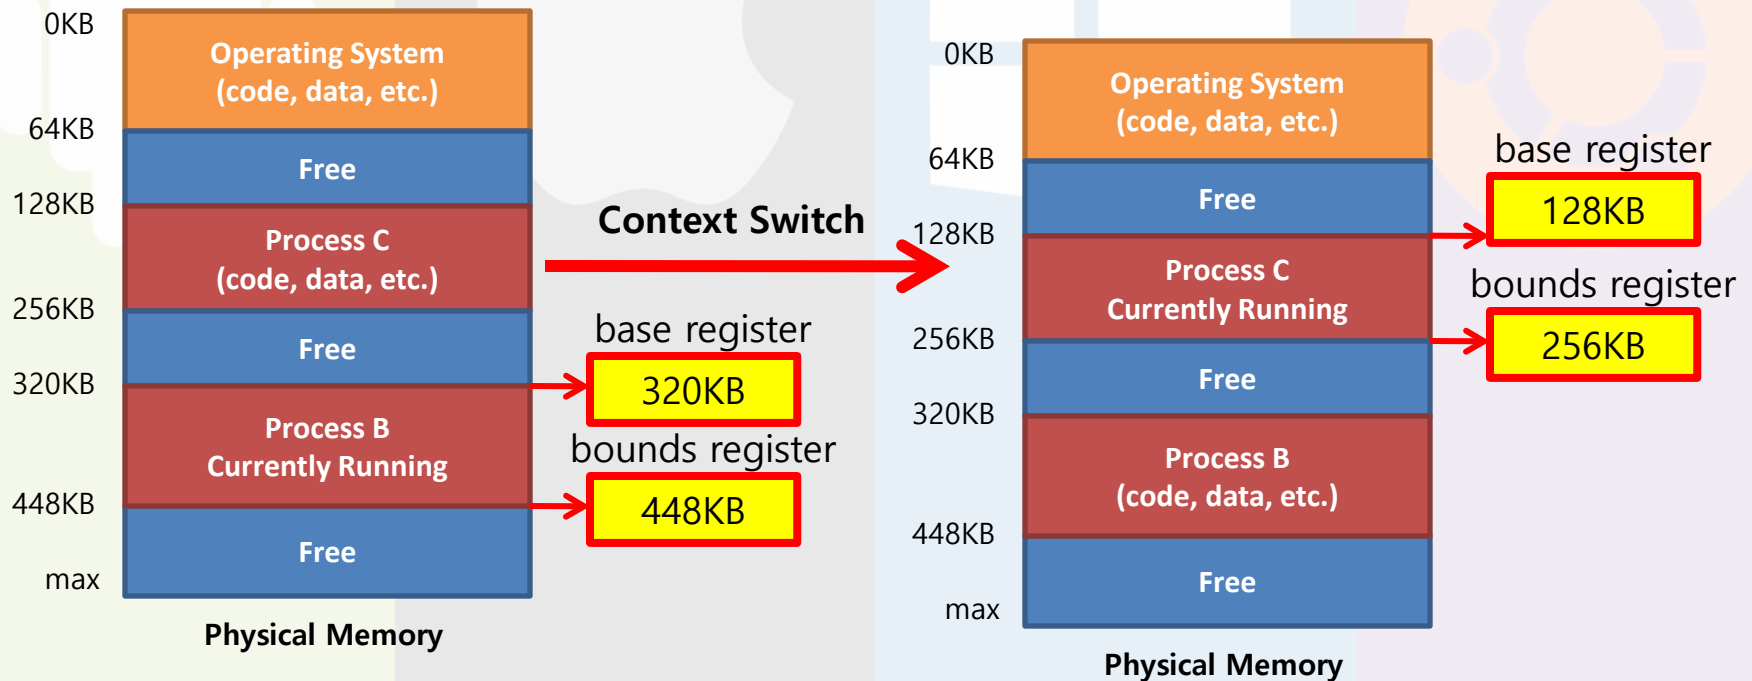
OS Issues: When a Process Is Terminated

- The OS must put the memory back on the free list



OS Issues: When Context Switch Occurs

- The OS must save and restore the base-and-bounds pair.
 - In process structure or **process control block(PCB)**



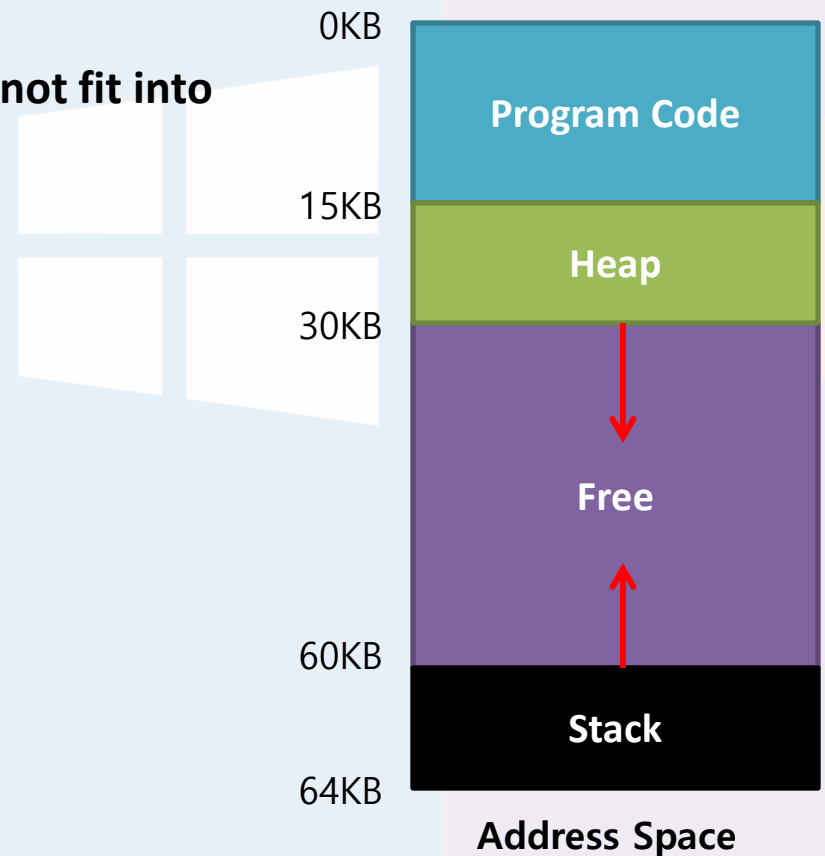
What if there is not enough space together to fit a program?

- How do we start a new program in this case



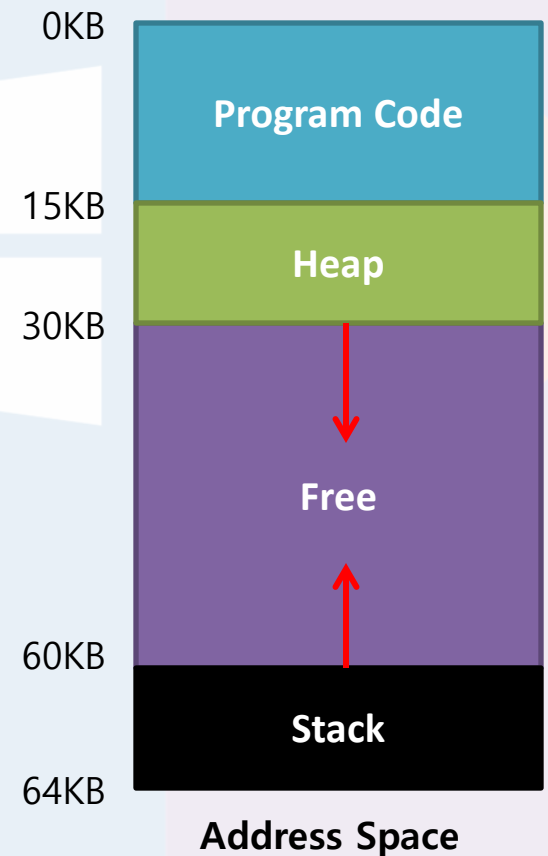
Inefficiency of the Base and Bound Approach

- Big chunk of “free” space
- “free” space takes up physical memory.
- Hard to run when an address space does not fit into physical memory



Segmentation

- Segment is just a contiguous portion of the address space of a particular length
 - Logically-different segment: code, stack, heap
- Each segment can be placed in different part of physical memory
 - Base and bounds exist per each segment




Segmentation

- Segment is just a contiguous portion of the address space of a particular length
 - Logically-different segment: code, stack, heap
- Each segment can be placed in different part of physical memory
 - Base and bounds exist per each segment

Segment Register

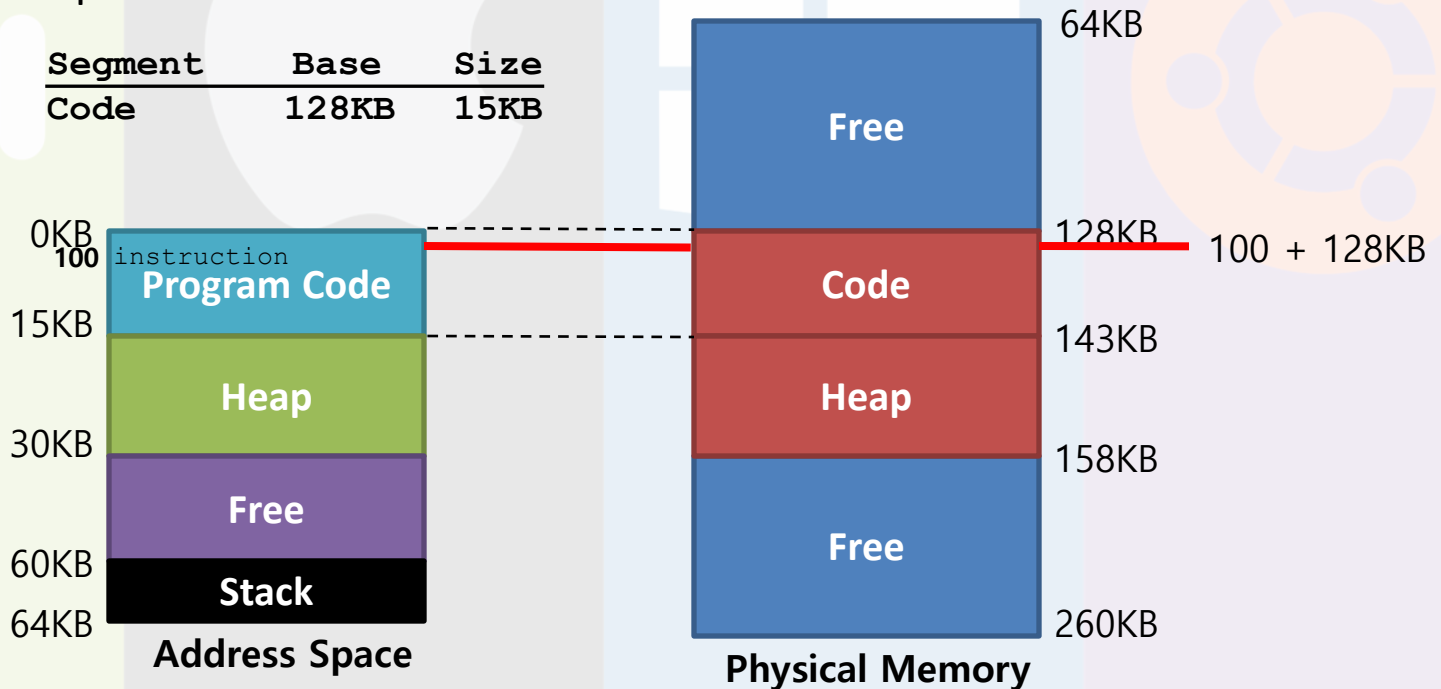
Segment	Base	Size
Code	256KB	4KB
Heap	260KB	60KB
Stack	128KB	64KB



Address Translation with Segmentation

Physical address = offset + base

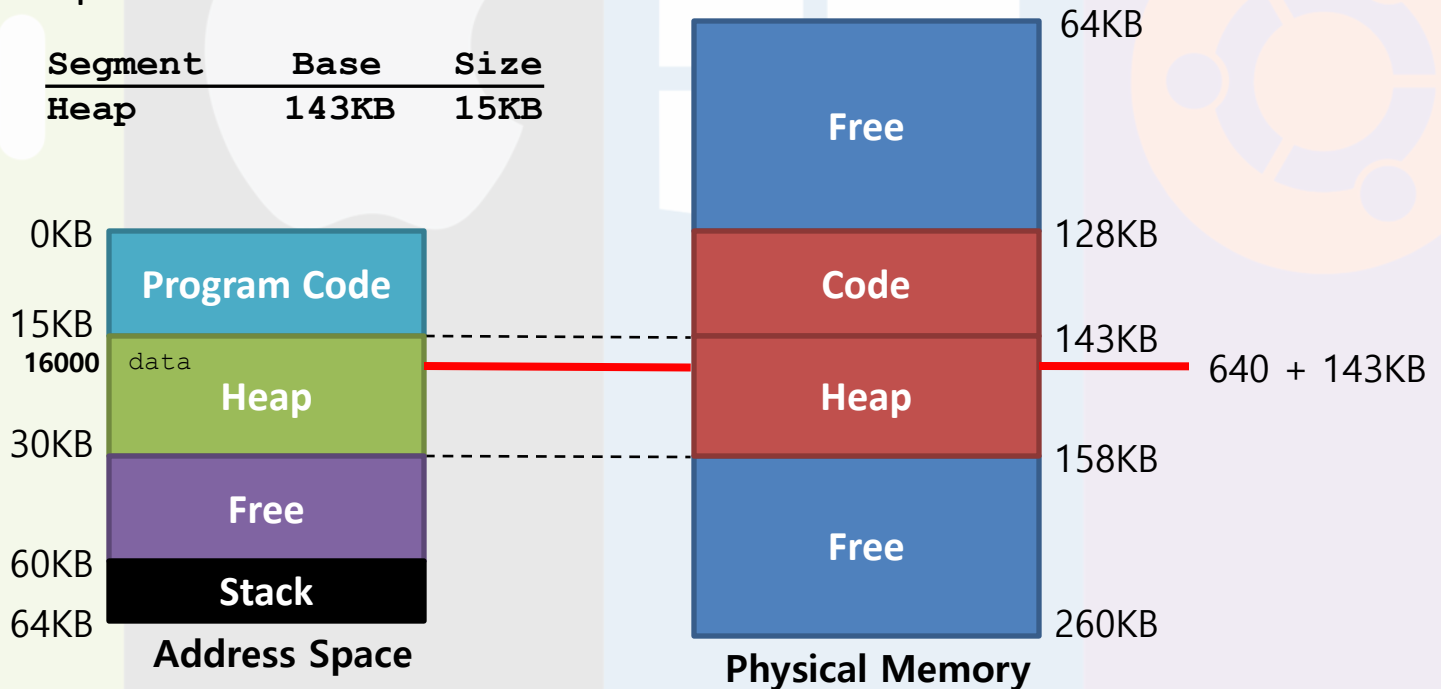
- The offset of virtual address 100 is 100.
- The code segment starts at virtual address 0 in address space.



Address Translation with Segmentation

Virtual address + base is not the correct physical address

- The offset of virtual address 16000 is 640.
- The code segment starts at virtual address 0 in address space.




What happens if an address is incorrectly referenced?

- Segmentation Fault occurs



Segmentation Faults


- If an illegal address such as 37KB which is beyond the end of heap is referenced, the OS occurs segmentation fault
 - The hardware detects that address is out of bounds



Referring to a Segment

- **Explicit approach**


- Chop up the address space into segments based on the top few bits of virtual address



Referring to a Segment


- **Explicit approach**

- Chop up the address space into segments based on the top few bits of virtual address



- **Example: virtual address 4200 (01000001101000)**

Segment	bits
Code	00
Heap	01
Stack	10
-	11




Referring to a Segment

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

- **SEG_MASK = 0x3000(11000000000000)**
- **SEG_SHIFT = 12**
- **OFFSET_MASK = 0xFFF (00111111111111)**
- VA = 01100011100011
- Segment = 01000000000000 >> 12 = 01
- Offset = 00100011100011

Referring to a Stack Segment

- **Stack grows backward**
- **Extra hardware support is need**
 - The hardware checks which way the segment grows
 - 1: positive direction
 - 0: negative direction



Segment Register(with Negative-Growth Support)

Segment	Base	Size	Grows	Positive?
Code	256KB	4KB		1
Heap	260KB	60KB		1
Stack	192KB	64KB		0

Support for Sharing

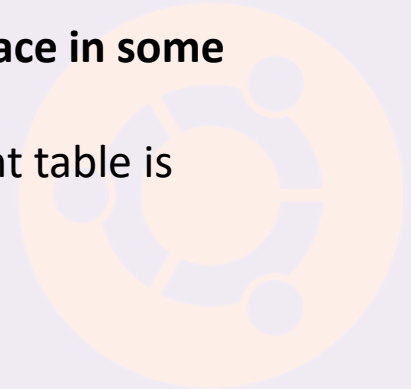
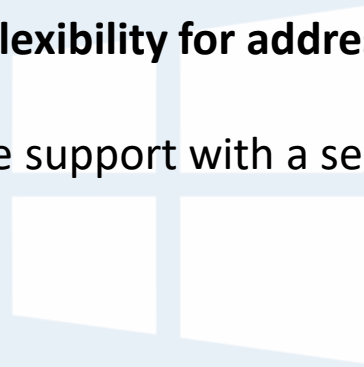
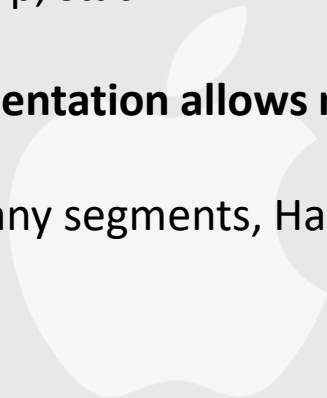
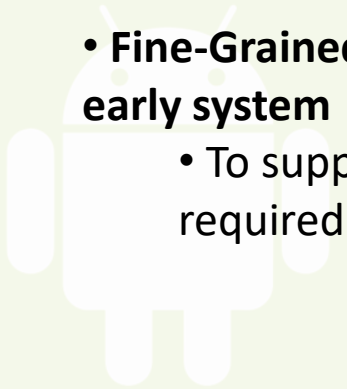
- **Segment can be shared between address space**
 - Code sharing is still in use in systems today
 - by extra hardware support
- **Extra hardware support is need for form of Protection bits**
 - A few more bits per segment to indicate permissions of read, write and execute

Segment Register(with Negative-Growth Support)

Segment	Base	Size	Grows	Positive?	Protection
Code	256KB	4KB		1	Read-Execute
Heap	260KB	60KB		1	Read-Write
Stack	192KB	64KB		0	Read-Write


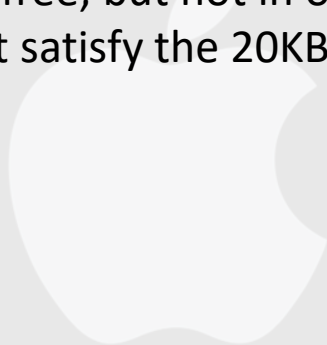

Fine-Grained and Coarse-Grained

- **Coarse-Grained means segmentation in a small number**
 - e.g., code, heap, stack
- **Fine-Grained segmentation allows more flexibility for address space in some early system**
 - To support many segments, Hardware support with a segment table is required



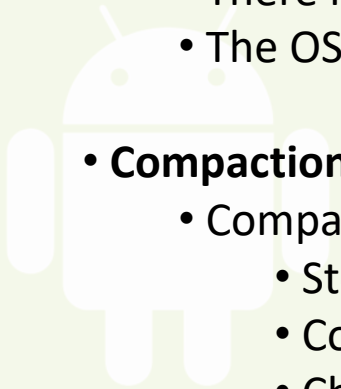
OS support: Fragmentation

- **External Fragmentation:** little holes of free space in physical memory that make difficulty to allocate new segments
 - There is 24KB free, but not in one contiguous segment
 - The OS cannot satisfy the 20KB request

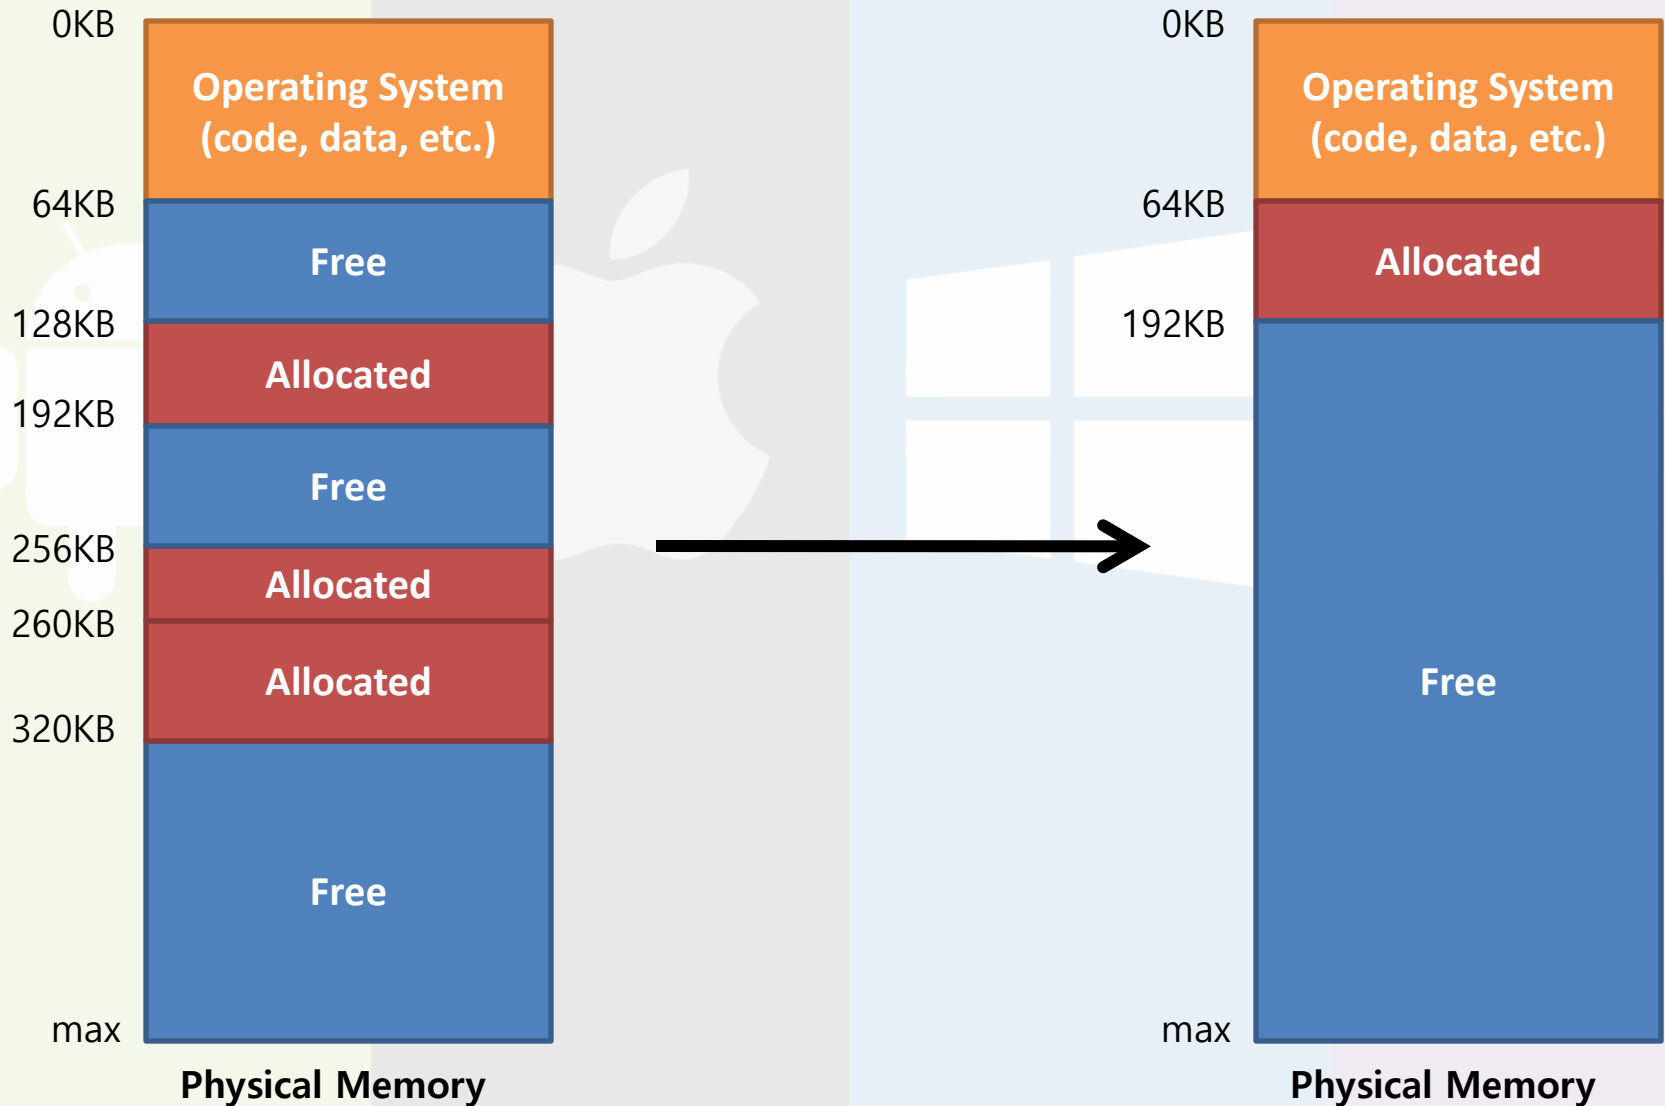


OS support: Fragmentation

- **External Fragmentation:** little holes of free space in physical memory that make difficulty to allocate new segments
 - There is 24KB free, but not in one contiguous segment
 - The OS cannot satisfy the 20KB request
- **Compaction:** rearranging the existing segments in physical memory
 - Compaction is costly
 - Stop running process
 - Copy data to somewhere
 - Change segment register value





Memory Compaction



Free-Space Management

- Splitting
- Coalescing
- Keeping Track of Memory Region Sizes
- Strategies for Managing Memory



Splitting

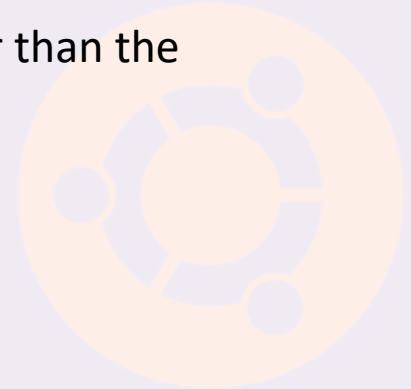

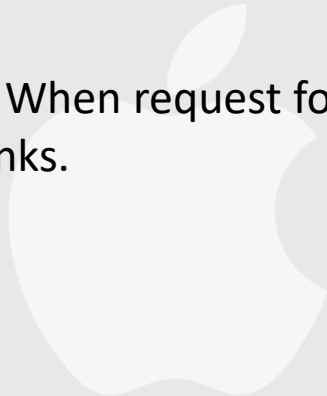

- Finding a free chunk of memory that can satisfy the request and splitting it into two



Splitting

- Finding a free chunk of memory that can satisfy the request and splitting it into two

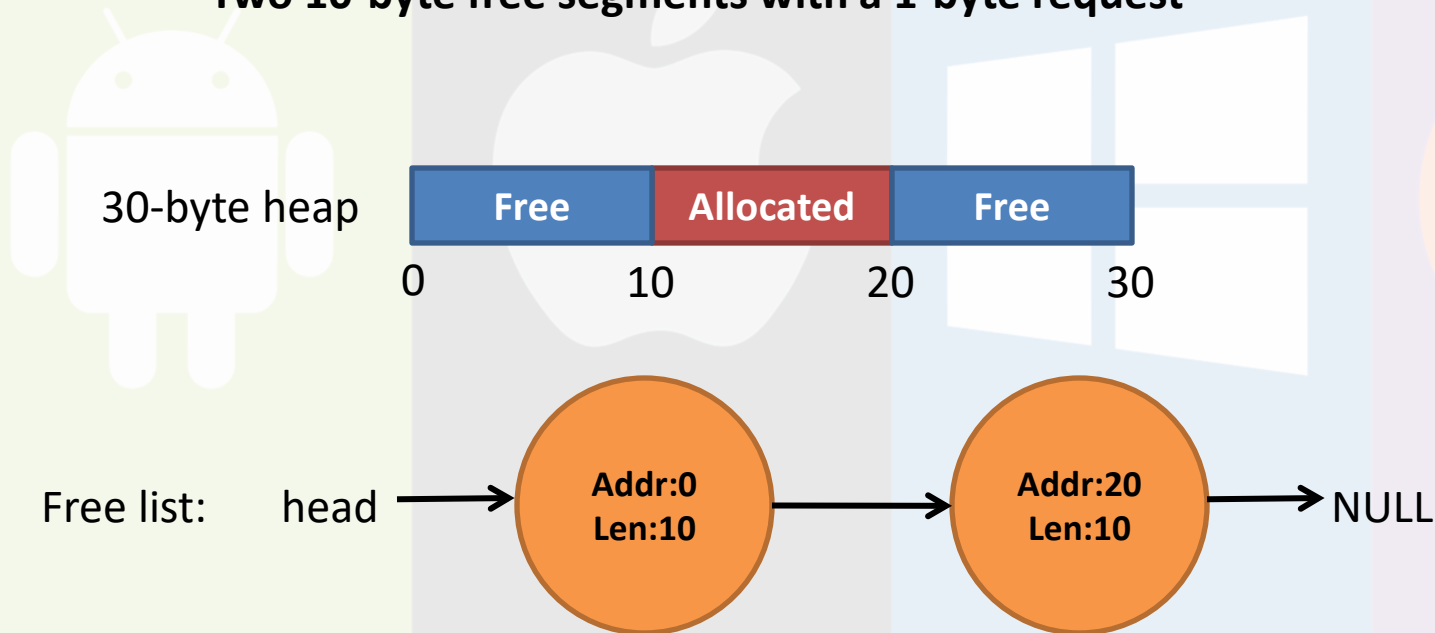
- **For Example** : When request for memory allocation is smaller than the size of free chunks.



Splitting

- **For Example** : When request for memory allocation is smaller than the size of free chunks.

- **Two 10-byte free segments with a 1-byte request**





Before Splitting

Splitting

- **For Example** : When request for memory allocation is smaller than the size of free chunks.

- **Two 10-byte free segments with a 1-byte request**



After Splitting 10-byte free segment

Coalescing

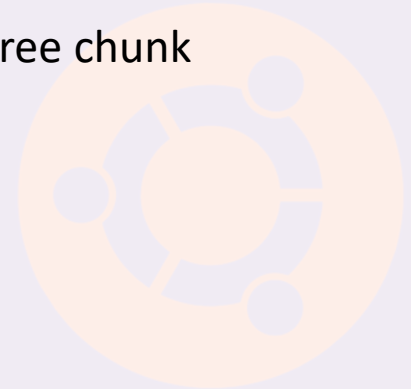

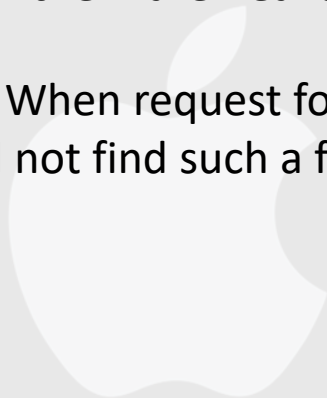

- Merge returning a free chunk with existing chunks into a large single free chunk if addresses of them are nearby



Coalescing

- **Merge** returning a free chunk with existing chunks into a large single free chunk if addresses of them are nearby

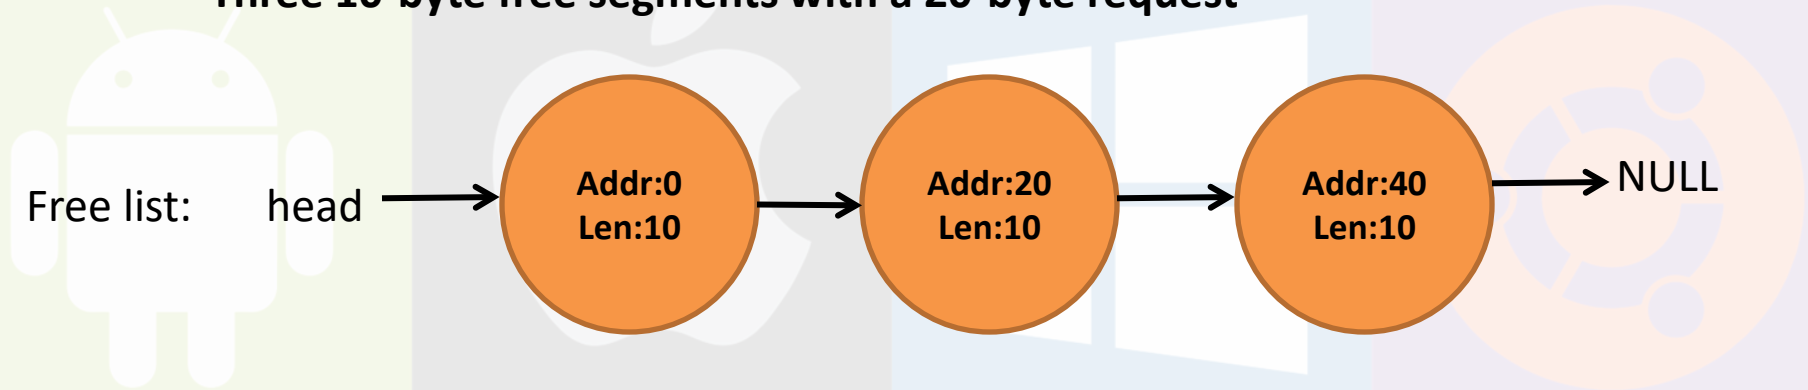
- **For Example** : When request for memory that is bigger than free chunk size, the list will not find such a free chunk




Coalescing

- **For Example** : When request for memory that is bigger than free chunk size, the list will not find such a free chunk

- **Three 10-byte free segments with a 20-byte request**



Before Coalescing

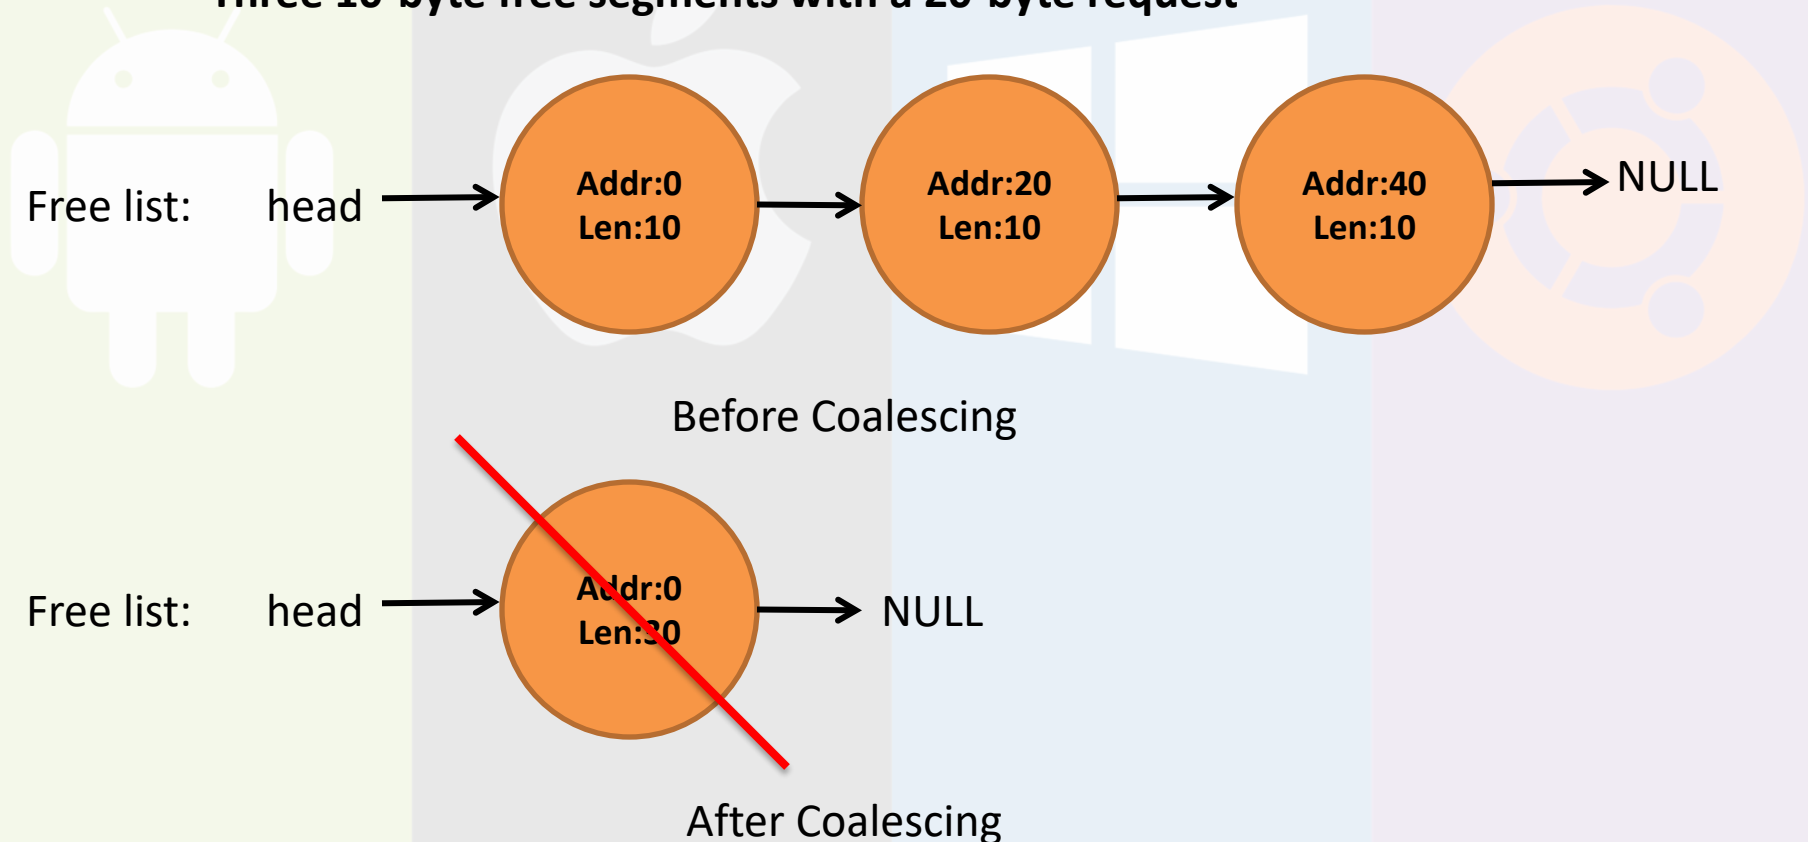


After Coalescing

Coalescing

- **For Example** : When request for memory that is bigger than free chunk size, the list will not find such a free chunk

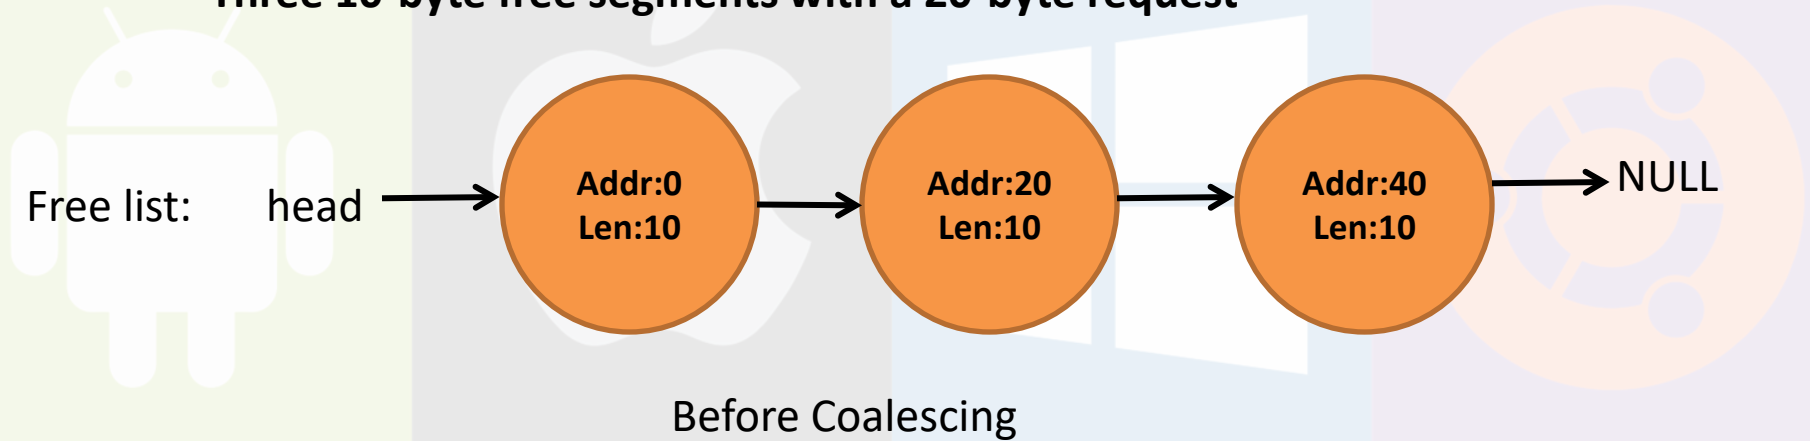
- **Three 10-byte free segments with a 20-byte request**



Coalescing

- **For Example :** When request for memory that is bigger than free chunk size, the list will not find such a free chunk

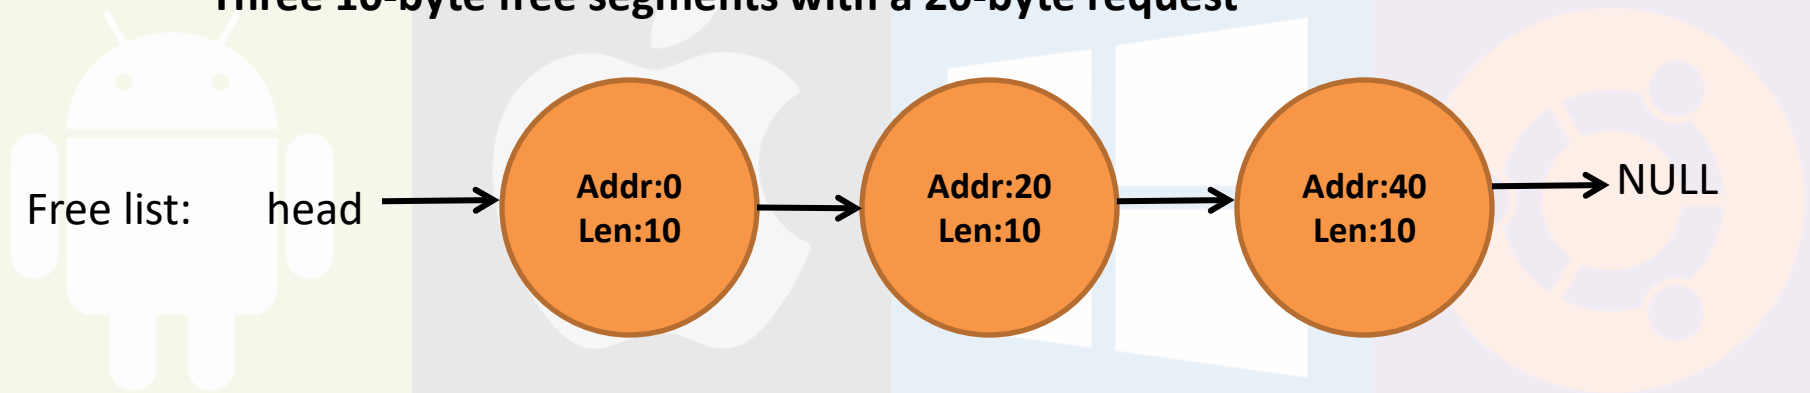
- **Three 10-byte free segments with a 20-byte request**




Coalescing

- **For Example :** When request for memory that is bigger than free chunk size, the list will not find such a free chunk

- **Three 10-byte free segments with a 20-byte request**



Before Coalescing

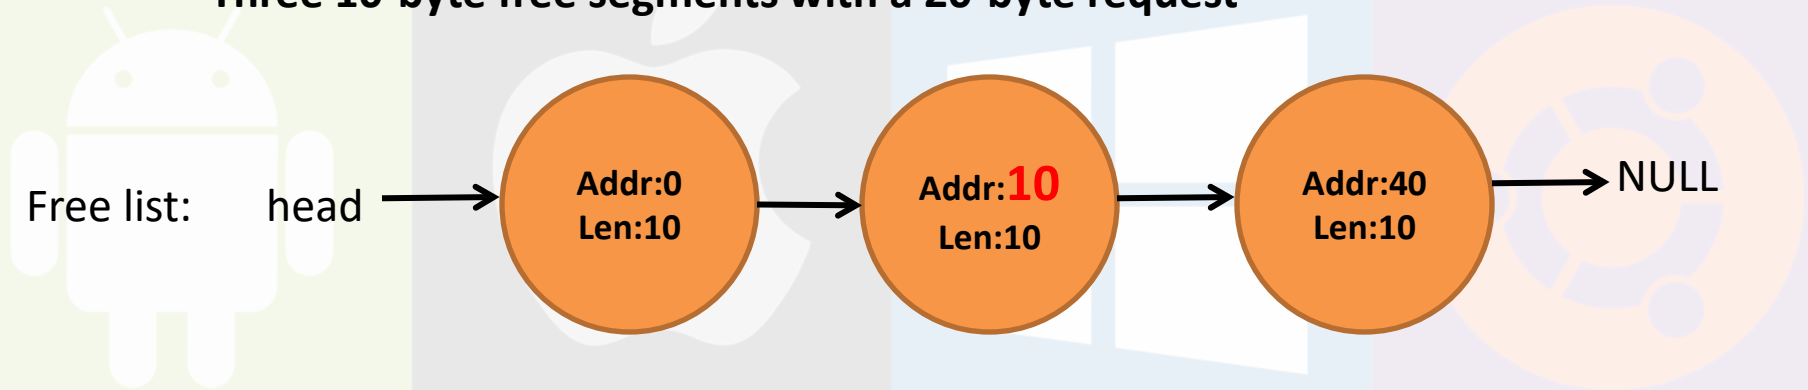


After Coalescing


Coalescing

- **For Example** : When request for memory that is bigger than free chunk size, the list will not find such a free chunk

- **Three 10-byte free segments with a 20-byte request**



Before Coalescing



After Coalescing



How do we do this reorganisation in an efficient way

- By tracking the size of allocated regions
- By embedding the free list



Tracking the size of allocated regions

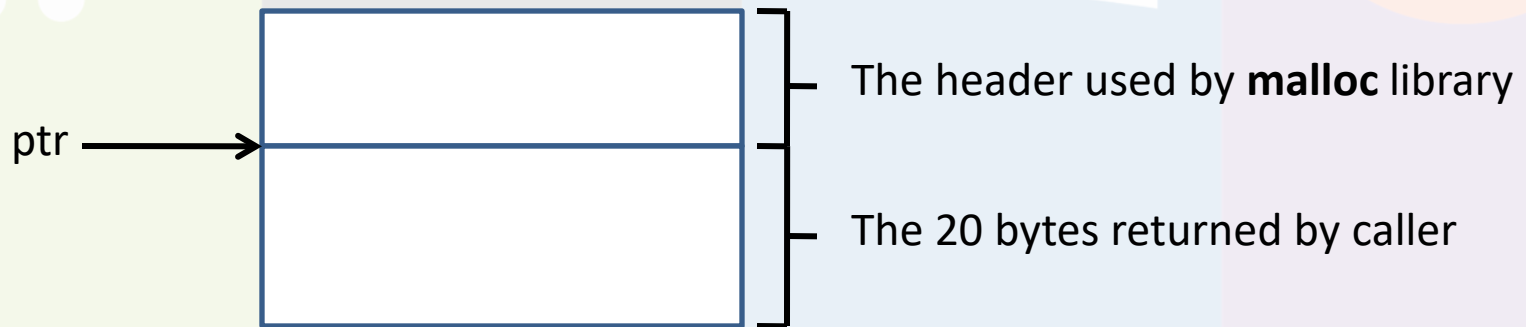
- The interface to `free(void *ptr)` does not take a size parameter.
- How does the library know the size of memory region that will be back into free list?



Tracking the size of allocated regions

- The interface to `free(void *ptr)` does not take a size parameter.
- How does the library know the size of memory region that will be back into free list?
 - Most allocators store **extra information** in a **header** block

```
ptr = malloc(20);
```



An Allocated Region Plus Header

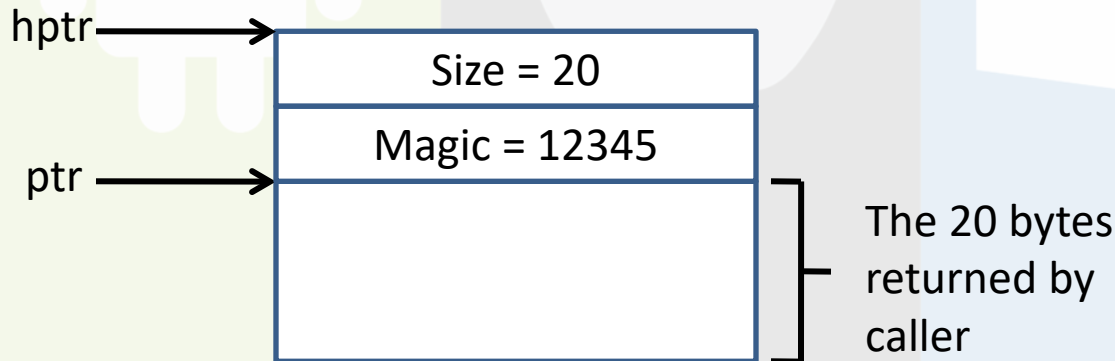
The header

- The header minimally contains the size of the allocated memory region
- The header may also contain

Additional pointers to speed up deallocation

A magic number for integrity checking

```
ptr = malloc(20);
```




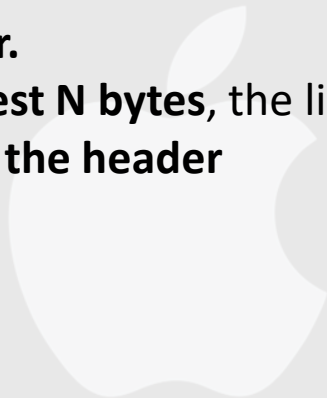

An Allocated Region Plus Header

```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

Simple Header

The header

- The size for required free region is the size of the header plus the size of the space allocated to the user.
- If a user request **N bytes**, the library searches for a free chunk of **size N plus the size of the header**



The header

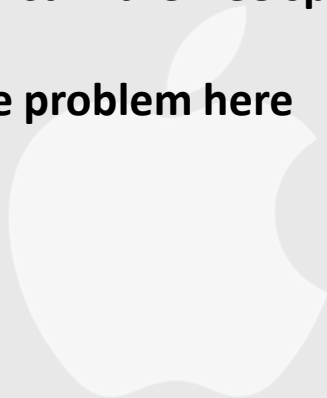

- The size for required free region is the size of the header plus the size of the space allocated to the user.
 - If a user request **N bytes**, the library searches for a free chunk of **size N plus the size of the header**
- Simple pointer arithmetic to find the header pointer

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```

Embedding a Free List

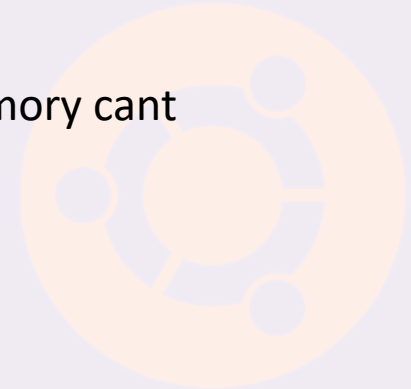
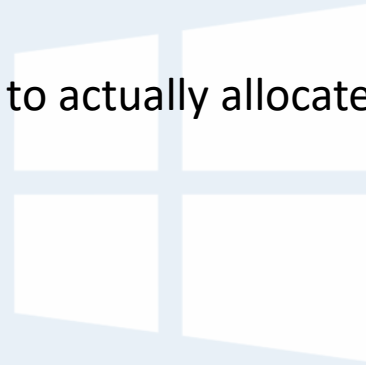
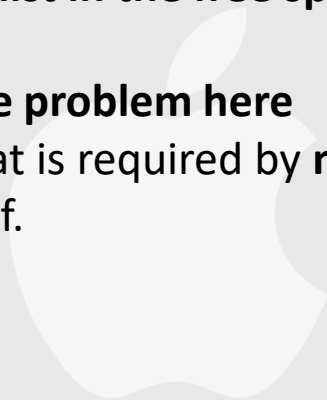
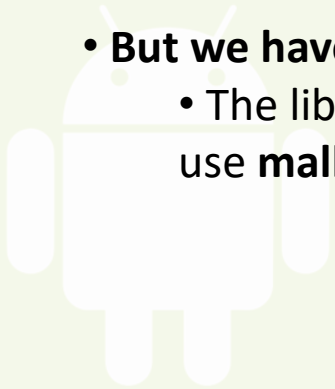
- The memory-allocation library initializes the heap and puts the first element of the free list in the free space

- But we have a little problem here



Embedding a Free List

- The memory-allocation library initializes the heap and puts the first element of the free list in the free space
- But we have a little problem here
 - The library that is required by **malloc** to actually allocate memory cant use **malloc** itself.



Embedding a Free List

- The memory-allocation library initializes the heap and puts the first element of the free list in the free space
- But we have a little problem here
 - The library that is required by **malloc** to actually allocate memory cant use **malloc** itself.
- Therefore we embed the free list directly in the heap

Embedding a Free List

- The memory-allocation library initializes the heap and puts the first element of the free list in the free space
- But we have a little problem here
 - The library that is required by **malloc** to actually allocate memory cant use **malloc** itself.
- Therefore we embed the free list directly in the heap

```
typedef struct __node_t {  
    int size;  
    struct __node_t *next;  
} nodet_t;
```

Description of a node of the free list

Embedding a Free List

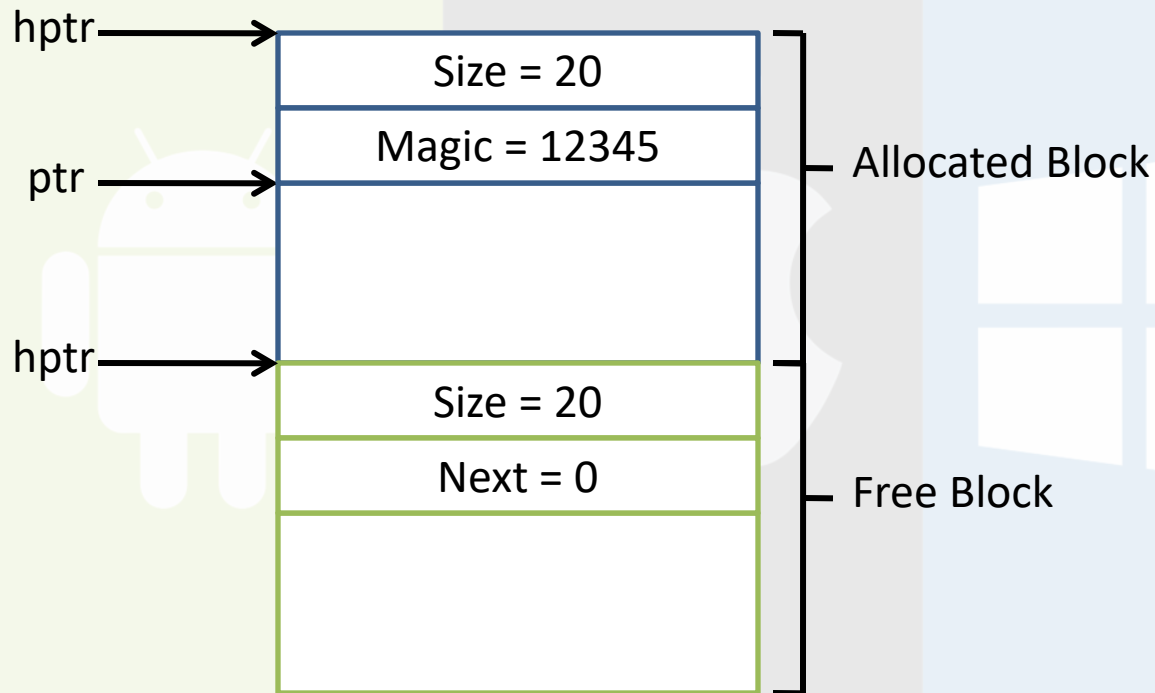
```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

Simple Header for an allocated
block

```
typedef struct __node_t {  
    int size;  
    struct __node_t *next;  
} node_t;
```


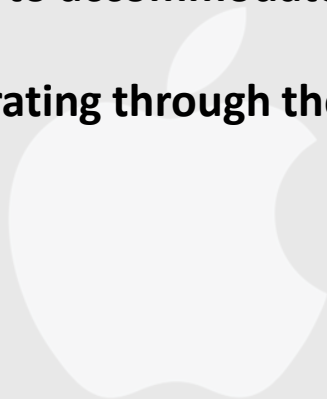

Description of a node of the
free list

Embedding a Free List



Embedding a Free List : Allocation

- If a chunk of memory is requested, the library will first find a chunk that is large enough to accommodate the request
- This is done by iterating through the free list and checking the size



Embedding a Free List : Allocation

- If a chunk of memory is requested, the library will first find a chunk that is large enough to accommodate the request
- This is done by iterating through the free list and checking the size
- The library will:
 - **Split** the large free chunk into two.
 - **One** for the **request** and the other being the **remaining** free chunk
 - **Shrink** the size of free chunk in the list


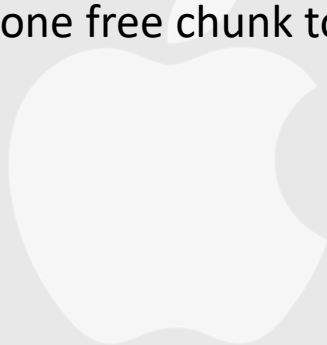

Embedding a Free List : Allocation

- **Example** a request for 100 bytes by `ptr = malloc(100)`



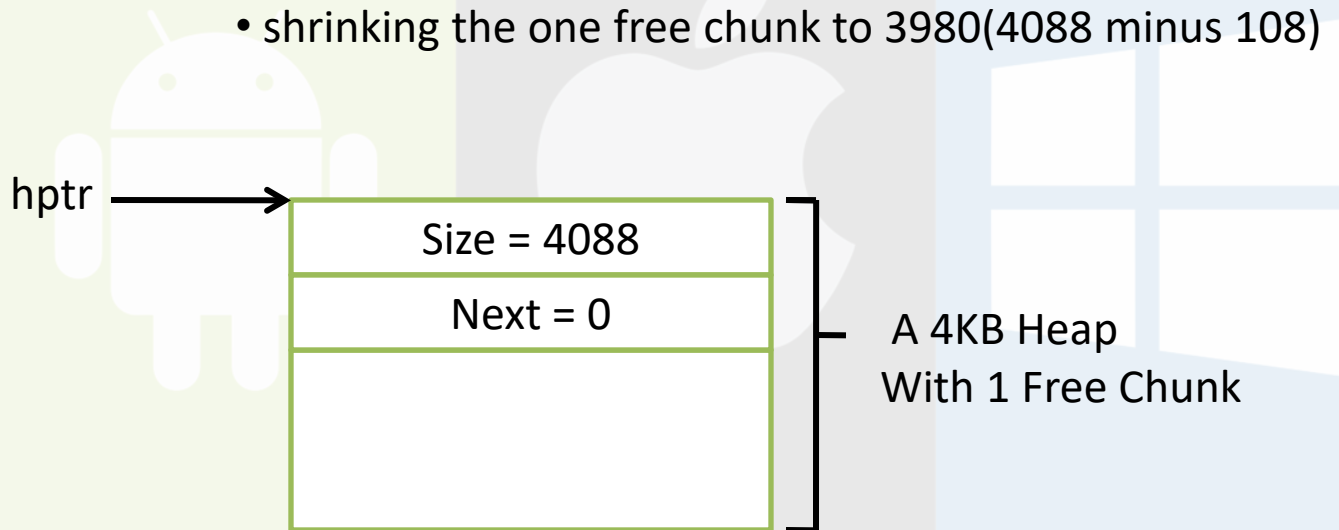
Embedding a Free List : Allocation

- **Example** a request for 100 bytes by `ptr = malloc(100)`
 - Allocating 108 bytes out of the existing one free chunk.
 - shrinking the one free chunk to 3980(4088 minus 108)



Embedding a Free List : Allocation

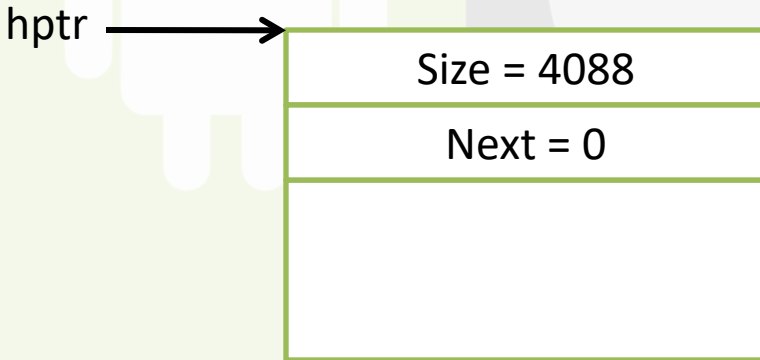
- **Example** a request for 100 bytes by `ptr = malloc(100)`
 - Allocating 108 bytes out of the existing one free chunk.
 - shrinking the one free chunk to 3980(4088 minus 108)



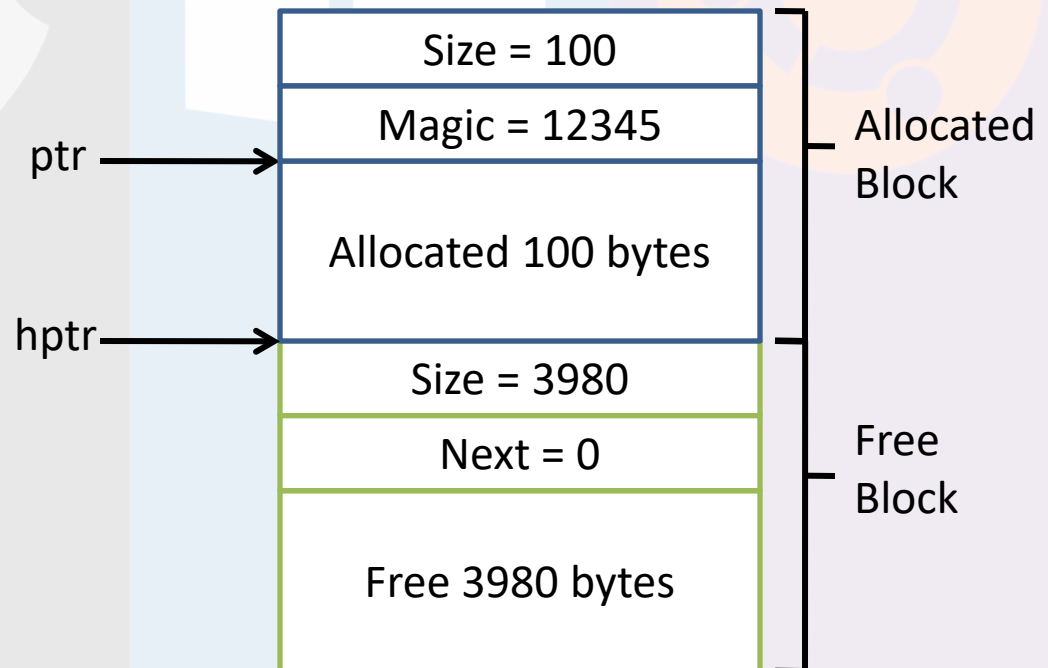
Embedding a Free List : Allocation

- **Example** a request for 100 bytes by `ptr = malloc(100)`
 - Allocating 108 bytes out of the existing one free chunk.
 - shrinking the one free chunk to 3980(4088 minus 108)

A 4KB Heap with 1 Free Chunk



A 4KB Heap after 1 Allocation



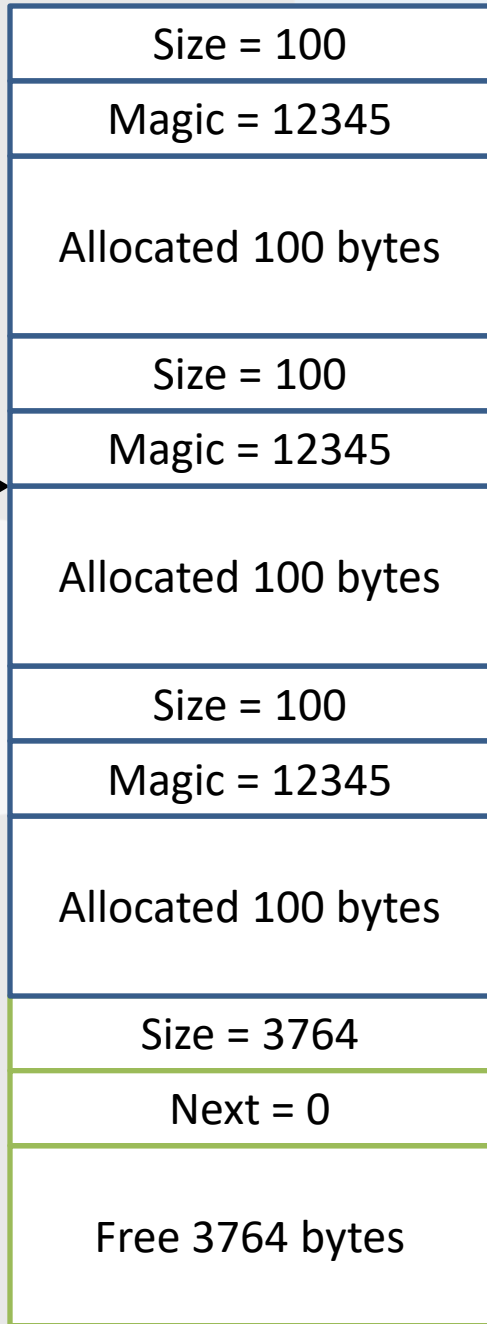
A 4KB Heap with 3 Allocated Chunks




sptr



hptr



100 allocated bytes
(but about to be freed)



Embedding a Free List : Free

- **Example** a request to free `sptr` : `free(sptr)`

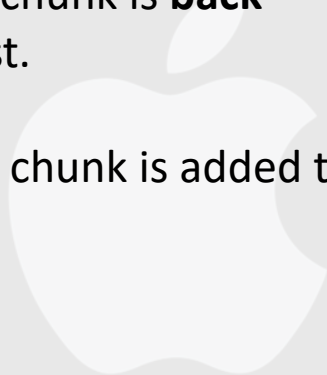



Embedding a Free List : Free

- **Example** a request to free `sptr` : `free(sptr)`

- The 100 byte chunk is **back into** the free list.

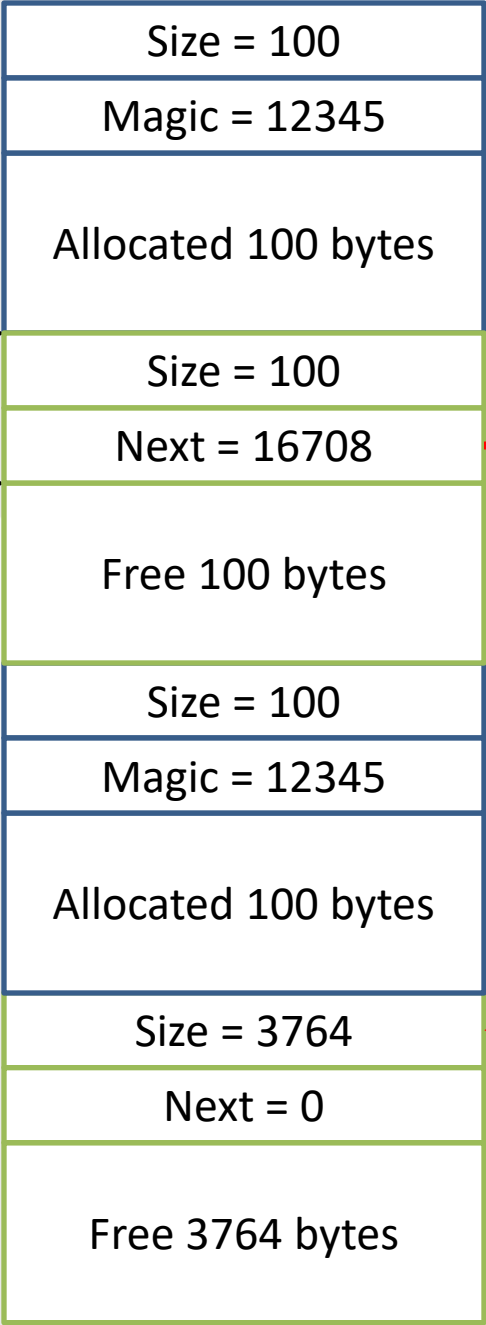
- The new Free chunk is added to the front of the list



A 4KB Heap with 2 Allocated
Chunks and 2 Free chunks



hptr →
sptr →



Embedding a Free List : Free

- **Example** Remaining chunks are freed


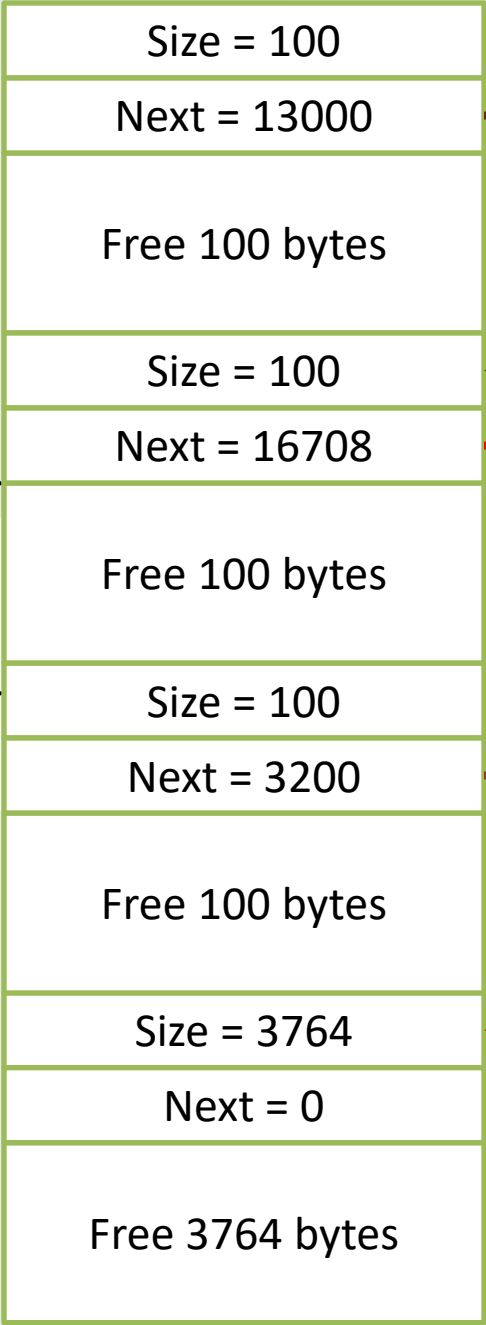


A 4KB Heap with 0 Allocated
Chunks and 4 Free chunks



sptr

hptr




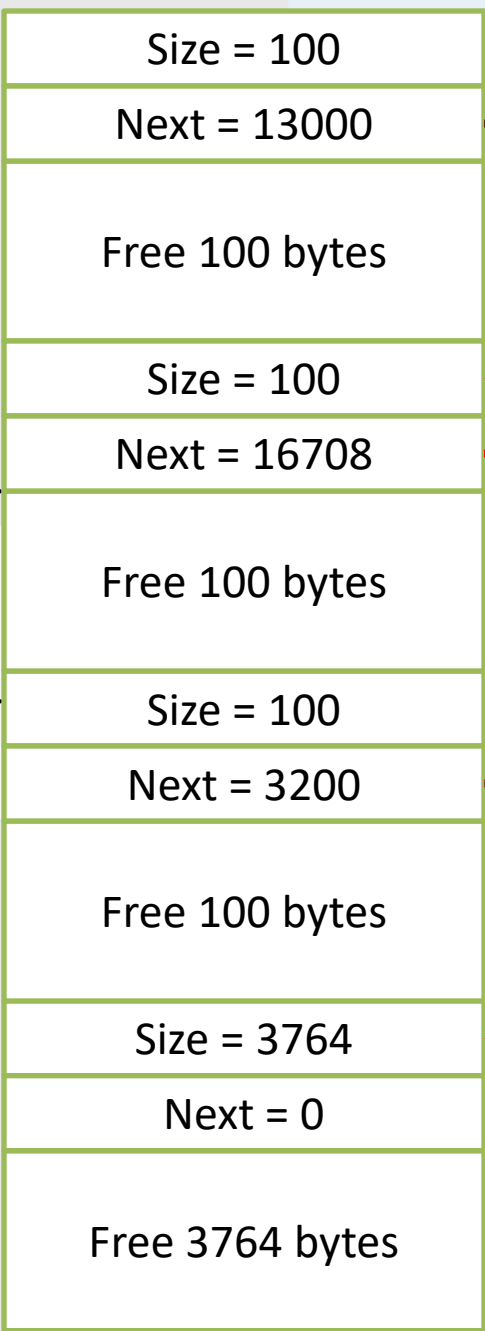
**A 4KB Heap with 0 Allocated
Chunks and 4 Free chunks**

**Fragmentation Occurs
Coalescing needed**



sptr →

hptr →



Managing Free Space : Basic Strategies

- **Best Fit:**

- Finding free chunks that are **big or bigger than the request**
- Returning the **one of smallest** chunks **in the group** of candidates



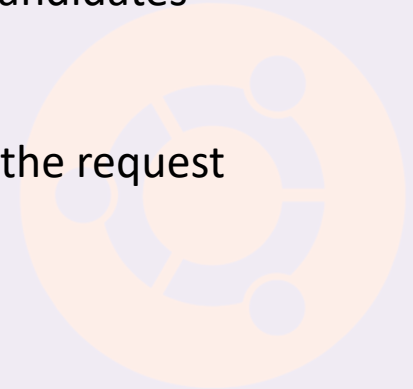

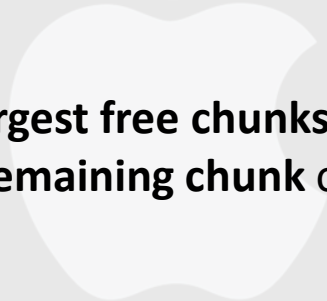

Managing Free Space : Basic Strategies

- **Best Fit:**

- Finding free chunks that are **equal or bigger than the request**
- Returning the **one of smallest** in the chunks **in the group** of candidates

- **Worst Fit:**


- Finding the **largest free chunks** and allocating the amount of the request
- **Keeping the remaining chunk** on the free list



Managing Free Space : Basic Strategies

- **First Fit:**

- Finding the **first chunk** that is **big enough** for the request
- Returning the requested amount and the remaining rest of the chunk




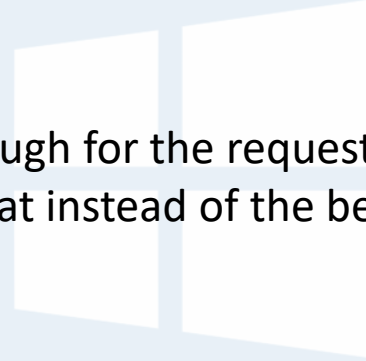
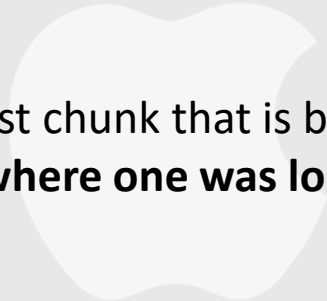

Managing Free Space : Basic Strategies

- **First Fit:**

- Finding the **first chunk** that is **big enough** for the request
- Returning the requested amount and the remaining rest of the chunk

- **Next Fit:**

- Finding the first chunk that is big enough for the request.
- Searching at **where one was looking** at instead of the beginning of the list



Basic Strategies - Examples

- Allocation Request Size 15





Free list: head →

Addr:0
Len:10

Addr:20
Len:30


Addr:60
Len:20

NULL




Basic Strategies - Examples

- Allocation Request Size 15




- Result of Best-fit




Basic Strategies - Examples


- Allocation Request Size 15



- Result of Best-fit



- Result of Worst-fit



Managing Free Space : Buddy Allocation


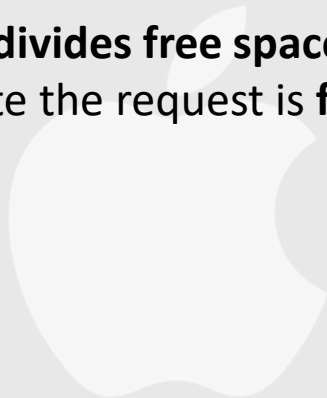

- Binary Buddy Allocation



Managing Free Space : Buddy Allocation

- Binary Buddy Allocation

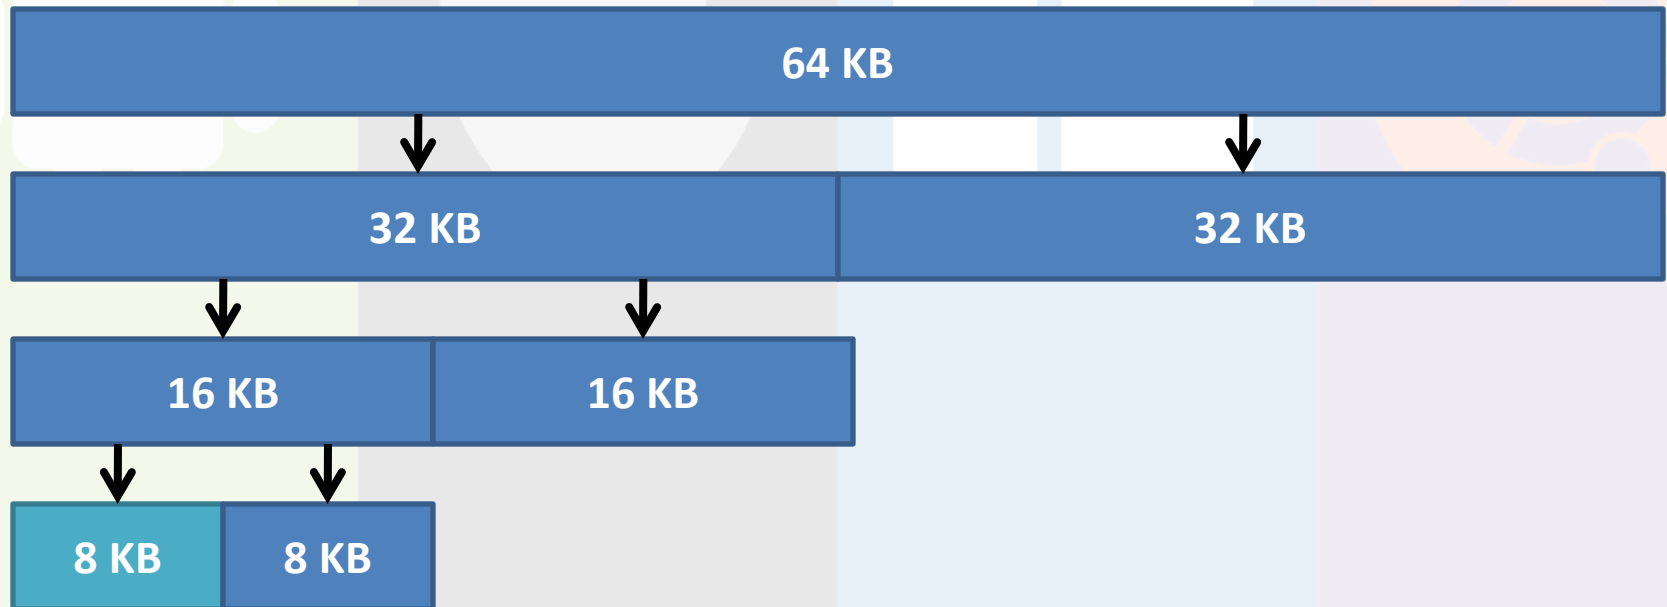
- The allocator **divides free space** by two **until a block** that is big enough to accommodate the request is **found**



Managing Free Space : Buddy Allocation

- Binary Buddy Allocation

- The allocator **divides free space** by two **until a block** that is big enough to accommodate the request is **found**



64KB free space for 7KB request