


Operating Systems

COMS(3010A)

Scheduling



Branden Ingram

branden.ingram@wits.ac.za


Recap

- **Swapping**





What is Scheduling

- Memory is oversubscribed and the memory demands of the set of running processes exceeds the available physical memory.
- Leads to a constant state of paging and page faults, inhibiting most application-level processing



What is Scheduling

- Memory is oversubscribed and the memory demands of the set of running processes exceeds the available physical memory.
- Leads to a constant state of paging and page faults, inhibiting most application-level processing
- So in order to remedy this we need to manage which processes we process and when



What is Scheduling

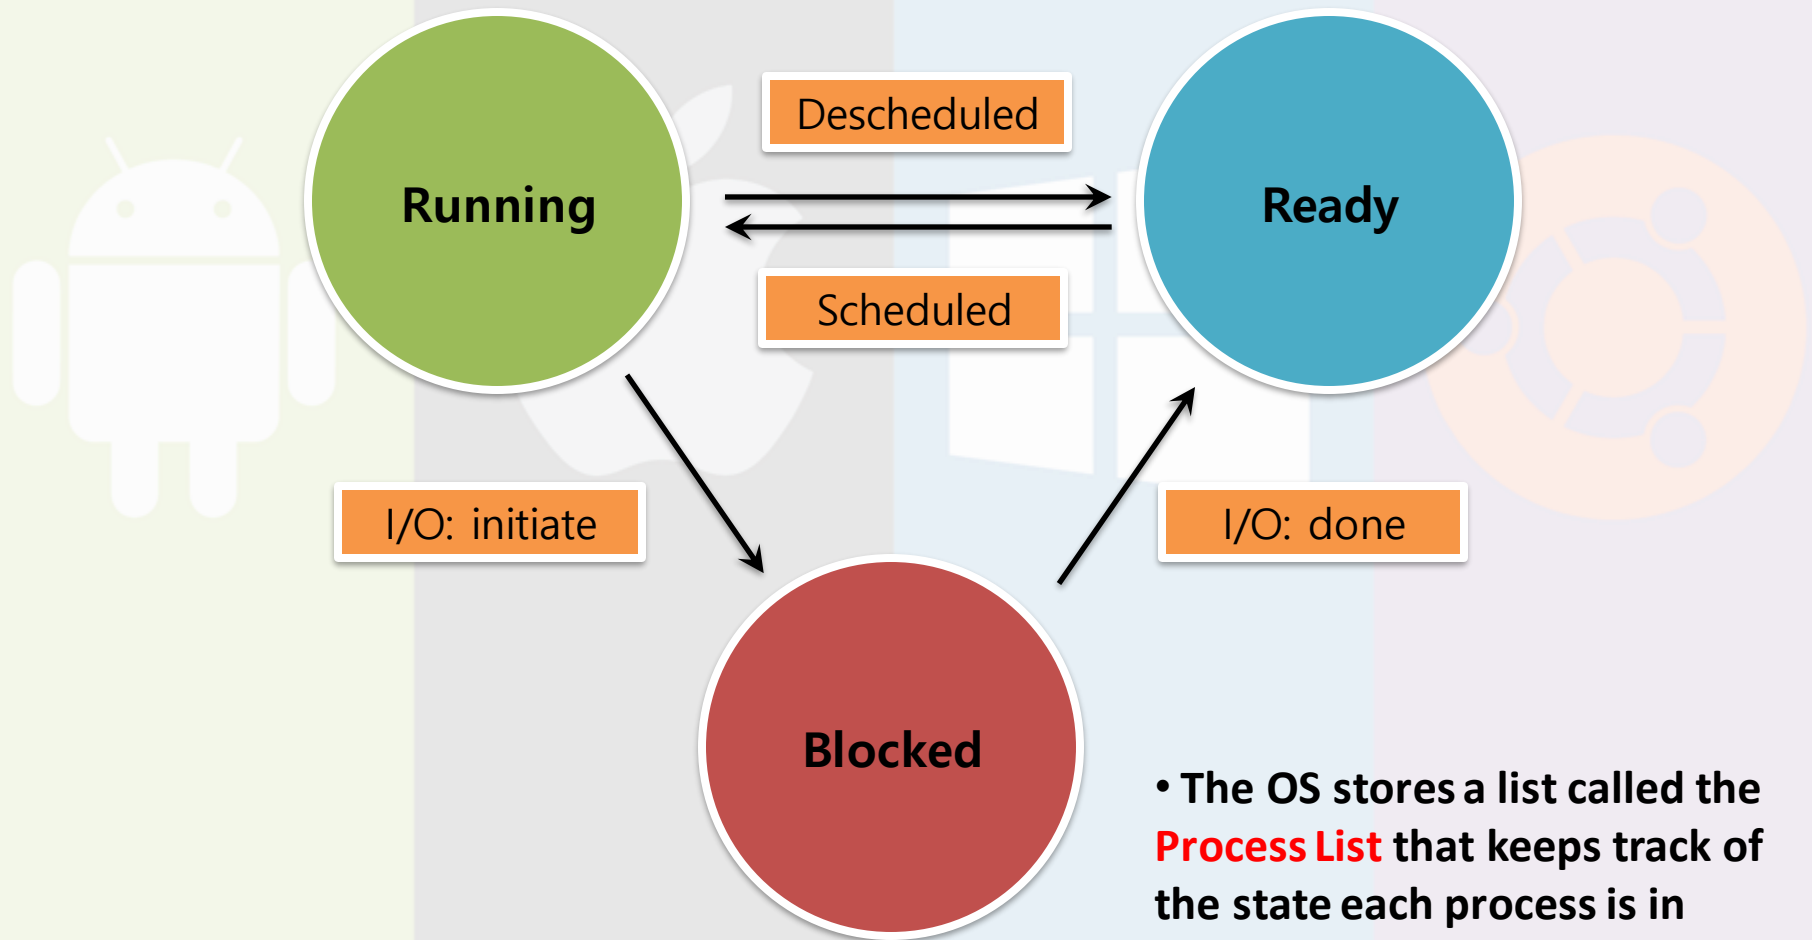
- Memory is oversubscribed and the memory demands of the set of running processes exceeds the available physical memory.
- Leads to a constant state of paging and page faults, inhibiting most application-level processing
- So in order to remedy this we need to manage which processes we process and when
- This role is performed by a **scheduler**

Scheduling

- **Process scheduling is the activity of the process manager that handles the removal of the running process from the CPU and the selection of another process on the basis of a particular strategy(scheme).**



Process State Transition



Metrics

- **Performance metric: Turnaround time**
 - The time at which the job completes minus the time at which the job arrived in the system.

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

First time it was ready NOT first time it was run!

Metrics

- Another metric is fairness – Response Time
 - Performance and fairness are often at odds in scheduling.


$$T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$$

Response and turnaround time are in contrast to each other, faster response = slower turn around time and vice versa


Use the metric to measure performance of different schemes

Schemes

- FIFO
- Shortest Job First (SJF)
- Round Robin
- Multi-Level Feedback Queue
- Proportional Share Scheduler



More modern




FIFO

- First Come, First Served (FCFS)
 - Very simple and easy to implement

- Example:

- A arrived just before B which arrived just before C.
 - Each job runs for 10 seconds.




Time (Second)

FIFO

- **First Come, First Served (FCFS)**
 - **Very simple and easy to implement**

- **Example:**

- A arrived just before B which arrived just before C. All arrive before $t=0$
- Each job runs for 10 seconds.




FIFO

- **First Come, First Served (FCFS)**
 - **Very simple and easy to implement**

- **Example:**

- A arrived just before B which arrived just before C. All arrive before $t=0$
- Each job runs for 10 seconds.




In a test he will ask which process will run when under a certain scheme

FIFO

- First Come, First Served (FCFS)
 - Very simple and easy to implement

- **Example:**

- A arrived just before B which arrived just before C. All arrive before t=0
- Each job runs for 10 seconds.




$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

FIFO - Weakness

Example:

A arrived just before B which arrived just before C. All arrive before $t=0$
A runs for 100 seconds, B and C run for 10 each.



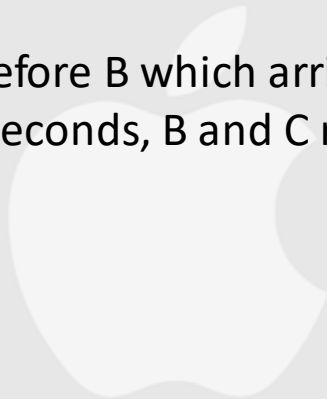

Convoy effect



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$


SJF - Shortest Job First

- Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- Example:
 - A arrived just before B which arrived just before C. All arrive before $t=0$
 - A runs for 100 seconds, B and C run for 10 each.




SJF - Shortest Job First

- Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- Example:
 - A arrived just before B which arrived just before C. All arrive before $t=0$
 - A runs for 100 seconds, B and C run for 10 each.




SJF - Shortest Job First

- Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- Example:
 - A arrived just before B which arrived just before C. All arrive before $t=0$
 - A runs for 100 seconds, B and C run for 10 each.



SJF - Shortest Job First

- Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- Example:
 - A arrived just before B which arrived just before C. All arrive before t=0
 - A runs for 100 seconds, B and C run for 10 each.




$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF - with Late Arrivals

- **Example:**

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds





$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$

Shortest Time-to-Completion First (STCF)

- **Add preemption to SJF**
 - Also known as **Preemptive Shortest Job First (PSJF)**
- **A new job enters the system:**
 - Determine of the remaining jobs and new job
 - Schedule the job which has the least time left


use a context switch to prevent late comers crunching the queue



Shortest Time-to-Completion First (STCF)

- **Example:**

- A arrives at $t=0$ and needs to run for 100 seconds.
- B and C arrive at $t=10$ and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

weakness: starvation - if we keep getting small tasks we will never process A

What's the catch

- STCF and related schemes are not particularly good for response time.

We want to prevent starvation (better response time)

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR)


- **Time slicing Scheduling**
 - **Run a job for a time slice and then switch to the next job in the run queue until the jobs are finished.**
 - Time slice is sometimes called a scheduling quantum.
- **It repeatedly does so until the jobs are finished.**
- **The length of a time slice must be a multiple of the timer-interrupt period**

RR is fair, but performs poorly on metrics such as turnaround time

We take much longer to complete a single process but we at least process everything

Round Robin (RR)

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.




SJF (Bad for Response Time)

This is a demo of SJF btw

$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5sec$$

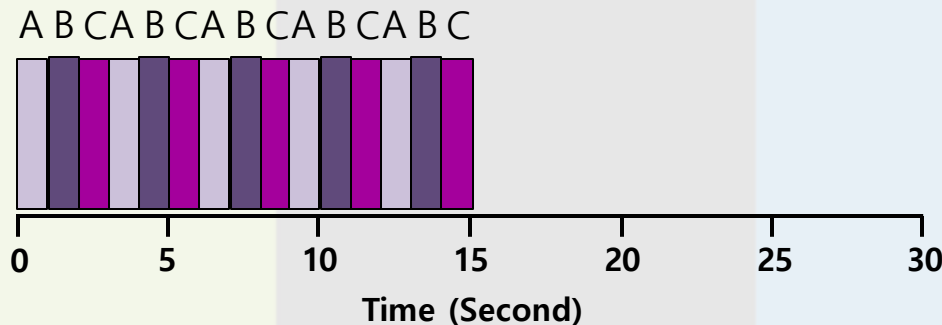
Round Robin (RR)

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



SJF (Bad for Response Time)

$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5sec$$



RR with a time-slice of 1sec (Good for Response Time)

$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1sec$$


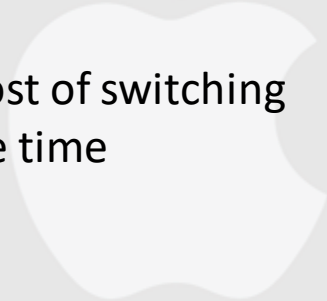

Time Slice Length is Critical

- **The shorter time slice**

- Better response time
- The cost of context switching will dominate overall performance.
above is really a big problem with RR, the cost is normally free cause its so fast but still needs to be considered

- **The longer time slice**



- Amortize the cost of switching
- Worse response time



Deciding on the length of the time slice presents
a **trade-off** to a system designer

Incorporating I/O




- When a job initiates an I/O request.
 - The job is blocked waiting for I/O completion.
 - The scheduler should schedule another job on the CPU.
- When the I/O completes
 - An interrupt is raised.
 - The OS moves the process from blocked back to the ready state.




Incorporating I/O

- **Example:**


- A and B need 50ms of CPU time each.
- A runs for 10ms and then issues an I/O request
- I/Os each take 10ms
- B simply uses the CPU for 50ms and performs no I/O
- The scheduler runs A first, then B after




Incorporating I/O



Incorporating I/O



Poor Use of Resources



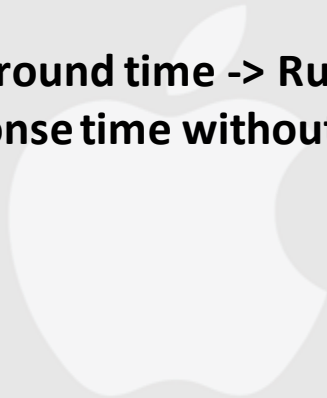



Overlap Allows Better Use of Resources

Maximize the
CPU utilization

Multi-Level Feedback Queue (MLFQ)

- A Scheduler that learns from the past to predict the future.
- Objective:
 - Optimize turnaround time -> Run shorter jobs first
 - Minimize response time without a prior knowledge of job length.



MLFQ – Basic Rules

- MLFQ has a number of distinct queues.
 - Each queue is assigned a different priority level.
- A job that is ready to run is on a single queue.
 - A job on a higher queue is chosen to run.
 - Use round-robin scheduling among jobs in the same queue

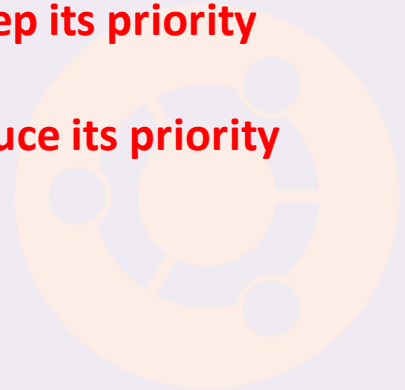



Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).

Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

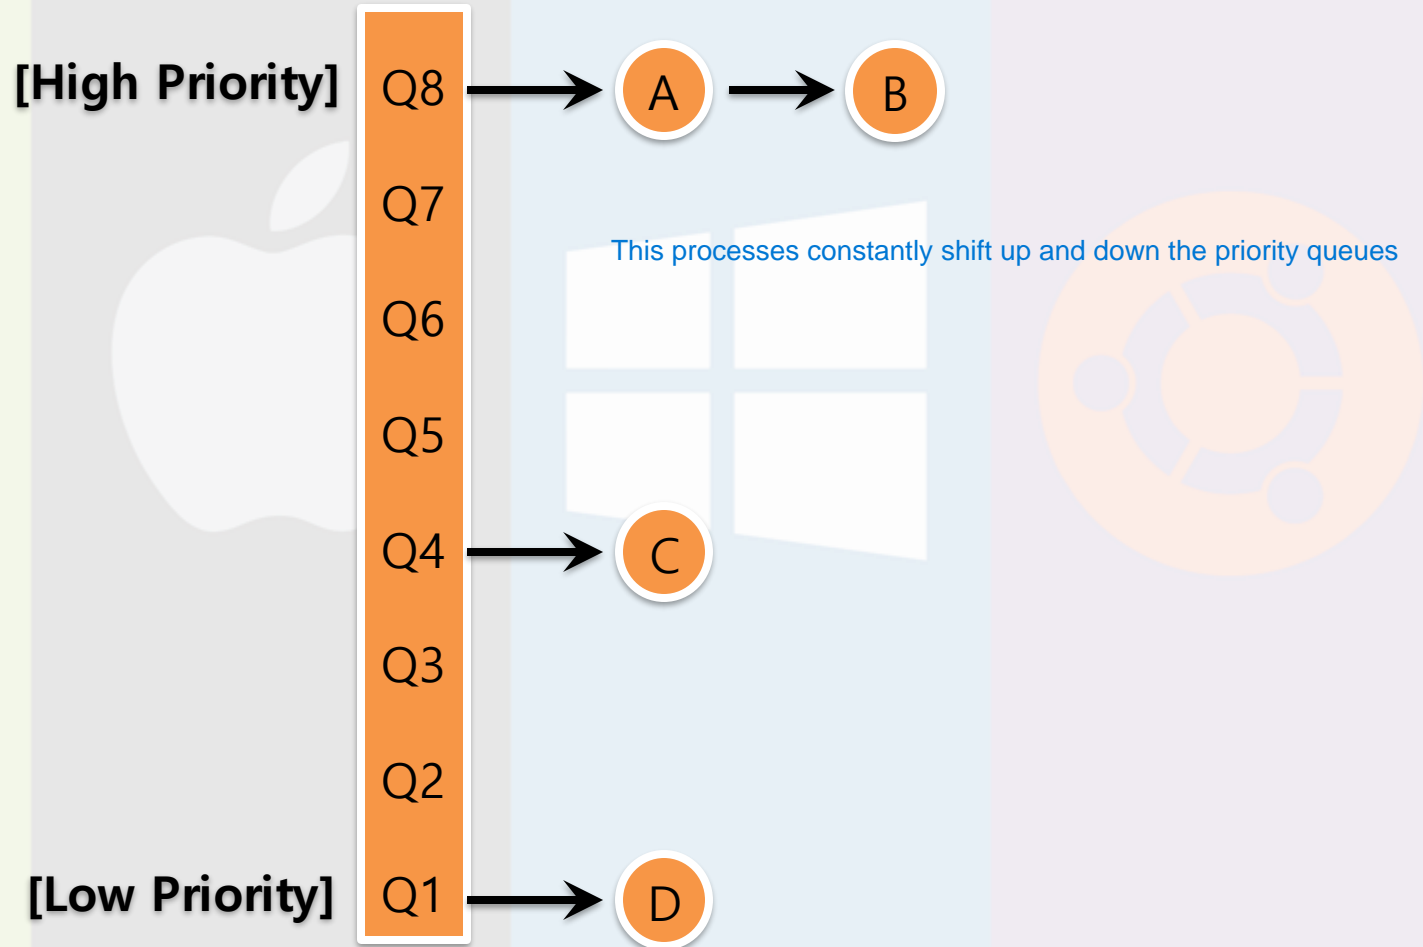
Same priorities just use RR

MLFQ – Basic Rules

- MLFQ varies the priority of a job based on its observed behaviour.
- Example:
 - A job repeatedly relinquishes the CPU while waiting IOs -> **Keep its priority high**
 - A job uses the CPU intensively for long periods of time -> **Reduce its priority**



MLFQ – Example



MLFQ – How Priority Changes


- **MLFQ priority adjustment algorithm:**

- **Rule 3:** When a job enters the system, it is placed at the highest priority
- **Rule 4a:** If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).
- **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level

In this manner, MLFQ approximates SJF

Example 1: A Single Long-Running Job

- A three-queue scheduler with time slice 10ms



Example 2: Along Came a Short Job


- **Assumption:**

- **Job A:** A long-running CPU-intensive job
- **Job B:** A short-running interactive job (20ms runtime)

- **A has been running for some time, and then B arrives at time T=100.**

jobs that run super fast end up as the top priority, whereas processes with a much longer run time would take a lower priority.
All processes start at the highest priority

at each time-slice we work down the priority queue



Along Came An Interactive Job (msec)



Example 3: What About I/O?

- **Assumption:**

- **Job A:** A long-running CPU-intensive job

- **Job B:** An interactive job that need the CPU only for 1ms before performing an I/O

since B never uses its full time slice for each of its operations it never moves down.
"It willingly gives up the CPU - we want this"



A: 
B: 

A Mixed I/O-intensive and CPU-intensive Workload (msec)

what if a process just gets stuck at the bottom and gets no time? We use a priority boost

Problems with the Basic MLFQ

- **Starvation**

- If there are “too many” interactive jobs in the system.
- Long-running jobs will never receive any CPU time.

- **Game the scheduler**

- After running 99% of a time slice, issue an I/O operation.
- The job gain a higher percentage of CPU time.

A process can out smart the scheduler by taking a higher priority by faking I/O operations to keep priority

- **A program may change its behaviour over time.**



- CPU bound process ->I/O bound process

The Priority Boost


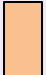

- Rule 5: After some time period S , move all the jobs in the system to the topmost queue.

- Example:

- A long-running job(A) with two short-running interactive job(B, C)



Without(Left) and With(Right) Priority Boost

A:  B:  C: 

Better Accounting



- How to prevent gaming of our scheduler?

- Solution:

- **Rule 4 (Rewrite Rules 4a and 4b):** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down on queue).

the percentage of time-slice you use is accumulative.

so if you use 1% of the time slice each time after 100 runs that operation will be moved down the queue




Without(Left) and With(Right) Gaming Tolerance

A:  B: 

Tuning MLFQ And Other Issues

- **The high-priority queues -> Short time slices**
 - E.g., 10 or fewer milliseconds
- **The Low-priority queue -> Longer time slices**
 - E.g., 100 milliseconds

Lower Priority, Longer Quanta


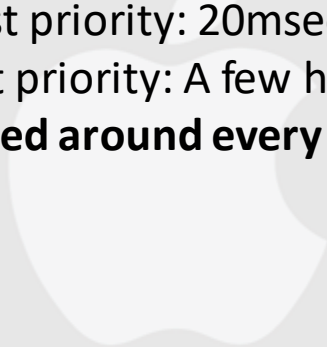



Example) 10ms for the highest queue, 20ms for the middle, 40ms for the lowest

A:  B: 

The Solaris MLFQ implementation

- For the Time-Sharing scheduling class (TS)
 - 60 Queues
 - Slowly increasing time-slice length
 - The highest priority: 20msec
 - The lowest priority: A few hundred milliseconds
 - Priorities boosted around every 1 second or so.



The Solaris MLFQ implementation

- **The refined set of MLFQ rules:**

- **Rule 1:** If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
- **Rule 2:** If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.
- **Rule 3:** When a job enters the system, it is placed at the highest priority.
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced(i.e., it moves down on queue).
- **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue.

YOU NEED TO MEMORISE THESE RULES NB!!!!


Proportional Share Scheduler

- Fair-share scheduler
 - Guarantee that each job obtain a certain percentage of CPU time.
 - Not optimized for turnaround or response time



Proportional Share Scheduler

- Fair-share scheduler
 - Guarantee that each job obtain a certain percentage of CPU time.
 - Not optimized for turnaround or response time



- Lottery Scheduling



Proportional Share Scheduler

- **Tickets**

- Represent the share of a resource that a process should receive
- The percent of tickets represents its share of the system resource in question.



Basic Concept


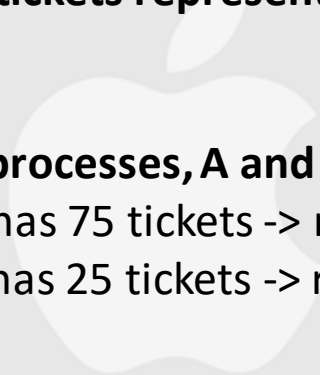

•Tickets

- Represent the share of a resource that a process should receive
- The percent of tickets represents its share of the system resource in question.

•Example

- There are two processes, A and B.
 - Process A has 75 tickets -> receive 75% of the CPU
 - Process B has 25 tickets -> receive 25% of the CPU

tickets issued according to a specific process



Lottery scheduling

- The scheduler picks a winning ticket.
- Load the state of that winning process and runs it.





Lottery scheduling

- The scheduler picks a winning ticket.
 - Load the state of that winning process and runs it.

- **Example**

- There are 100 tickets
- Process A has 75 tickets: 0 ~ 74
- Process B has 25 tickets: 75 ~ 99



Lottery scheduling

- The scheduler picks a winning ticket.
 - Load the state of that winning process and runs it.

- **Example**

- There are 100 tickets
- Process A has 75 tickets: 0 ~ 74
- Process B has 25 tickets: 75 ~ 99

Scheduler's winning tickets: 63 85 70 39 76 17 29 41 36 39 10 99 68 83 63

Resulting scheduler: A B A A B A A A A A A B A B A

The longer these two jobs compete,
The more likely they are to achieve the desired percentages.

Ticket Mechanisms

- **Ticket currency**

- A user allocates tickets among their own jobs in whatever currency they would like.
- The system converts the currency into the correct global value.

- **Example**

- There are 200 tickets (Global currency)
- Process A has 100 tickets
- Process B has 100 tickets

User A → 500 (A's currency) to A1 → 50 (global currency)
→ 500 (A's currency) to A2 → 50 (global currency)

User B → 10 (B's currency) to B1 → 100 (global currency)

Ticket Mechanisms

- **Example**



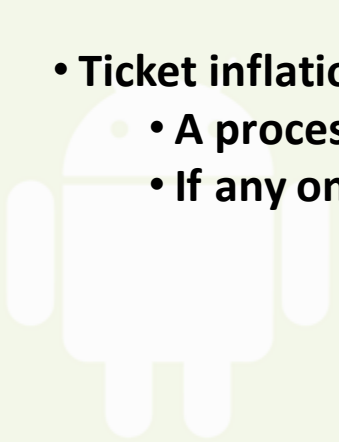
- There are 200 tickets (Global currency)
- Process A has 100 tickets
- Process B has 100 tickets

User A → 500 (A's currency) to A1 → 67 (global currency)
→ 250 (A's currency) to A2 → 33 (global currency)

User B → 10 (B's currency) to B1 → 100 (global currency)


Ticket Mechanisms

- **Ticket transfer**
 - A process can temporarily hand off its tickets to another process.
- **Ticket inflation**
 - A process can temporarily raise or lower the number of tickets it owns.
 - If any one process needs more CPU time, it can boost its tickets.



Ticket Mechanisms - Implementation

- Example: There are three processes, A, B, and C.
 - Keep the processes in a list:





```
1  // counter: used to track if we've found the winner yet
2  int counter = 0;
3
4  // winner: use some call to a random number generator to
5  // get a value, between 0 and the total # of tickets
6  int winner = getRandom(0, totaltickets);
7
8  // current: use this to walk through the list of jobs
9  node_t *current = head;
10
11 // loop until the sum of ticket values is > the winner
12 while (current) {
13     counter = counter + current->tickets;
14     if (counter > winner)
15         break; // found the winner
16     current = current->next;
17 }
18 // 'current' is the winner: schedule it...
```

we just check the bounds of each node, if the ticket is greater than that bound we got it


Ticket Mechanisms - Implementation

- **U** : unfairness metric
 - The time the first job completes divided by the time that the second job completes
- Example:
 - There are two jobs, each jobs has runtime 10
 - First job finishes at time 10
 - Second job finishes at time 20
 - **U** = $10 \div 20 = 0.5$
 - **U** will be close to 1 when both jobs finish at nearly the same time



Lottery Fairness Study

- There are two jobs.
- Each jobs has the same number of tickets (100).



When the job length is not very long, average unfairness can be **quite severe**.

Stride Scheduling

- Stride of each process
 - (A large number) / (the number of tickets of the process)
 - Example: A large number = 10,000
 - Process A has 100 tickets -> stride of A is 100
 - Process B has 50 tickets -> stride of B is 200
- A process runs, increment a counter(=pass value) for it by its stride.
- Pick the process to run that has the lowest pass value

bigger stride = lower priority

```
current = remove_min(queue);           // pick client with minimum pass
schedule(current);                     // use resource for quantum
current->pass += current->stride;       // compute next pass using stride
insert(queue, current);                // put back into the queue
```

A pseudo code implementation

Stride Scheduling - Example

Pass(A) (stride=100)	Pass(B) (stride=200)	Pass(C) (stride=40)	Who Runs?
0	0	0	A
100	0	0	B
100	200	0	C
100	200	40	C
100	200	80	C
100	200	120	C
200	200	120	A
200	200	160	C
200	200	200	C
			...

If new job enters with pass value 0,
It will **monopolize** the CPU!


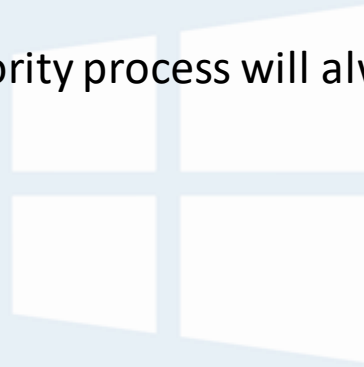


Windows Scheduling

- **Windows XP** scheduled processes using a **priority-based pre-emptive** scheduler with a flexible system of priority levels that includes **round robin** scheduling within each level





Windows Scheduling

- **Windows XP** scheduled processes using a **priority-based pre-emptive** scheduler with a flexible system of priority levels that includes **round robin** scheduling within each level
- The scheduler ensured that the **highest** priority process will always run




Windows Scheduling

- A process selected to run will run until:
 - **Pre-empted by a higher-priority process**
 - **Terminated**
 - **Time quantum ends**
 - **Calls a blocking system call such as I/O**




Windows Scheduling

- The scheduler uses a 32-level priority scheme to determine the order of execution
- **Priorities are divided into two classes:**
 - **Variable(Dynamic) Class : priorities 1-15**
 - **Real-time Class : priorities 16-31**
 - **0 is a special priorities for memory management**



Windows Scheduling

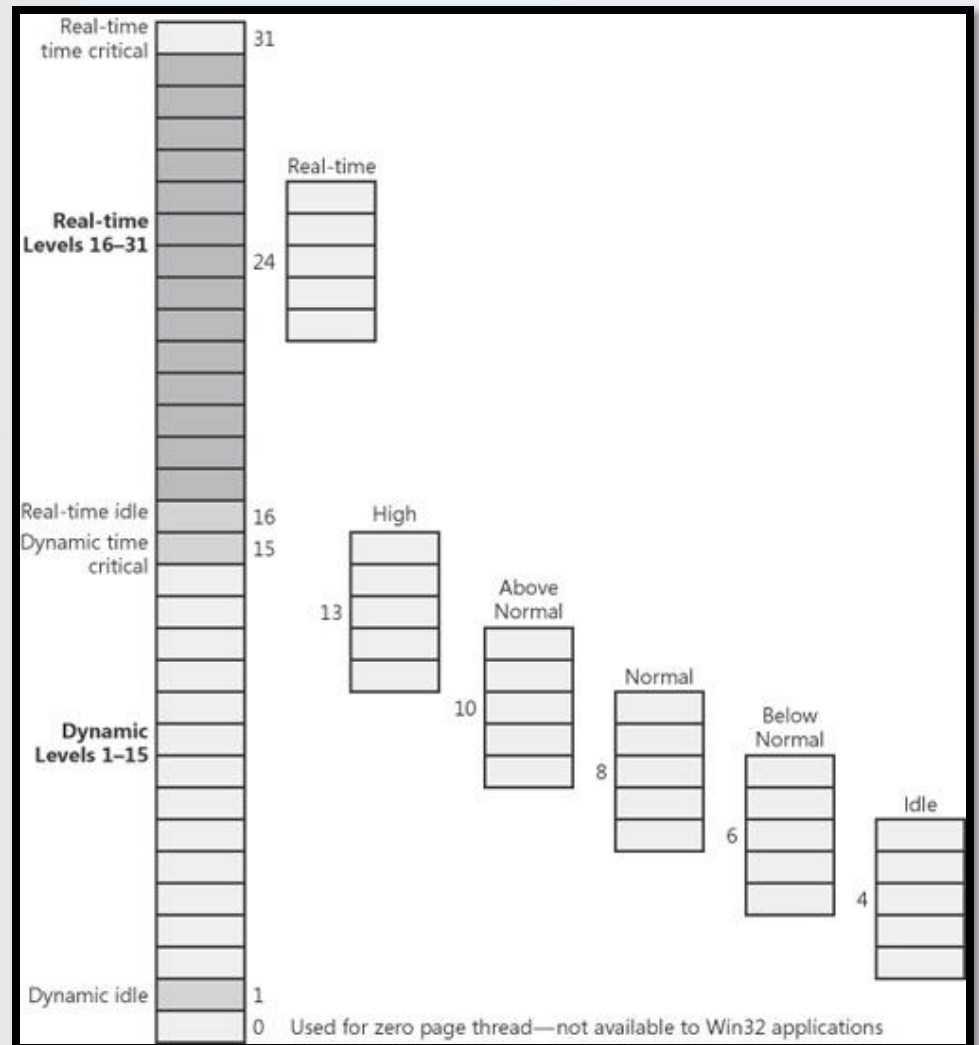
- Priorities in all the Dynamic classes are variable
- This means priorities can change
- So instead priorities here are defined by classes



Windows Scheduling


- Priorities in all the Dynamic classes are variable
- This means priorities can change
- So instead priorities here are defined by classes

- **Real-time**
- **High**
- **Above Normal**
- **Normal**
- **Below Normal**
- **Idle**



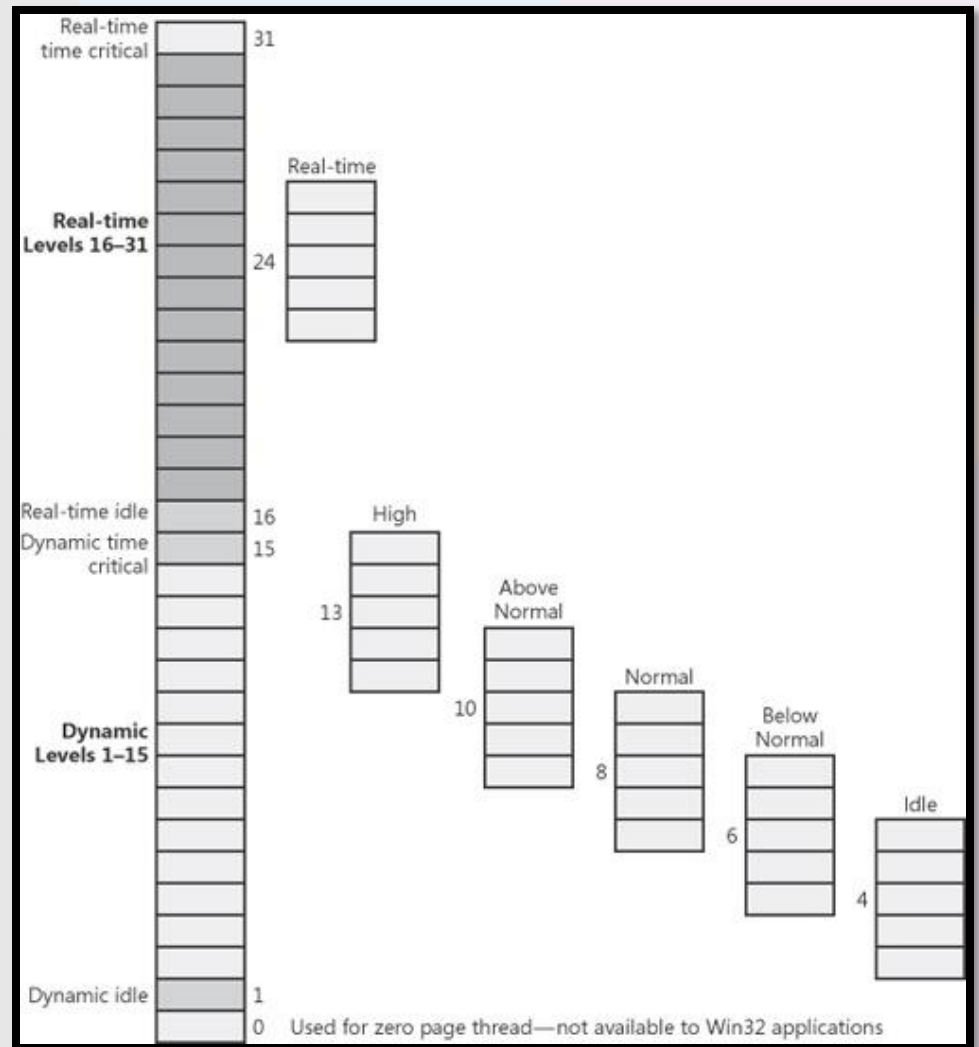
Windows Scheduling

- Priorities in all the Dynamic classes are variable
- This means priorities can change
- So instead priorities here are defined by classes
 - **Real-time**
 - **High**
 - **Above Normal**
 - **Normal**
 - **Below Normal**
 - **Idle**
- Therefore on creation a process is assigned a class and a **base priority** in that class.



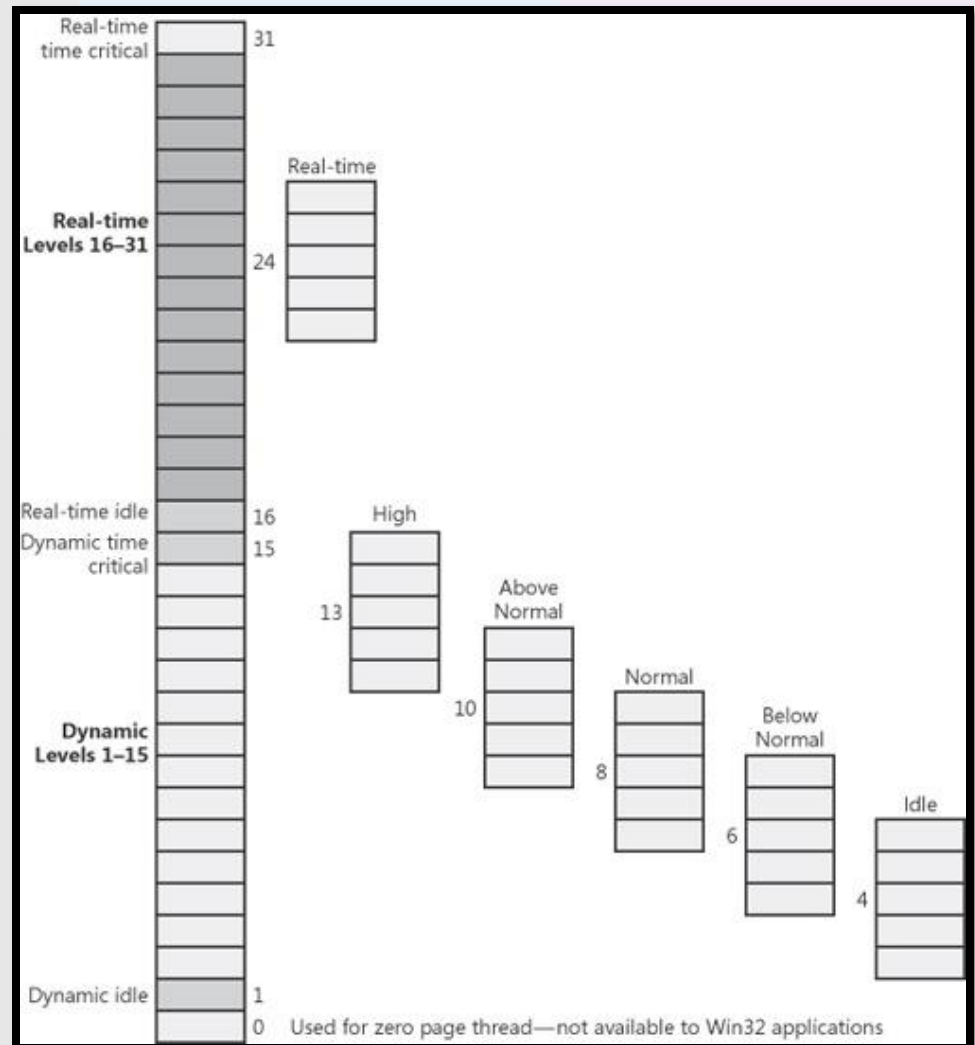
Windows Scheduling

- Processes are also each given a **base priority** within their priority class.
- When variable class processes **consume their entire time quanta**, then their **priority gets lowered**, but **not below their base priority**.



Windows Scheduling

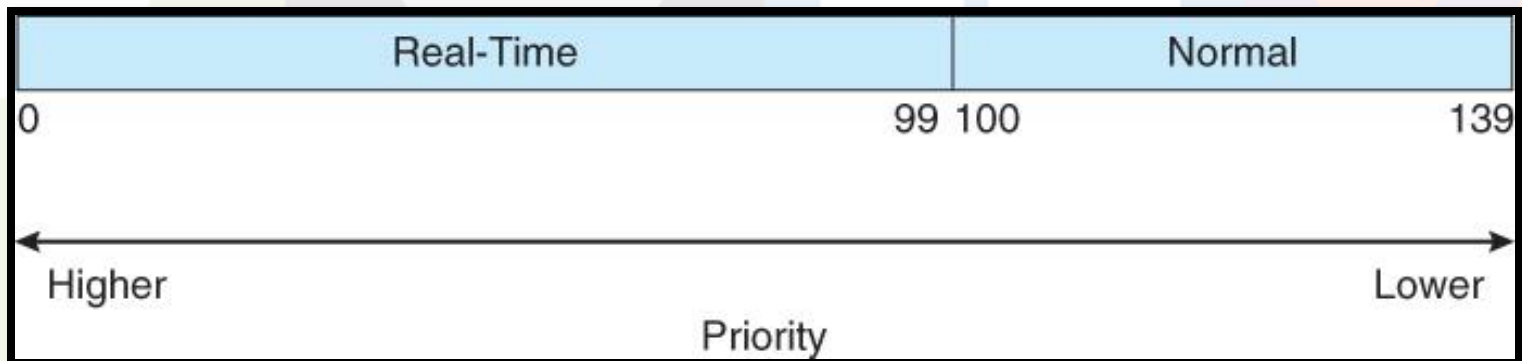
- Processes are also each given a **base priority** within their priority class.
- When variable class processes **consume their entire time quanta**, then their **priority gets lowered**, but **not below their base priority**.
- Processes in the **foreground** (active window) have their scheduling **quanta multiplied by 3**, to give **better response** to interactive processes in the foreground.



Linux Scheduling


- The Linux scheduler is a **preemptive priority-based** algorithm with two priority ranges

- **Real time** from 0 to 99
- **Normal** range from 100 to 140.




Linux Scheduling

- The Linux CFS (Completely Fair scheduler) provides an efficient algorithm for selecting which task to run next.
- Each runnable task is placed in a red-black tree—a balanced binary search tree whose key is based on the value of `vruntime`.



Linux Scheduling

- When a task becomes runnable, it is added to the tree.
- If a task on the tree is not runnable (for example, if it is blocked while waiting for I/O), it is removed



Linux Scheduling

- When a task becomes runnable, it is added to the tree.
- If a task on the tree is not runnable (for example, if it is blocked while waiting for I/O), it is removed
- **The scheduler runs the task with the lowest vruntime, the leftmost node**
- **vruntime is calculated with respect to priority and time spent processing**

