

Operating Systems

COMS(3010A)

Condition variables

and

Semaphores




Branden Ingram

branden.ingram@wits.ac.za

Office Number : ???

Recap

- Locks
- Synchronization Challenges
- Test and Set
- LLSC
- Ticket Locks
- Using Queues



Condition variable

- There are many cases where a thread wishes to check whether a condition is true before continuing its execution.



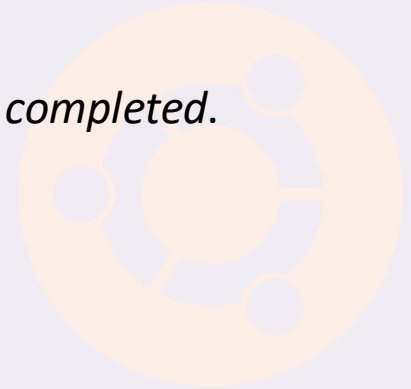
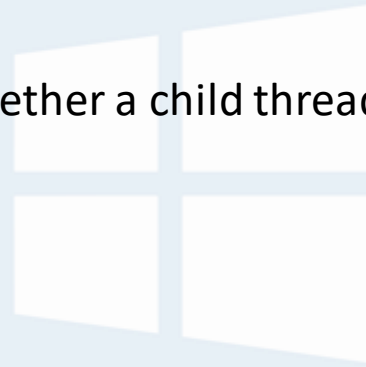

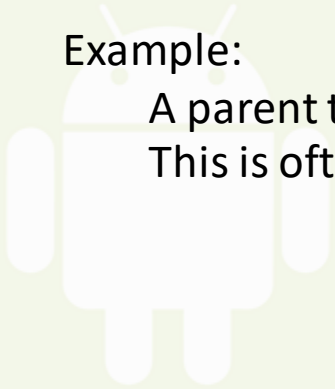
Condition variable

- There are many cases where a thread wishes to check whether a condition is true before continuing its execution.

Example:

A parent thread might wish to check whether a child thread has *completed*.

This is often called a `join()`.



Condition variable

- There are many cases where a thread wishes to check whether a condition is true before continuing its execution.

Example:

A parent thread might wish to check whether a child thread has *completed*. This is often called a `join()`.

- A Condition variable is a synchronization object that lets a thread efficiently wait for a change to shared state that is protected by a lock

Condition variable

A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3          // XXX how to indicate we are done?
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11         // XXX how to wait for child?
12         printf("parent: end\n");
13         return 0;
14     }
```

What we would like to see here is:

```
parent: begin
child
parent: end
```

Spin-based Approach

```
1      volatile int done = 0;
2      ^ volatile indicates to compiler its going to be used in multi-threaded systems
3      void *child(void *arg) {
4          printf("child\n");
5          done = 1;
6          return NULL;
7      }
8
9      int main(int argc, char *argv[]) {
10         printf("parent: begin\n");
11         pthread_t c;
12         Pthread_create(&c, NULL, child, NULL); // create child
13         while (done == 0)
14             ; // spin
15         printf("parent: end\n");
16         return 0;
17     }
```

- This is hugely inefficient as the parent spins and wastes CPU time

How do condition variables work

- A condition variable has two states

Waiting on the condition

An explicit queue that threads can put themselves on when some state of execution is not as desired.

Signaling on the condition

Some other thread, *when it changes said state*, can wake one of those waiting threads and allow them to continue.

How do condition variables work

- A condition variable has three methods

Function which puts itself to sleep

- **CV::wait(Lock *lock)** : This automatically releases the lock and suspends execution of the calling thread, placing it on the condition variables waiting list until the thread is re-enabled
- **CV::signal()** : This takes one thread off the condition variables waiting list and marks it eligible to run (puts it on the scheduler's ready list)
- **CV::broadcast()** : Takes all threads off the condition variables waiting list and marks them as eligible to run

How do condition variables work

- **Declare condition variable**

```
Pthread_cond_t c;
```

- Proper initialization is required.

- **Operation (the POSIX calls)**

Pass lock into wait so we release lock via wait

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);    // wait()  
pthread_cond_signal(pthread_cond_t *c);                      // signal()
```

- **The wait() call takes a mutex as a parameter.**

- The wait() call release the lock and put the calling thread to sleep.
- When the thread wakes up, it must re-acquire the lock.

Parent waiting for Child: Use a condition variable

```
1      int done = 0;
2      pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3      pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5      void thr_exit() {
6          Pthread_mutex_lock(&m);
7          done = 1;
8          Pthread_cond_signal(&c);
9          Pthread_mutex_unlock(&m);
10     }
11
12     void *child(void *arg) {
13         printf("child\n");
14         thr_exit();
15         return NULL;
16     }
17
18     void thr_join() {
19         Pthread_mutex_lock(&m);
20         while (done == 0)           we give up the lock here
21             Pthread_cond_wait(&c, &m);
22         Pthread_mutex_unlock(&m);
23     }
24
```

Parent waiting for Child: Use a condition variable

(cont.)

```
25     int main(int argc, char *argv[]) {
26         printf("parent: begin\n");
27         pthread_t p;
28         Pthread_create(&p, NULL, child, NULL);
29         thr_join();
30         printf("parent: end\n");
31         return 0;
32     }
```

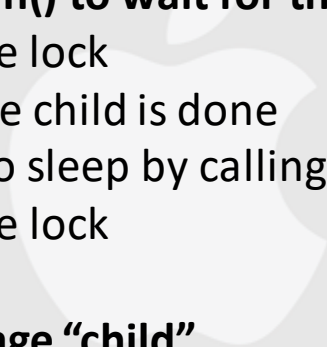

Parent waiting for Child: Use a condition variable

- **Parent:**

- **Create the child thread and continues running itself.**
- **Call into `thr_join()` to wait for the child thread to complete.**
 - Acquire the lock
 - Check if the child is done
 - Put itself to sleep by calling `wait()`
 - Release the lock

- **Child:**

- **Print the message “child”**
- **Call `thr_exit()` to wake the parent thread**
 - Grab the lock
 - Set the state variable done
 - Signal the parent thus waking it.



The importance of the state variable “done”

```
1      void thr_exit() {  
2          Pthread_mutex_lock(&m);  
3          Pthread_cond_signal(&c);  
4          Pthread_mutex_unlock(&m);  
5      }  
6  
7      void thr_join() {  
8          Pthread_mutex_lock(&m);  
9          Pthread_cond_wait(&c, &m);  
10         Pthread_mutex_unlock(&m);  
11     }
```

thr_exit() and thr_join() without variable done

- **Imagine the case where the child runs immediately.**
 - The child will signal, but there is no thread asleep on the condition.
 - When the parent runs, it will call wait and be stuck.
 - No thread will ever wake it.

Another poor implementation

```
1      void thr_exit() {
2          done = 1;
3          Pthread_cond_signal(&c);
4      }
5
6      void thr_join() {
7          if (done == 0)
8              Pthread_cond_wait(&c);
9      }
```

- **The issue here is a subtle race condition.**

- The parent calls thr_join().
 - The parent checks the value of done.
 - It will see that it is 0 and try to go to sleep.
 - Just before it calls wait to go to sleep, the parent is interrupted and the child runs.
- The child changes the state variable done to 1 and signals.
 - But no thread is waiting and thus no thread is woken.
 - When the parent runs again, it sleeps forever.


Another poor implementation

```
1      void thr_exit() {
2          done = 1;
3          Pthread_cond_signal(&c);
4      }
5
6      void thr_join() {
7          if (done == 0)
8              Pthread_cond_wait(&c);
9      }
```

- **The issue here is a subtle race condition.**

- The parent calls thr_join().
 - The parent checks the value of done.
 - It will see that it is 0 and try to go to sleep.
 - Just before it calls wait to go to sleep, the parent is interrupted and the child runs.
- The child changes the state variable done to 1 and signals.
 - But no thread is waiting and thus no thread is woken.
 - When the parent runs again, it sleeps forever.

The Producer / Consumer (Bound Buffer) Problem



- **Producer**

- Produce data items
- Wish to place data items in a buffer

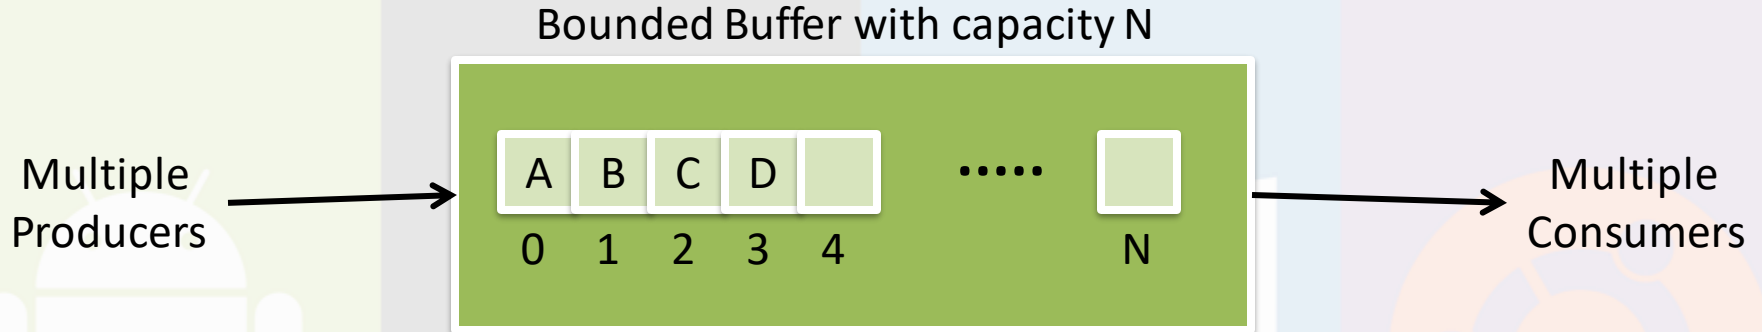
- **Consumer**

- Grab data items out of the buffer consume them in some way

- **Example: Multi-threaded web server**


- A producer puts HTTP requests in to a work queue
- Consumer threads take requests out of this queue and process them

The Producer / Consumer (Bound Buffer) Problem



- A bounded buffer is used when you pipe the output of one program into another.
- Example: `grep foo file.txt | wc -l`
 - The `grep` process is the producer.
 - The `wc` process is the consumer.
 - Between them is an in-kernel bounded buffer.
- Bounded buffer is Shared resource -> Synchronized access is required.

The Producer / Consumer (Bound Buffer) Problem



- **Without proper synchronization the following errors may occur**
 - The producers doesn't block when the buffer is full.
 - A Consumer consumes an empty slot in the buffer.
 - A consumer attempts to consume a slot that is only half-filled by a producer.
 - Two producers writes into the same slot.
 - Two consumers reads the same slot.

Simple Example

```
1      int buffer;  
2      int count = 0;    // initially, empty  
3  
4      void put(int value) {  
5          assert(count == 0);  
6          count = 1;  
7          buffer = value;  
8      }  
9  
10     int get() {  
11         assert(count == 1);  
12         count = 0;  
13         return buffer;  
14     }
```

- Only put data into the buffer when count is zero.
 - i.e., when the buffer is empty.
- Only get data from the buffer when count is one.
 - i.e., when the buffer is full.

Simple Example

```
1      void *producer(void *arg) {
2          int i;
3          int loops = (int) arg;
4          for (i = 0; i < loops; i++) {
5              put(i);
6          }
7      }
8
9      void *consumer(void *arg) {
10         int i;
11         while (1) {
12             int tmp = get();
13             printf("%d\n", tmp);
14         }
15     }
```

Producer puts an integer into the shared buffer loops number of times.
Consumer gets the data out of that shared buffer.

Simple Example (broken solution)

A single condition variable cond and associated lock mutex **with IF**

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              if (count == 1) if buffer is full, wait // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                                // p4
11             Pthread_cond_signal(&cond);            // p5
12             Pthread_mutex_unlock(&mutex);          // p6
13         }
14     }
15
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);            // c1
```

Simple Example (broken solution)

```
20      if (count == 0)    if buffer empty wait till its full           // c2
21          Pthread_cond_wait(&cond, &mutex);                          // c3
22      int tmp = get();                                           // c4
23      Pthread_cond_signal(&cond);                                  // c5
24      Pthread_mutex_unlock(&mutex);                               // c6
25      printf("%d\n", tmp);
26  }
27 }
```

With just a single producer and a single consumer, the code works.

If we have **more than** one of producer and consumer?

Simple Example (broken solution-Trace)

T _{c1}	State	T _{c2}	State	T _p	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	Tc1 awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	Tc2 sneaks in
	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	Tp awoken
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

Simple Example (broken solution 2)

A single condition variable `cond` and associated lock `mutex` **with While**

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);
8              while (count == 1)
9                  Pthread_cond_wait(&cond, &mutex);
10             put(i);
11             Pthread_cond_signal(&cond);
12             Pthread_mutex_unlock(&mutex);
13         }
14     }
15
```

// p1
// p2
// p3
// p4
// p5
// p6

Simple Example (broken solution 2)

A single condition variable `cond` and associated lock `mutex` **with While**

(Cont.)

```
16     void *consumer(void *arg) {  
17         int i;  
18         for (i = 0; i < loops; i++) {  
19             Pthread_mutex_lock(&mutex);  
20             while (count == 0)  
21                 Pthread_cond_wait(&cond, &mutex);  
22             int tmp = get();  
23             Pthread_cond_signal(&cond);  
24             Pthread_mutex_unlock(&mutex);  
25             printf("%d\n", tmp);  
26         }  
27     }
```

// c1
// c2
// c3
// c4
// c5
// c6

Simple Example (broken solution 2-Trace)

Consumer 1

Consumer 2

Producer

T _{c1}	State	T _{c2}	State	T _p	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	Tc1 awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	Tc1 grabs data
c5	Running		Ready		Sleep	0	Oops woke Tc2

c1 relates to
lines in code
above

Simple Example (broken solution 2-Trace)

T_{c1}	State	T_{c2}	State	T_p	State	Count	Comment
...	(cont.)
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	Everyone asleep ...

A consumer should not wake other consumers, only producers, and vice-versa

Simple Example (Solution 3)

- Use two condition variables and while

- Producer threads wait on the condition empty, and signals fill.
- Consumer threads wait on fill and signal empty.

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);
8              while (count == 1)
9                  Pthread_cond_wait(&empty, &mutex);
10             put(i);
11             Pthread_cond_signal(&fill);
12             Pthread_mutex_unlock(&mutex);
13         }
14     }
15
```

Simple Example (Solution 3)

- Use two condition variables and while

- Producer threads wait on the condition empty, and signals fill.
- Consumer threads wait on fill and signal empty.

(Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);
20             while (count == 0)
21                 Pthread_cond_wait(&fill, &mutex);
22             int tmp = get();
23             Pthread_cond_signal(&empty);
24             Pthread_mutex_unlock(&mutex);
25             printf("%d\n", tmp);
26         }
27     }
```

Simple Example (Final Solution 4)

- More concurrency and efficiency -> Add more buffer slots.
 - Allow concurrent production or consuming to take place.
 - Reduce context switches.

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4  int count = 0;
5
6  void put(int value) {
7      buffer[fill] = value;
8      fill = (fill + 1) % MAX;
9      count++;
10 }
11
12 int get() {
13     int tmp = buffer[use];
14     use = (use + 1) % MAX;
15     count--;
16     return tmp;
17 }
```

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == MAX)                 // p2
9                  Pthread_cond_wait(&empty, &mutex); // p3
10             put(i);                               // p4
11             Pthread_cond_signal(&fill);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                    // c2
21                 Pthread_cond_wait(&fill, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&empty);          // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```


Semaphore

- An object with an integer value



Semaphore

- An object with an integer value

- We can manipulate with two routines; `sem_wait()` and `sem_post()`.
- Initialization

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1); // initialize s to the value 1
```

- Declare a semaphore `s` and initialize it to the value 1
- The second argument, 0, indicates that the semaphore is shared between *threads in the same process*.

Semaphore - Interacting

- `sem_wait()`

```
1  int sem_wait(sem_t *s) {  
2      decrement the value of semaphore s by one  
3      wait if value of semaphore s is negative  
4  }
```

- If the value of the semaphore was *one* or *higher* when called `sem_wait()`, **return right away.**
- It will cause the caller to suspend execution waiting for a subsequent post.
- When negative, the value of the semaphore is equal to the number of waiting threads.

Semaphore - Interacting

- `sem_post()`

```
1  int sem_post(sem_t *s) {  
2      increment the value of semaphore s by one  
3      if there are one or more threads waiting, wake one  
4  }
```

- Simply increments the value of the semaphore.
- If there is a thread waiting to be woken, wakes one of them up

Binary Semaphores (Locks)

- What should X be?
 - The initial value should be 1.

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize semaphore to X; what should X be?  
3  
4  sem_wait(&m);  
5  //critical section here  
6  sem_post(&m);
```

Thread Trace: Single Thread Using A Semaphore

Value of Semaphore	Thread 0
1	
1	call sema_wait()
0	sem_wait() returns
0	(crit sect)
0	call sem_post()
1	sem_post() returns

Thread Trace: Two Threads Using A Semaphore

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() retruns	Running		Ready
0	(crit set: begin)	Running		Ready
0	Interrupt; Switch → T1	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0)→sleep	sleeping
-1		Running	Switch → T0	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	Interrupt; Switch → T1	Ready		Running
0		Ready	sem_wait() retruns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

Semaphore as a Conditional Variable

can use semaphore to simulate conditional variables

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

parent: begin
child
parent: end

The execution result

A Parent Waiting For Its Child

- What should X be?

Semaphore as a Conditional Variable

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

parent: begin
child
parent: end

The execution result

A Parent Waiting For Its Child

- What should X be?
 - The value of semaphore should be set to is 0.

Thread Trace: Parent Waiting For Child (Case 1)

The parent call `sem_wait()` before the child has called `sem_post()`.

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	(sem < 0) → sleep	sleeping		Ready
-1	Switch → Child	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake(Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	Interrupt; Switch → Parent	Ready
0	<code>sem_wait()</code> retruns	Running		Ready

Thread Trace: Parent Waiting For Child (Case 2)

The child runs to completion before the parent call `sem_wait()`.

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	Interrupt; switch→Child	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment sem	Running
1		Ready	wake(nobody)	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	Interrupt; Switch→Parent	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement sem	Running		Ready
0	(sem<0)→awake	Running		Ready
0	<code>sem_wait()</code> retruns	Running		Ready

The Producer/Consumer (Bounded-Buffer) Problem

- **Producer: put() interface**

- Wait for a buffer to become empty in order to put data into it.

- **Consumer: get() interface**

- Wait for a buffer to become filled before using it.

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;    // line g2
13     return tmp;
14 }
```

The Producer/Consumer (Bounded-Buffer) Problem

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);           // line P1
8          put(i);                     // line P2
9          sem_post(&full);            // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);             // line C1
17         tmp = get();                 // line C2
18         sem_post(&empty);            // line C3
19         printf("%d\n", tmp);
20     }
21 }
22 ...
```

First Attempt: Adding the Full and Empty Conditions

The Producer/Consumer (Bounded-Buffer) Problem

```
21  int main(int argc, char *argv[]) {
22      // ...
23      sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
24      sem_init(&full, 0, 0);    // ... and 0 are full
25      // ...
26  }
```

First Attempt: Adding the Full and Empty Conditions (Cont.)

- **Imagine that MAX is greater than 1 .**
 - If there are multiple producers, **race condition** can happen at line *f1*.
 - It means that the old data there is overwritten.
- **We've forgotten here is mutual exclusion.**
 - The filling of a buffer and incrementing of the index into the buffer is a **critical section**.

Adding Mutual Exclusion

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);           // line p0 (NEW LINE)
9          sem_wait(&empty);           // line p1
10         put(i);                     // line p2
11         sem_post(&full);             // line p3
12         sem_post(&mutex);           // line p4 (NEW LINE)
13     }
14 }
15
(Cont.)
```

Adding Mutual Exclusion (Incorrectly)

Adding Mutual Exclusion

(Cont.)

```
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&mutex);           // line c0 (NEW LINE)
20         sem_wait(&full);           // line c1
21         int tmp = get();           // line c2
22         sem_post(&empty);         // line c3
23         sem_post(&mutex);         // line c4 (NEW LINE)
24         printf("%d\n", tmp);
25     }
26 }
```

Adding Mutual Exclusion (Incorrectly)

Adding Mutual Exclusion

```
1  sem_t empty;  
2  sem_t full;  
3  sem_t mutex;  
4  
5  void *producer(void *arg) {  
6      int i;  
7      for (i = 0; i < loops; i++) {  
8          sem_wait(&empty);  
9          sem_wait(&mutex);  
10         put(i);  
11         sem_post(&mutex);  
12         sem_post(&full);  
13     }  
14 }  
15  
(Cont.)
```

// line p1
// line p1.5 (MOVED MUTEX HERE...)
// line p2
// line p2.5 (... AND HERE)
// line p3

Adding Mutual Exclusion (Correctly)

Adding Mutual Exclusion

(Cont.)

```
16  void *consumer(void *arg) {
17      int i;
18      for (i = 0; i < loops; i++) {
19          sem_wait(&full);
20          sem_wait(&mutex);
21          int tmp = get();
22          sem_post(&mutex);
23          sem_post(&empty);
24          printf("%d\n", tmp);
25      }
26  }
27
28  int main(int argc, char *argv[]) {
29      // ...
30      sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with ...
31      sem_init(&full, 0, 0);    // ... and 0 are full
32      sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock
33      // ...
34  }
```

// line c1
// line c1.5 (MOVED MUTEX HERE...)
// line c2
// line c2.5 (... AND HERE)
// line c3

Adding Mutual Exclusion (Correctly)

How To Implement Semaphores

```
1  typedef struct __Zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  } Zem_t;
6
7  // only one thread can call this
8  void Zem_init(Zem_t *s, int value) {
9      s->value = value;
10     Cond_init(&s->cond);
11     Mutex_init(&s->lock);
12 }
13
14 void Zem_wait(Zem_t *s) {
15     Mutex_lock(&s->lock);
16     while (s->value <= 0)
17         Cond_wait(&s->cond, &s->lock);
18     s->value--;
19     Mutex_unlock(&s->lock);
20 }
21 ...
```

Build our own version of semaphores called Zemaphores

How To Implement Semaphores

```
22  void Zem_post(Zem_t *s) {  
23      Mutex_lock(&s->lock);  
24      s->value++;  
25      Cond_signal(&s->cond);  
26      Mutex_unlock(&s->lock);  
27  }
```

Build our own version of semaphores called Zemaphores

- **Zemaphore don't maintain the invariant that the value of the semaphore.**
 - The value never be lower than zero.
 - This behavior is easier to implement and matches the current Linux implementation.